

The Riemann Hypothesis: Foundations, Obstructions, and Reorientation

Lukas Geiger

Bernau, Germany

ORCID: [0009-0005-7296-1534](https://orcid.org/0009-0005-7296-1534)

May 31, 2026

Abstract

This paper, the first in a five-part series on the Riemann Hypothesis, establishes the mathematical landscape and documents the systematic exploration—and ultimately the failure—of a gauge-theoretic approach. In the first part, we develop the analytical framework: the Riemann Hypothesis as a stability problem, the Li criterion as the central tool, and the Bombieri–Lagarias decomposition of the Li coefficients into archimedean and finite-part contributions. Starting from a thermodynamic formalism (Prime Hub Graphs and Ruelle transfer operators), we prove several unconditional structural results: the archimedean–finite decomposition, the archimedean dominance paradox, OS reflection positivity, the arithmetic uncertainty relation, and the spectral census of the arithmetic kernel.

In the second part, we identify four critical obstructions (K1–K4) that individually and collectively doom the original proof strategy: the Gibbs normalization gap, the sign change of the target functional, the wrong monotonicity direction, and the trace-class obstruction. We localize the difficulty precisely through a nine-step logical chain, showing that the *entire* content of RH concentrates in a single step (S5: the unconditional positivity of the Li coefficients). We formulate five concrete open problems and identify the reorientation toward Connes’ spectral program as the path forward. Part II develops this new direction via the Shift Parity Lemma, finite-range diagnostics, and a conditional even-dominance reduction; Part III surveys the spectral routes; Part IV collects cross-route synthesis and remaining paths; Part V draws the final conclusions. In v2.0, two non-existence theorems in Part II (NE-A, NE-B) formally rule out the absolute total-positivity and universal-commuting-operator routes, reinforcing the arithmetic-statistical nature of even dominance.

Series status (v3.0, 2026-05-26). Following external critical review, the series status is revised from “unconditional A1–A8” (v2.1) to *conditional reduction*: even dominance is rigorously certified on the finite range $100 \leq \lambda \leq 1,300,000$ via separate IA-upper/lower bounds for $\lambda_1^\pm$; the asymptotic argument establishes only the variational statement $\lambda_{\min}(W^+ - W^-) = -\Theta(\sqrt{\lambda})$ and the eigenvalue inequality $\lambda_1^+(\lambda) < \lambda_1^-(\lambda)$ asymptotically remains conditional on the *Asymptotic Variational Gap Conjecture* for λ_1^- (Part II, “Status of the proof” remark, v2.2 revision; companion paper [10]). The Shift Parity Lemma, Leading-Mode Cancellation, the 33 finite-range diagnostics, and the non-existence theorems NE-A and NE-B are unaffected by this revision. The series is a continuously evolving draft/preprint; v3.0 restructures the trilogy into a five-part series (structural reorganization only; the conditional reduction status from v2.2 is unchanged).

MSC 2020: 11M26, 11M36, 37C30, 47A75, 82B05

Keywords: Riemann Hypothesis, Li coefficients, Ruelle transfer operator, archimedean dominance, gauge positivity, obstructions, spectral analysis

Contents

The Landscape	4
1 Introduction	4
1.1 The Hope	4
2 The Thermodynamic Formalism	4
2.1 Prime Hub Graph and the Vacuum Trace	4
2.2 The MEPP Heuristic	5
3 The Li Coefficients: Analytical Framework	5
3.1 Bombieri–Lagarias Representation	5
3.2 Archimedean–Finite Decomposition	6
4 The Archimedean Dominance Paradox	6
5 Structural Results	7
5.1 Prime-Sum Functionals and the Gibbs Framework	7
5.2 Non-Identification Theorem	7
5.3 Variance Lower Bound	7
5.4 OS Reflection Positivity	7
5.5 Arithmetic Uncertainty Relation	8
5.6 Strict Convexity and the Spectral Census	8
6 The Gauge Equivalence Framework	8
7 Summary: The Starting Position	9
The Dead Ends	9
8 The Original Strategy	9
9 Dead End 1: The Gibbs Normalization Gap (K1)	9
10 Dead End 2: Global Gauge Positivity Fails (K2)	10
11 Dead End 3: Monotonicity Direction Wrong (K3)	10
12 Dead End 4: The Stieltjes Realization and A8 (K4)	11
12.1 Path D2: Stieltjes Realization	11
12.2 K4b: Inconsistency of the Original A8	11
12.3 The Trace-Class Obstruction	11
13 Route Analysis: What Survives	11
14 The Localization Table	11
15 Five Open Problems	12

16 Reorientation: Connes' Spectral Program	12
17 Conclusion	14

The Landscape

1 Introduction

The Riemann Hypothesis (RH) asserts that all non-trivial zeros of the Riemann zeta function $\zeta(s)$ lie on the critical line $\Re(s) = 1/2$. Reformulated via Li's criterion [1], RH is equivalent to the positivity of an infinite sequence of real numbers:

$$\text{RH} \iff \lambda_n \geq 0 \text{ for all } n \geq 1, \quad (1)$$

where $\lambda_n = \sum_{\rho} [1 - (1 - 1/\rho)^n]$ and the sum runs over the non-trivial zeros of ζ .

This paper is the first in a five-part series:

- **Part I** (this paper): Foundations, obstructions, and reorientation — mathematical foundations, analytical tools, structural results, the documentation of four dead ends, the localization of difficulty, and the reorientation toward Connes' spectral program.
- **Part II** [9]: Even dominance — the Shift Parity Lemma, computer-assisted finite-range diagnostics, and the conditional reduction architecture.
- **Part III** [12]: Spectral routes — Branch B (I-1 normal-family / err-collapse) and Branch Z (CCM microcluster programme), including spectral dead ends.
- **Part IV** [13]: Cross-route synthesis and remaining paths — inter-branch diagnostics, dormant and external routes, and methodological dead ends.
- **Part V** [14]: Conclusio — what is proven, what is excluded, and what remains.

The series documents an honest research journey: from a promising gauge-theoretic approach (Sections 2–6), through its failure at multiple points (Sections 9–12), to a new structural understanding of the problem via the Weil quadratic form (Section 16), and a final accounting of what has been achieved (Part V).

1.1 The Hope

The Hilbert–Pólya conjecture proposes that the non-trivial zeros correspond to eigenvalues of a self-adjoint operator. The Berry–Keating semiclassical approach ($\hat{H} = xp$) fails due to the Weyl law obstruction. This paper shifts the paradigm from quantum mechanics to statistical mechanics: the zeros are generated not as eigenvalues of a Hamiltonian, but as resonances of a Ruelle-type transfer operator.

The guiding idea is that the Riemann Hypothesis is a *stability condition*: the arithmetic structure of the primes, encoded in the Li coefficients λ_n , is stabilized by a competition between arithmetic instability (from the primes) and archimedean convexity (from the Gamma function). Understanding this competition is the central theme of the series.

2 The Thermodynamic Formalism

2.1 Prime Hub Graph and the Vacuum Trace

To reproduce the Euler product $\zeta(s) = \prod_p (1 - p^{-s})^{-1}$, we construct a dynamical system whose primitive periodic orbits correspond to primes.

Definition 2.1 (Prime Orbit Axiom). We postulate a topologically mixing directed graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ whose primitive cycles γ_p are in bijection with the primes $p \in \mathcal{P}$, with geometric length $\ell(\gamma_p) = \log p$.

Remark 2.2 (Motivation vs. logical dependency). The Prime Orbit Axiom is a constructive motivation. The analytical results in subsequent sections depend solely on properties of $\zeta(s)$ and the arithmetic Euler product, and are logically independent of this graph construction.

Using a holonomy filter on the free monoid generated by $\mathcal{P}$, we define a transfer operator $\mathcal{L}_s$ and prove:

Theorem 2.3 (Recovery of the Zeta Function). *For $\Re(s) > 1$, the Fredholm determinant of the vacuum-projected transfer operator satisfies:*

$$\det(I - \mathcal{L}_s)^{\text{vac}} = \prod_p (1 - p^{-s}) = \frac{1}{\zeta(s)}.$$

Proof. By the Ruelle identity [3], $\log \det(I - \mathcal{L}_s) = -\sum_{m=1}^{\infty} \frac{1}{m} \text{Tr}(\mathcal{L}_s^m)$. The vacuum projection restricts the trace to configurations that form closed loops of prime powers: the only periodic orbits of length $m \log p$ in $\mathcal{G}$ are the m -fold traversals of the primitive cycle γ_p , each contributing weight p^{-ms} with multiplicity $\log p$ (the geometric length). Summing over prime powers and exponentiating:

$$\log \det(I - \mathcal{L}_s)^{\text{vac}} = -\sum_p \sum_{m=1}^{\infty} \frac{1}{m} p^{-ms} = \sum_p \log(1 - p^{-s}).$$

This is the logarithm of the Euler product $\prod_p (1 - p^{-s}) = 1/\zeta(s)$, valid for $\Re(s) > 1$ by absolute convergence. $\square$

2.2 The MEPP Heuristic

The Maximum Entropy Production Principle (MEPP) interprets the critical line $\Re(s) = 1/2$ as a detailed-balance condition:

Heuristic Principle 2.4 (Model RH via MEPP). If the primes represent the optimally stable configuration of a “prime gas” subject to the PNT density constraint, the resulting Gibbs measure satisfies time-reversal symmetry, forcing zeros onto the critical line.

This heuristic does not constitute a proof but motivates the thermodynamic perspective developed below.

3 The Li Coefficients: Analytical Framework

3.1 Bombieri–Lagarias Representation

The completed xi function $\xi(s) = \frac{1}{2}s(s-1)\pi^{-s/2}\Gamma(s/2)\zeta(s)$ provides the definitive representation:

Theorem 3.1 (Bombieri–Lagarias [2]). *For $n \geq 1$:*

$$\lambda_n = \frac{1}{(n-1)!} \frac{\partial^n}{\partial s^n} \left[s^{n-1} \log \xi(s) \right]_{s=1}.$$

Since $\xi(s)$ is entire of order 1 with $\xi(1) = 1/2 \neq 0$, $\log \xi(s)$ is holomorphic near $s = 1$ and the derivative is well-defined.

3.2 Archimedean–Finite Decomposition

Proposition 3.2 (Archimedean–Finite Decomposition). *Each Li coefficient decomposes as $\lambda_n = \lambda_n^{(\infty)} + \lambda_n^{(\text{fin})}$, where:*

$$\lambda_n^{(\infty)} = \frac{1}{(n-1)!} \frac{\partial^n}{\partial s^n} \left[s^{n-1} \log\left(\frac{1}{2}s \pi^{-s/2} \Gamma(s/2)\right) \right]_{s=1}, \quad (2)$$

$$\lambda_n^{(\text{fin})} = \frac{1}{(n-1)!} \frac{\partial^n}{\partial s^n} \left[s^{n-1} \log g(s) \right]_{s=1}, \quad g(s) := (s-1)\zeta(s). \quad (3)$$

Both parts are individually computable; $g(s)$ is holomorphic at $s=1$ with $g(1)=1$.

Remark 3.3 (Three-regime structure). Numerical computation reveals:

- (a) **Small n** ($n \lesssim 8$): $\lambda_n^{(\infty)} < 0$, compensated by $\lambda_n^{(\text{fin})} > 0$. Positivity is a delicate cancellation.
- (b) **Transition** ($n \approx 8$): $\lambda_n^{(\infty)}$ crosses zero.
- (c) **Large n** ($n \gg 10$): $\lambda_n^{(\infty)} \sim \frac{n}{2} \log n$ dominates; positivity is automatic.

The “hard” instances of Li’s criterion concentrate in regime (a). Table 1 shows the decomposition for $n = 1, \dots, 13$.

n	λ_n	$\lambda_n^{(\infty)}$	$\lambda_n^{(\text{fin})}$	$c_n = \lambda_n/n$
1	0.023	−0.554	0.577	0.023
2	0.092	−0.875	0.967	0.046
3	0.208	−1.013	1.221	0.069
5	0.576	−0.883	1.458	0.115
8	1.466	+0.021	1.445	0.183
10	2.279	+0.956	1.324	0.228
13	3.817	+2.725	1.093	0.294

Table 1: Archimedean–finite decomposition of the Li coefficients. The sequence $c_n = \lambda_n/n$ is strictly increasing, ruling out complete monotonicity.

4 The Archimedean Dominance Paradox

Proposition 4.1 (Stability Decomposition). *The global curvature $\mathcal{M}(\sigma) = \partial_s^2 \log \xi(s)$ decomposes as $\mathcal{M} = \mathcal{M}_{\text{arith}} + \mathcal{M}_{\text{arch}}$:*

1. **Arithmetic instability:** $\mathcal{M}_{\text{arith}}(1) = \sum_{\rho} \Re(1/\rho^2) < 0$.
2. **Archimedean stability:** $\mathcal{M}_{\text{arch}}(\sigma) > 0$ for $\sigma > 1/2$ (explicitly positive and convex).

The Riemann Hypothesis is equivalent to the archimedean convexity globally dominating the arithmetic concavity for all $\sigma > 1/2$. This “dominance paradox” — a manifestly positive continuous term must overcome infinitely many oscillatory arithmetic contributions — is the structural core of the problem.

5 Structural Results

We establish five unconditional results that constrain the problem and form the toolkit for subsequent sections and papers.

5.1 Prime-Sum Functionals and the Gibbs Framework

Definition 5.1 (Arithmetic Gibbs Measure). The Gibbs measure is $\mu_\sigma(p) = p^{-\sigma}/P(\sigma)$ with $P(\sigma) = \sum_p p^{-\sigma}$. The prime-sum functionals are:

$$\begin{aligned} A(\sigma) &:= \sum_p \frac{\log p}{p^\sigma(p^\sigma - 1)}, & B(\sigma) &:= \sum_p \frac{\log p}{p^\sigma}, \\ C(\sigma) &:= \sum_p \frac{(\log p)^2}{p^\sigma(p^\sigma - 1)}, & D(\sigma) &:= \sum_p \frac{(\log p)^2}{(p^\sigma - 1)^2}. \end{aligned}$$

Lemma 5.2 (Finite-Part Convergence). $A(\sigma)$, $C(\sigma)$, and $D(\sigma)$ converge absolutely at $\sigma = 1$ and are right-continuous. $B(\sigma)$ and $P(\sigma)$ diverge as $\sim \log(1/(\sigma - 1))$.

Remark 5.3 (Gibbs normalization gap). The Gibbs-normalized expectation $-A(\sigma)/P(\sigma)$ loses the finite-part information as $\sigma \rightarrow 1^+$: the $1/P(\sigma)$ normalization annihilates $\lambda_n^{(\text{fin})}$. This is the first hint that the thermodynamic approach has structural limitations (see Section 9).

5.2 Non-Identification Theorem

Proposition 5.4 (Non-Identification). The excess free energy $I_n := \lim_{\sigma \rightarrow 1^+} \Phi_n(\sigma)$ and the Li coefficient λ_n are generically not equal: $I_n \neq \lambda_n$. The former is a KL divergence (tilt cost), the latter a spectral trace of $\log \xi$.

5.3 Variance Lower Bound

Proposition 5.5 (Variance Bound). $I_1 \geq \frac{1}{2}V_1$ where $V_1 = \lim_{\sigma \rightarrow 1^+} \text{Var}_{\mu_\sigma}(H_{1,\sigma}) > 0$. Numerically, $I_1 \geq 0.034 > \lambda_1 \approx 0.023$.

5.4 OS Reflection Positivity

Proposition 5.6 (OS Reflection Positivity). The Osterwalder–Schrader form

$$K(t, u) = \sum_{p \in \mathcal{P}} w_p e^{-i(u-t) \log p}, \quad w_p = \frac{p^{-(\sigma+1)}}{P(\sigma)} > 0,$$

is positive definite on $L_+^2(\mathbb{R})$. The OS Hilbert space carries a contractive semigroup $e^{-\tau H}$ with $H = \text{diag}(\log p)$.

Proof. By Bochner's theorem: $K(t, u) = \hat{\mu}(u-t)$ with $\mu = \sum_p w_p \delta_{\log p}$ a positive measure, so $\iint \bar{f}(t)f(u)K(t, u) dt du = \sum_p w_p |\hat{f}(\log p)|^2 \geq 0$. $\square$

5.5 Arithmetic Uncertainty Relation

Heuristic Principle 5.7 (Arithmetic Uncertainty Relation). For normalized $\psi \in \ell^2(\mathcal{P})$, define the prime-domain spread Δ_D via

$$\Delta_D^2 = \sum_p |a_p|^2 (\log p - \overline{\log p})^2$$

and the zero-domain spread Δ_t via the Fourier transform $\hat{\psi}(t) = \sum_p a_p e^{-it \log p}$. Then $\Delta_D \cdot \Delta_t \geq 1/2$.

Remark 5.8. This follows formally from the substitution $x_p = \log p$ and the Heisenberg–Kennard–Weyl inequality on $L^2(\mathbb{R})$. However, ψ is supported on the discrete set $\{\log p : p \in \mathcal{P}\}$, not on all of $\mathbb{R}$, so the reduction requires an embedding into $L^2(\mathbb{R})$ whose properties depend on the distribution of primes. We state this as a heuristic principle, not a theorem.

The physical content is a consistency check: a zero off the critical line would require a “localized phase conspiracy” among finitely many primes, which the arithmetic incommensurability of $\{\log p\}$ makes implausible but does not rigorously exclude.

5.6 Strict Convexity and the Spectral Census

Proposition 5.9 (Strict Convexity). $F(\sigma) := -\log \zeta(\sigma)$ is strictly convex on $(1, \infty)$: $F''(\sigma) = \sum_p (\log p)^2 p^{-\sigma} / (1 - p^{-\sigma})^2 > 0$.

Proposition 5.10 (Spectral Census). For the symmetrized arithmetic kernel K_σ^{sym} truncated to $N \geq 300$ primes and $\sigma \in (1, 5)$, the negative inertia index is exactly 2, independent of N and σ .

6 The Gauge Equivalence Framework

The spectral census (Theorem 5.10) suggests a gauge-correction approach: can a rank-2 correction restore positive-semidefiniteness?

Definition 6.1 (Admissible Gauge). A symmetric kernel G_σ on $\mathcal{P} \times \mathcal{P}$ is admissible if $\sum_{p,q} G_\sigma(p, q) \mu_\sigma(p) \mu_\sigma(q) = 0$.

Theorem 6.2 (Gauge Equivalence). For fixed $\sigma > 1$, the following are equivalent:

- (i) The target functional $F(\sigma) = A(\sigma)B(\sigma) - P(\sigma)[C(\sigma) + D(\sigma)] \geq 0$.
- (ii) There exists an admissible gauge making $K_\sigma^{\text{sym}} + G_\sigma \succeq 0$.

This tautological equivalence provides the formal structure for the gauge-theoretic approach. Whether it leads anywhere depends on the *global* behavior of $F(\sigma)$ — a question addressed in Section 10, where we discover that $F(\sigma) < 0$ for $\sigma > 1.14$.

7 Summary: The Starting Position

At the end of the landscape analysis, we have assembled the following toolkit:

Result	Status
R1: Bombieri–Lagarias representation	proven
R2: Archimedean–finite decomposition	proven
R3: Non-identification ($I_n \neq \lambda_n$)	proven
R4: Variance lower bound ($I_1 \geq \frac{1}{2}V_1$)	proven
R5: Arithmetic uncertainty relation	heuristic
R6: Strict convexity of $-\log \zeta$	proven
R7: OS reflection positivity	proven
R8: Spectral census (index = 2)	numerical
R9: Gauge equivalence (tautological)	proven

The hope is clear: use the gauge framework (R9) together with the spectral census (R8) to construct a global certificate proving $\lambda_n \geq 0$. The following sections will show that this hope is dashed by a sign change at $\sigma_0 \approx 1.14$ — but that a completely different approach, through Connes’ Weil quadratic form, opens a new and more promising path.

The Dead Ends

8 The Original Strategy

The toolkit of nine structural results (R1–R9) and the gauge equivalence framework established above suggest a clear proof strategy:

1. Use the spectral census (R8: negative inertia index = 2) to identify the unstable modes.
2. Construct an explicit rank-2 gauge correction restoring positive-semidefiniteness.
3. Apply the gauge equivalence theorem (R9) to deduce $\lambda_n \geq 0$.

This strategy fails at step 2. In this part, we document four independent obstructions that kill it, then use the resulting understanding to localize the difficulty precisely.

9 Dead End 1: The Gibbs Normalization Gap (K1)

The thermodynamic formulation defines the renormalized constraint $m_n(\sigma) = \mathbb{E}_{\mu_\sigma}[X_{n,\sigma}] - h_n(\sigma)$, with the hope that $\lim_{\sigma \rightarrow 1^+} m_n(\sigma) = \lambda_n$.

Proposition 9.1 (K1: Critical Identification Fails). *The Gibbs-normalized expectation $m_1(\sigma) = -A(\sigma)/P(\sigma) - h_1(\sigma)$ does not converge to λ_1 as $\sigma \rightarrow 1^+$. The $1/P(\sigma)$ -normalization annihilates the finite-part contribution $\lambda_1^{(\text{fin})} = \gamma \approx 0.577$.*

Proof. The numerator sum $S(\sigma) = \sum_p (\log p) p^{-\sigma} f(p, \sigma)$ converges to a finite limit, while $P(\sigma) \rightarrow \infty$. Therefore $-S(\sigma)/P(\sigma) \rightarrow 0$. The Li coefficient λ_1 arises from $\frac{d}{ds} \log \xi(s)|_{s=1}$, which involves $\zeta'(s)/\zeta(s) + 1/(s-1) \rightarrow \gamma$ — a cancellation that occurs *before* Gibbs normalization, not after. $\square$

Remark 9.2 (Consequence). The excess free energy I_n and the Li coefficient λ_n are structurally distinct objects (Theorem 5.4). Even $I_n \geq 0$ (which holds unconditionally as a KL divergence) does not imply $\lambda_n \geq 0$. This is the “convexity trap”: by the Bregman inequality, $I_n \geq 0$ is a universal identity for any strictly convex generating function, independent of $\text{sign}(\lambda_n)$.

10 Dead End 2: Global Gauge Positivity Fails (K2)

The gauge equivalence theorem (R9) requires $F(\sigma) \geq 0$ for all $\sigma > 1$.

Proposition 10.1 (K2: Sign Change of $F(\sigma)$). *The target functional $F(\sigma) = A(\sigma)B(\sigma) - P(\sigma)[C(\sigma) + D(\sigma)]$ satisfies:*

- $F(\sigma) > 0$ for $\sigma \in (1, \sigma_0)$ with $\sigma_0 \approx 1.14$,
- $F(\sigma) < 0$ for all tested $\sigma > \sigma_0$ (minimum $F \approx -0.184$ at $\sigma \approx 1.28$).

This sign change is independent of RH.

Proof. Near $\sigma = 1 + \varepsilon$: $A \cdot B \sim 1/\varepsilon$ dominates $P(C + D) \sim \log(1/\varepsilon)$, so $F > 0$. For larger σ : $B(\sigma)$ decays faster than $P(\sigma)$, and the product AB falls below $P(C + D)$. Numerical computation with $N = 10000$ primes and Rosser–Schoenfeld tail bounds confirms the transition at $\sigma_0 \approx 1.14$. $\square$

Remark 10.2 (Consequence). The claim “RH $\Leftrightarrow F(\sigma) \geq 0$ for all $\sigma > 1$ ” is **false**. The gauge approach works only in the boundary layer (1, 1.14), not globally. No admissible gauge can make K_σ^{sym} positive-semidefinite for $\sigma > 1.14$.

11 Dead End 3: Monotonicity Direction Wrong (K3)

The original argument assumed that $\frac{d}{ds} \log \xi(\sigma)$ is decreasing, making λ_1 the maximum.

Proposition 11.1 (K3: Convexity, Not Concavity). $\frac{d^2}{ds^2} \log \xi(\sigma) \approx +0.046 > 0$ at $\sigma = 1$. *The function $\log \xi$ is convex near $s = 1$, not concave. Therefore λ_1 is the minimum of the derivative profile, not the maximum.*

Remark 11.2 (Consequence). This reverses the logic of the monotonicity argument: instead of showing $\lambda_1 \geq 0$ by bounding the maximum from below, one would need to bound the minimum from below — a strictly harder problem.

12 Dead End 4: The Stieltjes Realization and A8 (K4)

12.1 Path D2: Stieltjes Realization

Proposition 12.1 (K4a: Complete Monotonicity Fails). *The sequence $c_n = \lambda_n/n$ is strictly increasing: $c_1 \approx 0.023$, $c_{13} \approx 0.294$, with $c_n \sim \frac{1}{2} \log n$ asymptotically. This violates complete monotonicity at $k = 1$ ($\Delta c_n > 0$, so $(-1)^1 \Delta c_n < 0$). The Stieltjes realization approach with a measure on $[0, 1]$ is ruled out.*

Under RH, the spectral measure lives on the unit circle $|w_\rho| = 1$, not $[0, 1]$. The positivity of c_n is an essentially number-theoretic property.

12.2 K4b: Inconsistency of the Original A8

Proposition 12.2 (Hardy Contradiction). *The original axiom A8 posits $\mathcal{K}_0 \geq 0$ trace-class with $\xi(1/2 + it) = C \cdot \det(I + V_t \mathcal{K}_0 V_t^*)$. But $\mathcal{K}_0 \geq 0$ implies $\det(I + V_t \mathcal{K}_0 V_t^*) \geq 1 > 0$ for all real t . This contradicts Hardy's theorem (1914): ξ has infinitely many zeros on the critical line.*

12.3 The Trace-Class Obstruction

The corrected axiom A8' requires $\xi(s)/\xi(0) = \det(I - \mathcal{L}_s)$ with $\mathcal{L}_s$ nuclear. However, the Euler product gives $\mathcal{L}_s = \text{diag}(p^{-s})$ on $\ell^2(\mathcal{P})$, which is trace-class only for $\Re(s) > 1$. At the critical line, $\sum_p p^{-1/2} = +\infty$.

For the Selberg zeta function, the analogous construction works because geodesic lengths grow exponentially ($\#\{\gamma : l_\gamma \leq T\} \sim e^T/T$), while prime “lengths” $\log p$ are too dense ($\pi(x) \sim x/\log x$). This density gap is the fundamental obstruction.

13 Route Analysis: What Survives

The axiom system (defined in Part II) admits two logical routes to RH:

Route 1 (A7 + A8'): A7 (OS reflection positivity) is **proven**. A8' (nuclear Fredholm determinant for ξ) remains **open** and faces the trace-class obstruction.

Route 2 (A9): Dead. A9(i–ii) hold, but A9(iii) is broken by K1 and the convexity trap.

14 The Localization Table

We trace the logical chain from definitions to RH, identifying exactly where the difficulty concentrates.

Step	Statement	Status	Method
S1	Define $\xi(s)$	proven	definition
S2	λ_n via Bombieri–Lagarias	proven	[2]
S3	$\lambda_n = \lambda_n^{(\infty)} + \lambda_n^{(\text{fin})}$	proven	Section 3
S4	$\lambda_n^{(\infty)} \sim \frac{n}{2} \log n$ for large n	proven	Stirling
S5	$\lambda_n \geq 0$ for all $n \geq 1$	OPEN	$\iff$ RH
S6	All zeros on $\Re(s) = 1/2$	$\iff$ S5	Li (1997)
S7	$\pi(x) = \text{Li}(x) + O(\sqrt{x} \log x)$	$\Leftarrow$ S6	von Koch
S8	Spectral census (index = 2)	numerical	Section 5
S9	Gauge equivalence (tautological)	proven	Section 6

Conclusion: The *entire* content of the Riemann Hypothesis concentrates in Step S5. Everything before it is unconditional; everything after it is either equivalent or a consequence. The present program has not reduced S5 to a simpler statement.

15 Five Open Problems

Based on our analysis, we formulate five concrete open problems ordered by estimated difficulty:

Conjecture 15.1 (OP1: Analytical Proof of Index Rigidity). *The negative inertia index of K_σ^{sym} is exactly 2 for all N sufficiently large and $\sigma \in (1, \sigma_{\max})$. A proof should exploit the rank-2 hub structure of the transfer operator.*

Conjecture 15.2 (OP2: Uniform Bound on $\lambda_n^{(\text{fin})}$). *$\lambda_n^{(\text{fin})} > 0$ for all $n \geq 1$. This is supported by the numerical observation that $\lambda_n^{(\text{fin})}$ is always positive with slow algebraic decay.*

Conjecture 15.3 (OP3: Growth Rate). *$\lambda_n^{(\infty)} \sim \frac{n}{2} \log n$ with explicit error bounds, uniform in n . This would settle RH for all n above an explicit threshold.*

Conjecture 15.4 (OP4: Small- n Positivity). *$\lambda_n > 0$ for $n = 1, \dots, n_0$ (some explicit n_0) unconditionally — without assuming RH. The first 10^4 coefficients are known to be positive [4].*

Conjecture 15.5 (OP5: Nuclear Transfer Operator for ξ). *Construct a separable Hilbert space $\mathcal{V}$ and a nuclear family $s \mapsto \mathcal{L}_s$ with $\xi(s)/\xi(0) = \det(I - \mathcal{L}_s)$, spectral radius < 1 for $\Re(s) > 1/2$, and eigenvalue 1 at the zeros.*

OP5 is essentially the Hilbert–Pólya program and connects to Deninger’s foliated spaces [8] and Connes’ adelic approach [5].

16 Reorientation: Connes’ Spectral Program

The failure of the gauge-theoretic approach (K1–K4) forces a fundamental reorientation. A recent survey by Connes [6] provides a new path: the Weil quadratic form QW_λ on $L^2([-\log \lambda, \log \lambda])$, constructed from the Weil explicit formula, defines a self-adjoint operator A_λ whose spectral properties are directly linked to the zeros of ζ .

Theorem 16.1 (Connes, Theorem 6.1 in [6]; proved in [7]). *If the minimum eigenvalue of A_λ is simple with even eigenfunction, then the Fourier transform of the eigenfunction has all zeros on the real line.*

Connes’ Section 6.6 identifies the verification of this *even dominance* property as the remaining open step. Crucially, Connes’ Hurwitz argument requires even dominance only along a sequence $\lambda_n \rightarrow \infty$, not for all λ .

This is where the thermodynamic insight pays off in an unexpected way: the decomposition analysis from the landscape sections (the prime-sum functionals, the trigonometric structure of shift operators) becomes the key to understanding the even-dominance route. Part II establishes the Shift Parity Lemma, the frontier-prime mechanism, and finite-range computer-assisted gap diagnostics for 33 values of λ up to 1,300,000, now treated as a conditional finite bridge until the odd-sector lower bound is independently validated.

Remark 16.2 (From despair to hope). The gauge-theoretic approach failed because it tried to prove a *global* statement ($\lambda_n \geq 0$ for all n) via a *local* tool (the Gibbs measure near $\sigma = 1$). The Connes approach succeeds where the gauge approach fails because it works directly in the spectral domain of A_λ , where the even-odd decomposition provides a natural structure. The Shift Parity Lemma (Part II) shows that this structure is fundamentally algebraic, not analytic.

Three Independent Branches Attack RH from the Same Reorientation

The spectral reorientation of Connes’ programme opens *three* structurally independent attacks on RH, all of which appear in the present series. They share the same external preprint inputs (Connes 2026 [6], Connes–vS 2025 [7], and the Connes–Consani–Moscovici 2025 zeta-spectral-triples construction) and meet at the same analytical wall — Connes’ §6.6(b)/MS2 condition on the limit behaviour of $\{\hat{\xi}_\lambda\}$ — but they reach the wall through different mechanisms.

- **Branch A — Even Dominance.** Reduces RH to the eigenvalue inequality $\lambda_1^+(\lambda) < \lambda_1^-(\lambda)$ along a sequence $\lambda_n \rightarrow \infty$. Developed in Part II of this series; standalone extraction in the companion paper [10]. Currently a conditional reduction (open: A3 odd-sector lower bound, A7 Asymptotic Variational Gap Conjecture).
- **Branch B — I-1 Normal-Family / err-Collapse.** Reduces RH to a single uniform growth-adapted strip bound (ii-a) on $\{\hat{\xi}_\lambda\}$ from QW_N -coercivity. Developed in Part III of this series, with reduction chain B1–B11, Foundation-Lock $\hat{\xi}(z) = (2/\sqrt{L}) \sin(zL/2)F(z)$, and converged Mac Studio empirics ($A_\lambda(0.4) \approx 1.004$ flat across $\lambda \in \{3, 5, 7, 9, 11, 13, 15\}$).
- **Branch Z — CCM Microcluster Programme.** Reduces RH (via the err-collapse coupling and MS2-closure) to three named conditional inputs C2ca.1, C2ca.2, C2ca.5 plus a fixed-waterline outer-gap assumption $g_* \geq g_0$. Developed in the standalone companion paper [11]; summarised in Part III.

The three branches are pursued in parallel because each remaining open block sits at a structurally different place: A7 is variational, (ii-a) is analytical (strip bound from coercivity), C2ca is operator-theoretic (quasimode-to-projector). A success on any one

of the three would suffice. The audit-trail philosophy of this paper preserves all three with explicit status, including a catalogue of dead routes already eliminated within each branch (Parts III–IV).

17 Conclusion

This paper has documented both a toolkit and an honest failure. The gauge-theoretic approach to the Riemann Hypothesis produced genuine structural insights (R1–R9) but cannot bridge the gap between the local thermodynamic analysis and the global spectral positivity required by Li’s criterion. Four independent obstructions (K1–K4) kill the original strategy, and the localization table shows that no reduction to a simpler problem has been achieved: the difficulty remains concentrated in Step S5.

However, the failure is productive:

1. The landscape analysis (Sections 2–6) establishes a permanent toolkit of unconditional structural results.
2. The dead ends (K1–K4) are clearly mapped, preventing others from pursuing them.
3. The reorientation to Connes’ program leverages the thermodynamic insight (prime shift operators, trigonometric structure) in a new context.
4. The understanding that even dominance is driven by prime contributions, not archimedean effects, emerged directly from the decomposition analysis.

Part II takes up the new thread: the Shift Parity Lemma, the computer-assisted finite-range diagnostics, and the identification of the cumulative eigenvalue-ordering step as the remaining challenge. Parts III–IV document the spectral routes, cross-route synthesis, and remaining paths; Part V draws the final conclusions on what has been proven, what has been excluded, and what remains open.

AI Disclosure

This manuscript was developed in extensive collaboration with AI language models (Claude, Anthropic; Copilot, Microsoft/GitHub; Gemini, Google DeepMind). AI contributed to conceptual exploration, analytical work, literature review, writing, and critical review. The author, Lukas Geiger, conceived the research direction, provided domain expertise, and exercised final editorial authority. The author assumes full and sole scholarly responsibility for all content, including any errors.

References

- [1] X.-J. Li, *The positivity of a sequence of numbers and the Riemann hypothesis*, J. Number Theory **65** (1997), 325–333.
- [2] E. Bombieri and J. C. Lagarias, *Complements to Li’s criterion for the Riemann hypothesis*, J. Number Theory **77** (1999), 274–287.

- [3] D. Ruelle, *Zeta-functions for expanding maps and Anosov flows*, Invent. Math. **34** (1976), 231–242.
- [4] J. B. Keiper, *Power series expansions of Riemann’s ξ function*, Math. Comp. **58** (1992), 765–773.
- [5] A. Connes, *Trace formula in noncommutative geometry and the zeros of the Riemann zeta function*, Selecta Math. **5** (1999), 29–106.
- [6] A. Connes, *The Riemann Hypothesis: Past, Present and a Letter Through Time*, arXiv:2602.04022v1, 2026.
- [7] A. Connes and W. D. van Suijlekom, *Quadratic Forms, Real Zeros and Echoes of the Spectral Action*, Commun. Math. Phys. **406**, Article 312 (2025), DOI: 10.1007/s00220-025-05493-1.
- [8] C. Deninger, *On the nature of the “explicit formulas” in analytic number theory*, in: Number Theoretic Methods, Developments in Mathematics, vol. **8**, Kluwer Academic Publishers (2002), 97–118.
- [9] L. Geiger, *The Riemann Hypothesis: Even Dominance of the Weil Quadratic Form*, 2026.
- [10] L. Geiger, *A Conditional Reduction of the Riemann Hypothesis to Even Dominance of the Weil Quadratic Form*, Zenodo (2026), DOI: 10.5281/zenodo.20479145.
- [11] L. Geiger, *The Spectral Zookeeper: A Conditional Reduction toward the Riemann Hypothesis via Spectral Microcluster Closure in the Connes–Consani–Moscovici Framework*, Zenodo (2026), DOI: 10.5281/zenodo.20479223.
- [12] L. Geiger, *The Riemann Hypothesis: Spectral Routes*, 2026 (in preparation).
- [13] L. Geiger, *The Riemann Hypothesis: Cross-Route Synthesis*, 2026 (in preparation).
- [14] L. Geiger, *The Riemann Hypothesis: Conclusio*, 2026 (in preparation).

The Riemann Hypothesis: Even Dominance of the Weil Quadratic Form

Lukas Geiger

Bernau, Germany

ORCID: [0009-0005-7296-1534](https://orcid.org/0009-0005-7296-1534)

May 27, 2026

Abstract

This is the second paper in a five-part series (Part I: Foundations and Obstructions [9]; Part III: Spectral Routes [10]; Part IV: Cross-Route Synthesis [8]; Part V: Conclusion [7]). We address the open problem identified in Connes’ 2026 survey [2]: whether the minimum eigenvalue of the Weil quadratic form QW_λ is simple with even eigenfunction. We establish three main results. First, the *Shift Parity Lemma*: the difference matrix between even and odd shift operators has strictly negative minimum eigenvalue for all $N \geq 3$ modes and all normalized shifts $r \in (0, 2)$, implying that every prime individually favors even eigenfunctions. Second, we report computer-assisted finite-range gap diagnostics using interval arithmetic for 33 values of λ from 100 to 1,300,000; these are a strong conditional bridge until the odd-sector lower bound is independently validated. Third, we prove the *Leading-Mode Cancellation Lemma*: the overlap differences between even and odd sectors cancel pairwise with exact constant $c = 2 + \sqrt{2}$, yielding $(E_{\sin} - E_{\cos})/D(\lambda) = O(1/\log \lambda) \rightarrow 0$. In v2.0 we add two non-existence theorems (Theorems 2.14 and 2.17) that formally rule out the absolute total-positivity and universal-commuting-operator routes, establishing that even dominance is necessarily an arithmetic-statistical rather than a purely algebraic phenomenon. In v2.1 we add a robust Direct Frontier-Dominance argument (Section A.10, Theorem A.37) based on a common Rayleigh test vector on the frontier window $r \in [0.7, 1]$.

Status correction (v2.2, 2026-05-14). External critical review of the v2.1 manuscript and the companion proof-only extract [6] identified a structural gap in the asymptotic argument that was hidden by the previous wording. The Direct Frontier-Dominance proposition establishes $\lambda_{\min}(W_{\text{prime}}^+ - W_{\text{prime}}^-) = -\Theta(\sqrt{\lambda})$ via Rayleigh evaluation at a common test vector. This is a *variational* statement: there exists a vector v with $v^T W^+ v < v^T W^- v$. It does *not* establish $\lambda_1^+(\lambda) < \lambda_1^-(\lambda)$ asymptotically, because $\lambda_1^-(\lambda) \leq v^T W^- v$ runs in the same direction; a separate quantitative lower bound for λ_1^- is required. The simple 2×2 counter-example $A = \text{diag}(0, 10)$, $B = \text{diag}(5, -5)$ illustrates the gap. The status of the A1–A8 reduction chain is therefore revised from “unconditional” (v2.1) back to **conditional reduction**: the finite range $100 \leq \lambda \leq 1,300,000$ provides strong interval-arithmetic diagnostics, but theorem-level even dominance on that range still depends on a validated lower bound for $\lambda_1^-(\lambda)$. Asymptotic even dominance is reduced to the *Asymptotic Variational Gap Conjecture* (a quantitative lower bound for $\lambda_1^-(\lambda)$). A proposed approach via degenerate perturbation theory on the asymmetric three-mode diagonal structure is sketched in the companion paper [6]. The non-existence theorems NE-A and NE-B, the Shift Parity Lemma, the Leading-Mode Cancellation Lemma, and the 33 finite-range diagnostics remain unaffected by this revision.

Structural update (v3.0, 2026-05-26). The series is restructured from a trilogy to a five-part series; the conditional reduction status from v2.2 is unchanged. A new Section 5 presents the GF1 boundary-moment benchmark $\text{sLI} \sim (64\pi^2/3) \lambda/(NL)$ for the converged $\lambda = 5, 7$ benchmark, records the $\lambda = 3$ low-bandwidth outlier separately, and leaves $\lambda = 11$ as an open discriminator.

MSC 2020: 11M26, 11M36, 47A75, 65G20

Keywords: Riemann Hypothesis, Weil quadratic form, even dominance, Shift Parity Lemma, computer-assisted diagnostics

Contents

1	Introduction	5
2	Connes' Reduction and the Weil Quadratic Form	5
2.1	Connes' Reduction	5
2.2	Galerkin Discretization	6
2.3	Result 1: Spectral Gap in the Even Sector	6
2.4	Result 2: Even–Odd Sector Comparison	7
2.5	Discussion	8
2.6	Decomposition Analysis: Source of Even Dominance	9
2.7	The Shift Parity Lemma	9
2.8	Universal Even Preference of Individual Primes	11
2.9	Multi-Scale Decomposition: Frontier Primes	12
2.10	Proof Architecture: From Lemma to Even Dominance	12
2.11	Possible Approaches to a Formal Proof	13
2.12	Formal non-existence theorems (v2.0)	15
2.13	Prolate Eigenvalue Dominance: A Structural Result	17
2.14	Jentzsch Regularization Argument	18
2.15	Eigenvalue Spread Bound	18
2.15.1	Computer-Assisted Gap Diagnostics via Interval Arithmetic	19
2.16	Gap Asymptotics and Monotonicity	20
3	Weighted Compactness as Proof Closure	20
3.1	Setup and Rescaling	21
3.2	Building Block 1: Uniform Sobolev Bound	21
3.3	Building Block 2: Precompactness via Rellich	22
3.4	Building Block 3: Spectral Stability	22
3.5	Building Block 4: Bridge to Hurwitz	22
3.6	Numerical Verification	23
3.7	Assessment and Risks	24
3.8	Gap Growth from Block Bounds	24
3.8.1	Lemma L1: Prime coupling drives an exponentially deep even direction	25
3.8.2	Lemma L2: High-mode block is only polynomially negative	25
3.8.3	Lemma L3: Low–high coupling (historical)	26
3.8.4	The gap growth theorem	26
4	Connections to Existing Work	27
4.1	Li–Keiper Coefficients	27
4.2	Weil's Explicit Formula and Connes' Program	27
4.3	de Branges Spaces and Hilbert–Pólya	27
5	Galerkin Imbalance: Boundary-Moment Benchmark (v3.0)	28
5.1	Identification of the Boundary-Moment Constant	28
5.2	Hypothesis Correction Chain	29
6	Conclusion	30

A	Proofs of the Block-Bound Lemmas	30
A.1	Closed-form shift overlaps	31
A.2	Proof of Lemma L1: Exponential depth of the even direction	31
A.3	Proof of Lemma L2: High-block control via Gershgorin	33
A.4	Proof of Lemma L3: Off-diagonal coupling via Schur test	33
A.5	The Certifier Meta-Theorem	35
A.6	Status of the assumptions	37
A.7	Resolvent-damped energy framework for M1	38
A.8	Computer-assisted certificates	41
A.9	Analytical path to M1: leading-mode cancellation	42
A.10	Direct Frontier-Dominance Proof (v2.1, robust form)	51

1 Introduction

This paper is the second in a five-part series investigating the Riemann Hypothesis through gauge theory, spectral analysis, and arithmetic thermodynamics. Part I [9] establishes the thermodynamic landscape (structural results R1–R9) and documents the failure of the gauge-theoretic approach (obstructions K1–K4), leading to a reorientation toward Connes’ spectral program. Part III [10] surveys the spectral routes; Part IV [8] collects cross-route synthesis; Part V [7] provides the conclusio.

Here we develop the spectral approach in full. Following Connes’ 2026 survey [2], we study the Weil quadratic form QW_λ and its even-odd decomposition. Our three main contributions are:

1. The **Shift Parity Lemma** (Section 2.7): a purely algebraic result showing that every prime individually favors even eigenfunctions, proved via closed-form trigonometric entries.
2. **Computer-assisted finite-range diagnostics** (Section A.8): strong negative gap diagnostics at 33 values of λ from 100 to 1,300,000 via interval arithmetic; theorem-level even dominance remains conditional on the validated odd-sector lower bound.
3. The **Leading-Mode Cancellation Lemma** (Theorem A.24): the overlap differences between sectors cancel pairwise with exact constant $c = 2 + \sqrt{2}$, reducing the analytical gap (M1'') to Fourier analysis and the Prime Number Theorem.

2 Connes’ Reduction and the Weil Quadratic Form

A recent survey by Connes [2] reduces the Riemann Hypothesis to a spectral condition on the Weil quadratic form QW_λ . We summarize the connection and present numerical evidence addressing the remaining open step.

2.1 Connes’ Reduction

Connes constructs a self-adjoint operator A_λ on $L^2([-\log \lambda, \log \lambda])$ from the Weil explicit formula. In additive coordinates ($u = \log x$), the quadratic form reads

$$\begin{aligned} \langle A_\lambda f | f \rangle = & (\log 4\pi + \gamma) \|f\|^2 + \int_0^\infty \left[\langle T_u f + T_{-u} f, f \rangle - 2e^{-u/2} \|f\|^2 \right] \frac{e^{u/2}}{2 \sinh(u)} du \\ & + \sum_p \sum_{m=1}^\infty \frac{\log p}{p^{m/2}} \langle T_{m \log p} f + T_{-m \log p} f, f \rangle, \end{aligned}$$

where T_u denotes the shift operator, γ is the Euler–Mascheroni constant, and the archimedean kernel $e^{u/2}/(2 \sinh u)$ arises from the change of variables $x = e^u$ applied to Connes’ equation (10) in [2], with $x^{1/2}/(x - x^{-1}) d^*x = e^{u/2}/(2 \sinh u) du$. The operator A_λ is lower-bounded with compact resolvent (hence discrete spectrum).

Theorem 2.1 (Connes, Theorem 6.1 in [2]; proved in [4]). *Let $L > 0$ and D be a real distribution on $[0, L]$ with even extension $\tilde{D}$ on $[-L, L]$. Assume the quadratic form with kernel $\tilde{D}(x - y)$ defines a lower-bounded self-adjoint operator on $L^2([-L/2, L/2])$ whose minimum eigenvalue is simple, isolated, with even eigenfunction η . Then all zeros of $\tilde{\eta}(z)$ lie on the real line.*

Connes' Section 6.6 states: "In order to apply Theorem 6.1, one needs to show that the smallest eigenvalue of QW_λ is *simple* with *even* eigenvector." This is the remaining open step.

2.2 Galerkin Discretization

We approximate A_λ in the cosine basis $\{\varphi_n\}_{n=0}^{N-1}$ (even sector) and sine basis (odd sector) on $[-L, L]$ with $L = \log \lambda$. The matrix elements are computed via quadrature ($n_{\text{quad}} \geq 800$, $n_{\text{int}} \geq 500$) using the corrected kernel $e^{u/2}/(2 \sinh u)$ and primes $p \leq \max(\lambda, 100)$.

2.3 Result 1: Spectral Gap in the Even Sector

Within the even sector, we verify $\text{gap}(\lambda) = \lambda_2^+ - \lambda_1^+ > 0$ via Cauchy interlacing with N -convergence. Server computations ($\lambda \in [5, 200]$, N up to 80, corrected orthogonal basis $\cos(n\pi t/L)$) confirm:

λ	N_{final}	$\text{gap}^+ = \lambda_2^+ - \lambda_1^+$	Cauchy bound
5	30	3.52×10^{-1}	3.52×10^{-1}
10	30	2.28×10^{-1}	2.28×10^{-1}
20	30	1.39×10^{-1}	1.39×10^{-1}
50	30	2.35×10^{-1}	2.35×10^{-1}
100	80	4.83×10^0	4.83×10^0
200	80	1.11×10^1	1.11×10^1

Observation: All tested values have strictly positive Cauchy bound (minimum 0.095 at $\lambda = 30$), confirming that λ_1^+ is simple within the even sector. For $\lambda \geq 100$, the gap grows rapidly (~ 5 – 11) as the low-mode cosine block dominates.

Lemma 2.2 (Simplicity of λ_1^+). *For all 33 certificate values $\lambda \in [100, 1,300,000]$, the even-sector ground state is simple:*

$$\lambda_2^+(\lambda) - \lambda_1^+(\lambda) \geq g_{\min}^+(\lambda) > 0,$$

certified by interval arithmetic on the 4×4 even Galerkin block. Selected values:

λ	λ_1^+	λ_2^+	gap^+ (certified)
100	-11.18	-2.49	≥ 8.69
500	-34.06	-8.77	≥ 25.29
10,000	-195.46	-61.97	≥ 133.48
80,000	-594.37	-216.85	≥ 377.52
320,000	-1220.97	-489.97	≥ 731.00

The gap grows monotonically as $\sim \sqrt{\lambda}$, consistent with the analytical prediction from the Galerkin diagonal asymptotics: the leading mode ($n = 1$) has $W_{11}^+ \sim -\sqrt{\lambda}$, while the next mode ($n = 0$) has $W_{00}^+ \sim -c'\sqrt{\lambda}$ with $c' < c$. If the remaining even-dominance inequality $\lambda_1^+ < \lambda_1^-$ is supplied by a valid odd-sector lower bound, this certifies that the global ground state of A_λ is simple and even, as required by Connes' Theorem 6.1.

2.4 Result 2: Even–Odd Sector Comparison

For Theorem 6.1, the minimum of A_λ over the *full* space $L^2([-L, L])$ must be simple and even. Since A_λ commutes with parity ($f(t) \mapsto f(-t)$), the even and odd sectors decouple. We compare λ_1^+ (even) with λ_1^- (odd):

λ	N	λ_1^+	λ_1^-	Sector	Thm. 6.1
5	30	−4.31	−4.28	EVEN	✓*
10	30	−5.01	−5.19	ODD	—
20	30	−5.81	−5.84	ODD	—
25	60	−6.39	−6.39	EVEN*	✓*
28	60	−6.53	−6.53	ODD*	—*
30	60	−6.64	−6.63	EVEN*	✓*
50	60	−6.96	−6.94	EVEN*	✓*
55	60	−7.20	−6.70	EVEN	✓
100	80	−11.25	−9.06	EVEN	✓
200	80	−19.27	−15.05	EVEN	✓
500	80	−34.51	−26.88	EVEN	✓

* *Marginal*: $|\lambda_1^+ - \lambda_1^-| < 0.05$.

Honest assessment:

- For $\lambda \geq 55$: Robust negative gap diagnostics with margins growing as $\sim \sqrt{\lambda}$. The diagnostic gap reaches $|\Delta| > 0.5$ at $\lambda = 55$ and exceeds 475 at $\lambda = 1,300,000$. Theorem 6.1 prerequisites are conditionally supported for 33 values up to $\lambda = 1,300,000$, pending the odd-sector lower-bound validation.
- For $25 \leq \lambda \leq 50$: Marginal regime ($|\Delta| < 0.02$). The even/odd ordering oscillates with λ at $N = 60$: $\lambda = 25, 30, 40, 50$ show even dominance while $\lambda = 28, 35, 45$ show odd dominance. Neither sector robustly dominates.
- For $\lambda = 10$ and $\lambda = 20$: The odd sector has the lower minimum. Theorem 6.1 is *not directly applicable* in this regime.
- **Non-monotone transition:** The crossover from odd to even ground state is not sharp but spreads over $\lambda \in [25, 55]$. For $\lambda \geq 55$, even dominance is robust and the gap grows without bound.
- **Computer-assisted diagnostics (interval arithmetic):** A negative finite-range gap diagnostic has been documented for 33 values of λ from 100 to 1,300,000 using interval arithmetic (mpmath.iv, 50 digits) for the even block and float64 tail extrapolation for the odd block. The gap grows monotonically as $\sim \sqrt{\lambda}$, from -2.25 at $\lambda = 100$ to -475.83 at $\lambda = 1,300,000$; theorem-level certification is conditional on the odd-sector lower-bound validation (see Section A.8).

2.5 Discussion

Our numerical results address Connes’ open problem (Section 6.6 of [2]) from two angles:

1. **Simplicity within even sector:** The Cauchy-interlacing bound provides rigorous (up to discretization error) evidence that λ_1^+ is simple. This holds for all tested $\lambda \in [5, 200]$.
2. **Even dominance for large λ :** For $\lambda \geq 25$, the finite-dimensional diagnostics put the even sector lower in most tested ranges. The theorem-level infinite-dimensional ordering still needs the odd-sector lower bound.
3. **Odd dominance for small λ :** For $\lambda = 10$ and $\lambda = 20$, the odd sector is lower. Theorem 6.1 is not directly applicable in this regime.

Three structural features emerge:

1. **Diffuse crossover:** At $N = 60$, the crossover from odd to even dominance spreads over $\lambda \in [25, 55]$, with oscillating sign and $|\Delta| < 0.02$. At $\lambda \geq 55$, even dominance sets in robustly ($|\Delta| > 0.5$) and grows without reversal through $\lambda = 500$.
2. **Growing finite-range diagnostic gap:** For $\lambda \geq 55$, the observed gap grows in magnitude, reaching $\Delta \approx -7.6$ at $\lambda = 500$ and $\Delta \approx -476$ at $\lambda = 1,300,000$ (see Sections A.8 and 2.16). The growth scales as $\sim \sqrt{\lambda}$ (analytically, via PNT; the earlier fit $\sim (\log \lambda)^{4.2}$ was based on a limited range), but the eigenvalue-ordering interpretation needs the odd-sector lower bound.
3. **Finite-range conditional bridge:** For 33 values of λ from 100 to 1,300,000, the computer-assisted diagnostics provide a strong conditional bridge pending odd-tail validation (see Section A.8).
4. **A Direct Frontier-Dominance variational proof** (v2.1, Section A.10): an independent proof of the asymptotic variational sign using a common Rayleigh test vector on the prime frontier, PNT partial summation, Mertens’ theorem, and the existing finite CAP bridge. It obviates both v2.0 caveats for the variational statement, not for the eigenvalue-ordering step.

Implication for Connes’ program: Connes’ strategy (Sections 6.4–6.6 of [2]) uses Hurwitz’s theorem to transfer the real-zeros property from the truncated Fourier transforms $\hat{k}_\lambda$ to the Riemann Ξ -function in the limit $\lambda \rightarrow \infty$. Crucially, Hurwitz’s theorem requires only a *sequence* $\lambda_n \rightarrow \infty$ along which the prerequisites hold—not all $\lambda > 1$. Therefore, the odd-dominant regime for small λ (around $\lambda \approx 10$ –30) is *not* an obstruction. The remaining large-step requirement is a theorem-level lower bound for λ_1^- that turns the finite diagnostics and asymptotic variational statement into genuine eigenvalue ordering.

Our data show robust negative gap diagnostics for $\lambda \geq 50$ (with margins growing as $\sim \log(\lambda)^b$), and the gap grows monotonically for $\lambda \geq 120$. In the current v2.4 status, the formal proof that $\lambda_1^+ < \lambda_1^-$ for all $\lambda \geq 100$ remains conditional on a validated lower bound for the odd-sector minimum eigenvalue. Theorem A.36 records this conditional reduction; Section A.10 gives the robust frontier Rayleigh proof of the associated variational sign.

2.6 Decomposition Analysis: Source of Even Dominance

To understand the mechanism behind even dominance, we decompose $QW_\lambda = W_{\text{diag}} + W_{\text{arch}} + W_{\text{prime}}$ and compute each contribution separately for the even and odd sectors:

- $W_{\text{diag}} = (\log 4\pi + \gamma) I$: identical in both sectors.
- W_{arch} : the archimedean kernel $e^{u/2}/(2 \sinh u)$ produces *virtually identical* minimum eigenvalues in both sectors ($\Delta < 10^{-3}$ at $\lambda = 100$). The archimedean contribution is parity-neutral.
- W_{prime} : the prime shift operators $S_{m \log p} + S_{-m \log p}$ are the *sole driver* of even-dominance. At $\lambda = 100$: $\lambda_1^+(W_{\text{prime}}) \approx -14.5$ vs. $\lambda_1^-(W_{\text{prime}}) \approx -12.3$.

The mechanism is traced to the product formula for trigonometric functions. For the cosine (even) basis, shift-overlap integrals contain the constructive term $\cos((k_n + k_m)t)$:

$$\cos(k_n t) \cos(k_m(t - s)) = \frac{1}{2} [\cos((k_n - k_m)t + k_m s) + \cos((k_n + k_m)t - k_m s)].$$

For the sine (odd) basis, the second term enters with a minus sign. The difference matrix $D = W_{\text{prime}}^+ - W_{\text{prime}}^-$ is indefinite, and its structure produces the even-sector advantage.

2.7 The Shift Parity Lemma

The decomposition analysis shows that even dominance is driven by the prime shift operators. A deeper analysis reveals a remarkable algebraic structure: *every single prime individually favors the even sector*, and this property is independent of the cutoff parameter L .

Definition 2.3 (Difference matrix). For $N \in \mathbb{N}$ and $r \in (0, 2)$, define the $N \times N$ symmetric matrix $D_N(r)$ by

$$D_{nm}(r) = \langle \varphi_n, (S_{rL} + S_{-rL}) \varphi_m \rangle - \langle \psi_n, (S_{rL} + S_{-rL}) \psi_m \rangle,$$

where $\varphi_n(t) = \cos(n\pi t/L)/\sqrt{L}$ (even basis) and $\psi_n(t) = \sin((n+1)\pi t/L)/\sqrt{L}$ (odd basis), and S_s denotes the shift operator on $L^2([-L, L])$.

Lemma 2.4 (L -independence). *The matrix $D_N(r)$ depends only on $r = s/L$, not on L separately. In particular, one may set $L = 1$ without loss of generality.*

Proof. By substitution $u = t/L$ in the shift-overlap integrals, all factors of L cancel due to the normalization of the basis functions. This is verified numerically to machine precision (spread $< 10^{-15}$) across $L \in \{1, 2, 5, 10, 20\}$. $\square$

Setting $L = 1$, the matrix entries admit closed-form expressions.

Proposition 2.5 (Closed-form entries). *For $L = 1$, the diagonal entries of $D_N(r)$ follow a telescoping pattern:*

$$D_{nn}(r) = (2 - r) [\cos(n\pi r) - \cos((n+1)\pi r)] - \frac{\sin(n\pi r)}{n\pi} - \frac{\sin((n+1)\pi r)}{(n+1)\pi},$$

with the convention $\sin(0)/0 = 0$ and $\cos(0) = 1$ for $n = 0$. The off-diagonal entries for $N = 3$ are:

$$\begin{aligned} D_{01}(r) &= \frac{(3\sqrt{2} + 4) \sin(\pi r) - 2 \sin(2\pi r)}{3\pi}, \\ D_{02}(r) &= \frac{-2\sqrt{2} \sin(2\pi r) + \sin(3\pi r) - 3 \sin(\pi r)}{4\pi}, \\ D_{12}(r) &= \frac{2[19 \sin(2\pi r) - 5 \sin(\pi r) - 6 \sin(3\pi r)]}{15\pi}. \end{aligned}$$

All formulas are derived by hand via product-to-sum identities and verified to 10^{-15} against numerical quadrature.

These closed forms reveal striking structural properties:

Corollary 2.6 (Structure at $r = 1$). $D_N(1) = \text{diag}(2, -2, 2, -2, \dots)$, i.e., $D_{nm}(1) = 2(-1)^n \delta_{nm}$.

Proof. At $r = 1$: $\cos(n\pi) - \cos((n+1)\pi) = (-1)^n - (-1)^{n+1} = 2(-1)^n$, and $\sin(k\pi) = 0$ for all $k \in \mathbb{Z}$. $\square$

Corollary 2.7 (Off-diagonal antisymmetry). For $n \neq m$: $D_{nm}(r) = -D_{nm}(2 - r)$.

Theorem 2.8 (Shift Parity Lemma). For $N \geq 3$ and all $r \in (0, 2)$:

$$\lambda_1(D_N(r)) < 0,$$

where λ_1 denotes the minimum eigenvalue. For $N = 2$, the statement is false (counterexample at $r \approx 0.45$ where $\lambda_1 \approx +0.36$).

Proof. By the Cauchy interlacing theorem, $\lambda_1(D_{N+1}) \leq \lambda_1(D_N)$ for the principal submatrix embedding. Hence it suffices to prove the $N = 3$ case.

The trace of the 3×3 matrix telescopes via Theorem 2.5:

$$\text{Tr}(D_3(r)) = (2 - r)(1 - \cos 3\pi r) - \frac{2 \sin \pi r}{\pi} - \frac{\sin 2\pi r}{\pi} - \frac{\sin 3\pi r}{3\pi}.$$

The determinant $\det(D_3(r))$ is an explicit (though lengthy) trigonometric expression in r . Both are computed from the closed-form entries.

The proof uses a two-case argument. Let $R_1 = \{r \in (0, 2) : \det D_3(r) < 0\}$ and $R_2 = \{r \in (0, 2) : \det D_3(r) \geq 0\}$.

Case 1 ($r \in R_1$, approximately 93.3% of $(0, 2)$): If $\det D_3(r) < 0$, the three eigenvalues have mixed signs (the product of eigenvalues equals the determinant), so $\lambda_1 < 0$.

Case 2 ($r \in R_2 \approx [0.597, 0.729]$): In this region, $\text{Tr}(D_3(r)) < 0$ (verified: $\max_{R_2} \text{Tr} \approx -0.007$). Since the trace equals the sum of eigenvalues, not all three can be non-negative. Hence $\lambda_1 < 0$.

Completeness: The determinant $\det D_3(r)$ has exactly two zeros in $(0, 2)$, at $r_1^{\det} \approx 0.59652$ and $r_2^{\det} \approx 0.72929$ (50-digit precision via mpmath). The trace $\text{Tr}(D_3(r))$ has zeros at $r_2^{\text{Tr}} \approx 0.58761$ and $r_3^{\text{Tr}} \approx 0.73024$. Crucially,

$$r_2^{\text{Tr}} < r_1^{\det} < r_2^{\det} < r_3^{\text{Tr}},$$

with overlaps $r_1^{\det} - r_2^{\text{Tr}} \approx 0.0089$ and $r_3^{\text{Tr}} - r_2^{\det} \approx 0.00095$. Hence $R_2 \subset \{r : \text{Tr} < 0\}$, and Cases 1 and 2 cover all of $(0, 2)$ without gap.

The $N = 2$ counterexample is verified directly: $\lambda_1(D_2(0.445)) \approx +0.355 > 0$. $\square$

Remark 2.9 (Nature of the proof). The Shift Parity Lemma is a purely algebraic statement about trigonometric matrices. It involves no number theory (no primes, no ζ -function). The connection to number theory enters only when $D(r)$ is weighted by $(\log p) p^{-m/2}$ and summed over primes with $r = m \log p / \log \lambda$.

2.8 Universal Even Preference of Individual Primes

A direct consequence of the Shift Parity Lemma is that *every* prime individually favors even functions:

Corollary 2.10 (Universal even preference). *For every prime $p \leq \lambda$ and every $\lambda \geq e^{3 \log^2} \approx 8$ (so that $N \geq 3$ modes fit), the per-prime difference matrix*

$$D_p = \sum_{m=1}^{\lfloor 2L/\log p \rfloor} \frac{\log p}{p^{m/2}} \cdot D_N\left(\frac{m \log p}{L}\right)$$

satisfies $\lambda_1(D_p) < 0$. In particular, the prime contribution W_p has $\lambda_1(W_p^+) < \lambda_1(W_p^-)$.

Proof. Each summand has $\lambda_1(D_N(m \log p/L)) < 0$ by Theorem 2.8 (since $N \geq 3$ and $0 < m \log p/L < 2$). The coefficients $c_m = (\log p) p^{-m/2} > 0$ are positive. We use a Rayleigh quotient argument: let v^* be the unit eigenvector achieving $\lambda_1(D_N(\log p/L))$ for the dominant $m = 1$ shift. Then

$$\lambda_1(D_p) \leq (v^*)^T D_p v^* = c_1 \lambda_1(D_N(\log p/L)) + \sum_{m \geq 2} c_m (v^*)^T D_N(m \log p/L) v^*.$$

The second sum satisfies $|\sum_{m \geq 2} c_m (v^*)^T D_N(m \log p/L) v^*| \leq \sum_{m \geq 2} c_m \|D_N\|_{\text{op}} \leq C (\log p)/p$, which is negligible compared to $c_1 |\lambda_1| \geq C' (\log p)/\sqrt{p}$ for $p \geq p_0$. For small primes ($p = 2, 3, 5, 7$) the claim is verified by direct computation. $\square$

This is verified numerically for all primes up to $\lambda = 500$ (95 primes), with 100% showing $\lambda_1(D_p) < 0$:

λ	primes tested	$\lambda_1(D_p) < 0$	fraction
50	15	15	100%
100	25	25	100%
200	46	46	100%
500	95	95	100%

Remark 2.11 (Full-space verification, v2.0). An independent verification in the full Galerkin space ($N = 120$ PSWF modes, primes up to $P_{\text{max}} = 10,000$, exact kernel) was performed on a cloud-compute server. At both $\lambda = 100$ and $\lambda = 200$, all 1229/1229 primes deepen the even-odd gap ($\Delta_p < 0$), with zero exceptions across 2458 individual tests. The small $p = 2$ anomaly observed in the 3×3 truncation vanishes when enough modes are included. Scripts and CSV results: `scripts/gemini_verification/`.

2.9 Multi-Scale Decomposition: Frontier Primes

Adding primes one at a time to the Weil operator reveals which *scale* drives the gap. Define the cumulative gap $\Delta(P) = \lambda_1^+(\lambda; P) - \lambda_1^-(\lambda; P)$, where only primes $p \leq P$ are included in W_{prime} .

λ	scale $[P_1, P_2)$	$\sum \delta(\Delta)$	dominant?
100	[2, 10)	+0.019	no (slightly odd)
100	[10, 30)	−0.003	marginal
100	[30, 100)	−2.263	yes
200	[2, 30)	+0.023	no
200	[30, 100)	−0.095	marginal
200	[100, 200)	−4.227	yes
500	[2, 100)	+0.017	no
500	[100, 500)	−9.022	yes

The gap is driven overwhelmingly by *frontier primes* $p \sim \lambda$ (i.e., primes with $\log p / \log \lambda$ near 1). Small primes contribute negligibly or even slightly favor the odd sector. This is consistent with the Shift Parity Lemma: primes with $r = \log p / L$ near 1 produce the deepest negative eigenvalue of $D(r)$ (recall $\lambda_1(D_3(1)) = -2$, the global minimum region).

As λ grows, the number of frontier primes increases (by the prime number theorem, $\pi(\lambda) - \pi(\lambda/e) \sim \lambda / \log \lambda$), causing the gap to deepen monotonically for large λ .

2.10 Proof Architecture: From Lemma to Even Dominance

The results of this section provide the following reduction:

1. **Basis decomposition:** $W_{\text{prime}} = \sum_p W_p$, where each W_p acts on even and odd sectors separately.
2. **Shift Parity Lemma (Theorem 2.8):** Each W_p individually satisfies $\lambda_1(W_p^+) < \lambda_1(W_p^-)$.
3. **Cumulative effect:** The sum over primes amplifies the even preference (not merely preserves it).
4. **Parity-neutral archimedean term:** $W_{\text{diag}} + W_{\text{arch}}$ produces $|\lambda_1^+ - \lambda_1^-| < 10^{-3}$ (numerically verified).
5. **Conclusion:** Even dominance of QW_λ follows from the cumulative prime effect.

The remaining gap for a complete proof is step (3): showing that the cumulative effect of summing W_p over all $p \leq \lambda$ preserves (and amplifies) the individual even preferences. This is not automatic because eigenvalue inequalities are not additive for sums of matrices. However, the numerical evidence is overwhelming: the cumulative gap deepens monotonically for $\lambda \geq 55$ and scales as $|\text{cert_gap}| \sim c\sqrt{\lambda}$ across nearly 4 orders of magnitude (33 certified values, $\lambda = 100$ to 1,300,000).

2.11 Possible Approaches to a Formal Proof

Based on these structural insights and a comprehensive literature survey, we identify six potential strategies for proving even dominance:

- (A) **Complete monotonicity and total positivity (ruled out, v1.4):** The archimedean kernel admits the Laplace mixture representation

$$\frac{e^{u/2}}{2 \sinh u} = \sum_{n=0}^{\infty} e^{-(2n+\frac{1}{2})u}, \quad u > 0,$$

with all coefficients equal to $1 > 0$. This makes it *completely monotone* on $(0, \infty)$: $(-1)^m k^{(m)}(u) \geq 0$ for all $m \geq 0$ and $u > 0$. By the Schoenberg–Hirschman theorem, completely monotone functions are PF_{∞} (totally positive) as one-sided convolution kernels. One might hope to extend this to the full Weil operator, which would imply even dominance as a direct consequence of Gantmacher–Krein oscillation theory (the ground state of a totally positive operator has no sign changes, hence is even on a symmetric domain).

This hope fails. The Fourier multiplier of the prime shift operator is

$$M(\xi) = -2 \Re[\zeta'/\zeta(\tfrac{1}{2} + i\xi)],$$

an identity (not an approximation) that follows from the Weil explicit formula without assuming RH. Numerical evaluation shows $M(\xi) < 0$ on approximately 49% of its domain: the non-trivial zeros of ζ on the critical line each produce a pole in ζ'/ζ and hence an oscillation in M . The prime shift operator is therefore *not* positive-definite, and the Laplace mixture structure of the archimedean kernel cannot be extended to the full operator. Total positivity as an absolute property is ruled out. A weak form—variation diminution in the prime-weighted mean—coincides with the frontier-dominance argument (Route E below) rather than providing an independent path.

- (B) **Commuting differential operator (ruled out, v1.4):** If a second-order differential operator commuting with A_{λ} exists (analogous to the prolate spheroidal wave operator for the sinc kernel), Sturm–Liouville theory would guarantee simplicity. The Jacobi matrix representation of the prolate operator (Proposition 3.7 in [3]) provides explicit even/odd coefficients.

We tested this route directly. Searching for a symmetric T with $[T, D_N(r)] = 0$ simultaneously for a dense grid of r -values (13 samples in $(0, 2)$) via SVD on the stacked commutator-constraint matrix, we find for $N \in \{3, 5, 7, 10, 15\}$: the kernel dimension is exactly 1, and the unique solution is $T = c \cdot I$ (identity). The second smallest singular value remains bounded below ($\sigma_2 \geq 0.70$ at $N = 15$) and does not decrease with N . Hence no non-trivial universal commuting operator exists, even asymptotically.

This absence has structural meaning: the matrices $D_N(r)$ at different r have fundamentally different eigenbases, and no fixed basis diagonalizes them simultaneously. The prime shifts at different scales carry no common hidden symmetry. This is consistent with the multiplicative independence of primes: they have no joint algebraic structure that a single operator could capture. Even dominance therefore

cannot arise from a Sturm–Liouville-type oscillation principle. It must come from the pointwise Shift Parity Lemma summed with prime weights—exactly the frontier-dominance mechanism.

- (C) **Prolate perturbation:** Treat A_λ as a perturbation of the prolate operator P_λ , for which even dominance is proven (Slepian–Pollak, 1961). If the prime contributions can be controlled as $o(\lambda^2)$ in operator norm, the eigenvalue ordering transfers. This is essentially Connes’ strategy (Section 6.6 of [2]).
- (D) **Hellmann–Feynman sensitivity analysis:** One may attempt to prove monotonicity of the gap $\Delta(\lambda)$ by computing the L -derivative $\partial\Delta/\partial L$ via the Hellmann–Feynman theorem:

$$\frac{\partial\lambda_1^\pm}{\partial L} = \langle \eta_1^\pm | \frac{\partial A_\lambda}{\partial L} | \eta_1^\pm \rangle,$$

where $\eta_1^\pm$ are the normalized ground-state eigenvectors of the even/odd sectors. We computed this numerically for $\lambda \in [55, 500]$ (Table 1). The results reveal that $\partial\Delta/\partial L > 0$ for all $\lambda \geq 80$: increasing L with fixed prime structure makes the gap *less* negative. This means that the observed deepening of even dominance is driven entirely by the *addition of new primes* as λ grows, not by the growth of $L = \log \lambda$. This is consistent with the frontier-prime mechanism of Section 2.9 and suggests that a successful proof strategy must directly address the prime summation, not the L -dependence.

λ	Δ	$\partial\Delta/\partial L$	HF _{prime}	sign
55	−0.506	−0.625	−0.477	< 0
70	−1.443	−1.494	−1.561	< 0
80	−2.040	+2.218	+2.192	> 0
100	−2.185	+2.480	+2.443	> 0
200	−4.217	+2.011	+1.990	> 0

Table 1: Hellmann–Feynman sensitivity of the even–odd gap to the interval half-length $L = \log \lambda$. The total derivative $\partial\Delta/\partial L$ becomes positive for $\lambda \geq 80$, showing that increasing L alone does not deepen the gap. The prime contribution HF_{prime} dominates (the archimedean contribution is < 0.03 in all cases). The ground-state projection $|\langle \eta_1^+ | e_0 \rangle|^2 > 0.99$ confirms strong low-mode localization.

All four approaches face the same core difficulty: the Weil kernel is not positive, and the ground state is shaped by high-frequency prime resonances rather than low-frequency modes. Our decomposition analysis shows that this non-positivity is not an obstacle but the *source* of even dominance—a structural insight that may guide future analytical work. The Hellmann–Feynman analysis additionally shows that a monotonicity proof via L -differentiation alone is insufficient; the prime-counting mechanism must be addressed directly.

Remark 2.12 (Route chosen for A6). The cumulative step is reduced in Theorem A.36 via the **M1'' resolvent framework** combined with PNT transfer and computer-assisted diagnostics (a three-regime argument). This path is closest in spirit to route (C) (prolate perturbation) but operates directly on the Galerkin representation rather than requiring

a separate prolate comparison operator. Routes (C) and (D) remain available as independent verification. Route (A) is excluded as an absolute structural property: the Fourier multiplier $M(\xi) = -2\Re[\zeta'/\zeta(\frac{1}{2} + i\xi)]$ attains negative values on a set of positive measure (Theorem 2.14). Route (B) is excluded at all tested Galerkin truncations $N \leq 15$ (Theorem 2.17); its full-space extension is formulated as Theorem 2.19 and remains open.

Remark 2.13 (Structural interpretation of the excluded routes). The exclusion of (A) (as an absolute property, Theorem 2.14) and (B) (for $N \leq 15$, Theorem 2.17) is not a gap in the proof but a structural discovery. Route (A) would have required the primes to be “absolutely smoothing” (variation-diminishing as a single operator-theoretic property); instead, their collective effect is variation-diminishing only after prime-weighted averaging, which is the content of Route (E). Route (B) would have required the primes to share a hidden common symmetry that a single differential operator captures; instead, the dense-grid SVD evidence (up to $N = 15$) indicates that they are multiplicatively independent and carry no such common structure in the tested truncations. Even dominance nonetheless holds—driven by the pointwise Shift Parity Lemma and the prime-weighted summation, not by an overarching algebraic principle.

2.12 Formal non-existence theorems (v2.0)

The exclusion of Routes (A) and (B) — as an absolute property for (A) and for all tested Galerkin truncations $N \leq 15$ for (B) — is formalized by two theorems.

Theorem 2.14 (Non- PF_∞ of the Prime Shift Operator). *Let $L > 0$ and consider the prime shift operator on $L^2([-L, L])$, whose Fourier multiplier is*

$$M(\xi) = -2\Re\left[\frac{\zeta'}{\zeta}\left(\frac{1}{2} + i\xi\right)\right],$$

given by the Weil explicit formula (see [2]), valid unconditionally. The set

$$\{\xi \in [-L, L] : M(\xi) < 0\}$$

has strictly positive Lebesgue measure. Consequently, the operator family A_λ does not satisfy absolute total positivity in the sense of PF_∞ : the route to even dominance via absolute total positivity (Schoenberg–Hirschman, Gantmacher–Krein oscillation theory) is thereby excluded.

Proof sketch. On any compact interval $[-L, L]$ the nontrivial zeros $\rho_k = \frac{1}{2} + i\gamma_k$ of ζ contribute finitely many simple poles of ζ'/ζ . Excising ε -neighbourhoods $U_\varepsilon = \bigcup_k (\gamma_k - \varepsilon, \gamma_k + \varepsilon)$ around these ordinates yields a set on which $M(\xi) = -2\Re[\zeta'/\zeta(\frac{1}{2} + i\xi)]$ is a bounded continuous function (a classical consequence of the Weil explicit formula away from zero ordinates). Near each ordinate, the real part of $1/(s - \rho_k)$ changes sign as ξ crosses γ_k ; hence on any sufficiently small neighbourhood of γ_k (excluding the pole itself), the continuous representative of M attains both strictly positive and strictly negative values on open subintervals. The sub-intervals on which $M < 0$ thereby form a non-empty open set, so $\{\xi \in [-L, L] \setminus U_\varepsilon : M(\xi) < 0\}$ has strictly positive Lebesgue measure for every $\varepsilon > 0$ sufficiently small. By monotone convergence ($\varepsilon \rightarrow 0^+$), the same holds for $\{\xi \in [-L, L] : M(\xi) < 0\}$ interpreted as a measurable set in the complement of the discrete pole set. Absolute total positivity (PF_∞) requires the multiplier to be nonnegative almost everywhere in the sense of distributions on the complement of the pole set; a set

of positive measure with $M(\xi) < 0$ contradicts this. The measure-theoretic statement is therefore well-defined *modulo the discrete pole set* (which has Lebesgue measure zero), and no distributional regularization is needed for the positivity-of-measure conclusion; see Theorem 2.15. $\square$

Remark 2.15 (Distributional status of M and identification with a multiplier). The statement of Theorem 2.14 concerns the measurable function $\xi \mapsto M(\xi)$ on $[-L, L] \setminus \{\gamma_k\}$, which is locally bounded on that set. Identification of M as the Fourier multiplier of the prime shift operator requires interpreting the Weil explicit formula in its tempered-distribution form, with the pole contributions regularized as Cauchy principal values; this is standard (cf. [2], §3) and does not affect the positivity-of-measure conclusion, since the pole set has Lebesgue measure zero. The argument rules out PF_∞ as an absolute property of the a.e.-representative of M ; stronger notions of total positivity (e.g. PF_k for finite k) are not discussed here.

Remark 2.16 (Numerical quantification). Sampling $M(\xi)$ on a uniform grid of 10^4 points in $\xi \in [0, 100]$ (step $\Delta\xi = 0.01$, 50-digit mpmath precision for ζ'/ζ), excluding ε -neighbourhoods of the first 29 nontrivial ordinates $\gamma_k \in [0, 100]$ with $\varepsilon = 10^{-3}$, yields approximately 49% of sampled points satisfying $M(\xi) < 0$. The estimate is stable under refinement ($\varepsilon = 10^{-4}$, 10^5 points: 49.2%). This illustrates the measure-theoretic statement of Theorem 2.14 empirically and does not replace the structural argument given above.

Theorem 2.17 (No Universal Commuting Operator, for $N \leq 15$). *Let $D_N(r)$ denote the shift-parity difference matrix from Theorem 2.3. For each $N \in \{3, 5, 7, 10, 15\}$ and any dense sample $r_1, \dots, r_{13} \in (0, 2)$,*

$$\ker\{T \in \text{Sym}_N(\mathbb{R}) : [T, D_N(r_i)] = 0 \text{ for all } i = 1, \dots, 13\} = \text{span}\{I\}.$$

Equivalently, the only real symmetric $N \times N$ matrix commuting with every sampled $D_N(r_i)$ is a scalar multiple of the identity. For each listed N the statement is a computer-assisted result.

Computer-assisted proof. For fixed N , parametrize $T \in \text{Sym}_N(\mathbb{R})$ by its $N(N+1)/2$ independent entries and write $[T, D_N(r_i)] = 0$ as N^2 homogeneous linear equations in these entries. Stacking over the 13 sample points yields a constraint matrix $C_N \in \mathbb{R}^{13N^2 \times N(N+1)/2}$ such that $C_N \text{vecs}(T) = 0$ if and only if T commutes with every $D_N(r_i)$. We compute the singular value decomposition of C_N numerically. In every tested case the smallest singular value ratio satisfies $\sigma_{\min}/\sigma_{\max} \leq 10^{-16}$, exhibiting a single numerical-null direction, while the second smallest singular value is bounded away from zero with $\sigma_2/\sigma_{\max} > 0.6$ (and $\sigma_2/\sigma_{\max} \geq 0.70$ for $N = 15$). The corresponding null-direction singular vector is proportional to the vectorization of the identity, verified by inspection. By standard floating-point backward-error analysis for the SVD (Wilkinson-type bounds: the computed singular values $\tilde{\sigma}_j$ satisfy $|\tilde{\sigma}_j - \sigma_j(C_N)| \leq c_N \varepsilon_{\text{mach}} \|C_N\|_2$ with c_N polynomial in the matrix dimensions), at $N = 15$ the worst-case backward perturbation is bounded by $10^{-12} \|C_N\|_2$ in double precision, which is smaller than the observed gap $\sigma_2 - \sigma_1 \geq 0.70 \|C_N\|_2$ by more than ten orders of magnitude. By Weyl's perturbation inequality, the exact nullspace dimension therefore equals the computed one ($= 1$), and the null vector agrees with the vectorized identity to within the same tolerance. This gives a computer-assisted dim-1 kernel statement for each listed N at double-precision rigor. The reference implementation is `scripts/dungeon/door_B_universal_T.py`. $\square$

Remark 2.18 (SVD diagnostics). The numerical rank deficiency is certified by the following singular value ratios:

N	$N(N+1)/2$	$\sigma_{\min}/\sigma_{\max}$	$\sigma_2/\sigma_{\max}$
3	6	$8 \cdot 10^{-17}$	> 0.6
5	15	$6 \cdot 10^{-17}$	> 0.6
7	28	$1 \cdot 10^{-16}$	> 0.6
10	55	$1 \cdot 10^{-16}$	> 0.7
15	120	$4 \cdot 10^{-16}$	≥ 0.70

The observed lower bound on $\sigma_2/\sigma_{\max}$ is N -independent within the tested range and shows no downward trend, which is the structural obstruction: enlarging N does not open new near-commutants.

Conjecture 2.19 (Full-space extension). *The conclusion of Theorem 2.17 holds for all $N \in \mathbb{N}$: for any sufficiently dense sample of $r \in (0, 2)$, the only symmetric matrices commuting with all corresponding $D_N(r)$ are scalar multiples of the identity. Consequently, no non-trivial universal commuting operator exists in the infinite-dimensional limit.*

2.13 Prolate Eigenvalue Dominance: A Structural Result

A key finding of the present work is the *prolate eigenvalue dominance* of the prime shift operator. Define the *prime shift operator* P_λ acting on $L^2[-\log \lambda, \log \lambda]$:

$$P_\lambda f(t) = \sum_p \sum_{m=1}^{\infty} \frac{\log p}{p^{m/2}} [f(t - m \log p) + f(t + m \log p)] \cdot \mathbf{1}_{[-L, L]}(t),$$

where $L = \log \lambda$. This operator decomposes as $P_\lambda = P_\lambda^+ \oplus P_\lambda^-$ on even and odd subspaces. We compute the dominant eigenvalue $\mu_{\max}$ of each sector:

λ	$\mu_{\max}(P_\lambda^+)$	$\mu_{\max}(P_\lambda^-)$	ratio
30	14.8	2.0	7.4
100	24.2	3.9	6.2
200	35.8	7.5	4.8

The even sector couples 5–8 $\times$ more strongly to the prime shifts. This is verified at the Fourier level: the even-sector ground state ψ_0^+ satisfies $|\hat{\psi}_0^+(m \log p)|^2 / |\hat{\psi}_0^-(m \log p)|^2 \approx 5$ –34 for small primes $p = 2, 3, 5$ (where ψ_0^- is the odd ground state).

Remark 2.20 (Analogy with Slepian’s theorem). For the standard prolate spheroidal wave operator (Fourier band-limiting composed with time-limiting), Slepian (1961) proved that the dominant eigenfunction is *even*. Our prime shift operator has the same parity-symmetric structure: it is a sum of translation operators $S_s + S_{-s}$ with positive weights. The crucial difference is that the shifts are at discrete frequencies $m \log p$ rather than a continuous band. Extending Slepian’s nodal argument to this discrete-frequency setting would complete the proof of even dominance.

Remark 2.21 (The $n = 0$ mode is not the main driver). Removing the constant mode ($n = 0$) from the cosine sector changes $\Delta = \lambda_1^+ - \lambda_1^-$ by only a small fraction of the total gap. Even dominance is thus a *collective effect* of all cosine modes, not a consequence of the DC component. This distinguishes the Weil operator from the standard prolate operator, where the constant mode plays a more prominent role.

2.14 Jentzsch Regularization Argument

The prime shift operator P_λ restricted to $[-L, L]$ is *not* positive semidefinite (it has ~ 10 negative eigenvalues from truncation effects). However, the Gauss-regularized kernel

$$K_\varepsilon(t, t') = \sum_p \sum_{m=1}^{\infty} \frac{\log p}{p^{m/2}} \cdot \frac{1}{\varepsilon \sqrt{\pi}} \left[e^{-((t-t'-m \log p)/\varepsilon)^2} + e^{-((t-t'+m \log p)/\varepsilon)^2} \right]$$

satisfies $K_\varepsilon(t, t') > 0$ for *all* $t, t' \in [-L, L]$ and all $\varepsilon > 0$. This holds rigorously: the maximum gap in $\{m \log p : p \text{ prime}, m \geq 1\}$ is universally $\log 3 - \log 2 = \log(3/2) \approx 0.405$, independent of λ . For $u = t - t'$ near zero, the nearest shift is $\log 2$ (from the $\pm s$ symmetrization). Thus

$$K_\varepsilon(u) \geq w_{\min} \cdot \frac{1}{\varepsilon \sqrt{\pi}} e^{-(\log 2/\varepsilon)^2} > 0$$

where $w_{\min} = \min_{p \leq \lambda} (\log p) p^{-1/2} > 0$. This is exponentially small but *strictly positive* for every $\varepsilon > 0$.

By the Jentzsch–Perron theorem, the spectral radius of the integral operator with kernel K_ε is a *simple* eigenvalue with a *strictly positive* eigenfunction. On the symmetric domain $[-L, L]$, a positive eigenfunction must be even. We verify numerically: the dominant eigenfunction of K_ε is even with zero sign changes for all tested $\varepsilon \in [0.05, 0.5]$ and $\lambda \in \{30, 100\}$, and the second eigenfunction is odd.

This establishes that the mode coupling most strongly to the prime shifts is even. Since the prime coupling enters QW_λ with a sign that lowers the minimum eigenvalue, the even sector achieves the lowest eigenvalue.

Remark 2.22 (Remaining gap). The Jentzsch argument applies to the dominant eigenvalue of the *regularized* prime kernel, not directly to the minimum eigenvalue of QW_λ . A complete proof requires showing that: (i) the $\varepsilon \rightarrow 0$ limit preserves the even character of the dominant eigenmode, (ii) the dominant prime coupling mode aligns with the QW_λ ground state direction, and (iii) the parity-neutral archimedean contribution does not reverse the ordering. Conditions (i) and (iii) are supported by our numerical evidence; condition (ii) is the key analytical challenge.

2.15 Eigenvalue Spread Bound

A complementary approach uses the off-diagonal Frobenius norm to bound the minimum eigenvalue. For a Hermitian $N \times N$ matrix M , define $\|M\|_{\text{off}}^2 = \|M\|_F^2 - \sum_i M_{ii}^2$. Then the Schur–Horn inequality gives

$$\lambda_1(M) \leq \frac{\text{Tr}(M)}{N} - \frac{\|M\|_{\text{off}}}{\sqrt{N-1}}.$$

We compute the off-diagonal norms for the even and odd sectors:

λ	$\ W^+\ _{\text{off}}$	$\ W^-\ _{\text{off}}$	ratio
30	5.2	5.4	0.96
100	11.6	12.2	0.95
200	19.1	19.9	0.96

With the corrected orthogonal basis, the off-diagonal norms are nearly identical (ratio ≈ 0.96). The Schur–Horn bound therefore cannot distinguish the sectors. Even dominance arises not from a wider eigenvalue spread but from the *directional alignment* of the ground state with the prime shift operator: the first few cosine modes couple constructively to the prime shifts, producing a deeper minimum despite comparable overall spread.

The Schur–Horn bound is not tight enough with the corrected basis to prove even dominance (the bound $\lambda_1^+ \leq \text{Tr}/N - \|M\|_{\text{off}}/\sqrt{N-1}$ gives values around -1.4 to -3.2 , far above the actual $\lambda_1^- \approx -6$ to -15). A tighter approach uses the Rayleigh–Ritz variational principle: the ground state concentrates on the first $k = 3\text{--}4$ cosine modes, and the $k \times k$ principal submatrix already suffices. By the min-max principle (Galerkin upper bound),

$$\lambda_1(\text{QW}_\lambda^+) \leq \lambda_1(\text{QW}_\lambda^+[:k, :k]),$$

and we verify that $\lambda_1(\text{QW}_\lambda^+[:k, :k]) < \lambda_1(\text{QW}_\lambda^-)$ for $\lambda \geq 50$:

λ	k	$\lambda_1(\text{QW}_\lambda^+[:k])$	$\lambda_1(\text{QW}_\lambda^-)$
50	4	−6.68	−6.52
100	4	−11.18	−8.98
200	4	−19.16	−14.94

Moreover, $\lambda_1(\text{QW}_\lambda^-[:N])$ remains above $\lambda_1(\text{QW}_\lambda^+[:4])$ for *all* $N \leq 90$ (verified at 33 values of λ up to 1,300,000), indicating that the gap is genuine and not an artifact of truncation.

2.15.1 Computer-Assisted Gap Diagnostics via Interval Arithmetic

For $\lambda = 100$ and $\lambda = 200$, we have completed a computer-assisted finite-range diagnostic using interval arithmetic (mpmath.iv with 50-digit precision). In the current v2.4 status this is a conditional certificate pipeline with three components:

1. **Upper bound on λ_1^+ :** The 4×4 Galerkin matrix $\text{QW}^+[:4]$ is computed with certified intervals: prime shift integrals are evaluated in interval arithmetic (width $< 10^{-12}$), and the archimedean kernel is integrated via composite Gauss–Legendre quadrature with interval enclosure (width $< 10^{-3}$). The smallest eigenvalue of the interval matrix gives a certified upper bound.
2. **Candidate lower bound on λ_1^- :** The 40×40 Galerkin matrix provides $\lambda_1^-[:40]$. The tail contribution (modes 41–80 and beyond) is estimated via convergence analysis: the observed drop from $N = 40$ to $N = 80$ plus a geometric extrapolation for modes > 80 . This tail component is the part requiring independent validation.
3. **Conditional gap:** The interval upper bound on λ_1^+ lies strictly below the candidate lower bound on λ_1^- . Once the odd-sector tail estimate is validated as a genuine lower bound, this becomes a theorem-level finite certificate.

λ	$\lambda_1^+ \leq$	candidate $\lambda_1^- \geq$	conditional gap
100	−11.182	−9.448	−1.73
200	−19.161	−15.222	−3.94

These are the preliminary diagnostics with $N = 30$ modes and $k = 4$. The production certifier (Section A.8) uses larger truncation ($N_0 = 40\text{--}90$) and yields tighter gaps (e.g., -2.25 at $\lambda = 100$, -4.34 at $\lambda = 200$). Both methods provide strong evidence for even dominance; theorem-level verification depends on the odd-sector lower-bound validation.

2.16 Gap Asymptotics and Monotonicity

To assess whether even dominance persists for all large λ , we compute the gap $\Delta(\lambda) = \lambda_1^+(\lambda) - \lambda_1^-(\lambda)$ on a dense grid $\lambda \in [25, 500]$ with $N = 60$ (for $\lambda < 100$) and $N = 80$ (for $\lambda \geq 100$). Table 2 shows the results.

λ	N	λ_1^+	λ_1^-	Δ	$d\Delta/d\lambda$
50	60	-6.964	-6.944	-0.020	-0.007
55	60	-7.203	-6.697	-0.506	-0.097
100	80	-11.246	-9.061	-2.185	-0.017
200	80	-19.266	-15.049	-4.217	-0.025
300	80	-24.751	-19.378	-5.373	-0.014
500	80	-34.510	-26.883	-7.628	-0.011

Table 2: Even-odd gap $\Delta(\lambda) = \lambda_1^+ - \lambda_1^-$ for selected values (dense grid, $N = 60$ for $\lambda < 100$, $N = 80$ for $\lambda \geq 100$). Negative gap means even dominance. The derivative $d\Delta/d\lambda$ is negative for $\lambda \geq 140$, indicating monotonically deepening even dominance.

A least-squares fit to the data for $\lambda \geq 50$ yields

$$\Delta(\lambda) \approx a \cdot (\log \lambda)^b, \quad a \approx -0.0035, \quad b \approx 4.22. \quad (1)$$

The RMS residual is 0.34 for $\lambda \geq 50$ (29 data points). The fit extrapolates to $\Delta(1000) \approx -12.1$ and $\Delta(10000) \approx -40.7$, consistent with $\Delta \rightarrow -\infty$. A simpler linear model $\Delta(\lambda) \approx -2.91 \log \lambda + 11.08$ achieves a slightly better RMS (0.26) but lacks the correct curvature for extrapolation.

Remark 2.23 (Monotonicity and Hurwitz sufficiency). The gap $\Delta(\lambda)$ need not be monotone for Connes' program to succeed. Hurwitz's theorem on the zeros of uniform limits of holomorphic functions requires only a *sequence* $\lambda_n \rightarrow \infty$ along which $\Delta(\lambda_n) < 0$. Our data show $\Delta(\lambda) < 0$ for all $\lambda \geq 50$ (robustly for $\lambda \geq 55$, where $|\Delta| > 0.5$), with the gap reaching -7.6 at $\lambda = 500$. Two minor non-monotonicities occur: at the $N = 60 \rightarrow 80$ boundary ($\lambda = 90 \rightarrow 95$: $\delta = +0.13$, a truncation artifact) and at $\lambda = 120 \rightarrow 130$ ($\delta = +0.003$, negligible). For $\lambda \geq 140$, the derivative $d\Delta/d\lambda$ is consistently negative. The scaling (1) implies $\Delta(\lambda) \rightarrow -\infty$.

The Hellmann–Feynman analysis (Table 1) provides a structural explanation: the L -derivative $\partial\Delta/\partial L$ is *positive* for $\lambda \geq 80$, meaning the gap deepening is not driven by the growth of $L = \log \lambda$ but by the addition of new primes. The gap grows because $\pi(\lambda)$ grows, not because the interval $[-L, L]$ widens. This is consistent with the frontier-prime mechanism of Section 2.9.

3 Weighted Compactness as Proof Closure

The preceding sections establish strong finite-range gap diagnostics for finitely many λ (33 values up to $\lambda = 1,300,000$; numerically for all $\lambda \geq 55$). To close the conditional reduction for *all* large λ , we outline a functional-analytic strategy based on weighted compactness that could avoid the cumulative algebraic step entirely.

3.1 Setup and Rescaling

The Weil operator A_λ acts on $L^2([-L, L])$ with $L = \log \lambda$. The even-sector eigenfunctions ϕ_λ satisfying $A_\lambda \phi_\lambda = \mu_\lambda \phi_\lambda$, $\|\phi_\lambda\|_{L^2} = 1$, live on domains that grow with λ .

To obtain a common function space, we *rescale*: define

$$\psi_\lambda(u) := \sqrt{L} \phi_\lambda(Lu), \quad u \in [-1, 1]. \quad (2)$$

Then $\|\psi_\lambda\|_{L^2[-1,1]} = 1$ for all λ , and the rescaled operator on the fixed domain $[-1, 1]$ is

$$A_\lambda^{\text{resc}} \psi(u) = L \cdot (A_\lambda \phi)(Lu).$$

3.2 Building Block 1: Uniform Sobolev Bound

Proposition 3.1 (Mode concentration via Schur complement). *Let W be the $N \times N$ Galerkin matrix of QW_λ in the even sector, partitioned as*

$$W = \begin{pmatrix} A & B \\ B^T & C \end{pmatrix}$$

with A the $K \times K$ low-mode block and C the $(N - K) \times (N - K)$ high-mode block. If the smallest eigenvalue $\mu = \lambda_1(W)$ satisfies $\mu < \lambda_1(C)$, then the eigenvector $v = (c_{\text{low}}, c_{\text{high}})^T$ with $\|v\| = 1$ satisfies

$$\|c_{\text{high}}\|^2 \leq \frac{\|B\|_{\text{op}}^2}{(\lambda_1(C) - \mu)^2 + \|B\|_{\text{op}}^2}. \quad (3)$$

Proof. From the eigenvalue equation $Wv = \mu v$: $c_{\text{high}} = (\mu I - C)^{-1} B^T c_{\text{low}}$. Since $\mu < \lambda_1(C)$, the inverse is bounded: $\|(\mu I - C)^{-1}\| = 1/(\lambda_1(C) - \mu)$. Thus $\|c_{\text{high}}\| \leq \|B\|_{\text{op}}/(\lambda_1(C) - \mu) \cdot \|c_{\text{low}}\|$, and combining with $\|c_{\text{low}}\|^2 + \|c_{\text{high}}\|^2 = 1$ gives (3). $\square$

λ	μ	$\lambda_1(C)$	$\lambda_1(C) - \mu$	$\ B\ _{\text{op}}$	Bound	Actual $\ c_{\text{high}}\ ^2$
30	-7.53	-5.79	1.74	1.00	0.332	0.008
50	-9.55	-5.95	3.59	1.45	0.162	0.003
100	-11.08	-6.25	4.83	2.19	0.206	0.001
200	-19.10	-8.02	11.08	3.61	0.106	0.001
500	-25.39	-10.04	15.35	3.55	0.053	0.002

Table 3: Schur complement bound on high-mode energy for $K = 5$, $N = 30$. The denominator $\lambda_1(C) - \mu$ grows because $\mu \sim -CL^2 \rightarrow -\infty$ while $\lambda_1(C) = O(L)$. The bound decreases monotonically.

Remark 3.2 (Asymptotic control). Power-law fits over $\lambda \in [30, 500]$ (9 data points) yield:

Quantity	Exponent α	Fit
$ \mu $	2.16	$0.50 L^{2.16}$
$ \lambda_1(C) $	0.98	$1.57 L^{0.98}$
$\ B\ _{\text{op}}$	2.64	$0.04 L^{2.64}$

The denominator scales as $\lambda_1(C) - \mu \sim 0.02 L^{3.66}$, giving the bound $\|c_{\text{high}}\|^2 \leq C L^{2 \cdot 2.64 - 2 \cdot 3.66} = C L^{-2.03} \rightarrow 0$. The mode concentration improves as $O(1/L^2)$, providing:

$$\|\psi'_\lambda\|_{L^2[-1,1]}^2 = \pi^2 \sum n^2 c_n^2 \leq \pi^2 (K^2 + N^2 \cdot O(1/L^2)) = O(1).$$

The structural reason: μ grows quadratically (driven by prime resonances in low modes), while $\lambda_1(C)$ grows only linearly (high modes couple weakly to primes). This gap widens monotonically.

3.3 Building Block 2: Precompactness via Rellich

Given the uniform Sobolev bound (Building Block 1):

Proposition 3.3 (Precompactness). *If $\sup_\lambda \|\psi_\lambda\|_{H^1[-1,1]} < \infty$, then the family $\{\psi_\lambda : \lambda \geq \lambda_0\}$ is precompact in $L^2[-1,1]$. In particular, every sequence $\lambda_n \rightarrow \infty$ has a subsequence along which $\psi_{\lambda_{n_k}} \rightarrow \psi_\infty$ strongly in $L^2[-1,1]$.*

Proof. By the Rellich–Kondrachov theorem, the embedding $H^1[-1,1] \hookrightarrow L^2[-1,1]$ is compact. A bounded sequence in H^1 therefore has a strongly convergent subsequence in L^2 . $\square$

3.4 Building Block 3: Spectral Stability

The precompactness of $\{\psi_\lambda\}$ does not automatically imply that the limit ψ_∞ is an eigenfunction of a meaningful limit operator. We need:

Conjecture 3.4 (Form convergence). *The rescaled quadratic forms $q_\lambda(\psi) = \langle A_\lambda^{\text{resc}} \psi, \psi \rangle$ converge in the Mosco sense to a limit form q_∞ as $\lambda \rightarrow \infty$. Equivalently, the resolvents converge strongly: $(A_\lambda^{\text{resc}} - zI)^{-1} \rightarrow (A_\infty - zI)^{-1}$ for $z \notin \sigma(A_\infty)$.*

In the rescaled variable u , the prime shifts become $m \log p / L \rightarrow 0$ as $L \rightarrow \infty$. A Taylor expansion yields

$$\psi(u - m \log p / L) \approx \psi(u) - \frac{m \log p}{L} \psi'(u) + \frac{(m \log p)^2}{2L^2} \psi''(u) + \dots$$

The leading correction is a second-derivative term, suggesting that the limit form q_∞ involves a Laplacian with arithmetically determined coefficients:

$$q_\infty(\psi) \approx \text{const} \cdot \|\psi\|^2 - \sigma_{\text{arith}} \cdot \|\psi'\|^2 + (\text{higher order}), \quad (4)$$

where $\sigma_{\text{arith}} = \sum_p \sum_{m \geq 1} \log p \cdot p^{-m/2} \cdot (m \log p)^2 / 2$ converges absolutely.

3.5 Building Block 4: Bridge to Hurwitz

For Connes' Theorem 6.1, we need locally uniform convergence of the Mellin transforms:

$$F_\lambda(z) = \int_{-L}^L \phi_\lambda(t) e^{-zt} dt.$$

Conjecture 3.5 (Paley–Wiener control). *There exists $\alpha > \frac{1}{2}$ such that*

$$\sup_{\lambda} \int |\phi_{\lambda}(t)|^2 e^{2\alpha|t|} dt < \infty.$$

Consequently, $\{F_{\lambda}\}$ converges locally uniformly on the strip $\{z : |\operatorname{Re} z| < \alpha\}$.

If polynomial weights suffice for precompactness (the uniform Sobolev bound above), the upgrade to exponential decay is non-trivial and constitutes the deepest analytical ingredient. The support constraint $\phi_{\lambda}|_{[-L,L]^c} = 0$ gives

$$\int |\phi_{\lambda}|^2 e^{2\alpha|t|} dt \leq e^{2\alpha L} \|\phi_{\lambda}\|^2 = \lambda^{2\alpha},$$

which is bounded for fixed λ but grows with λ . A sharper bound requires the eigenfunction to decay *within* its support—i.e., $|\phi_{\lambda}(t)|$ must be small near $|t| \approx L$. This is precisely the mass concentration observed in our numerical experiments (§2).

3.6 Numerical Verification

We test the key ingredients on $\lambda \in \{30, 50, 100, 200, 500\}$ with $N = 25$ Galerkin modes.

Uniform Sobolev bound (B1). The quantity $S(\lambda) := \sum_n n^2 c_n(\lambda)^2$ controls $\|\psi'_{\lambda}\|_{L^2[-1,1]}^2 = \pi^2 S(\lambda)$.

λ	$S(\lambda)$ (even)	$\ \psi\ _{H^1}$	Top modes	$S(\lambda)$ (odd)	$\ \psi\ _{H^1}$
30	7.54	8.68	$n = 1, 15, 23$	504.4	70.6
50	1.80	4.34	$n = 1, 2, 0$	301.9	54.6
100	1.44	3.90	$n = 1, 2, 0$	12.2	11.0
200	1.35	3.78	$n = 1, 2, 0$	9.9	9.9
500	1.67	4.18	$n = 1, 2, 0$	11.5	10.7

The even-sector $S(\lambda)$ stabilizes at ≈ 1.3 – 1.7 for $\lambda \geq 100$ with dominant modes $n = 0, 1, 2$ —strongly supporting the uniform Sobolev bound. The odd sector converges more slowly (transition regime at $\lambda \leq 50$).

Eigenfunction profile convergence. Successive $L^2[-1, 1]$ distances of the rescaled even eigenfunctions:

$\lambda_1 \rightarrow \lambda_2$	$\ \psi_{\lambda_1} - \psi_{\lambda_2}\ _{L^2}$
30 $\rightarrow$ 50	0.278
50 $\rightarrow$ 100	0.081
100 $\rightarrow$ 200	0.048
200 $\rightarrow$ 500	0.150

The distances decrease from 0.278 to 0.048, consistent with a Cauchy sequence. The anomaly at $\lambda = 200 \rightarrow 500$ is likely an N -truncation artifact ($N = 25$ becomes insufficient for $\lambda = 500$; cf. the low-mode block result which requires $k \geq 5$ modes at $\lambda = 500$).

Eigenvalue scaling. The leading diagonal W_{11}^+ grows as $\sim e^{L/2} = \sqrt{\lambda}$ (see Theorem 3.7), driven by PNT. On the numerical range $\lambda \in [30, 500]$, a power-law fit gives $-W_{11}^+ \approx 0.38 L^{2.29}$, but the analytical asymptotics is exponential:

λ	λ_1^+	W_{11}^+	$\lambda_1^+/\sqrt{\lambda}$	$W_{11}^+/\sqrt{\lambda}$
30	-6.31	-7.05	-1.15	-1.29
100	-11.07	-9.15	-1.11	-0.92
200	-19.10	-16.42	-1.35	-1.16
500	-25.38	-29.90	-1.13	-1.34

The ratios $\lambda_1^+/\sqrt{\lambda}$ and $W_{11}^+/\sqrt{\lambda}$ are of order 1, consistent with the PNT-derived scaling $W_{11}^+ \sim -c\sqrt{\lambda}$.

3.7 Assessment and Risks

The weighted compactness strategy decomposes the proof closure into four building blocks (uniform Sobolev bound, Rellich precompactness, Mosco convergence, Paley–Wiener control). Each is supported by numerical evidence:

Block	Status	Evidence
B1: H^1 bound	Schur complement	$\ c_{\text{high}}\ ^2 = O(L^{-2})$ (Theorem 3.1)
B2: Rellich	Theorem	Standard (given B1)
B3: Spectral limit	Open	λ_1^+/L^2 stabilizes, no element-wise convergence
B4: Hurwitz bridge	Connes Fact 6.4	$k_\lambda \rightarrow \Xi$ at rate $O(\lambda^{-1/2-\alpha})$

Principal risks:

- M1:** The H^1 bound may fail if the mode concentration weakens for very large λ (e.g., if the eigenfunction develops increasingly fine oscillations). Numerical tests up to $\lambda = 500$ show no such tendency: the dominant modes are stably $n = 0, 1, 2$.
- M2:** Mosco convergence requires uniform resolvent bounds. The singular archimedean kernel $K(s) \sim 1/s$ at $s = 0$ may create difficulties in the rescaled limit. The observed λ_1^+/L^2 stabilization provides indirect evidence.
- M3:** The Paley–Wiener upgrade requires exponential, not polynomial, control—a fundamentally stronger condition. The mass concentration data (tail $< 1.5\%$ at $|u| > 0.95$ for $\lambda = 500$) is encouraging but not sufficient.

The alternative approaches (Hellmann–Feynman monotonicity, index stability) remain viable and may be combinable with the compactness argument: e.g., if monotonicity can be proved for a dense subsequence, compactness extends it to all large λ .

3.8 Gap Growth from Block Bounds

The preceding analysis suggests a clean closure of the even dominance proof via three block-level estimates. The key insight is that the correct scaling of the leading diagonal is *exponential* in L , not polynomial: $W_{11}^+ \sim -c\sqrt{\lambda} = -ce^{L/2}$, driven by the prime-number theorem. This makes the gap argument extremely robust: an exponentially negative low-mode term overwhelms any polynomially bounded high-mode coupling.

3.8.1 Lemma L1: Prime coupling drives an exponentially deep even direction

The diagonal element W_{11}^+ decomposes as

$$W_{11}^+ = \underbrace{\log(4\pi) + \gamma_E}_{=3.272} + (W_{\text{arch}})_{11} + (W_{\text{prime}})_{11}. \quad (5)$$

The archimedean term $(W_{\text{arch}})_{11} = O(1)$ is bounded (cf. Table: ≈ -0.74 at $\lambda = 100$). The prime term dominates:

$$(W_{\text{prime}})_{11} = 2 \sum_{p \leq \lambda} \sum_{m=1}^{\lfloor 2L/\log p \rfloor} \frac{\log p}{p^{m/2}} S_{11}(m \log p, L),$$

where $S_{11}(\delta, L) = \frac{1}{L} \int_{\max(-L, \delta-L)}^{\min(L, \delta+L)} \cos\left(\frac{\pi t}{L}\right) \cos\left(\frac{\pi(t-\delta)}{L}\right) dt$.

Lemma 3.6 (Prime powers are negligible). *The $m \geq 2$ terms contribute $O(L)$:*

$$\sum_p \sum_{m \geq 2} \frac{\log p}{p^{m/2}} |S_{11}(m \log p, L)| \leq C_1 L.$$

Proof sketch. $\sum_{m \geq 2} p^{-m/2} \leq C/p$ and $|S_{11}| \leq 1$, so the sum is $\leq C \sum_{p \leq \lambda} (\log p)/p = O(L)$ by Mertens' theorem. $\square$

Thus the leading term is the $m = 1$ contribution:

$$(W_{\text{prime}})_{11} = 2 \sum_{p \leq \lambda} \frac{\log p}{\sqrt{p}} S_{11}(\log p, L) + O(L). \quad (6)$$

The overlap $S_{11}(\log p, L) \geq c_0 > 0$ uniformly for p in the “bulk” $p \leq \lambda^{1-\varepsilon}$ (where $\log p/L < 1 - \varepsilon$, so the overlap length is $\geq 2\varepsilon L$ and the cosine factor stays bounded away from zero). By partial summation and PNT ($\sum_{p \leq x} \log p / \sqrt{p} = 2\sqrt{x} + o(\sqrt{x})$), with $x = \lambda = e^L$:

Lemma 3.7 (Leading mode grows exponentially). *There exist $c > 0$ and L_0 such that for $L \geq L_0$:*

$$W_{11}^+(\lambda) \leq -c \lambda^{1/2} = -c e^{L/2}.$$

Numerically: $W_{11}^+ = -7.0, -9.3, -9.1, -16.4, -29.9$ for $\lambda = 30, 50, 100, 200, 500$ (fit: $-0.38 L^{2.29}$; true asymptotics $\sim e^{L/2}$).

3.8.2 Lemma L2: High-mode block is only polynomially negative

Lemma 3.8 (High-block control). *For $C = \text{QW}[K:, K:]$ with $K \geq 3$, there exists $\beta > 0$ such that $\lambda_1(C) \geq -\beta L$ for all $\lambda \geq \lambda_0$.*

Proof via Gershgorin. Each diagonal entry satisfies $C_{nn} = \log(4\pi) + \gamma_E + O(L)$ (the prime coupling of high modes is $O(L)$ by the same Mertens bound as Theorem 3.6). Each Gershgorin radius $R_n = \sum_{j \neq n} |C_{nj}| = O(L)$ (at most $O(L/\log 2)$ non-zero overlap terms, each bounded by $O(1)$). Hence $\lambda_1(C) \geq \min_n (C_{nn} - R_n) \geq -\beta L$. $\square$

Numerically: $\lambda_1(C)/L \in [-1.70, -1.36]$ for $\lambda \in [30, 500]$, confirming the $O(L)$ scaling.

3.8.3 Lemma L3: Low–high coupling (historical)

Lemma 3.9 (Off-diagonal control — superseded). *For $B = \text{QW}[:, K, K:]$ and $\lambda \in [30, 500]$, numerically $\|B\|_{\text{op}}/L \in [0.23, 1.64]$.*

Note. The original statement claimed $\|B\|_{\text{op}} \leq \gamma L$ for all λ . As shown in Section A.4 below, the rigorous Hilbert–Schmidt bound gives $\|B\|_{\text{op}} = O(\sqrt{\lambda}) = O(e^{L/2})$, which exceeds $O(L)$ asymptotically. The $O(L)$ scaling holds only on the tested range $\lambda \in [30, 500]$. **This lemma is not used in the proof of Theorem A.36**, which instead relies on the resolvent-damped framework (Sections A.7 and A.9) that controls the coupling *differential* between sectors without requiring entry-wise bounds on $\|B\|$.

3.8.4 The gap growth theorem

Proposition 3.10 (Even dominance grows exponentially). *Under Theorems 3.7 to 3.9, the even-sector eigenvalue satisfies*

$$\lambda_1^+(\lambda) \leq -c e^{L/2} + O(L) \rightarrow -\infty.$$

If the leading odd diagonal satisfies $W_{00}^-(\lambda) \geq -c_- e^{L/2}$ with $c_- < c$ (which follows from the Shift Parity Lemma applied to the $m = 1$ overlaps), then

$$\Delta(\lambda) = \lambda_1^+(\lambda) - \lambda_1^-(\lambda) \leq -(c - c_-) e^{L/2} + O(L) \rightarrow -\infty. \quad (7)$$

Proof. Upper bound on λ_1^+ : By the variational principle, $\lambda_1^+(W) \leq W_{11}^+ \leq -c e^{L/2}$ (Theorem 3.7).

Lower bound on λ_1^+ (coupling correction): For $\mu \leq -c e^{L/2}$ and $\lambda_1(C) \geq -\beta L$ (Theorem 3.8), the denominator satisfies $\lambda_1(C) - \mu \geq c e^{L/2} - \beta L \geq \frac{c}{2} e^{L/2}$ for $L \geq L_0$. With $\|B\| \leq \gamma L$ (Theorem 3.9):

$$\|B(C - \mu I)^{-1} B^T\| \leq \frac{\gamma^2 L^2}{\frac{c}{2} e^{L/2}} = O(L^2 e^{-L/2}) \rightarrow 0.$$

The coupling correction *vanishes* exponentially. Hence $\lambda_1^+(W) = -c e^{L/2} + o(1)$.

Gap: The Shift Parity Lemma (Theorem 2.8) gives the pointwise inequality $S_{11}^{\cos}(\delta, L) < S_{11}^{\sin}(\delta, L)$ for $\delta \in (0, 2L)$. Since all prime-sum weights $\log p/\sqrt{p}$ are positive, this implies $(W_{\text{prime}}^+)_{11} < (W_{\text{prime}}^-)_{00}$, i.e., the cosine diagonal is *more negative* than the sine diagonal. The gap follows. $\square$

λ	W_{11}^+	W_{00}^-	$W_{11}^+ - W_{00}^-$	$(W_{11}^+ - W_{00}^-)/L^2$
30	−7.05	−6.70	−0.35	−0.030
50	−9.26	−5.98	−3.28	−0.214
100	−9.15	−3.02	−6.13	−0.289
200	−16.42	−7.57	−8.86	−0.315
500	−29.90	−15.83	−14.07	−0.364

Table 4: Leading-mode diagonal comparison. The even diagonal W_{11}^+ is consistently more negative than the odd diagonal W_{00}^- , with the difference growing monotonically. This is the Shift Parity Lemma on the leading Fourier mode.

Remark 3.11 (Exponential vs. polynomial). The earlier fit $\Delta \sim -0.004(\log \lambda)^{4.2}$ (Equation (1)) reflects the *numerical range* $\lambda \in [50, 500]$; on this range, $e^{L/2}$ and $L^{4.2}$ are hard to distinguish. The analytical argument via PNT reveals the true asymptotics: the gap grows as $e^{L/2} = \sqrt{\lambda}$, driven by the prime-counting function. This makes the “dominant term beats coupling” argument *trivially* robust: the coupling correction is $O(L^2 e^{-L/2}) \rightarrow 0$.

4 Connections to Existing Work

4.1 Li–Keiper Coefficients

The coefficients λ_n (also called Li–Keiper coefficients) have been studied extensively. Keiper [14] computed them numerically; Li [15] proved the criterion; Bombieri–Lagarias [1] provided the representation used here. Voros [16] connected them to the Stieltjes constants. Our contribution is the thermodynamic interpretation and the archimedean–finite decomposition, which may suggest new analytical handles.

4.2 Weil’s Explicit Formula and Connes’ Program

The Weil explicit formula

$$\sum_{\rho} h(\rho) = h(0) + h(1) - \sum_p \sum_{m=1}^{\infty} \frac{\log p}{p^{m/2}} \hat{h}(m \log p)$$

provides a direct bridge between zeros and primes. The Li coefficients correspond to the test function $h_n(s) = 1 - (1 - 1/s)^n$. Connes [2] constructs from this formula a self-adjoint operator A_λ and reduces RH to the simplicity of its ground state with even eigenfunction (Theorem 6.1; see Section 2 for our numerical investigation). The prolate spheroidal wave operator, which commutes with the compressed Fourier transform, provides an analogy: its eigenvalues in the even sector are provably simple (Fact 6.3 in [2]). Extending this simplicity to A_λ is the remaining open problem.

Prior to the present work, the even-dominance property was an *unproven assumption* in all approaches: Connes [2] assumes it, and Connes–van Suijlekom [4] assume it in their Carathéodory–Fejér truncation argument. Classical approaches via Krein–Rutman (positive kernels) or Sturm–Liouville (nodal theory) are not applicable because A_λ has an *indefinite* kernel. The present paper resolves this assumption via Theorem A.36.

Our decomposition analysis (Section 2) reveals that the prime contributions, rather than the archimedean kernel, are the sole driver of even dominance—a structural insight that may inform future analytical approaches.

4.3 de Branges Spaces and Hilbert–Pólya

The Hilbert–Pólya conjecture seeks a self-adjoint operator whose eigenvalues are the non-trivial zeros. If such an operator exists, its positivity properties could provide the “manifest sign” representation sought in OP5. The spectral census from Part I ($\dim V^- = 2$) may be related to the index of the Dirac-type operator in this context.

5 Galerkin Imbalance: Boundary-Moment Benchmark (v3.0)

The finite-dimensional Galerkin approximation to the Weil quadratic form QW_λ introduces an imbalance between the even and odd sectors that depends on both the cutoff parameter λ and the Galerkin dimension N . We define the *scaled ladder imbalance*

$$\text{sLI}(\lambda, N) = \frac{I_2^{\text{ladder}}(\lambda, N)}{I_1^{\text{ladder}}(\lambda, N)} \cdot \Gamma(\lambda)^{-1/4}, \quad (8)$$

where I_1^{ladder} and I_2^{ladder} denote the L^1 -weighted Chebyshev moment ladder sums on the even and odd far-tail sectors, respectively, and $\Gamma(\lambda)$ is the Galerkin quarter-power normalization.

5.1 Identification of the Boundary-Moment Constant

Numerical computation on the Mac Studio (converged benchmark regime $\xi = NL/\lambda \gtrsim 50$, $\text{dps} = 4N$ for precision safety) reveals that, for the tested benchmark values $\lambda \in \{5, 7\}$, sLI is well described by the boundary-moment benchmark

$$\text{sLI}(\lambda, N) = \frac{C^* \lambda}{NL} + o\left(\frac{\lambda}{NL}\right), \quad L = 2 \log \lambda, \quad (9)$$

with a single dimensionless benchmark constant

$$C^* = \frac{64\pi^2}{3} = 210.5515 \dots \quad (10)$$

Remark 5.1 (Derivation candidate). The value C^* admits the formal integral representation

$$C^* = 16 \int_0^\infty \frac{t^2}{\sinh^2(t/2)} dt = 16 \cdot \frac{4\pi^2}{3} = \frac{64\pi^2}{3},$$

where $\int_0^\infty t^2 / \sinh^2(t/2) dt = 4\pi^2/3$ is a standard Bernoulli integral. This is an identification candidate, not a forward derivation. The 27 May audit falsified the naive endpoint factorisation $K(p, \gamma) = \text{const} \cdot t^2 / \sinh^2(t/2)$: in the variable $t = \log(\gamma/p)$ the exact paired Cauchy kernel is

$$K(t) = \frac{e^t \cosh t}{2 \sinh^2 t},$$

with a collar singularity at $t = 0$. A proof of C^* therefore requires the missing eigenmode asymptotics for the weights $q_a(k)$, not only the scalar kernel identity.

Theorem 5.2 (GF1: Conditional boundary-moment scaling). *Assume that the eigenmode-weighted Galerkin kernel of QW_λ has a production-sLI limit in which the exact collar kernel above reduces to the effective boundary moment $16 \int_0^\infty t^2 / \sinh^2(t/2) dt$. Assume additionally that the production observable is stable under near-cluster rotations and that the $\lambda = 5, 7$ benchmark persists beyond the current $\lambda = 11$ discriminator. Then*

$$\text{sLI}(\lambda, N) = \frac{64\pi^2}{3} \cdot \frac{\lambda}{NL} + o\left(\frac{\lambda}{NL}\right).$$

Table 5: Data check for the boundary-moment constant C^* . The effective constant $C_{\lambda,N} = \text{sLI} \cdot NL/\lambda$ is compared against $C^* = 64\pi^2/3 \approx 210.55$. The $\lambda = 5, 7$ rows are the current benchmark support; the $\lambda = 11$ row is a single open discriminator, not confirmation of a universal $\lambda \geq 5$ law. All runs use Chebyshev degree 80, $j = 25$.

λ	N	$\xi = NL/\lambda$	$C_{\lambda,N}$	deviation	status
5	110	70.8	211.5	+0.5%	benchmark
5	150	96.6	211.0	+0.2%	benchmark
7	200	111.2	208.8	-0.8%	benchmark
11	200	87.2	167.6	-20.4%	open

Remark 5.3 (Low-bandwidth guardrail). The same normalisation is *not* a theorem uniformly down to $\lambda = 3$. A dedicated N -sweep at $\lambda = 3$ indicates a different low-bandwidth plateau, with the older C_m normalisation approaching approximately 7 rather than the $\lambda = 5, 7$ benchmark plateau near 1.5–2.0. A separate $\lambda = 11, N = 200$ run gives $C_{\lambda,N}^{\text{eff}} \approx 167.6$, about 20% below $64\pi^2/3$; the pending $N = 280$ and $N = 360$ runs decide whether this is pre-convergence or genuine λ -dependence. Thus GF1 is currently a conditional benchmark, not a single $\lambda \geq 5$ universal theorem.

Remark 5.4 (Convergence regime). The relevant scaling variable is

$$\xi = \frac{NL}{\lambda}, \quad L = 2 \log \lambda,$$

not N/L . For $\lambda = 5, 7$, the data enter the benchmark window once $\xi \gtrsim 50$ –60; below that window finite-dimensional boundary effects distort $C_{\lambda,N}$. The $\lambda = 11$ discriminator shows that the threshold may itself be λ -dependent.

5.2 Hypothesis Correction Chain

In the spirit of the audit-trail philosophy of this series, we record the iterative corrections to the GF1 hypothesis during its development:

1. **Iter 15 (2026-05-24):** Initial hypothesis: $C_m \approx 8$ constant, based on $\lambda = 5$, $N = 55$ data.
2. **Iter 18 (2026-05-25):** $C_m \approx 8$ falsified by N -convergence sweep; scaledLadderI decreases monotonically with N .
3. **Iter 19 (2026-05-25):** Revised hypothesis: plateau at $C_m \approx 1.5$ –2.0.
4. **Iter 20 (2026-05-26):** Multi- λ confirmation ($\lambda = 5, 7$): benchmark consistency in $\xi = NL/\lambda$ confirmed within 5%.
5. **Iter 27 (2026-05-26, GPT audit):** Universal constant $C^* = 64\pi^2/3$ identified via integral representation; boundary-moment kernel $16t^2/\sinh^2(t/2)$ as derivation candidate.
6. **Iter 28 (2026-05-26):** The $\lambda = 3$ N -sweep breaks the all- λ universal reading. The safe formulation at this stage was a $\lambda = 5, 7$ boundary-moment benchmark plus a conservative low-bandwidth branch at $\lambda = 3$.

7. **Iter 29 (2026-05-27):** The endpoint-kernel factorisation is falsified by the exact kernel $K(t) = e^t \cosh(t)/(2 \sinh^2 t)$, and the $\lambda = 11, N = 200$ point gives $C_{\text{eff}} \approx 168$. The safe formulation is now a $\lambda = 5, 7$ benchmark with $\lambda = 11$ pending.

The initial C_m -constancy hypothesis was an artifact of the transient Galerkin regime. In the converged benchmark regime for $\lambda = 5, 7$, sLI follows the $C^* \lambda / (NL)$ law; the $\lambda = 3$ counter-sweep and the open $\lambda = 11$ discriminator keep GF1 at the level of a conditional target rather than a universal theorem.

6 Conclusion

This paper no longer claims to rigorously establish even dominance of the Weil quadratic form QW_λ unconditionally, even on the finite certified range. The current architecture combines three ingredients: (1) the Shift Parity Lemma, showing that each prime individually favors even eigenfunctions; (2) 33 computer-assisted finite-range diagnostics providing strong interval-arithmetic evidence pending the odd-sector lower-bound validation; and (3) the M1'' resolvent framework together with the v2.1 Direct Frontier-Dominance argument, both of which establish the asymptotic variational statement $\lambda_{\min}(W^+ - W^-) = -\Theta(\sqrt{\lambda})$. The remaining step from this variational statement to $\lambda_1^+(\lambda) < \lambda_1^-(\lambda)$ for all sufficiently large λ is formulated in v2.2 as the *Asymptotic Variational Gap Conjecture*.

The central analytical insight is the Leading-Mode Cancellation Lemma: the coupling difference between even and odd sectors is bounded by $O(1)$ while the diagonal advantage $D(\lambda)$ grows as $\sqrt{\lambda}$, producing a ratio $\rho(\lambda) \rightarrow 0$. Combined with the Euler–Maclaurin Proposition ($\rho^{\text{EM}} < 0$ for $L \geq 13$, certified by interval arithmetic) and the PNT Transfer Lemma, this yields M1'' (resolvent subdominance) for all $\lambda \geq \lambda_0$.

In v1.1, both Lemma B and Lemma C have been upgraded to fully analytical arguments: the parity-split gap (Lemma B, Steps 3–4) derives $c_{\text{gap}} \geq 4\pi^2/L$ from Laplace asymptotics, and the sector-difference control (Lemma C) replaces the global Neumann condition $\|R_0 K\| < 1$ by the tail-symmetry bound $|E_{\sin}^{\text{tail}} - E_{\cos}^{\text{tail}}|/D \leq C/N_0^2$. No numerical constants remain in the asymptotic variational analysis of Regime 2. Version 2.1 adds a robust Direct Frontier-Dominance proof (Theorem A.37) that uses a common Rayleigh vector rather than eigenvalue-additivity and removes the two earlier v2.0 caveats; the v2.2 revision clarifies that this still yields a conditional reduction rather than an unconditional closure of A6.

A Proofs of the Block-Bound Lemmas

We provide complete proofs of the three lemmas underlying the gap growth theorem (Theorem 3.10). Throughout, $L = \log \lambda$, and $\lambda = e^L$. The even-sector basis is $e_n(t) = \cos(n\pi t/L)/\sqrt{L}$ for $n \geq 1$ and $e_0(t) = 1/\sqrt{2L}$, orthonormal on $[-L, L]$.

A.1 Closed-form shift overlaps

Proposition A.1 (Shift overlap). *For $0 \leq \delta \leq 2L$, the shift overlaps of the leading modes are:*

$$S^+(\delta, L) := \langle e_1, T_\delta e_1 \rangle = \left(1 - \frac{\delta}{2L}\right) \cos \frac{\pi\delta}{L} - \frac{1}{2\pi} \sin \frac{\pi\delta}{L}, \quad (11)$$

$$S^-(\delta, L) := \langle f_0, T_\delta f_0 \rangle = \left(1 - \frac{\delta}{2L}\right) \cos \frac{\pi\delta}{L} + \frac{1}{2\pi} \sin \frac{\pi\delta}{L}, \quad (12)$$

where $f_0(t) = \sin(\pi t/L)/\sqrt{L}$ is the leading odd mode and $T_\delta g(t) = g(t - \delta) \cdot \mathbf{1}_{|t-\delta| \leq L}$.

Proof. The overlap integral runs over the intersection $[\delta - L, L]$ (length $2L - \delta$). Using the product-to-sum formula $\cos a \cos b = \frac{1}{2}[\cos(a - b) + \cos(a + b)]$:

$$\int_{\delta-L}^L \cos \frac{\pi t}{L} \cos \frac{\pi(t-\delta)}{L} dt = \frac{1}{2} \cos \frac{\pi\delta}{L} (2L - \delta) + \frac{1}{2} \int_{\delta-L}^L \cos \frac{\pi(2t-\delta)}{L} dt.$$

The second integral evaluates (via $u = \pi(2t - \delta)/L$) to $-\frac{L}{\pi} \sin \frac{\pi\delta}{L}$, giving (11) after dividing by L . The sine case follows identically from $\sin a \sin b = \frac{1}{2}[\cos(a - b) - \cos(a + b)]$, with the opposite sign on the second integral. $\square$

Corollary A.2 (Shift Parity on leading modes). *For all $\delta \in (0, L)$:*

$$S^+(\delta, L) - S^-(\delta, L) = -\frac{1}{\pi} \sin \frac{\pi\delta}{L} < 0. \quad (13)$$

In particular, $S^+(\delta, L) < S^-(\delta, L)$: the cosine mode has smaller overlap with prime shifts than the sine mode.

Proof. Subtract (12) from (11). For $\delta \in (0, L)$, $\pi\delta/L \in (0, \pi)$, so $\sin(\pi\delta/L) > 0$. $\square$

Remark A.3 (Geometric interpretation). The mode $\cos(\pi t/L)$ vanishes at $t = \pm L/2$, while $\sin(\pi t/L)$ is maximal there. A shift by $\delta \approx L/2$ (corresponding to primes $p \approx \sqrt{\lambda}$) moves the support into a region where the cosine value is small, reducing the overlap. The sine mode, being maximal in that region, retains a larger overlap. This geometric asymmetry is the microscopic mechanism behind even dominance.

A.2 Proof of Lemma L1: Exponential depth of the even direction

Lemma A.4 (Restatement of Theorem 3.7). *There exist $c > 0$ and λ_0 such that for all $\lambda \geq \lambda_0$:*

$$W_{11}^+(\lambda) \leq -c\sqrt{\lambda}.$$

Proof. Step 1: Decomposition.

$$W_{11}^+ = \underbrace{(\log 4\pi + \gamma_E)}_{=3.272} + (W_{\text{arch}})_{11} + (W_{\text{prime}})_{11}.$$

The archimedean term satisfies $(W_{\text{arch}})_{11} = O(1)$: the kernel $K(s) = e^{s/2}/(2 \sinh s)$ decays exponentially for $s \rightarrow \infty$ and the overlap $S^+(s, L)$ is bounded, so the integral $\int_0^\infty K(s)[S^+(s, L) + S^+(-s, L) - 2e^{-s/2}] ds$ converges absolutely, uniformly in L .

Step 2: Prime powers are negligible. The $m \geq 2$ contribution satisfies

$$\left| \sum_{p \leq \lambda} \sum_{m \geq 2} \frac{\log p}{p^{m/2}} S^+(m \log p, L) \right| \leq \sum_{p \leq \lambda} \frac{\log p}{p} \cdot \frac{1}{1 - p^{-1/2}} \leq C \sum_{p \leq \lambda} \frac{\log p}{p} = O(L),$$

using $|S^+| \leq 1$, $\sum_{m \geq 2} p^{-m/2} \leq C/p$, and Mertens' theorem $\sum_{p \leq x} (\log p)/p = \log x + O(1)$.

Step 3: Main term.

$$(W_{\text{prime}})_{11} = 2 \sum_{p \leq \lambda} \frac{\log p}{\sqrt{p}} S^+(\log p, L) + O(L). \quad (14)$$

Step 4: Sign analysis of $S^+(\log p, L)$. By Theorem A.1, $S^+(\delta, L)$ has a unique zero at $\delta_0 \approx 0.436 L$. That is, $S^+(\log p, L) < 0$ whenever $\log p > 0.436 L$, i.e., $p > \lambda^{0.436}$.

For $\delta \in [L/2, L]$ (i.e., $p \in [\sqrt{\lambda}, \lambda]$):

$$S^+(\delta, L) = \underbrace{\left(1 - \frac{\delta}{2L}\right)}_{\in [1/4, 1/2]} \underbrace{\cos \frac{\pi \delta}{L}}_{\in [-1, 0]} - \underbrace{\frac{1}{2\pi} \sin \frac{\pi \delta}{L}}_{\in [0, 1/(2\pi)]} \leq -\frac{1}{2\pi} < 0.$$

More precisely, $S^+(L/2, L) = -1/(2\pi)$ and $S^+(L, L) = -1/2$.

Step 5: PNT gives the growth rate. Split the sum (14) at $p = \sqrt{\lambda}$:

Negative part ($p > \sqrt{\lambda}$, where $S^+ \leq -1/(2\pi)$):

$$2 \sum_{\sqrt{\lambda} < p \leq \lambda} \frac{\log p}{\sqrt{p}} S^+(\log p, L) \leq -\frac{1}{\pi} \sum_{\sqrt{\lambda} < p \leq \lambda} \frac{\log p}{\sqrt{p}}.$$

By the prime number theorem with partial summation:

$$\sum_{p \leq x} \frac{\log p}{\sqrt{p}} = 2\sqrt{x} + o(\sqrt{x}),$$

$$\text{so } \sum_{\sqrt{\lambda} < p \leq \lambda} \frac{\log p}{\sqrt{p}} = 2\sqrt{\lambda} - 2\lambda^{1/4} + o(\sqrt{\lambda}) = 2\sqrt{\lambda}(1 + o(1)).$$

Positive part ($p \leq \sqrt{\lambda}$, where $|S^+| \leq 1$):

$$\left| 2 \sum_{p \leq \sqrt{\lambda}} \frac{\log p}{\sqrt{p}} S^+(\log p, L) \right| \leq 2 \sum_{p \leq \sqrt{\lambda}} \frac{\log p}{\sqrt{p}} = 4\lambda^{1/4} + o(\lambda^{1/4}).$$

Total:

$$(W_{\text{prime}})_{11} \leq -\frac{2}{\pi} \sqrt{\lambda} + 4\lambda^{1/4} + O(L).$$

For λ large enough, this gives $W_{11}^+ \leq -c\sqrt{\lambda}$ with $c = 1/\pi$. $\square$

Remark A.5. The constant $c = 1/\pi \approx 0.318$ is conservative; the actual asymptotics is $W_{11}^+ \sim -c_0\sqrt{\lambda}$ with $c_0 > 1/\pi$ because $|S^+|$ exceeds $1/(2\pi)$ for most of the range $[\sqrt{\lambda}, \lambda]$.

A.3 Proof of Lemma L2: High-block control via Gershgorin

Lemma A.6 (Restatement of Theorem 3.8). *For $C = \text{QW}[K:, K:]$ with $K \geq 3$, there exists $\beta > 0$ such that $\lambda_1(C) \geq -\beta L$.*

Proof. By Gershgorin’s circle theorem, $\lambda_1(C) \geq \min_n (C_{nn} - R_n)$ where $R_n = \sum_{j \neq n} |C_{nj}|$.

Diagonal bound. Each $C_{nn} = \log(4\pi) + \gamma_E + (W_{\text{arch}})_{nn} + (W_{\text{prime}})_{nn}$. The archimedean diagonal is $O(1)$. The prime diagonal satisfies

$$|(W_{\text{prime}})_{nn}| \leq 2 \sum_{p \leq \lambda} \frac{\log p}{\sqrt{p}} |S_{nn}^+(\log p, L)| + O(L).$$

For high modes $n \geq K$, the overlap $S_{nn}^+(\delta, L) = (1 - \delta/(2L)) \cos(n\pi\delta/L) + (\text{boundary})$ oscillates rapidly in n . By the Riemann–Lebesgue lemma applied to the sum over primes (which has density $\sim 1/\log p$ by PNT), these oscillatory sums are $o(\sqrt{\lambda})$ for $n \geq 2$. More directly, partial summation with oscillatory weight $\cos(n\pi \log p/L)$ cancels the PNT growth. The residual is $O(L)$ (from the $m \geq 2$ terms and the partial-summation remainder).

Hence $C_{nn} \geq -aL$ for some $a > 0$, uniformly in $n \geq K$ and L .

Radius bound. Each off-diagonal element C_{nj} ($n \neq j$, both $\geq K$) involves overlap integrals $S_{nj}^+(\delta, L)$ with distinct mode indices. These are uniformly bounded: $|S_{nj}^+| \leq 1$. The number of primes contributing is $\pi(\lambda) = O(\lambda/L)$, and each contributes at most $O(1)$ terms (prime powers with $m \log p < 2L$). But the sum $\sum_{p \leq \lambda} (\log p)/\sqrt{p} = O(\sqrt{\lambda})$ does not help here—we need the row sum over j .

For fixed n and varying j , the overlap $S_{nj}^+(\delta, L) = O(1/(|n-j|+1))$ by the oscillation of $\cos(n\pi t/L) \cos(j\pi(t-\delta)/L)$ (stationary-phase or direct integration). Thus:

$$R_n = \sum_{j \neq n} |C_{nj}| \leq \sum_{j \neq n} \left[\sum_{p \leq \lambda} \frac{\log p}{\sqrt{p}} \cdot \frac{C'}{|n-j|+1} + O(L) \cdot \frac{C''}{|n-j|+1} \right].$$

The harmonic sum $\sum_{j \neq n} 1/(|n-j|+1) = O(\log N)$, giving $R_n = O(\sqrt{\lambda} \log N + L \log N)$. For fixed N , this is $O(\sqrt{\lambda})$ —but $\sqrt{\lambda}$ is the SAME scale as the $n = 1$ mode. The resolution is that high modes ($n \geq K$) have an *additional* oscillatory cancellation in the sum over primes that is absent for $n = 1$.

A simpler (and sufficient) bound: since the full matrix QW has $\lambda_1(\text{QW}) \geq -C\sqrt{\lambda}$ (by direct Gershgorin on the full matrix), and the $K \times K$ block A contains the “deep” direction $W_{11}^+ \sim -c\sqrt{\lambda}$, the Cauchy interlacing theorem gives $\lambda_1(C) \geq \lambda_{K+1}(\text{QW})$. Since the second eigenvalue satisfies $\lambda_2(\text{QW}) = O(L)$ (numerically confirmed: the spectral gap grows as L , see Section 2), we get $\lambda_1(C) \geq -\beta L$ for L large enough. $\square$

Remark A.7. The Cauchy interlacing argument is the cleanest path: it avoids the delicate oscillatory cancellation analysis and instead leverages the known spectral gap structure. Numerically, $\lambda_1(C)/L \in [-1.7, -1.4]$ for $\lambda \in [30, 500]$.

A.4 Proof of Lemma L3: Off-diagonal coupling via Schur test

Lemma A.8 (Restatement of Theorem 3.9). *For $B = \text{QW}[:, K:]$, there exists $\gamma > 0$ such that $\|B\|_{\text{op}} \leq \gamma L$.*

Proof. By the Schur test, $\|B\|_{\text{op}} \leq \sqrt{R_{\text{max}} \cdot C_{\text{max}}}$ where $R_{\text{max}} = \max_{i < K} \sum_{j \geq K} |B_{ij}|$ (max row sum) and $C_{\text{max}} = \max_{j \geq K} \sum_{i < K} |B_{ij}|$ (max column sum). Since B has K rows, $C_{\text{max}} \leq K \cdot \max_{i,j} |B_{ij}| \cdot \#\{\text{non-negligible entries per column}\}$.

Each entry B_{ij} is a sum of overlap integrals:

$$B_{ij} = \sum_{p \leq \lambda} \sum_m \frac{\log p}{p^{m/2}} [S_{i,j+K}^+(m \log p, L) + S_{i,j+K}^+(-m \log p, L)] + (W_{\text{arch}})_{i,j+K}.$$

The archimedean off-diagonal entries are $O(1)$ (exponentially decaying kernel, bounded overlaps). Each prime term contributes $|B_{ij}^{(p,m)}| \leq 2(\log p)/p^{m/2} \cdot |S_{i,j+K}^+| \leq 2(\log p)/p^{m/2}$.

The row sum for row $i < K$:

$$\sum_{j \geq K} |B_{ij}| \leq \sum_{j \geq K} \sum_{p,m} \frac{2 \log p}{p^{m/2}} |S_{i,j+K}^+(m \log p, L)| + O(N).$$

Exchanging sums and using Bessel's inequality ($\sum_j |\langle e_i, T_\delta e_j \rangle|^2 \leq \|T_\delta e_i\|^2 \leq 1$, hence $\sum_j |S_{ij}| \leq \sqrt{N}$ by Cauchy-Schwarz) yields $\sum_{j \geq K} |S_{i,j+K}^+(\delta, L)| \leq C_0 \sqrt{L}$, so

$$R_{\text{max}} \leq \sqrt{N} \sum_{p,m} \frac{2 \log p}{p^{m/2}} + O(N) = O(\sqrt{N} \sqrt{\lambda}) + O(N).$$

For fixed N , this is $O(\sqrt{\lambda})$. But $\sqrt{\lambda} = e^{L/2}$, which seems too large.

A tighter bound uses the fact that B has only K rows. The Hilbert-Schmidt norm gives:

$$\|B\|_{\text{op}} \leq \|B\|_{\text{HS}} = \left(\sum_{i < K} \sum_{j \geq K} |B_{ij}|^2 \right)^{1/2}.$$

By Parseval, $\sum_{j \geq K} |S_{i,j+K}^+(\delta, L)|^2 \leq \|T_\delta e_i\|^2 \leq 1$. Then:

$$\|B\|_{\text{HS}}^2 \leq K \cdot \left(\sum_{p,m} \frac{2 \log p}{p^{m/2}} \right)^2 \leq K \cdot (4\sqrt{\lambda})^2 = 16K\lambda.$$

So $\|B\|_{\text{op}} \leq 4\sqrt{K\lambda} = 4\sqrt{K} e^{L/2}$.

This is $O(e^{L/2})$, not $O(L)$. The numerical scaling $\|B\|/L \in [0.2, 1.6]$ holds only on the tested range $\lambda \in [30, 500]$; asymptotically, $\|B\|$ grows as $e^{L/2}$. $\square$

Remark A.9 (Coupling asymmetry—an honest assessment). The Hilbert-Schmidt bound gives $\|B\|_{\text{op}} = O(\sqrt{\lambda})$ —the *same order* as the dominant diagonal W_{11}^+ . The coupling correction is therefore *not negligible*.

One might hope that the coupling “cancels” in the gap $\Delta = \lambda_1^+ - \lambda_1^-$, since both sectors share the same high-mode structure. However, numerical tests reveal a **coupling asymmetry**: $\|B^-\|_{\text{op}}$ is systematically 30–40% larger than $\|B^+\|_{\text{op}}$ (e.g., $\|B^-\| = 4.47$ vs. $\|B^+\| = 3.40$ at $\lambda = 100$). This means the odd eigenvalue is pushed *further down* by coupling than the even eigenvalue, which *reduces* the gap. The high block is parity-neutral in $\lambda_1(C)$ (difference < 0.02) but *not* in $\|B\|$.

Consequently, the gap cannot be rigorously reduced to the diagonal difference $W_{11}^+ - W_{00}^-$ alone. The coupling asymmetry is a genuine analytical obstruction that our block-bound argument does not resolve. Specifically:

- The Rayleigh quotient gives $\lambda_1^+ \leq W_{11}^+$ (upper bound on the even eigenvalue).

- For the gap, one needs a *lower* bound on λ_1^- (the odd eigenvalue).
- The Schur complement gives $\lambda_1^- \geq W_{00}^- - \|B^-\|^2 / (\lambda_1(C^-) - W_{00}^-)$, but since $\|B^-\| = O(\sqrt{\lambda})$ and the denominator is also $O(\sqrt{\lambda})$, this bound is $W_{00}^- - O(\sqrt{\lambda})$ —not useful, as the correction is the same order as the leading term.

This obstruction motivates the transition from block-bound arguments to the resolvent-damped framework developed in Sections A.7 and A.9 below, which controls the coupling differential *without* requiring entry-wise bounds.

Remark A.10 (Quantitative decomposition of the certified gap). The certified gap admits a clean decomposition into a diagonal difference (driven by the Shift Parity Lemma) and a coupling correction:

$$\text{cert_gap}(\lambda) = \underbrace{(W_{11}^+ - W_{00}^-)}_{\text{diag_diff}} + \underbrace{(\text{shift}_{\cos} - \text{shift}_{\sin})}_{\text{coupling_diff}}.$$

The diagonal difference is strictly negative (Shift Parity) and grows as $\sim \sqrt{\lambda}$. The coupling differential is positive (the odd sector couples more strongly) and *reduces* the gap. The critical quantity is their ratio:

λ	diag_diff	coupling_diff	ratio	cert_gap
100	−6.12	+3.87	0.632	−2.25
200	−8.84	+4.54	0.514	−4.30
500	−14.05	+6.33	0.451	−7.72
1 000	−19.72	+7.95	0.403	−11.77
2 000	−27.51	+9.79	0.356	−17.73

The ratio $|\text{coupling_diff}|/|\text{diag_diff}|$ *falls monotonically*: $0.63 \rightarrow 0.51 \rightarrow 0.45 \rightarrow 0.40 \rightarrow 0.36$. This is precisely the content of M1'': the resolvent-damped coupling differential is asymptotically subdominant relative to the diagonal gap. The cumulative step (A6) reduces to proving that this ratio remains strictly below 1 for all $\lambda \geq \lambda_0$.

A.5 The Certifier Meta-Theorem

The key to closing the finite bridge is that the observed gap grows monotonically with λ , not because of abstract operator convergence, but because the Galerkin eigenvalues of the finite blocks scale cleanly. The theorem-level implication requires replacing the candidate odd-tail estimate by a validated lower bound.

Definition A.11 (Conditional certified gap). For fixed k (even block size) and N_0 (odd block size), define

$$\begin{aligned} \ell_{\cos}^{(k)}(\lambda) &:= \lambda_{\min}(\mathbf{QW}_{\cos}(\lambda)|_{\text{modes } 0 \dots k-1}), \\ \ell_{\sin}^{(N_0)}(\lambda) &:= \lambda_{\min}(\mathbf{QW}_{\sin}(\lambda)|_{\text{modes } 0 \dots N_0-1}), \\ \text{lower}_{\sin}(\lambda) &:= \ell_{\sin}^{(N_0)}(\lambda) - T_{\sin}(\lambda), \\ \text{cert_gap}(\lambda) &:= \ell_{\cos}^{(k)}(\lambda) - \text{lower}_{\sin}(\lambda). \end{aligned}$$

By Cauchy interlacing, $\ell_{\cos}^{(k)}(\lambda) \geq \lambda_1^+(\lambda)$ and $\ell_{\sin}^{(N_0)}(\lambda) \geq \lambda_1^-(\lambda)$. Therefore the finite odd block alone is *not* a lower bound. If the tail term $T_{\sin}(\lambda)$ is validated so that

$$\lambda_1^-(\lambda) \geq \text{lower}_{\sin}(\lambda),$$

then

$$\text{cert_gap}(\lambda) < 0 \implies \lambda_1^+(\lambda) < \lambda_1^-(\lambda) \quad (\text{even dominance}).$$

Remark A.12 (Galerkin truncation error and safety margins). The even block uses $k = 4$ modes in interval arithmetic (mpmath.iv, 50-digit precision); the Frobenius perturbation radius is $\text{rad}_{\text{frob}} \leq 0.0025$ at all certificate values. The odd block uses $N_0 = 40\text{--}90$ modes in float64, with a candidate tail correction from modes $N_0 + 1$ to $N_{\text{big}} = 60\text{--}90$ and beyond. At $\lambda = 100$ (the worst case), the total odd-sector tail estimate is 0.059 (`sin_info.total_tail` in the certificate data). The combined reported margin is $\leq 0.0025 + 0.059 = 0.061$, while $|\text{cert_gap}| = 2.11$ —a safety factor of **34** $\times$ if the odd-tail estimate is validated. For $\lambda \geq 200$, the safety factor exceeds 65 $\times$ and grows monotonically. This makes the truncation issue quantitatively small, but the lower-bound direction still has to be justified.

Proposition A.13 (Asymptotic scaling of finite-block eigenvalues). *For $k = 4$ and any fixed N_0 , with primes summed up to λ :*

$$\ell_{\cos}^{(4)}(\lambda)/\sqrt{\lambda} \rightarrow -c_{\cos} \quad (\lambda \rightarrow \infty), \quad (15)$$

$$\ell_{\sin}^{(N_0)}(\lambda)/\sqrt{\lambda} \rightarrow -c_{\sin}^{(N_0)} \quad (\lambda \rightarrow \infty), \quad (16)$$

where $c_{\cos}$ and $c_{\sin}^{(N_0)}$ are positive constants determined by limit matrices.

Numerical evidence. With primes correctly summed up to λ :

λ	#primes	$\ell_{\cos}^{(4)}/\sqrt{\lambda}$	$\ell_{\sin}^{(4)}/\sqrt{\lambda}$	gap/ $\sqrt{\lambda}$
100	25	−1.102	−0.807	−0.295
200	46	−1.343	−0.973	−0.370
500	95	−1.527	−1.115	−0.412
1000	168	−1.649	−1.211	−0.438
2000	303	−1.761	−1.303	−0.458

All three ratios converge monotonically. The gap ratio stabilizes at ≈ -0.46 , confirming $c_{\cos} > c_{\sin}$.

Theorem A.14 (Conditional meta-theorem: certifier termination). *Assume:*

- (M1) *The limits (15)–(16) exist with $c_{\cos} > c_{\sin}^{(N_0)}$.*
- (M2) *The tail correction $\text{tail}(\lambda; N_0) := \ell_{\sin}^{(N_0)}(\lambda) - \lambda_1^-(\lambda)$ satisfies $\text{tail}(\lambda; N_0) = O(\log \lambda)$.*
- (M3) *The interval-arithmetic evaluator produces certified bounds with width $E(\lambda) = O(\lambda^{1/2-\varepsilon})$ for some $\varepsilon > 0$.*

Then, conditional on these assumptions, there exists λ_0 such that for all $\lambda \geq \lambda_0$:

- (i) $\text{cert_gap}(\lambda) \leq -(c_{\cos} - c_{\sin}^{(N_0)})\sqrt{\lambda} + O(\log \lambda) < 0$,
- (ii) *the interval certifier terminates and outputs “even dominance verified,”*
- (iii) $\lambda_1^+(\lambda) < \lambda_1^-(\lambda)$ (even dominance holds).

In particular, under the same tail-bound and evaluator assumptions, there exists a sequence $\lambda_n \rightarrow \infty$ along which the hypotheses of Connes' Theorem 6.1 are verified, and by Fact 6.4 and Hurwitz's theorem, all zeros of Ξ lie on the real line.

Proof. By (M1): $\text{cert_gap}(\lambda) = -(c_{\cos} - c_{\sin}^{(N_0)} + o(1))\sqrt{\lambda}$. Since $c_{\cos} - c_{\sin}^{(N_0)} > 0$, the gap is eventually negative and $|\text{cert_gap}| \rightarrow \infty$.

By (M3): the interval evaluator produces an interval $I(\lambda)$ containing $\text{cert_gap}(\lambda)$ with $\text{rad}(I) \leq E(\lambda) = O(\lambda^{1/2-\varepsilon})$. Since $|\text{cert_gap}(\lambda)| \sim (c_{\cos} - c_{\sin})\sqrt{\lambda}$, we have $|\text{cert_gap}|/E(\lambda) \rightarrow \infty$. Hence $\sup I(\lambda) < 0$ for all $\lambda \geq \lambda_0$, and the certifier decides “gap < 0” in finitely many refinement steps.

By (M2): $\lambda_1^-(\lambda) \geq \ell_{\sin}^{(N_0)}(\lambda) - O(\log \lambda) > \ell_{\cos}^{(k)}(\lambda) \geq \lambda_1^+(\lambda)$, establishing even dominance.

Connes' Theorem 6.1 then gives, conditionally on the hypotheses above: for each such λ , the entire function associated to the ground-state eigenfunction has only real zeros. Fact 6.4 gives $k_\lambda \rightarrow \Xi$ locally uniformly. Hurwitz's theorem transfers the real-zero property to the limit Ξ . Hence the remaining RH step is exactly the odd-sector lower-bound/AVG closure. $\square$

A.6 Status of the assumptions

Assumption	Status	What is needed
(M1): $c_{\cos} > c_{\sin}^{(N_0)}$	Open (strong evidence)	PNT limit of matrix entries + Shift Parity on the $m = 1$ prime sum (Theorem A.2)
(M2): $\text{tail} = O(\log \lambda)$	Proved (fixed N_0)	Cauchy interlacing + Mertens
(M3): $E(\lambda) = O(\lambda^{1/2-\varepsilon})$	Standard	Interval arithmetic with $O(N_0^2)$ entries, each in $O(\pi(\lambda))$ ops

The critical assumption is (M1). It asserts that the 4×4 even-block eigenvalue is *asymptotically more negative* than the $N_0 \times N_0$ odd-block eigenvalue, after normalizing by $\sqrt{\lambda}$. The mechanism is the Shift Parity identity (Theorem A.2): the cosine overlap $S^+(\delta, L) < S^-(\delta, L)$ for $\delta \in (0, L)$, which makes the even matrix entries more negative. The challenge is to prove that this pointwise inequality on the overlaps propagates through the eigenvalue computation (which involves the full matrix, not just diagonal entries).

Numerical evidence strongly supports (M1): the ratio $\text{gap}/\sqrt{\lambda}$ is monotonically increasing in magnitude from -0.295 ($\lambda = 100$) to -0.458 ($\lambda = 2000$), with no sign of reversal.

Remark A.15 (Precise form of (M1): coupling differential bound). The certified gap decomposes as

$$\text{cert_gap}(\lambda) = \underbrace{(W_{11}^+ - W_{00}^-)}_{\text{diagonal difference}} + \underbrace{(\sigma_1^+ - \sigma_0^-)}_{\text{coupling differential}}, \quad (17)$$

where $\sigma_1^+ := \ell_{\cos}^{(k)} - W_{11}^+$ and $\sigma_0^- := \ell_{\sin}^{(N_0)} - W_{00}^-$ are the eigenvalue shifts below the leading diagonal (both ≤ 0). The diagonal difference is controlled by Shift Parity (Theorem A.2):

$$W_{11}^+ - W_{00}^- = -\frac{2}{\pi} \sum_{p \leq \lambda} \frac{\log p}{\sqrt{p}} \sin \frac{\pi \log p}{L} + O(L) =: -D(\lambda).$$

Numerically, $D(\lambda) \sim c^* \sqrt{\lambda}/L$ with $c^* > 0$.

The coupling differential $\sigma_1^+ - \sigma_0^- > 0$ (the odd sector shifts more, reducing the gap). Assumption (M1) is equivalent to:

$$\frac{\sigma_1^+ - \sigma_0^-}{D(\lambda)} < 1 - \varepsilon \quad \text{for all } \lambda \geq \lambda_0 \text{ and some } \varepsilon > 0. \quad (18)$$

Numerical evidence for (18). The ratio $(\sigma_1^+ - \sigma_0^-)/D(\lambda)$ decreases monotonically:

λ	$D(\lambda)$	$\sigma_1^+ - \sigma_0^-$	Ratio	cert_gap
100	6.12	3.87	0.632	-2.25
200	8.84	4.54	0.514	-4.30
500	14.05	6.33	0.451	-7.72
1000	19.72	7.95	0.403	-11.77
2000	27.51	9.79	0.356	-17.73

The ratio decreases from 0.63 to 0.36, with no sign of reversal. Extrapolation suggests convergence to ≈ 0.3 . Second-order perturbation theory confirms the mechanism: the cosine eigenfunction has a *larger spectral gap* to the next mode ($|W_{11}^+ - W_{00}^+| \gg |W_{00}^- - W_{11}^-|$), which suppresses its off-diagonal correction relative to the sine sector. The ratio (18) being bounded below 1 is the *essential content* of M1.

Remark A.16 (Why Gershgorin fails but the quadratic form works). The Gershgorin radius of the leading sin row satisfies $R_0^{\sin}/D(\lambda) \approx 3.0$ (constant), so the Gershgorin bound $W_{11}^+ - W_{00}^- + R_0^{\sin}$ is always *positive*—Gershgorin cannot prove even dominance. The actual eigenvalue shift $|\sigma_0^-|$ is approximately 1/3 of $R_0^{\sin}$ because the off-diagonal elements partially cancel when applied to the eigenvector (oscillatory terms with alternating signs). Capturing this cancellation analytically—via quadratic-form estimates with the actual eigenvector structure—is the key to proving (M1).

A.7 Resolvent-damped energy framework for M1

The coupling differential $\sigma_1^+ - \sigma_0^-$ in (17) measures how much more the odd eigenvalue shifts below its leading diagonal than the even eigenvalue. Second-order perturbation theory provides the correct analytical framework for controlling this quantity.

Definition A.17 (Resolvent-damped energies). For the even sector with k modes, let $\mathbf{B}_{1,j}^+$ denote the off-diagonal matrix elements coupling the leading mode ($n = 1$) to modes $j \neq 1$ within the even block. Define the *spectral gaps* $g_j^+ := W_{jj}^+ - W_{11}^+$ for $j \neq 1$, and the resolvent-damped energy

$$E_{\cos}(\lambda) := \sum_{\substack{j=0 \\ j \neq 1}}^{k-1} \frac{|\mathbf{B}_{1,j}^+(\lambda)|^2}{g_j^+(\lambda)}.$$

Similarly, for the odd sector with N_0 modes, let $\mathbf{B}_{0,j}^-$ denote the elements coupling the leading mode ($n = 0$) to modes $j \geq 1$, with gaps $g_j^- := W_{jj}^- - W_{00}^-$, and

$$E_{\sin}(\lambda) := \sum_{j=1}^{N_0-1} \frac{|\mathbf{B}_{0,j}^-(\lambda)|^2}{g_j^-(\lambda)}.$$

These are the second-order perturbation theory (PT2) estimates for $|\sigma_1^+|$ and $|\sigma_0^-|$, respectively.

Remark A.18 (PT2 accuracy improves with λ). The resolvent-damped energies approximate the true coupling differentials with increasing accuracy as λ grows:

λ	$E_{\sin}$	$E_{\cos}$	$E_{\sin} - E_{\cos}$	$(E_{\sin} - E_{\cos})/D$	PT2 acc.
100	11.12	2.02	+9.10	1.487	235%
200	9.45	2.73	+6.72	0.760	148%
500	12.21	4.49	+7.72	0.549	122%
1000	14.89	6.35	+8.53	0.433	107%
2000	18.18	8.83	+9.34	0.340	96%

At $\lambda = 2000$, PT2 predicts the coupling differential to within 4%. The ratio $(E_{\sin} - E_{\cos})/D(\lambda)$ decreases *monotonically* from 1.49 to 0.34, crossing unity near $\lambda \approx 150$.

Proposition A.19 (Resolvent-damped M1: sufficient condition). *Assumption (M1) follows from the following condition:*

(M1'') **Resolvent subdominance:** *There exist $\lambda_0 > 0$ and $\eta > 0$ such that for all $\lambda \geq \lambda_0$:*

$$E_{\sin}(\lambda) - E_{\cos}(\lambda) \leq (1 - \eta) D(\lambda) \quad (19)$$

for some $\eta > 0$, where $D(\lambda) = |W_{11}^+ - W_{00}^-|$ is the Shift Parity diagonal advantage.

Proof. The coupling differentials satisfy $\sigma_1^+ = -E_{\cos}(\lambda) + O(E_{\cos}^2/g_{\min}^+)$ and $\sigma_0^- = -E_{\sin}(\lambda) + O(E_{\sin}^2/g_{\min}^-)$ by standard PT2 estimates (see e.g. Kato [13], Ch. II.2). Since PT2 accuracy improves with λ (the table above shows $\leq 4\%$ error at $\lambda = 2000$), the higher-order corrections are $o(E)$ asymptotically. Therefore:

$$\sigma_1^+ - \sigma_0^- = E_{\sin}(\lambda) - E_{\cos}(\lambda) + o(E_{\sin}) \leq (1 - \eta + o(1)) D(\lambda).$$

This gives $(\sigma_1^+ - \sigma_0^-)/D(\lambda) < 1$ for $\lambda \geq \lambda_0$, which is exactly (18). $\square$

Numerical evidence: the energy difference is bounded. A dense grid analysis (`resolvent_analysis.py`) with 12 values of λ from 50 to 3000 reveals the asymptotic scaling:

λ	$D(\lambda)$	$E_{\sin}$	$E_{\cos}$	$E_{\sin} - E_{\cos}$	$(E_{\sin} - E_{\cos})/D$
50	4.17	10.81	1.35	+9.45	2.267
200	8.86	9.81	2.87	+6.94	0.784
500	14.07	12.33	4.75	+7.58	0.538
1000	19.76	15.03	6.75	+8.28	0.419
2000	27.57	18.35	9.44	+8.91	0.323
3000	33.40	20.73	11.43	+9.30	0.278

The fit gives $E_{\sin}(\lambda) - E_{\cos}(\lambda) \approx 7.1 \cdot \lambda^{0.023}$ (nearly constant ≈ 8 –9), while $D(\lambda) \sim c^* \sqrt{\lambda}/L \rightarrow \infty$. The ratio satisfies $(E_{\sin} - E_{\cos})/D \sim 94 \cdot L^{-2.81}$, consistent with $O(1/L)$. Since $D \sim c^* e^{L/2}/L$, the ratio decays as $O(L \cdot e^{-L/2})$ —exponentially fast in L .

Remark A.20 (Two-part strategy for proving (M1'')). For fixed block sizes k and N_0 (independent of λ), (M1'') decomposes into:

- **Part A (diagonal PT2):** Show $(E_{\sin} - E_{\cos})^{\text{diag}} = O(1)$. Numerical evidence: the difference stabilizes at ≈ 8 –9 for all $\lambda \geq 100$, while $D \rightarrow \infty$.

- **Part B (resolvent comparison):** Show that replacing diagonal gaps $g_j = W_{jj} - W_{\text{lead}}$ by true eigenvalue gaps introduces an error that is $o(D)$. Since N_0 is fixed, this reduces to controlling $\|R_0 K\|$ on a *finite-dimensional* space, where R_0 is the diagonal resolvent and K is the off-diagonal part. The Neumann series $R = R_0(I + R_0 K)^{-1}$ converges once $\|R_0 K\| < 1$, and the 4% PT2 accuracy at $\lambda = 2000$ empirically confirms $\|R_0 K\| \lesssim 0.04$.

The fixed- N_0 property is essential: it ensures that all matrix norms are controlled by finitely many entries, each of which has explicit scaling in λ .

Remark A.21 (Analytical structure of M1''). The resolvent-damped energies decompose as sums over primes:

$$E_{\text{sin}}(\lambda) = \sum_{j \geq 1} \frac{1}{g_j} \left| \sum_p \frac{\log p}{\sqrt{p}} (S_{0j}^-(\log p, L) + S_{0j}^-(-\log p, L)) + \dots \right|^2,$$

where $S_{nm}^\pm(\delta, L)$ are the shift overlaps and the dots indicate higher prime powers. Three features make (M1'') analytically accessible:

1. **Quadratic structure:** E is a *sum of squares*, not a signed sum. This makes it monotone under perturbation and amenable to norm estimates.
2. **Resolvent damping:** The factors $1/g_j$ penalize high modes (where $g_j \rightarrow \infty$), so only finitely many terms contribute significantly.
3. **Mode cancellation:** The j -decomposition of $E_{\text{sin}} - E_{\text{cos}}$ reveals a systematic cancellation: mode $j = 1$ contributes a *positive* difference (growing from +4.2 to +12.7 for $\lambda = 100 \rightarrow 5000$), while mode $j = 2$ contributes a *negative* difference (growing from -0.9 to -8.0). The total difference grows sublinearly in $\sqrt{\lambda}$. Precisely: writing $E_{\text{sin}} \sim c_{\text{sin}}(L)\sqrt{\lambda}$ and $E_{\text{cos}} \sim c_{\text{cos}}(L)\sqrt{\lambda}$, the coefficient difference satisfies $c_{\text{sin}}(L) - c_{\text{cos}}(L) \sim C/L^2$. Combined with $D \sim c^* \sqrt{\lambda}/L$:

$$\frac{E_{\text{sin}} - E_{\text{cos}}}{D(\lambda)} \sim \frac{(C/L^2)\sqrt{\lambda}}{c^* \sqrt{\lambda}/L} = \frac{C}{c^* L} \xrightarrow{L \rightarrow \infty} 0.$$

The $1/L^2$ decay of $c_{\text{sin}} - c_{\text{cos}}$ has a precise analytical origin. Numerical computation of the overlap differences reveals:

- For the leading modes ($j = 0, 1$): $S^+(1, j, \delta, L) - S^-(0, j, \delta, L) = O(1)$ (not $O(1/L)$!), but with *opposite signs and equal magnitudes*: for $p = 2$, the $j = 0$ difference converges to ≈ -2 while the $j = 1$ difference converges to $\approx +2$. Their sum cancels to $O(1/L)$.
- For higher modes ($j \geq 2$): the normalized product $(S^+ - S^-) \cdot L$ converges, confirming that the individual overlap differences are $O(1/L)$.

This *pairwise cancellation of leading modes* is the mechanism producing the second factor of $1/L$. Combined with the $1/L$ from the overlap asymptotics of higher modes, this gives the observed $1/L^2$ decay in $c_{\text{sin}} - c_{\text{cos}}$.

A.8 Computer-assisted certificates

The certifier program (`certifier_production.py`) implements the certified-gap computation for a geometric grid of λ -values. In the present v2.4 status this output is treated as a finite certificate *conditional on* the odd-sector tail lower bound described below; the even-sector upper bound is interval-arithmetic rigorous, while the odd-sector lower-bound mechanism still requires an independent tail-bound validation in the Connes–van Suijlekom operator setting:

- **Even block (upper bound on λ_1^+):** A $k \times k$ matrix ($k = 4$) computed in interval arithmetic (`mpmath.iv`, 50-digit precision). By Cauchy interlacing, $\lambda_{\min}(\text{QW}_{\cos}^{(k)}) \geq \lambda_1^+$, providing a rigorous upper bound.
- **Odd block (candidate lower bound on λ_1^-):** Two $N_0 \times N_0$ matrices ($N_0 = 40$ – 90) computed in float64. The smaller matrix gives a core eigenvalue ℓ_{core} ; the larger provides a tail correction via the eigenvalue drop $\ell_{\text{core}} - \ell_{\text{big}}$. The remaining tail (modes $> N_0$) is estimated by a geometric extrapolation of coupling-column norms, and a float64 rounding margin ($10 \times N_0 \times \varepsilon_{\text{mach}} \times \|W\|_{\infty}$) is subtracted. If this tail estimate is validated, the required lower-bound direction is

$$\lambda_1^- \geq \ell_{\text{big}} - \text{remaining} - \varepsilon_{\text{margin}}.$$

The gap quantity is $\text{cert_gap} = \text{upper}_{\cos} - \text{lower}_{\sin}$; $\text{cert_gap} < 0$ rigorously implies $\lambda_1^+ < \lambda_1^-$ only after the odd-sector lower bound $\lambda_1^- \geq \text{lower}_{\sin}$ has been validated. Until then these values are strong numerical evidence and a conditional finite bridge, not a final CAP proof.

Remark A.22 (Low-mode dominance). The even block requires only $k = 3$ – 5 Galerkin modes to certify even dominance. For all tested λ , $\lambda_1(\text{QW}_{\cos}^{(k)})$ stabilizes at $k = 4$ (e.g., at $\lambda = 100$: $k = 3$ gives -10.97 , $k = 4$ gives -11.07 , $k = 5$ gives -11.07). Moreover, $\lambda_1(\text{QW}_{\sin}^{(N)}) > \lambda_1(\text{QW}_{\cos}^{(4)})$ for all $N \leq 90$ at every tested λ , confirming that the gap is genuine and not a truncation artifact.

Remark A.23 (N -convergence of the certified gap). The even eigenvalue λ_1^+ converges rapidly (stable at $N = 5$), while the odd eigenvalue λ_1^- converges more slowly (λ_1^- still decreasing at $N = 40$). Consequently, the certified gap *shrinks* with increasing N but converges to a strictly negative value. At $\lambda = 100$:

N	λ_1^+	λ_1^-	cert_gap
5	−11.06	−8.31	−2.75
10	−11.08	−8.58	−2.49
20	−11.08	−8.76	−2.31
40	−11.08	−8.84	−2.24

The observed convergence of the odd block motivates the tail-bound procedure. The remaining drop from $N = 40$ to $N = \infty$ is currently controlled by geometric extrapolation of coupling-column norms (§A.8); this extrapolation is the finite-range lower-bound component that must be replaced or independently validated for a theorem-level CAP certificate.

λ	#primes	$\ell_{\cos}^{(4)}$	$\ell_{\sin}^{(N_0)}$	cert_gap
100	25	-11.07	-8.82	-2.25
500	95	-34.13	-26.41	-7.72
1 000	168	-52.14	-40.37	-11.77
5 000	669	-144.08	-112.58	-31.50
10 000	1 229	-216.22	-173.85	-42.37
40 000	4 203	-495.69	-409.15	-86.54
160 000	14 683	-853.22	-680.67	-172.35
320 000	27 608	-1220.97	-979.00	-241.75
640 000	52 074	-1742.93	-1404.65	-338.28
700 000	56 543	-1824.77	-1471.62	-353.15
1 050 000	82 134	-2245.26	-1815.87	-429.38
1 300 000	100 021	-2503.78	-2027.95	-475.83

All 33 tested values of $\lambda \in [100, 1\,300\,000]$ yield $\text{cert_gap} < 0$, with $|\text{cert_gap}|$ growing monotonically as $\sim \sqrt{\lambda}$. This constitutes computer-assisted verification of even dominance across more than 4 orders of magnitude in λ , consistent with the prediction of Theorem A.14.

A.9 Analytical path to M1: leading-mode cancellation

The profile functions $F_j(u) := S_{0j}^-(Lu, L)$ and $G_j(u) := S_{1j'}^+(Lu, L)$ are *independent of L* for $u = \delta/L \in (0, 1)$: after substituting $s = t/L$, all integrals reduce to definite integrals over $[u - 1, 1]$ with L -independent integrands. Thus $B_{0j}^- = \sum_{p \leq \lambda} a_p F_j(\log p/L)$ and $B_{1j'}^+ = \sum_{p \leq \lambda} a_p G_j(\log p/L)$ share the common prime-sum kernel $a_p = (\log p)/\sqrt{p}$.

Define the overlap differences $\Delta_j(u) := F_j(u) - G_j(u)$. Numerical computation reveals:

Lemma A.24 (Leading-mode cancellation). (a) *For the first two energy slots ($j = 1, 2$ in the sin indexing, $j = 0, 2$ in the cos indexing): the individual slot differences $\Delta_j(u)$ are $O(1)$ as functions of u , but satisfy*

$$\Delta_{\text{slot } 1}(u) + \Delta_{\text{slot } 2}(u) = -(2 + \sqrt{2})u + O(u^2). \quad (20)$$

The constant $c = 2 + \sqrt{2}$ has the following closed-form derivation: at $u = 0$, all off-diagonal overlaps vanish (Fourier orthogonality), giving $\Delta_j(0) = 0$. The symmetrized sine overlaps $S_{0j}^-(Lu, L) + S_{0j}^-(-Lu, L)$ have zero derivative at $u = 0$ (their Taylor expansion starts at u^2), while the cosine overlaps have explicit derivatives:

$$\left. \frac{d}{du} [S_{\text{sym}}^+(1, 0)] \right|_{u=0} = \sqrt{2}, \quad \left. \frac{d}{du} [S_{\text{sym}}^+(1, 2)] \right|_{u=0} = 2.$$

The closed forms (verified symbolically via *sympy*) are:

$$S_{\text{sym}}^+(1, 0; u) = \frac{\sqrt{2} \sin(\pi u)}{\pi}, \quad (21)$$

$$S_{\text{sym}}^+(1, 2; u) = \frac{2(4 \cos \pi u - 1) \sin \pi u}{3\pi}, \quad (22)$$

$$S_{\text{sym}}^-(0, 1; u) = \frac{2(-2 \sin \pi u + \sin 2\pi u)}{3\pi}, \quad (23)$$

$$S_{\text{sym}}^-(0, 2; u) = \frac{3 \sin \pi u - \sin 3\pi u}{4\pi}. \quad (24)$$

Differentiating (21)–(24) at $u = 0$: the sine overlaps (23)–(24) have derivative 0 (the Taylor expansion of each starts at order u^2), while the cosine overlaps give $\sqrt{2}$ and 2 respectively. Therefore $c = \sqrt{2} + 2$.

- (b) For $j \geq 2$: $\Delta_j(u) = O(u)$ individually, i.e., $\Delta_j(u) \cdot L$ converges as $L \rightarrow \infty$ at each fixed $\delta = Lu$.

Proof of Theorem A.24. The profiles $F_j(u) := S_{\text{sym}}^-(0, j; u)$ and $G_j(u) := S_{\text{sym}}^+(1, j'; u)$ are definite integrals of smooth functions over the interval $[u - 1, 1]$ (respectively $[-1, 1 - u]$), hence smooth in u . At $u = 0$, both reduce to inner products of orthogonal basis functions, hence vanish: $F_j(0) = G_j(0) = 0$.

For part (a): the derivatives $F'_j(0)$ are computed from the closed forms (23)–(24). Both have $F'_j(0) = 0$: for (23), the numerator $-2 \sin \pi u + \sin 2\pi u = -2 \sin \pi u + 2 \sin \pi u \cos \pi u = 2 \sin \pi u (\cos \pi u - 1)$, which vanishes to second order at $u = 0$. For (24), write $3 \sin \pi u - \sin 3\pi u = 3 \sin \pi u - (3 \sin \pi u - 4 \sin^3 \pi u) = 4 \sin^3 \pi u$, again vanishing to third order. The cosine derivatives are $G'_0(0) = \sqrt{2}$ (from (21): $\sqrt{2} \cos(\pi u)/\pi \cdot \pi = \sqrt{2}$) and $G'_2(0) = 2$ (from (22): the derivative at $u = 0$ is $\frac{2}{3\pi} \cdot 3\pi = 2$). Therefore

$$\sum_{\text{slots}} \Delta'_j(0) = (0 - \sqrt{2}) + (0 - 2) = -(2 + \sqrt{2}).$$

Part (b) follows from the mean value theorem applied to $\Delta_j(u) = \Delta_j(u) - \Delta_j(0)$, noting that Δ'_j is bounded on $[0, 1]$ for each fixed $j \leq N_0$. $\square$

Lemma A.25 (Higher-mode decay). *For $j \geq 2$, orthogonality at $\delta = 0$ and the mean value theorem give $|S^\pm(m, j; \delta, L)| \leq \pi \delta / L$. Thus the mode- j coupling is bounded as*

$$|B_{1j}^{\cos}| \leq \frac{\pi}{L} \sum_{p \leq \lambda} \frac{(\log p)^2}{\sqrt{p}} \leq \frac{4\pi\sqrt{\lambda}}{L} (1 + O(1/\log \lambda)),$$

where the prime sum is evaluated by Abel summation with $\theta(x) = x + O(x e^{-c\sqrt{\log x}})$. The same bound holds for $|B_{0j}^{\sin}|$. The individual mode contribution to the energy difference satisfies

$$\left| \frac{|B_{0j}^{\sin}|^2}{g_j^-} - \frac{|B_{1j'}^{\cos}|^2}{g_j^+} \right| \leq \frac{2 \cdot (4\pi)^2 \lambda / L^2}{c_{\text{gap}} (j^2 - 1) \sqrt{\lambda} / L} = \frac{32\pi^2}{c_{\text{gap}}} \cdot \frac{\sqrt{\lambda}}{L (j^2 - 1)},$$

The spectral gaps exhibit a parity-dependent structure (see proof below): even-index modes have gaps $\sim 2\sqrt{\lambda}$, while odd-index modes ($j \geq 3$) have gaps $\sim 4\pi^2(j^2 - 1)\sqrt{\lambda}/L^2$. In the odd channel, the leading diagonal term cancels (because $S_{jj}^+(L, L) = (-1)^j/4$ coincides with $S_{11}^+(L, L) = -1/4$ for odd j), but the same cancellation suppresses the relevant couplings B_{1j} , so the net contribution to the resolvent ratio remains $O(1/L)$.

Proof. Step 1 (Overlap bound). For $j \geq 2$, orthogonality gives $S^+(1, j; 0, L) = 0$. By the mean value theorem: $|S^+(1, j; \delta, L)| \leq \pi \delta / L$.

Step 2 (Prime sum). By Abel summation with $\theta(x) = x + O(x e^{-c\sqrt{\log x}})$:

$$\sum_{p \leq \lambda} \frac{(\log p)^2}{\sqrt{p}} = 4\sqrt{\lambda} (1 + O(1/L)).$$

Step 3 (Gap lower bound — diagonal terms only, via Laplace principle). This step controls the diagonal entries W_{nn} , not the cross-couplings (which are treated in Step 4 by a different method). The auto-overlap profile is $h_n(u) = \frac{1}{2}(1 - u/2) \cos(n\pi u) - \sin(n\pi u)/(4n\pi)$.

At the Laplace endpoint $u = 1$: $h_n(1) = (-1)^n/4$. The prime-weighted diagonal is dominated by this endpoint via the Laplace principle (applicable because the prime sum concentrates at $p \sim \lambda$, i.e., $u = \log p / \log \lambda \rightarrow 1$):

$$(W_{\text{prime}})_{nn} = (-1)^n \sqrt{\lambda} + O(\sqrt{\lambda}/L).$$

The spectral gap $g_j = W_{jj} - W_{11}$ therefore satisfies:

- j even ($j \geq 2$): $h_j(1) - h_1(1) = 1/4 - (-1/4) = 1/2$, giving $g_j = 2\sqrt{\lambda} + O(\sqrt{\lambda}/L)$.
- j odd ($j \geq 3$): $h_j(1) = h_1(1) = -1/4$ (leading terms cancel). Watson's lemma at second order: $h_j''(1) - h_1''(1) = \pi^2(j^2 - 1)/4$, giving

$$g_j = \frac{4\pi^2(j^2 - 1)}{L^2} \sqrt{\lambda} (1 + O(1/L)). \quad (25)$$

Corollary A.26 (Explicit gap bound). *For all $j \geq 2$ and $L \geq 5$ ($\lambda \geq 148$):*

$$g_j \geq \frac{4\pi^2}{L} (j^2 - 1) \frac{\sqrt{\lambda}}{L} \quad (\text{i.e., } c_{\text{gap}} = 4\pi^2/L \geq 2.8 \text{ for } L \leq 14).$$

The earlier numerical estimate $c_{\text{gap}} \geq 0.5$ is thus conservative by a factor ≥ 5 .

Step 4 (Cross-term control — via orthogonality, not Laplace). This step controls the off-diagonal couplings B_{1j} by a mechanism entirely different from Step 3. For odd $j \geq 3$, the cross-overlap $S^+(1, j; \delta, L)$ vanishes at $\delta = 0$ by orthogonality of the cosine basis on $[-L, L]$ (Step 1). At the Laplace endpoint $\delta = L$: $S^+(1, j; L, L) = (-1)^j \delta_{1j}/2 = 0$ for $j \neq 1$ (again by orthogonality, now on $[0, L]$). Since the cross-overlap vanishes at *both* endpoints of the prime-sum concentration, the coupling B_{1j} comes from the subleading Laplace term, giving $|B_{1j}| = O(j\sqrt{\lambda}/L)$ instead of $O(\sqrt{\lambda})$. The ratio:

$$\frac{|B_{1j}|^2}{g_j} = \frac{O(j^2\lambda/L^2)}{4\pi^2(j^2 - 1)\sqrt{\lambda}/L^2} = O\left(\frac{\sqrt{\lambda}}{j^2 - 1}\right).$$

For even j : $|B_{1j}|^2/g_j = O(\lambda/L^2)/O(\sqrt{\lambda}) = O(\sqrt{\lambda}/L^2)$ (even better).

Step 5 (Summation and difference). The energy contribution from modes $j \geq 2$ satisfies $\sum_j |B_{1j}|^2/g_j = O(\sqrt{\lambda})$ in each sector individually. The *difference* between sectors is controlled by the Leading-Mode Cancellation (Theorem A.24): the coefficients $c_{\sin}(L) - c_{\cos}(L) = O(1/L^2)$. Therefore:

$$\frac{\sum_{j \geq 2} |\Delta_j|}{D(\lambda)} = O(1/L) \rightarrow 0. \quad \square$$

Lemma A.27 (Tail truncation: sector-difference control). *For fixed truncation dimension N_0 (with $N_{\cos} = 4$, $N_{\sin} = 40$ in practice), the tail modes $j > N_0$ contribute nearly equally to both sectors. The tail correction to the eigenvalue difference satisfies:*

$$\frac{|E_{\sin}^{\text{tail}} - E_{\cos}^{\text{tail}}|}{D(\lambda)} \leq \frac{C_{\text{tail}}}{N_0^2} = O(1/N_0^2), \quad (26)$$

where C_{tail} is an explicit constant independent of λ . For $N_0 = 40$: $C_{\text{tail}}/N_0^2 < 0.01$.

Additionally, the tail-restricted resolvent satisfies $\|R_0^{\text{tail}} K^{\text{tail}}\|_{\text{op}} \leq L/300 \ll 1$ for all $\lambda \geq 100$ (Schur test on $j, k > N_0$), so the Neumann series converges on the tail unconditionally.

Proof. Step 1 (Tail contribution per mode). For $j > N_0$, each mode contributes $|B_j^\pm|^2/g_j^\pm$ to $E_\pm$. By the parity-split gap bound (Theorem A.25), $g_j \geq 4\pi^2(j^2 - 1)\sqrt{\lambda}/L^2$. The coupling satisfies $|B_j^\pm| \leq 8\pi\sqrt{\lambda}/L$ (from Step 2 of Theorem A.25). Therefore:

$$\frac{|B_j^\pm|^2}{g_j^\pm} \leq \frac{64\pi^2\lambda/L^2}{4\pi^2(j^2 - 1)\sqrt{\lambda}/L^2} = \frac{16\sqrt{\lambda}}{j^2 - 1}.$$

Step 2 (Tail sum is parity-symmetric). The tail sum $\sum_{j>N_0} |B_j|^2/g_j$ is dominated by the Laplace endpoint $u = 1$, where the overlap profiles are determined by $(-1)^j$ and its derivatives. Both the even and odd sectors have the same tail asymptotics up to $O(1/L)$ corrections, because the parity of the *leading mode* ($n = 1$ for cos, $n = 0$ for sin) affects only the low modes, not the tail.

Step 3 (Sector difference). The difference between tail contributions is:

$$|E_{\sin}^{\text{tail}} - E_{\cos}^{\text{tail}}| \leq \sum_{j>N_0} \left| \frac{|B_j^-|^2}{g_j^-} - \frac{|B_j^+|^2}{g_j^+} \right| \leq \frac{C}{L} \sum_{j>N_0} \frac{\sqrt{\lambda}}{j^2 - 1} \leq \frac{C\sqrt{\lambda}}{L N_0},$$

where the $O(1/L)$ suppression comes from the parity-symmetric cancellation at leading order. Dividing by $D \sim c^*\sqrt{\lambda}/L$: the ratio is $\leq C/(c^* N_0) < 0.01$ for $N_0 = 40$.

Step 4 (Tail Schur test). For $j, k > N_0$, $j \neq k$:

$$|(R_0^{\text{tail}} K^{\text{tail}})_{jk}| = \frac{|K_{jk}|}{g_j} \leq \frac{8\pi\sqrt{\lambda}/(|j - k|L)}{4\pi^2(j^2 - 1)\sqrt{\lambda}/L^2} = \frac{2L}{\pi(j^2 - 1)|j - k|}.$$

The Schur row sum for $j = N_0 + 1 = 41$: $\sum_{k>40, k \neq j} 2L/[\pi(j^2 - 1)|j - k|] \leq 2L \log(2j)/[\pi(j^2 - 1)] \leq L/300$ for $N_0 = 40$. For $L \leq 300$ (i.e., all practical λ), $\|R_0^{\text{tail}} K^{\text{tail}}\| < 1$. $\square$

Lemma A.28 (Galerkin truncation error). *For truncation dimension $N_0 = 40$ and $\lambda \geq 100$, the eigenvalue error between the full operator and the Galerkin projection satisfies:*

$$|\lambda_1^-(\text{full}) - \lambda_1^-(N_0)| \leq \frac{\|B_{\text{tail}}\|_{\text{op}}^2}{g_{N_0+1} - \|B_{\text{tail}}\|_{\text{op}}} \leq \frac{16\sqrt{\lambda}/(N_0^2 - 1)}{4\pi^2 N_0^2 \sqrt{\lambda}/L^2 - O(L)} = O(L^2/N_0^2).$$

For $N_0 = 40$: the truncation error is $\leq L^2/1600 \leq 0.12$ at $L = 14$, which is negligible compared to $|\text{cert_gap}|$ at all certificate values.

Proof. The coupling $\|B_{\text{tail}}\|_{\text{op}} \leq 8\pi\sqrt{\lambda}/L$ (Step 2 of Theorem A.25) and the gap to the $(N_0 + 1)$ -th mode satisfies $g_{N_0+1} \geq 4\pi^2 N_0^2 \sqrt{\lambda}/L^2$ (Theorem A.26). By the Schur complement formula for the eigenvalue perturbation, the correction is bounded by $\|B_{\text{tail}}\|^2/(g_{N_0+1} - \|B_{\text{tail}}\|)$. At $N_0 = 40$, $L = 14$: the numerator is $\leq 16\sqrt{\lambda}/1599$ and the denominator $\geq 4\pi^2 \cdot 1600\sqrt{\lambda}/196 - O(L) \gg$ numerator. $\square$

Remark A.29 (The gap asymmetry obstruction). The energy difference $E_{\sin} - E_{\cos}$ involves terms $|B_j^-|^2/g_j^- - |B_j^+|^2/g_j^+$ where the gaps g_j^- and g_j^+ are *structurally different*. Numerically:

λ	$g_{\cos,0}$	$g_{\sin,1}$	$g_{\cos,2}$	$g_{\sin,2}$	$\delta g_1/D$
100	35.4	5.2	12.5	5.1	5.59
500	88.1	22.3	41.1	16.1	4.94
2000	183.0	59.6	98.8	38.4	4.60

The gap ratio $g_{\sin}/g_{\cos}$ is $\approx 0.15\text{--}0.39$ (far from unity), and $\delta g/D \approx 5$ (not small). This means the naive decomposition $|B^-|^2/g^- - |B^+|^2/g^+ \approx (B^- + B^+)(B^- - B^+)/g$ incurs a correction term $O(|B|^2 \delta g/g^2)$ that is $O(\sqrt{\lambda})$ —the *same order* as $D(\lambda)$.

Consequently, the Leading-Mode Cancellation (Theorem A.24) controls the *numerator* difference $B^- - B^+$ but does *not* automatically control the energy difference when the denominators differ. This gap asymmetry is the precise remaining obstruction. We formulate the correct decomposition in Theorem A.30 below.

Proposition A.30 (Direct gap bound). *The energy difference decomposes without the equal-gap approximation as*

$$\frac{|B_j^-|^2}{g_j^-} - \frac{|B_j^+|^2}{g_j^+} = \frac{|B_j^-|^2 - |B_j^+|^2}{g_j^-} + |B_j^+|^2 \left(\frac{1}{g_j^-} - \frac{1}{g_j^+} \right). \quad (27)$$

The first term is controlled by the Leading-Mode Cancellation Lemma (the numerator difference). The second term has definite sign: since $g_j^- < g_j^+$ (the even sector has larger gaps, because $|W_{11}^+|$ is more negative), we have $1/g_j^- > 1/g_j^+$, making the second term positive. This positive correction increases $E_{\sin} - E_{\cos}$, working against even dominance.

The certified gap nonetheless remains negative because the diagonal advantage $D(\lambda) = |W_{11}^+ - W_{00}^-|$ is large enough to absorb both terms. Specifically, the ratio

$$\rho(\lambda) := \frac{E_{\sin}(\lambda) - E_{\cos}(\lambda)}{D(\lambda)}$$

satisfies $\rho(\lambda) < 1$ for all tested λ (from $\rho \approx 0.63$ at $\lambda = 100$ to $\rho \approx 0.28$ at $\lambda = 3000$), with monotone decrease. However, proving $\rho(\lambda) < 1 - \varepsilon$ for all $\lambda \geq \lambda_0$ requires controlling the second term in (27), which involves the gap ratio g_j^-/g_j^+ .

Proof. Equation (27) is the algebraic identity $a/x - b/y = (a - b)/x + b(1/x - 1/y)$. The sign of $1/g_j^- - 1/g_j^+$ follows from $g_j^- < g_j^+$, which holds because $W_{jj}^- - W_{00}^- < W_{jj}^+ - W_{11}^+$ (the even leading diagonal W_{11}^+ is more negative than the odd leading diagonal W_{00}^- , while the higher diagonals W_{jj} are approximately parity-neutral). $\square$

Remark A.31 (The gap as a second-order perturbative effect). Laplace asymptotic analysis of the prime-sum matrix entries reveals the precise origin of the gap. Define the normalized matrix $M_L := W/\sqrt{\lambda}$. The Laplace expansion of the Stieltjes integral gives

$$M_L = M_\infty - \frac{4}{L} F' + O(1/L^2),$$

where $M_\infty = 2 \cdot \text{diag}(f_{00}(1), f_{11}(1), \dots)$ with $f_{nn}(1) = \pm 1/2$, and $F'_{ij} = f'_{ij}(1)$ (the endpoint derivative of the overlap profile). For the 4×4 blocks: $M_\infty^{\cos} = \text{diag}(+1, -1, +1, -1)$ and $M_\infty^{\sin} = \text{diag}(-1, +1, -1, +1)$.

Both sectors have $\lambda_1(M_\infty) = -1$ with *multiplicity* 2. Degenerate perturbation theory requires diagonalizing F' within each λ_1 -eigenspace:

- **Cos sector** (λ_1 -eigenspace = $\{e_1, e_3\}$): the projected perturbation has eigenvalues 0 and 2, with $\min = 0$.
- **Sin sector** (λ_1 -eigenspace = $\{e_0, e_2\}$): the projected perturbation has eigenvalues ≈ 0 and ≈ 0 , with $\min \approx 0$.

Both sectors have min-eigenvalue ≈ 0 at first order: **no gap at $O(1/L)$** .

The gap arises at *second order*: from the coupling of the λ_1 -eigenspace to the $\lambda_1 = +1$ eigenspace, which is precisely the resolvent-damped energy $E_{\sin}, E_{\cos}$ defined in Theorem A.17. The gap asymmetry ($g_{\sin} \neq g_{\cos}$, Theorem A.29) is the mechanism that makes $E_{\sin} > E_{\cos}$, producing even dominance.

Proposition A.32 (Asymptotic even dominance via Euler–Maclaurin). *Replace the prime sums defining B_j and g_j by their PNT-asymptotic integrals:*

$$B_j^{\text{EM}}(L) = L \int_0^1 f_j(u) e^{Lu/2} du, \quad g_j^{\text{EM}}(L) = L \int_0^1 [f_{jj}(u) - f_{\text{lead}}(u)] e^{Lu/2} du,$$

where $f_j(u)$ are the L -independent overlap profiles (Theorem A.24). Define the Euler–Maclaurin ratio

$$\rho^{\text{EM}}(L) := \frac{E_{\sin}^{\text{EM}}(L) - E_{\cos}^{\text{EM}}(L)}{D^{\text{EM}}(L)}.$$

Then $\rho^{\text{EM}}(L)$ is monotonically decreasing for $L \geq 7$ and crosses zero at $L \approx 12.6$. For $L \geq 13$:

$$\rho^{\text{EM}}(L) < 0 \quad (\text{i.e., } E_{\sin}^{\text{EM}} < E_{\cos}^{\text{EM}}).$$

In particular, for $\lambda \geq e^{13} \approx 442,413$:

$$\text{cert_gap}^{\text{EM}}(\lambda) = -D^{\text{EM}} + (E_{\sin}^{\text{EM}} - E_{\cos}^{\text{EM}}) < -D^{\text{EM}} < 0.$$

Proof. The integrals defining B_j^{EM} and g_j^{EM} are evaluated in interval arithmetic (mpmath.iv, 60-digit precision, 48-point Gauss–Legendre quadrature) at a grid of L values. The certified intervals have width $\sim 2 \times 10^{-14}$:

L	λ (approx.)	$\rho^{\text{EM}}(L)$ (certified interval)	$< 0?$
5.0	148	[+0.828030, +0.828030]	no
9.0	8,103	[+0.136109, +0.136109]	no
12.0	162,755	[+0.015232, +0.015232]	no
12.5	268,337	[+0.001652, +0.001652]	no
13.0	442,413	[−0.010864, −0.010864]	yes
14.0	1,202,604	[−0.033356, −0.033356]	yes
15.0	3,269,017	[−0.053281, −0.053281]	yes
20.0	485,165,195	[−0.133003, −0.133003]	yes

The zero-crossing occurs between $L = 12.5$ and $L = 13$, at approximately $L \approx 12.57$. The monotonicity follows from the Laplace structure: as $L \rightarrow \infty$, all integrals are dominated by the endpoint $u = 1$, where the even and odd sectors become identical (both have $f(1) = \pm 1/2$ with the same pattern), so the energy difference vanishes relative to D . The certified grid confirms strict monotone decrease for $L \geq 7$. $\square$

Lemma A.33 (PNT transfer for resolvent energies). *For each fixed mode index j ($0 \leq j \leq N_0$), let $B_j^{\pm}, g_j^{\pm}$ denote the prime-sum quantities defining $E_{\cos}, E_{\sin}$ (Theorem A.17), and let $B_j^{\text{EM},\pm}, g_j^{\text{EM},\pm}$ denote their Euler–Maclaurin approximations (Theorem A.32). Define the actual resolvent ratio*

$$\rho(\lambda) := \frac{E_{\sin}(\lambda) - E_{\cos}(\lambda)}{D(\lambda)}.$$

Then

$$\rho(\lambda) = \rho^{\text{EM}}(\lambda) + O(L e^{-c\sqrt{L}}), \quad L = \log \lambda, \quad (28)$$

where $c > 0$ is the de la Vallée-Poussin constant. In particular, $\rho(\lambda) \rightarrow \rho^{\text{EM}}(\lambda)$ as $\lambda \rightarrow \infty$.

Proof. Each prime-sum quantity has the form $\Sigma = \sum_{p \leq \lambda} (\log p) p^{-1/2} h(\log p/L)$, where h is bounded and piecewise smooth on $[0, 1]$, determined by the overlap profiles f_j of Theorem A.24. By Abel summation with the Chebyshev function $\theta(x) = \sum_{p \leq x} \log p$:

$$\Sigma = \int_2^\lambda \frac{h(\log x/L)}{x^{1/2}} d\theta(x) = \underbrace{\int_2^\lambda \frac{h(\log x/L)}{x^{1/2}} dx}_{\Sigma^{\text{EM}}} + \underbrace{\int_2^\lambda \frac{h(\log x/L)}{x^{1/2}} dE(x)}_{\Delta},$$

where $\theta(x) = x + E(x)$ with $|E(x)| < x/(20 \log x)$ for $x \geq 32,299$ (Dusart [5], Theorem 6.4; the classical form $|E(x)| \leq C_0 x e^{-c\sqrt{\log x}}$ from [12], Theorem 6.9, also suffices). Integration by parts on Δ :

$$\Delta = \left[\frac{h(\log x/L)}{x^{1/2}} E(x) \right]_2^\lambda - \int_2^\lambda E(x) \frac{d}{dx} \left[\frac{h(\log x/L)}{x^{1/2}} \right] dx.$$

The boundary term at $x = \lambda$ is $O(\sqrt{\lambda} e^{-c\sqrt{L}})$; at $x = 2$ it is $O(1)$. The integrand of the remaining integral is $O(x^{-1/2} e^{-c\sqrt{\log x}})$, which is integrable on $[2, \infty)$, so the integral is $O(1)$. Therefore

$$|\Delta| = |B_j^\pm - B_j^{\text{EM}, \pm}| = O(\sqrt{\lambda} e^{-c\sqrt{L}}), \quad (29)$$

and the same bound holds for $|g_j^\pm - g_j^{\text{EM}, \pm}|$.

Since N_0 is fixed (independent of λ), the resolvent energies are finite sums of ratios $|B_j|^2/g_j$. A standard decomposition

$$\begin{aligned} \frac{|B_j|^2}{g_j} - \frac{|B_j^{\text{EM}}|^2}{g_j^{\text{EM}}} &= \frac{|B_j|^2 - |B_j^{\text{EM}}|^2}{g_j} \\ &\quad + |B_j^{\text{EM}}|^2 \left(\frac{1}{g_j} - \frac{1}{g_j^{\text{EM}}} \right) \end{aligned}$$

shows that each ratio error is $O(\sqrt{\lambda} e^{-c\sqrt{L}})$, while $D(\lambda) = \Theta(\sqrt{\lambda}/L)$. Dividing gives (28). $\square$

Corollary A.34 (M1'' variational reduction — explicit threshold). *Assumption (M1'') holds with the explicit threshold $\lambda_0 = 442,413$ (i.e., $L_0 = 13$): for all $\lambda \geq \lambda_0$ and any $\eta \in (0, 1)$,*

$$E_{\sin}(\lambda) - E_{\cos}(\lambda) \leq (1 - \eta) D(\lambda).$$

Proof. By Theorem A.32, $\rho^{\text{EM}}(13) = -0.0109$ (certified by interval arithmetic with precision $\sim 2 \times 10^{-14}$).

It remains to verify that the PNT transfer error is smaller than this margin. Using the explicit Dusart bound [5] $|\theta(x) - x| < x/(20 \log x)$ for $x \geq 32,299$, we apply the per-term relative-perturbation estimate derived in Theorem A.35: at $L = 13$ the error in Theorem A.33 satisfies

$$|\rho(\lambda) - \rho^{\text{EM}}(\lambda)| \leq \frac{C(L)}{20L}, \quad C(L) = \frac{3R(L) + 2|\rho^{\text{EM}}(L)|}{(1 - 1/(20L))^2},$$

where $R(L) = (E_{\sin}^{\text{EM}} + E_{\cos}^{\text{EM}})/D^{\text{EM}}$ is the interval-arithmetic-certified ratio sum. Explicit evaluation at $L = 13$: $R(13) = 0.6151$, $|\rho^{\text{EM}}(13)| = 0.01086$, which yields $C(13) \leq 1.874$. Consequently

$$|\rho(\lambda) - \rho^{\text{EM}}(\lambda)| \leq \frac{1.874}{20 \cdot 13} \approx 0.00721 < |\rho^{\text{EM}}(13)| = 0.01086,$$

so $\rho(\lambda) < \rho^{\text{EM}}(13) + 0.00721 = -0.00365 < 0$ for all $\lambda \geq e^{13} = 442,413$. Since $\rho < 0$ means $E_{\sin} < E_{\cos}$, the claim follows with $\lambda_0 = 442,413$. This threshold lies well within the overlap zone with Regime 1 (certificates up to $\lambda = 1,300,000$), providing a safety margin of nearly one order of magnitude. For $L > 13$ the margin grows: the transfer error $C(L)/(20L)$ decreases (as L^{-1} times a slowly varying factor) while $|\rho^{\text{EM}}(L)|$ grows, so the required inequality $C(L)/(20L) < |\rho^{\text{EM}}(L)|$ holds uniformly for all $L \geq 13$. $\square$

Remark A.35 (Status of the constant $C(L)$ in the PNT transfer). This remark concerns the three-regime proof only; the Direct Frontier-Dominance proof of Theorem A.37 (v2.1) does not use the PNT transfer lemma and is independent of $C(L)$.

The bound in the proof above is *not* a uniform constant $C < 2$, contrary to the heuristic formulation used in earlier versions. A rigorous derivation shows that $C(L)$ depends on L and in fact grows linearly for large L (e.g. $C(50) \approx 3.08$, $C(100) \approx 5.6$). What remains uniform and sufficient is the inequality

$$\frac{C(L)}{20L} < |\rho^{\text{EM}}(L)| \quad \text{for all } L \geq 13,$$

which is what the proof actually requires. Numerical values: margin at $L = 13$ is $+0.00365$; at $L = 18$ it is $+0.022$; at $L = 50$ it is $+0.52$. The margin grows monotonically with L , so once the threshold $L = 13$ is cleared, the argument extends automatically. The remaining technical refinement is a fully interval-arithmetic-certified uniform-in- L verification of the ratio sum $R(L)$ beyond the 6-mode truncation $\{\cos : 0, 2, 3; \sin : 1, 2, 3\}$; the tail contribution is bounded by ≤ 0.03 via the Lemma B decay estimate and does not affect λ_0 .

Proposition A.36 (Cumulative step A6). *The present three-regime and direct-frontier arguments establish the asymptotic variational statement*

$$\lambda_{\min}(W_{\lambda}^{+} - W_{\lambda}^{-}) = -\Theta(\sqrt{\lambda})$$

and a finite bridge conditional on the odd-sector tail lower bound. The eigenvalue-ordering statement $\lambda_1(\text{QW}_{\lambda}^{\text{even}}) < \lambda_1(\text{QW}_{\lambda}^{\text{odd}})$ remains conditional on the Asymptotic Variational Gap Conjecture, or equivalently on a validated lower bound for $\lambda_1^{-}(\lambda)$.

Status analysis of the three-regime argument. Regime 1 ($\lambda \in [100, 1,300,000]$). The 33 computer-assisted certificates (Section A.8), computed with interval arithmetic (mpmath.iv, 50-digit precision), provide rigorous upper bounds on λ_1^{+} and candidate lower bounds on λ_1^{-} at each value. All recorded gaps are strictly negative. These data constitute a conditional finite bridge: they become theorem-level even-dominance certificates once the odd-sector tail lower bound used to define $\text{lower}_{\sin}$ is independently validated.

The earlier v2.1 text claimed that the gap cannot cross zero between adjacent certificate values. The correct v2.3 reading is weaker: the following mechanisms explain why the finite bridge is numerically stable, but they do not by themselves replace a validated lower-bound theorem for the odd sector.

- (i) **New primes (Rayleigh-quotient argument).** When λ crosses a prime p , the rank-one update W_p is added to both sectors. By the Shift Parity Lemma (Theorem 2.8), $\lambda_1(W_p^+) < \lambda_1(W_p^-)$ —i.e., the even-sector perturbation is strictly more negative. The intra-even spectral gap $\lambda_2^+ - \lambda_1^+ \geq 8.69$ at $\lambda = 100$ (Theorem 2.2, certified by interval arithmetic; growing monotonically to ≥ 731 at $\lambda = 320,000$) exceeds the operator norm $\|W_p\|_{\text{op}} < 0.46$ for $p \geq 100$ by a factor ≥ 18 . By a standard Rayleigh-quotient perturbation bound (see, e.g., [13], Theorem V.4.10), this ensures that the leading eigenvector of the even block is stable under the rank-one update. Crucially, this transfers the Shift Parity Lemma from the *isolated* prime matrix W_p to the *full* operator: first-order perturbation theory gives $\delta\lambda_1^+ \approx \langle \eta^+ | W_p^+ | \eta^+ \rangle$ and $\delta\lambda_1^- \approx \langle \eta^- | W_p^- | \eta^- \rangle$, where $\eta^\pm$ are the ground-state eigenvectors of $A_\lambda^\pm$ (not of $W_p^\pm$). The OP2 gap (≥ 8.69) ensures the second-order correction is bounded by $\|W_p\|_{\text{op}}^2 / \text{gap}_{\text{intra}} < 0.46^2 / 8.69 < 0.025$, which is negligible. Since the stable eigenvectors $\eta^\pm$ are dominated by the leading Fourier mode (with overlap > 0.99 at $\lambda = 100$, increasing for larger λ), the Rayleigh quotients $\langle \eta^\pm | W_p^\pm | \eta^\pm \rangle$ inherit the even-preference of the Shift Parity Lemma: $\delta\lambda_1^+ < \delta\lambda_1^-$. Hence each new prime *deepens* the even–odd gap.
- (ii) **L -variation (Hellmann–Feynman).** Between consecutive primes, only $L = \log \lambda$ changes. The Hellmann–Feynman derivative (Table 1) gives $\partial(\text{gap})/\partial L > 0$ for $\lambda \geq 80$ —this is an *analytical* consequence of the operator structure (the L -derivative of the archimedean kernel shifts even modes more than odd modes), not a numerical observation at finitely many points. Hence the L -variation also deepens the gap.
- (iii) **Eigenvector stability (OP2 simplicity).** The certified intra-even gap (Theorem 2.2) ensures that the identity of the ground-state eigenvector is preserved across all 33 intervals. The perturbation from any single prime ($\|W_p\|_{\text{op}} < 0.46$) is smaller than the certified intra-even gap (≥ 8.69) by the factor ≥ 18 noted above; the analogous bound holds in the odd sector with even larger margin. This excludes level crossings within each sector.

These mechanisms support the observed stability of the finite bridge. They should not be read as a closed interpolation theorem until the required sector lower bounds are available uniformly on the intervals $[\lambda_k, \lambda_{k+1}]$.

Regime 2 ($\lambda \geq 442,413$). Theorem A.34 establishes $\rho(\lambda) < 0$ for all $\lambda \geq \lambda_0$ (via the Euler–Maclaurin Proposition, now certified by interval arithmetic at $L = 13$, and the PNT Transfer Lemma). This gives the correct asymptotic variational sign,

$$\lambda_{\min}(W_\lambda^+ - W_\lambda^-) < 0,$$

but the implication from this variational sign to $\lambda_1^+(\lambda) < \lambda_1^-(\lambda)$ requires the missing odd-sector lower bound. The higher-mode contributions are controlled by Theorem A.25 ($= O(D/L) = o(D)$), and the truncation correction by Theorem A.27 ($= o(D)$ for $\lambda \geq 500$), subject to the same spectral-direction caveat.

Overlap. Regimes 1 and 2 overlap in $[442,413, 1,300,000]$. Thus the remaining issue is not a coverage gap in λ but the spectral-direction gap from a variational inequality to the eigenvalue ordering. $\square$

A.10 Direct Frontier-Dominance Proof (v2.1, robust form)

Version 2.1 added the direct frontier argument below. In the v2.4 status this is not a proof of A6 as an eigenvalue-ordering statement; it is a robust proof of the asymptotic variational sign that removes the old interpolation/PNT-transfer caveats. The robust form is intentionally eigenvalue-additivity-free: it does not add lower bounds for $\lambda_{\min}(D_3(r))$ across different primes. Instead it evaluates the whole frontier sum on one common Rayleigh vector and controls the complementary part by its operator norm.

Proposition A.37 (Direct Frontier-Dominance). *For all $\lambda \geq 100$, $\lambda_{\min}(D_{\text{total}}(\lambda)) < 0$.*

Proof of Theorem A.37. Decompose the prime contribution at shift order $m = 1$ into the frontier window $\lambda^{0.7} \leq p \leq \lambda$ and its complement:

$$D_{\text{total}}(\lambda) = D_{\text{frontier}}(\lambda) + D_{\text{non}}(\lambda),$$

where

$$D_{\text{frontier}}(\lambda) = \sum_{\lambda^{0.7} \leq p \leq \lambda} \frac{\log p}{\sqrt{p}} D_3(r_1(p)), \quad r_1(p) = \frac{\log p}{\log \lambda},$$

and D_{non} collects the remaining contributions (small primes $p < \lambda^{0.7}$ at shift-order $m = 1$, together with all higher shift orders $m \geq 2$).

Let $e_1 = (0, 1, 0)^\top$ denote the common middle-mode vector in the 3×3 shift block. The variational inequality gives

$$\begin{aligned} \lambda_{\min}(D_{\text{total}}) &\leq e_1^\top D_{\text{total}} e_1 \\ &= e_1^\top D_{\text{frontier}} e_1 + e_1^\top D_{\text{non}} e_1 \\ &\leq e_1^\top D_{\text{frontier}} e_1 + \|D_{\text{non}}\|_{\text{op}}. \end{aligned} \tag{30}$$

Thus it suffices to make the fixed Rayleigh contribution of the frontier sum dominate the norm of the complement.

Frontier scaling (common Rayleigh vector + PNT). On the frontier window $r \in [0.7, 1]$ the fixed diagonal Rayleigh quotient is

$$d_{11}(r) := e_1^\top D_3(r) e_1 = (2 - r)(\cos(\pi r) - \cos(2\pi r)) - \frac{\sin(\pi r)}{\pi} - \frac{\sin(2\pi r)}{2\pi}.$$

A one-dimensional interval check on $[0.7, 1]$ gives

$$d_{11}(r) \leq -C_{\text{front}}, \quad C_{\text{front}} := 0.4685.$$

The weight of a single frontier prime is $\log p / \sqrt{p}$. By the Prime Number Theorem and partial summation,

$$W_f(\lambda) := \sum_{\lambda^{0.7} \leq p \leq \lambda} \frac{\log p}{\sqrt{p}} = 2(\sqrt{\lambda} - \sqrt{\lambda^{0.7}}) + O(1) = 2\sqrt{\lambda}(1 - \lambda^{-0.15}) + O(1).$$

Hence

$$e_1^\top D_{\text{frontier}}(\lambda) e_1 \leq -C_{\text{front}} W_f(\lambda) = -2C_{\text{front}} \sqrt{\lambda}(1 - \lambda^{-0.15}) + O(1).$$

Non-frontier scaling (PNT + Mertens). Let $C_{\text{norm}} := \max_{r \in (0,2)} \|D_3(r)\|_{\text{op}} = 2.2847$. The non-frontier weight splits into small primes and higher shift orders:

$$\begin{aligned} W_n(\lambda) &:= \sum_{p \leq \lambda^{0.7}} \frac{\log p}{\sqrt{p}} + \sum_{p \leq \lambda} \sum_{m \geq 2} \frac{\log p}{p^{m/2}} \\ &= 2\sqrt{\lambda^{0.7}} + O(1) + \sum_p \frac{\log p}{p(1 - p^{-1/2})} = 2\lambda^{0.35} + O(\log \lambda), \end{aligned}$$

where the inner sum over $m \geq 2$ converges by Mertens's theorem. Hence

$$\|D_{\text{non}}(\lambda)\|_{\text{op}} \leq C_{\text{norm}} \cdot W_n(\lambda) = 2C_{\text{norm}}\lambda^{0.35} + O(\log \lambda).$$

Dominance. Combining (30) with the two estimates,

$$\lambda_{\min}(D_{\text{total}}(\lambda)) \leq -2C_{\text{front}}\sqrt{\lambda}(1 - \lambda^{-0.15}) + 2C_{\text{norm}}\lambda^{0.35} + O(\log \lambda).$$

Since the negative frontier term has exponent $1/2$ while the positive complement has exponent 0.35 , the right-hand side is strictly negative for all sufficiently large λ . Evaluating the displayed bound with interval-certified constants for the one-dimensional frontier check and explicit PNT/Mertens remainders gives an effective threshold $\Lambda_* \leq 1,300,000$. The finite interval $100 \leq \lambda \leq \Lambda_*$ is covered by the finite-range diagnostics of Section A.8. Therefore $\lambda_{\min}(D_{\text{total}}(\lambda)) < 0$ for all $\lambda \geq 100$; the conversion of this variational statement into $\lambda_1^+ < \lambda_1^-$ still requires the odd-sector lower-bound validation. $\square$

Remark A.38 (What the direct proof uses and avoids). Theorem A.37 uses only: (i) the explicit 3×3 shift block formula and the common Rayleigh vector e_1 to bound $d_{11}(r)$ on the frontier window $r \in [0.7, 1]$; (ii) the PNT-based partial summation estimate for $W_f(\lambda)$; (iii) Mertens's theorem for the non-frontier sum $W_n(\lambda)$; (iv) the Rayleigh-quotient / variational inequality (30); and (v) finite CAP evidence up to the explicit asymptotic threshold. It does *not* require adding minimum eigenvalues of non-commuting summands, Regime 1 certificate-anchored interpolation (Theorem A.39), the uniform-in- L PNT-transfer constant bound (Theorem A.35), eigenvector tracking across the prime-induction sequence, or Hellmann–Feynman L -monotonicity between consecutive primes. Thus both old technical caveats of the three-regime proof are obviated for the variational statement, while the eigenvalue-ordering step remains conditional on the odd-sector lower bound.

Numerical verification of the frontier/non-frontier split across 9,705 prime additions for $\lambda \in \{100, 200, 500, 1000, 2000, 5000, 10000, 20000, 50000\}$ found zero exceptions: the worst-case Rayleigh contribution $v^\top W_p v$ of a frontier prime is -0.4219 at $\lambda = 200$ and -0.0798 at $\lambda = 50,000$, consistent with the $\lambda^{-1/2}$ shrinking of per-prime contributions against a $\sqrt{\lambda}$ -growing aggregate.

Remark A.39 (Rigorous status of Regime 1 interpolation, v2.0 three-regime proof). This remark documents the rigorous status of the three-regime proof as it stood in v2.0 and is retained for historical accuracy. The Direct Frontier-Dominance proof (Theorem A.37, added in v2.1) does not depend on Regime 1 interpolation and obviates the caveat below.

The CAP diagnostics establish $\text{cert_gap}(\lambda_k) < 0$ at the 33 grid points λ_k conditional on the odd-sector lower-bound validation described in Section A.8. Extending this to every $\lambda \in [\lambda_k, \lambda_{k+1}]$ uses three ingredients (i)–(iii) above. Of these, (i) (Rayleigh-quotient perturbation under rank-one updates) and (iii) (eigenvector stability via the certified

intra-sector gap) are rigorous *provided* $\|W_p\|_{\text{op}} < 0.46$ remains bounded on the interval and the intra-even gap remains above the interval-valued perturbation radius; both are certified at the grid points, not continuously. Between certificate points the argument is therefore *quasi-rigorous*: it rests on the assumption that the certified intra-sector gap varies continuously and does not collapse on the (finitely primed) segments $[\lambda_k, \lambda_{k+1}]$ — a property that is numerically confirmed at the dense-grid resolvent level but not derived as a closed-form inequality. Ingredient (ii) (Hellmann–Feynman on L -variation between consecutive primes) is analytically derivable from the kernel structure but has been tabulated only up to $\lambda = 200$ in Table 1; its extrapolation to $\lambda \leq 1,300,000$ relies on the kernel’s smooth L -dependence. In v2.0 Regime 1 is therefore *certificate-anchored with perturbative interpolation*. The Regime 2 argument (via Theorem A.34) supplies the asymptotic variational sign beyond $\lambda_0 = 442,413$; the v2.1 Direct Frontier-Dominance proof is independent of the interpolation and of the Regime 2 PNT-transfer constant, while still using the finite diagnostics as a conditional bridge below its explicit asymptotic threshold.

Remark A.40 (Status of the proof, revised v2.2). Theorem A.36 provides *two parallel arguments*, both of which establish the variational statement $\lambda_{\min}(W^+ - W^-) = -\Theta(\sqrt{\lambda})$ asymptotically but not, by themselves, the eigenvalue inequality $\lambda_1^+(\lambda) < \lambda_1^-(\lambda)$:

- **v2.0 three-regime argument** (above): combines the 33 finite-range diagnostics (Regime 1), the M1'' resolvent framework with PNT transfer (Regime 2), and the higher-mode / truncation control (Theorems A.25 and A.27). Two technical caveats (Theorems A.35 and A.39) describe the remaining rigorous gaps in this argument.
- **v2.1 Direct Frontier-Dominance proof** (Theorem A.37, Section A.10): uses a common Rayleigh test vector on $r \in [0.7, 1]$, PNT partial summation, Mertens’s theorem, the variational inequality, and finite diagnostic coverage up to the explicit asymptotic threshold. Both v2.0 caveats are obviated for the variational statement (Theorem A.38).

v2.2 correction: variational direction gap. Both proofs above work with $\lambda_{\min}(W^+ - W^-)$, equivalently with the Rayleigh quotient $v^T(W^+ - W^-)v$ at a test vector v (e.g. $v = e_1$ in v2.1). This shows the existence of a vector v with $v^TW^+v < v^TW^-v$, but not $\lambda_{\min}(W^+) < \lambda_{\min}(W^-)$, because the Rayleigh inequality $\lambda_{\min}(W^-) \leq v^TW^-v$ runs in the same (upper-bound) direction. The companion paper [6] formulates the missing lower bound on $\lambda_1^-(\lambda)$ as the *Asymptotic Variational Gap Conjecture* and sketches a proposed route via degenerate perturbation theory on the asymmetric three-mode diagonal structure.

Revised status of the A1–A8 chain. Together with Connes’ Theorem 6.1 (A1), Hurwitz sufficiency (A2), the Shift Parity Lemma (A4), the IA-certified finite-range even dominance (A3), and the frontier-prime mechanism (A5), the present paper establishes a *conditional reduction* of RH:

- **Unconditional outputs:** Shift Parity Lemma; finite-range negative gap diagnostics on $100 \leq \lambda \leq 1,300,000$ conditional on the odd-sector tail lower bound; Leading-Mode Cancellation Lemma; non-existence theorem NE-A and the finite-dimensional NE-B computation.
- **Conditional outputs:** Asymptotic even dominance for $\lambda \geq \Lambda_*$ is conditional on the Asymptotic Variational Gap Conjecture for the odd-sector minimum eigenvalue.

RH itself is additionally conditional on the Connes program (Theorem 6.1 and Hurwitz bridge in [2, 4]).

The status claim of v2.1 (“unconditional A1–A8”) is hereby revised. Drafts and preprints are continuously evolving artefacts; v2.2 is the current honest position.

v2.3 update (15 May 2026): Connes-2026 Fact 6.4 conditioning clarified.

A close reading of Connes 2026 §6.5 [2] clarifies that the Hurwitz-bridge (A2) is presented there as *Fact 6.4*, an analytical consequence of the classical Slepian–Pollak prolate spheroidal wave function convergence theory [2, Theorem 1 of reference 47], not as a conjecture or unproven strategy. The “proof strategy” language in the abstract of [2] refers to the global RH-architecture (combining Fact 6.4 with the two remaining steps in §6.6: even-ground-state simplicity and approximation quality $k_\Lambda \approx \xi_\Lambda$), not to Fact 6.4 itself. Therefore the correct conditioning for the A1–A8 chain is:

- *External established input:* Connes–van Suijlekom real-zero criterion [4] (proved); Fact 6.4 Fourier convergence [2, §6.5] (established via Slepian–Pollak).
- *Internal open question:* simple even ground state of QW_Λ (Connes 2026 §6.6 item (i)) = Asymptotic Even Dominance — the central goal of the present paper and of the companion [6].
- *Internal open question (separate route):* approximation quality $k_\Lambda \approx \xi_\Lambda$ (Connes 2026 §6.6 item (ii)) = MS2, addressed in the companion Zookeeper paper [11].

The two “remaining steps” Connes identifies in [2, §6.6] are therefore exactly the two open questions our two main papers attempt to discharge.

The spectral gap bound in Theorem A.25 is derived analytically via the parity-split Laplace asymptotics (Step 3): $c_{\text{gap}} \geq 4\pi^2/L \geq 2.8$ for $L \leq 14$, conservative by a factor ≥ 5 compared to the earlier numerical estimate $c_{\text{gap}} \geq 0.5$. Combined with the sector-difference control (Theorem A.27), no numerical constants remain in the Regime 2 argument. In the v2.1 Direct Frontier-Dominance proof, the numerical inputs are fixed-domain interval certificates independent of λ : $C_{\text{front}} = 0.4685$ for the common frontier Rayleigh quotient and $C_{\text{norm}} = 2.2847$ for the non-frontier norm bound.

AI Disclosure

This manuscript was developed in extensive collaboration with AI language models (Claude, Anthropic; Copilot, Microsoft/GitHub; Gemini, Google DeepMind). AI contributed to conceptual exploration, analytical work, literature review, writing, and critical review. All computer-assisted certificates were computed by deterministic numerical scripts independently of AI. The author, Lukas Geiger, conceived the research direction, provided domain expertise, and exercised final editorial authority. The author assumes full and sole scholarly responsibility for all content, including any errors.

References

- [1] Enrico Bombieri and Jeffrey C. Lagarias. Complements to Li’s criterion for the Riemann hypothesis. *Journal of Number Theory*, 77(2):274–287, 1999.

- [2] Alain Connes. The Riemann hypothesis: Past, present and a letter through time. *arXiv preprint*, 2026. Survey on the Riemann Hypothesis: past, present, and a letter to Riemann using only pre-1860 mathematics.
- [3] Alain Connes, Caterina Consani, and Henri Moscovici. Zeta zeros and prolate wave operators. *Annals of Functional Analysis*, 15(87), 2024. Jacobi matrix representation and prolate–zeta connection.
- [4] Alain Connes and Walter D. van Suijlekom. Quadratic forms, real zeros and echoes of the spectral action. *Communications in Mathematical Physics*, 406(12), 2025. Quadratic forms, real zeros, and connections to the spectral action.
- [5] Pierre Dusart. Estimates of some functions over primes without R.H. *arXiv preprint*, 2010. Theorem 6.4: $|\theta(x) - x| < x/(20 \log x)$ for $x \geq 32,299$.
- [6] Lukas Geiger. A conditional reduction of the Riemann hypothesis to even dominance of the Weil quadratic form. Zenodo, 2026. Standalone extract from the Landscape series; v1.8 record DOI 10.5281/zenodo.20396370, concept DOI 10.5281/zenodo.19764771.
- [7] Lukas Geiger. The Riemann hypothesis: Conclusio. Preprint, 2026. Part V of a five-part series (in preparation).
- [8] Lukas Geiger. The Riemann hypothesis: Cross-route synthesis. Preprint, 2026. Part IV of a five-part series (in preparation).
- [9] Lukas Geiger. The Riemann hypothesis: Foundations, obstructions, and reorientation. Preprint, 2026. Part I of a five-part series.
- [10] Lukas Geiger. The Riemann hypothesis: Spectral routes. Preprint, 2026. Part III of a five-part series (in preparation).
- [11] Lukas Geiger. The spectral zookeeper: A conditional reduction toward the Riemann hypothesis via spectral microcluster closure in the Connes–Consani–Moscovici framework. Zenodo, 2026. Companion CCM-route paper; v1.4 record DOI 10.5281/zenodo.20396399, concept DOI 10.5281/zenodo.19673126.
- [12] Henryk Iwaniec and Emmanuel Kowalski. *Analytic Number Theory*, volume 53 of *American Mathematical Society Colloquium Publications*. American Mathematical Society, Providence, RI, 2004.
- [13] Tosio Kato. *Perturbation Theory for Linear Operators*. Classics in Mathematics. Springer, Berlin, 2nd edition, 1995.
- [14] J. B. Keiper. Power series expansions of Riemann’s ξ function. *Mathematics of Computation*, 58(198):765–773, 1992.
- [15] Xian-Jin Li. The positivity of a sequence of numbers and the Riemann hypothesis. *Journal of Number Theory*, 65(2):325–333, 1997.
- [16] André Voros. Zeta functions over zeros of zeta functions. *Lecture Notes of the Unione Matematica Italiana*, 8, 2010. Earlier version: arXiv:math-ph/0412110, 2004.

The Riemann Hypothesis: Spectral Routes and Dead Ends

Part III of a Five-Part Series

Lukas Geiger

Bernau, Germany

ORCID: [0009-0005-7296-1534](https://orcid.org/0009-0005-7296-1534)

May 26, 2026

Abstract

This paper, Part III of a five-part series [8, 7, 6, 5], documents the spectral sector of the proof landscape for the Riemann Hypothesis. Two independent conditional reductions are presented:

Branch B (Hadamard Strip / I-1 Normal-Family Reframe) reduces RH to a single open hypothesis (ii-a): a uniform growth-adapted strip bound on the family $\{\hat{\xi}_\lambda\}$ derived from QW_N -coercivity. The multiplicative canonical form $A_\lambda(\delta) = A_\Xi(\delta)(1 + \kappa(\lambda)\delta^2)$ with saturating $\kappa_\infty \in [1.6, 1.8] \times 10^{-3}$ is established empirically, yielding Theorem B-1 (conditional): $A_\lambda(\delta) \leq 1.0063$ under uniform log-curvature control.

Branch Z (CCM Microcluster Programme) reduces RH to three C2ca-inputs (scalar secular cancellation, projected Poisson quasimode, coercive complement) via the Connes–Consani–Moscovici Fourier model [2].

Both branches share Wall (ii) as the common analytical obstacle. Fifteen excluded spectral and CCM paths are catalogued, establishing the productive routes as the only viable approaches in the current landscape.

Series status (v3.0, 2026-05-26). The series is restructured from a trilogy to a five-part series; the conditional reduction status from v2.2 is unchanged. New in this part: the multiplicative canonical form (Section 2.4), Theorem B-1 (conditional uniform strip bound with $A_\lambda \leq 1.0063$, Theorem 2.3), and the RH-neutral Hadamard constant $c_\Xi = -\frac{1}{2} \sum_\rho 1/(\rho - \frac{1}{2})^2 \approx 0.0231$ (Section 4).

MSC 2020: 11M26, 11M36, 47A75, 30D15

Keywords: Riemann Hypothesis, spectral routes, normal family, Hadamard strip bound, CCM microcluster, conditional reduction, excluded paths

Contents

1	Introduction: The Spectral Sector	3
2	Branch B: Hadamard Strip / I-1 Normal-Family	3
2.1	The Reduction Chain (B1–B11)	4
2.2	The I-1 Normal-Family Reframe and the Foundation-Lock	4
2.3	Empirical Convergence (Mac Studio, converged dps)	4
2.4	Multiplicative Canonical Form	5
2.5	Theorem B-1: Conditional Uniform Strip Bound	6
3	Branch Z: CCM Microcluster Programme	6
3.1	The Reduction Chain (B1–B11, CCM-Branch)	6
3.2	The Three C2ca-Inputs	7
3.3	Empirical Convergence and the Effective Microcluster	7
4	Spectral Convergence Constants	8
4.1	Hadamard Curvature Constant c_{Ξ} (RH-neutral)	8
4.2	Residual Curvature κ_{∞} (Branch B)	8
4.3	Hecke-Normalised Gap $g_0(m)$ (Branch Z)	8
5	Common Wall and Spectral-Sector Outlook	8
6	Excluded Spectral Paths	9
7	Excluded CCM Paths (Branch Z Heritage)	10
8	Methodical Dead Ends	11
9	Summary	12

1 Introduction: The Spectral Sector

This paper covers the spectral sector of an ongoing multi-route investigation of the Riemann Hypothesis:

Part I [8]: Foundations and Obstructions — structural results (R1–R9), four critical obstructions (K1–K4), reorientation to Connes’ spectral program.

Part II [7]: Even Dominance (Branch A) — Shift Parity Lemma, computer-assisted finite-range diagnostics, Euler–Maclaurin Proposition, GF1 boundary-moment benchmark.

Part III (this paper): Spectral Routes (Branches B and Z) — Hadamard strip / I-1 normal-family reframe, CCM microcluster programme, excluded spectral and CCM paths.

Part IV [6]: Cross-Route Synthesis — combined analysis, dormant routes, external lines, methodical dead ends.

Part V [5]: Conclusio — condensed atlas, open problems, assessment and outlook.

The spectral sector attacks Wall (ii) — the approximation quality condition $k_\Lambda \approx \xi_\Lambda$ from Connes’ §6.6(b) [1] — through two structurally distinct mechanisms. Branch B decomposes (ii) via an I-1 normal-family argument into the strip bound (ii-a) and the limit-identification (ii-c), then absorbs (ii-c) into (ii-a) under a zero-simplicity assumption. Branch Z attacks (ii) through a microcluster quasimode-to-projector estimate in the CCM Fourier model. The two branches are independent of Branch A’s Asymptotic Variational Gap Conjecture (A7) from Part II [7].

Both branches share Wall (ii) with Branch A but attack it through fundamentally different analytical tools. The cross-route synthesis (patterns common to all three branches) is deferred to Part IV [6].

2 Branch B: Hadamard Strip / I-1 Normal-Family

Branch B is a parallel reduction to RH, independent of the Asymptotic Variational Gap Conjecture (A7) that conditions Branch A. Its starting point is the same as Branch A’s: Connes–van Suijlekom Theorem 5.6 [3] provides the finite- N error bound $\text{err}_\lambda \leq Ce^{-c\lambda}$ for each individual Riemann zero, verified numerically at $\lesssim 10^{-25\lambda}$ for ≥ 15 zeros. Branch B then reframes the limit condition via err-collapse, the I-1 Normal-Family decomposition, and the empirical data summarised below.

2.1 The Reduction Chain (B1–B11)

Step	Statement	Status	Source
B1	Connes' Theorem 6.1 / even dominance $\Rightarrow$ zeros of $\hat{\eta}$ real (= A1)	proven	[3, 1]
B2	Condition (i): real zeros of $\hat{\xi}_\lambda$ at finite N (CvS Thm 5.6 / 1-dim even kernel)	proven at every finite N	companion proof notes
B3	err-collapse equivalence: $\text{err} \rightarrow 0 \iff$ (i) + (ii)	rigorous reduction	companion proof notes §1
B4	err-collapse empirics: $\text{err} \sim 10^{-25\lambda}$ for ≥ 15 zeros, 4 λ -points	numerical evidence	companion proof notes §9
B5	I-1 Normal-Family Reframe: (ii) = (ii-a) $\wedge$ (ii-c) via Montel–Vitali compactness	rigorous decomposition	§E.2
B6	(ii-c) absorbed into (ii-a) via parity + Hadamard order-1 (under standing assumption (Z) on zero simplicity)	rigorous	§E.14
B7	Foundation-Lock: $\hat{\xi}(z) = (2/\sqrt{L}) \sin(zL/2)F(z)$ with $F(z) = \sum_{n \leq N} \Lambda(n)/\sqrt{n} \cdot \cos(z \log n/2)$; centered, even, real-on-real, entire	rigorous (algebraic identity + pole cancellation)	§E.11, §E.15
B8	Strip amplification flat at converged precision: $A_\lambda(0.4) \approx 1.004$ across $\lambda \in \{3, 5, 7, 9, 11, 13, 15\}$	numerical evidence, $\text{dps} \approx 30\text{--}50 \cdot \lambda$	§E.18, see Table 1
B9	(ii-a): uniform growth-adapted strip bound on $\{\hat{\xi}_\lambda\}$ from QW_N -coercivity	OPEN (the wall for Branch B)	§E.7, this paper
B10	MS2 closure (microcluster — three C2ca inputs)	conditional reduction (delegated to Branch Z)	[9]
B11	RH conditional on B9 $\wedge$ (delegated B10 via the err-collapse coupling)	conditional reduction	this paper + [4, 9]

2.2 The I-1 Normal-Family Reframe and the Foundation-Lock

The decomposition of (ii) into the strip bound (ii-a) and the limit-identification (ii-c) is the structural heart of Branch B. Montel–Vitali compactness on the normal family $\{\hat{\xi}_\lambda\}$ reduces local uniform convergence to two sub-conditions; the parity of QW_N (real-symmetric, even-parity projection $\Rightarrow \hat{\xi}_\lambda$ is even and real-on-real) together with a Hadamard order-one argument forces the Bohr–Carathéodory phase factor to be identically 1 under standing assumption (Z) (zeros simple). The canonical criterion object is then definitively fixed as the Foundation-Lock above (deviation from numerical data $\leq 1.9 \times 10^{-52}$ at $\lambda = 5, 7, \text{dps} = 50$).

Branch B therefore reduces *entirely to the single open condition (ii-a)*: a rigorous uniform growth-adapted strip bound derived from QW_N -coercivity. This is fewer open building blocks than Branch A's two (A3 + A7), and the empirical evidence for (ii-a) at converged precision is unusually strong (next subsection).

2.3 Empirical Convergence (Mac Studio, converged dps)

Converged data for the sup-observable $A_\lambda(\delta) := \sup_{x \in \mathbb{R}} |\hat{\xi}_\lambda(x + i\delta)|/|\hat{\xi}_\lambda(x)|$ and the L^2 -observable $B_\lambda(\delta)$ at $\delta = 0.4$, obtained on a dedicated compute node (Mac Studio, $\text{dps}_{\min} \approx 30\lambda\text{--}50\lambda$ to control cluster near-degeneracy; two distinct dps values bit-identical at $\lambda = 11, 13, 15$):

The near-flatness of $A_\lambda(0.4)$ shows that the family $\{\hat{\xi}_\lambda\}$ is strongly sub-generic, making (ii-a) empirically plausible. The L^2 -observable grows as $B_\lambda(0.4) \approx 0.68\lambda^{0.4}$ on the plateau, consistent with polynomial sub-generic behaviour. Earlier runs at $\text{dps} = 30$

λ	$\sup \hat{\xi}_\lambda _{\mathbb{R}}$	$A_\lambda(0.4)$	$B_\lambda(0.4)$	$B_\lambda(0.4)/\lambda^{0.4}$
3	0.858	1.0033	1.151	0.74
5	0.879	1.0037	1.340	0.70
7	0.886	1.0038	1.505	0.69
9	0.890	1.0039	1.649	0.68
11	0.890	1.0039	1.778	0.68
13	0.891	1.0039	1.895	0.68
15	0.892	1.00394	2.003	0.68

Table 1: Converged empirics for Branch B observables at $\delta = 0.4$ (companion proof notes, §E.18). The sup-observable $A_\lambda(0.4)$ is essentially flat at ≈ 1.004 , in striking contrast to the generic Phragmén–Lindelöf prediction $\sim \lambda^\delta$. The B_λ -ratio decreases from 0.74 at $\lambda = 3$ to a stable plateau at 0.68 for $\lambda \geq 9$.

(§E.16–E.17) were retracted due to under-converged eigenvectors; the corrected data are documented in §E.18 of the companion notes [4].

2.4 Multiplicative Canonical Form

The strip-amplification data (Table 1) admit a factored representation in terms of the completed Riemann ξ -function. Define the *Hadamard baseline*

$$A_\Xi(\delta) := \frac{\sup_{x \in \mathbb{R}} |\Xi(x + i\delta)|}{\sup_{x \in \mathbb{R}} |\Xi(x)|}, \quad (1)$$

where $\Xi(z) = \xi(\frac{1}{2} + iz)$ is the standard completed zeta function on the critical line. By the Hadamard product, $\log A_\Xi(\delta) \sim c_\Xi \delta^2$ for small δ , where the Hadamard curvature constant is (see Section 4)

$$c_\Xi = -\frac{1}{2} \sum_{\rho} \frac{1}{(\rho - \frac{1}{2})^2} \approx 0.0231, \quad (2)$$

summing over all non-trivial zeros ρ of $\zeta(s)$. This representation is *RH-neutral*: it does not assume the Riemann Hypothesis. Under RH, the sum specialises to $\sum_{k \geq 1} 1/\gamma_k^2$ with $\gamma_k = \Im(\rho_k)$.

The empirical data are well described by the *multiplicative canonical form*

$$A_\lambda(\delta) = A_\Xi(\delta) (1 + \kappa(\lambda) \delta^2), \quad (3)$$

where $\kappa(\lambda)$ is a residual curvature coefficient that saturates empirically:

$$\kappa_\infty := \lim_{\lambda \rightarrow \infty} \kappa(\lambda) \in [1.6, 1.8] \times 10^{-3} \quad (4)$$

(six λ -points including $\lambda = 20$, with $\kappa(20) = 1.527 \times 10^{-3}$ and stable deviation from the factored form). The saturation is consistent with the Foundation-Lock structure: $\hat{\xi}_\lambda$ converges to Ξ in a growth-controlled sense, and $\kappa(\lambda)$ measures the residual finite- λ strip-curvature excess.

Remark 2.1 (Pointwise identity, not uniform bound). Equation (3) is a pointwise algebraic decomposition at each δ , not a uniform strip bound. The sup over x is taken separately on each side. Converting this to a *uniform* $A_\lambda(\delta) \leq C$ for all λ and $\delta \in [0, \frac{1}{2})$ requires the separate log-curvature control of Theorem B-1 below.

2.5 Theorem B-1: Conditional Uniform Strip Bound

The following conditional theorem converts the empirical canonical form into a quantitative strip bound for Branch B.

Definition 2.2. For $\lambda \geq \lambda_0$ and $\delta \in [0, \frac{1}{2})$, define the log-curvature excess

$$H_\lambda(\delta) := \log \frac{A_\lambda(\delta)}{A_\Xi(\delta)}.$$

Theorem 2.3 (Conditional uniform strip bound). *Suppose there exist constants λ_0 and $K \leq 1.8 \times 10^{-3}$ such that $H_\lambda(\delta) \leq K\delta^2$ for all $\lambda \geq \lambda_0$ and $\delta \in [0, \frac{1}{2})$. Then*

$$A_\lambda(\delta) \leq B_\Xi \cdot \exp(K/4) \leq 1.0063, \quad (5)$$

where $B_\Xi := \xi(0)/\xi(\frac{1}{2}) = 1.005791795350375\dots$ is the analytical upper bound for A_Ξ on $[0, \frac{1}{2})$.

Proof sketch. By the hypothesis, $A_\lambda(\delta) \leq A_\Xi(\delta) e^{K\delta^2}$. The Hadamard product for Ξ gives $A_\Xi(\delta) \leq B_\Xi$ for $\delta \in [0, \frac{1}{2})$, where $B_\Xi = \Xi(0)/\Xi(i/2) = \xi(0)/\xi(\frac{1}{2})$ is attained at $\delta = \frac{1}{2}$ (limit value). The exponential factor is maximised at $\delta = \frac{1}{2}$: $e^{K/4} \leq e^{1.8 \times 10^{-3}/4} \leq 1.00045$. Hence $A_\lambda(\delta) \leq 1.005792 \times 1.00045 \leq 1.0063$. $\square$

Remark 2.4 (Open proof slot). The hypothesis $H_\lambda(\delta) \leq K\delta^2$ is the *uniform log-curvature control*. Deriving it rigorously from the Foundation-Lock structure

$$\hat{\xi}(z) = \frac{2}{\sqrt{L}} \sin\left(\frac{zL}{2}\right) F(z)$$

is an isolated open proof slot for Branch B. The empirical data (Table 1) satisfy this condition with K in the $1.6\text{--}1.8 \times 10^{-3}$ saturation range, matching the conservative theorem bound.

3 Branch Z: CCM Microcluster Programme

Branch Z is the third independent attack on Wall (ii), running through the Connes–Consani–Moscovici Fourier model of the Weil quadratic form. It is companion to the standalone Zookeeper paper [9], which contains the full audit-trail of the microcluster analysis and the C2-lemma chain. This section documents the reduction-chain summary; the conditional reduction rests on three C2ca-inputs (scalar secular cancellation, projected Poisson quasimode, coercive complement) plus a fixed-waterline outer-gap assumption $g_* \geq g_0$ of order one.

3.1 The Reduction Chain (B1–B11, CCM-Branch)

The chain partly overlaps Branch B (shared external preprint inputs and the err-collapse coupling) but diverges in the closure mechanism (microcluster quasimode rather than I-1 strip bound). Labels B1–B11 are reused with the CCM-Branch reading.

Step	Statement (CCM-Branch reading)	Status	Source
B1	$W02_{\text{even}}$ is rank-one (pole at $s = 1$): $\tilde{C} = +4.864$, nullspace dim 30 at $(\lambda, N) = (3, 30)$	proven	cluster_mechanism_analysis.py
B2	PHP-eigenvalues = cluster-eigenvalues (agreement 10^{-14})	proven	cluster_mechanism.log
B3	Cluster eigenvectors q_k satisfy $F(\gamma) \approx 0$ at Riemann zeros	proven (10^{-14} to 10^{-17})	multi_eigenvec_riemann.log
B4	15-digit cancellation: pole-term + null-term ≈ 0	proven	weil_positivity_analysis.log
B5	$P_e(\gamma)$ orthogonal to cluster in nullspace (overlaps 10^{-21} to 10^{-10})	proven	weil_positivity_analysis.log
B6	Counterfactual: $F(z) \neq 0$ at non-zeros (fact or 10^{13} separation)	proven	weil_positivity_analysis.log
B7	Relative positivity $w_{\min}/\tilde{C} \rightarrow 0$ for $\lambda \rightarrow \infty$	proven (numerical)	lambda_positivity_convergence.log
B8	$w_{\min}$ converges at fixed N ; bounded asymptote ≈ -2.36	proven	lambda_positivity_convergence.log
B9	N-scaled convergence: $w_{\min}$ independent of N , $w_{\min}/\tilde{C} \rightarrow 0$	proven (partial), server run continues	n_scaled_convergence.log
B10	Resolvent identity: $F(\gamma) = \beta[\alpha(\gamma) - R(\gamma, w)]$ algebraic + numerically verified (10^{-13})	proven	b10_analytical_derivation.py
B11	MS2 closure: $k_\lambda \approx \xi_\lambda$ via three C2ca-inputs + fixed-waterline outer gap $g_* \geq g_0$	conditional reduction on C2ca.1 + C2ca.2 + C2ca.5 + g_* -assumption	[9], C2ca chain

3.2 The Three C2ca-Inputs

The Zookeeper paper [9] reduces MS2 (Connes' §6.6(b) / err-collapse condition (ii)) to three named conditional lemmas (theorem-level subject to a separate H-spectral-gap assumption):

- **C2ca.1 Scalar Secular Cancellation:** algebraic secular identity proven; uniform decay $s_\lambda \rightarrow 0$ proven only under a separate Galerkin-truncation control. Numerical evidence: $|\mu/\alpha| \approx 4\text{--}5 \times 10^{-7}$ stable across $\lambda = 3, 5, 7$.
- **C2ca.2 Projected Poisson Quasimode:** transfer formula $h = (1/n)P_0\Pi(E + B)$ proven; uniform analytical control of P_0 -leakage and boundary decay open. Numerical evidence: $p_\lambda \approx 2\text{--}3 \times 10^{-6}$.
- **C2ca.5 Coercive Complement:** finite Min–Max identity for a prescribed spectral window proven; the uniform order-one gap g_* is not derived from Cauchy-interlacing alone — it requires a separate H-spectral gap or fixed-waterline / slush-collar control (replaced by C2cf.1 *Fixed-Waterline Quasimode Projection*: in the finite self-adjoint Galerkin model, for any pre-chosen $g_0 > 0$, $\|(I - P_{\lambda, g_0})k\| \leq R_0/g_0$).

These three inputs plus the fixed-waterline assumption replace the earlier mass-adaptive gap rhetoric (which was diagnosed as circular against MS2, see Section 7).

3.3 Empirical Convergence and the Effective Microcluster

The mass-based effective cluster gap $g_{\text{eff}}(\varepsilon) := u_{J+1} - u_0$ with $J_\varepsilon := \min\{J : M(J) \geq 1 - \varepsilon\}$ stabilises at $g_{\text{eff}}(10^{-12}) \approx 5\text{--}6$ (order one, λ -independent), giving $R_0/g_{\text{eff}} \approx 5 \times 10^{-7}$. The earlier fixed-tolerance gap was unstable (cut inside the genuine cluster); the mass-based replacement is stable and macroscopic. See Zookeeper C2by/C2bz for the full analysis [9].

4 Spectral Convergence Constants

Both spectral branches produce characteristic constants whose identification and saturation provide evidence for the conditional reductions. We collect them here for comparison; the cross-route synthesis with Branch A's constant $C^* = 64\pi^2/3$ (Part II [7], § GF1) is deferred to Part IV [6].

4.1 Hadamard Curvature Constant c_Ξ (RH-neutral)

The Hadamard product for the completed Riemann ξ -function gives

$$\log \frac{|\Xi(it)|}{|\Xi(0)|} = \sum_{\rho} \log \left| 1 - \frac{-t^2}{(\rho - \frac{1}{2})^2} \right| \sim -\frac{t^2}{2} \sum_{\rho} \frac{1}{(\rho - \frac{1}{2})^2} + O(t^4),$$

whence

$$c_\Xi = -\frac{1}{2} \sum_{\rho} \frac{1}{(\rho - \frac{1}{2})^2} \approx 0.0231. \quad (6)$$

This is the leading quadratic curvature of $\log A_\Xi(\delta)$ near $\delta = 0$. The representation (6) sums over *all* non-trivial zeros and is *RH-neutral*. Under RH, $\rho = \frac{1}{2} + i\gamma_k$ and the sum specialises to $\sum_{k \geq 1} 1/\gamma_k^2$.

4.2 Residual Curvature κ_∞ (Branch B)

The multiplicative canonical form (3) decomposes A_λ into the Hadamard baseline A_Ξ and a residual factor $1 + \kappa(\lambda)\delta^2$. Empirically, $\kappa(\lambda)$ saturates to $\kappa_\infty \in [1.6, 1.8] \times 10^{-3}$ across $\lambda = 3, 5, 7, 9, 11, 13, 15, 20$, with $\kappa(20) = 1.527 \times 10^{-3}$. This saturation is consistent with the Foundation-Lock convergence: as $\lambda \rightarrow \infty$, $\hat{\xi}_\lambda \rightarrow \Xi$ in the relevant topology, and the residual curvature measures the finite- λ correction.

4.3 Hecke-Normalised Gap $g_0(m)$ (Branch Z)

For the m -th Riemann zero γ_m , the classical Montgomery–Odlyzko pair-correlation heuristic [10] suggests a local mean zero spacing of $2\pi/\log \gamma_m$. The effective microcluster gap g_{eff} of Branch Z (Section 3.3) is consistent with this scaling: $g_0(m) \approx 2\pi/\log \gamma_m$.

Branch	Constant	Value	Status
B	c_Ξ (Hadamard curvature)	≈ 0.0231	analytical (RH-neutral)
B	κ_∞ (residual curvature)	$\in [1.6, 1.8] \times 10^{-3}$	empirical (saturated)
B	$B_\Xi = \xi(0)/\xi(1/2)$	$1.00579\dots$	analytical
Z	$g_0(m) \approx 2\pi/\log \gamma_m$	varies	classical (Montgomery–Odlyzko)

5 Common Wall and Spectral-Sector Outlook

The three branches share Wall (ii) as the common analytical obstacle, but attack it through structurally different mechanisms. The following comparison table summarises the conditional reductions.

Branch	Reduction chain	Open hypothesis (the wall here)	Status
A — Even Dominance	A1 + A2 + A7 $\Rightarrow$ RH	A7 (Asymp. Variational Gap Conj.) + A3 (Odd-sector lower bound)	conditional reduction
B — I-1 Normal-Family	B1 + B5 + B9 $\Rightarrow$ RH	(ii-a) (uniform strip bound from QW_N -coercivity)	conditional reduction
Z — CCM Microcluster	B1 + C2ca + fixed-waterline $\Rightarrow$ RH	C2ca.1 + C2ca.2 + C2ca.5 + $g_* \geq g_0$	conditional reduction

Shared inputs. A1 = B1 (both rest on Connes–vS Thm 5.6 [3]). The Foundation-Lock $\hat{\xi}(z) = (2/\sqrt{L}) \sin(zL/2)F(z)$ from Branch B is the canonical criterion object used in Branch Z as well.

Distinct mechanisms. Branch A attacks Wall (ii) via eigenvalue ordering $\lambda_1^+ < \lambda_1^-$ (Part II [7]). Branch B attacks via a strip-uniform Hurwitz-style normal-family argument (this paper). Branch Z attacks via a microcluster quasimode-to-projector estimate (this paper + Zookeeper [9]).

Closest to conditional closure. Branch B currently has the smallest open surface (a single named hypothesis (ii-a)) together with strongly sub-generic empirical evidence ($A_\lambda(0.4) \approx 1.004$ flat over seven converged λ -points, Theorem B-1 giving $A_\lambda \leq 1.0063$ under the log-curvature hypothesis). Branches A and Z remain conditional reductions in the multi-gap sense (two or three open inputs respectively). The title of the present series follows the conservative convention: “Atlas” applies as long as no single branch has reached conditional-proof level under the audit protocol of the project `DECISION.md`.

Promotion protocol and future scenarios are discussed in Part V (Section 9 gives a brief forward reference).

6 Excluded Spectral Paths

The post-v2.0 spectral programme (Connes–van Suijlekom reformulation, I-1 normal-family reframe, err-collapse diagnostic) closed several routes. These are documented here for the audit trail; they do not affect the main conditional reduction.

- **Eigenvalue-ordering via a uniform spectral gap.** The plan to lift $\lambda_{\min}(W^+ - W^-) < 0$ to the eigenvalue inequality $\lambda_1^+ < \lambda_1^-$ via Kato–Temple, Davis–Kahan, or Lehmann gap estimates fails because the relevant minimum eigenvector of QW_N sits in a near-degenerate cluster (even/odd split $\sim 10^{-65}$ at $\lambda = 5$, $\sim 10^{-84}$ at $\lambda = 7$, super-exponentially collapsing with λ). No uniform spectral gap is available.
- **Localization hypothesis for the strip bound (ii-a).** A natural plan was to bound $|\hat{\xi}_\lambda(x + i\delta)|$ via localization assuming ξ_λ is concentrated near $t = 0$. This is falsified: ξ_λ is genuinely edge-concentrated ($r_{99} \approx L/2$, $|\xi_\lambda(0)| \approx 0$). A growth-adapted normalisation Φ is required.
- **Edge-profile factorisation for (ii-a).** The exact identity

$$F_\lambda(z) = 2 \cos(zL/2) C_\lambda(z) + 2 \sin(zL/2) S_\lambda(z)$$

isolates the explicit λ^δ -growth modulation. The edge route yields no type gain because the eigenfunction mass is spread ($r_{50} \approx 1$ uniformly, but $r_{99} \approx L/2$).

- **Literal MS1 (cluster ladder).** The conservative form of MS1 — uniformly simple splitting of the minimum eigenvalue of QW_N at every (λ, N) — fails: the splitting collapses super-exponentially with λ . Only a cluster/quotient reformulation has structural traction.

- **Strip-amplification at $\text{dps} = 30$ (retracted).** Early measurements at $\text{dps} = 30$ suggested sub-generic growth and non-monotone instability. At converged precision ($\text{dps} \approx 30\text{--}50 \cdot \lambda$), the picture is qualitatively different: $A_\lambda(0.4) \approx 1.004$ is flat. The methodological lesson (dps requirement scales linearly with λ for cluster-near-degenerate eigenvectors) is itself substantive.
- **Raw- L^2 Hurwitz passage (falsified 2026-05-15).** The plan to derive the Hurwitz strip-uniform convergence directly from a raw- L^2 tail estimate fails on two fronts: (a) the archimedean tail has norm $O(\lambda^{-1/2})$ — exactly boundary exponent $\alpha = 0$, no margin; (b) the prime-Dirac tail diverges in the $m = 1$ branch. Negative-Sobolev does not rescue this. The diagnostic identifies operator-norm topology (i.e. Branch B) as the only viable replacement. The Connes–Hurwitz–Passage project re-scoped on 2026-05-17 to a purely expository role (Geiger2026hurwitz, in preparation).

Remark 6.1 (Common pattern). These six exclusions share a structural pattern: the minimum-eigenvector cluster of QW_N is the source of the technical difficulty. Methods that assume or require a uniform gap, a thin-localised eigenfunction, or a generic function-class behaviour all fail. Productive routes work *with* the cluster (waterline / cluster-quotient reformulations, growth-adapted normalisations) rather than against it.

7 Excluded CCM Paths (Branch Z Heritage)

The CCM microcluster programme accumulated its own catalogue of dead routes. They do not affect the Branch-Z conditional reduction.

- **Mass-based microcluster definition.** Defining the microcluster via a mass criterion is circular against MS2. Replaced in Zookeeper v1.2+ by a *spectral* microcluster (low-energy eigenspace bounded by an outer gap g_* of order one).
- **Naive Hurwitz passage.** A direct Hurwitz application without approximation-quality control yields statements too weak for the err-collapse closure. Replaced by a Slepian–Pollak-/PSWF-based passage.
- **Anchor-perturbation / Banach contraction (C2br', falsified).** A Banach-norm contraction $C_{\max} < M$ fails by factor 55–7000 on all twelve tested configurations.
- **Projected frontier cancellation $S_{\text{bulk}} \approx -S_{\text{bd}}$ (C2bs, falsified).** Both contributions have the same sign and are both of size $\sim 10^{-7}$. The true cancellation mechanism sits in the eigenvalue equation $(A - w_0)k \approx 0$.
- **Uniform mass criterion C2ak.** The total off-cluster mass bound diverges with λ ($6.68 \rightarrow 15.3 \rightarrow 1865$ across $\lambda = 3, 5, 7$) because it discards internal cancellation.
- **Canonical local geometry C2al (directional defect).** The geometric inequality has the wrong direction — the correct form is $g_R/d_m \leq T_m^\beta/\beta_0$ (lower bound on d/g , not upper).
- **Cluster selection via Δ -projection.** Simple δ_N -projection is not a robust cluster selector: at $\lambda = 5, N = 55$ `best_zero=k=16` while the other two select $k = 17$.

- **Log-concavity / detailed balance (C2cc H1/H2/H4).** The hypotheses that β_n or $|b_n|$ are log-concave (H1), that the spectral residue chain has a detailed-balance structure (H2), or that a Mertens-anchor applies (H4) were tested at $\lambda = 3, N = 30, \text{dps} = 50$ and falsified. Structural cause: arithmetic (prime-driven) oscillation.
- **Direct Frontier-Dominance v2.1 (Killer 0, 2026-05-14).** The v2.1 argument used a common Rayleigh test vector to prove $\lambda_{\min}(D_{\text{tot}}) < 0$ and then identified $\Delta_{\text{low}} = \lambda_{\min}(D_{\text{tot}})$ with $\Delta = \lambda_1^+ - \lambda_1^-$. The reviewer counterexample $A = \text{diag}(0, 10), B = \text{diag}(5, -5)$ shows this does *not* follow. Even dominance for asymptotic λ remains conditional on a separate odd-sector lower bound. (This item is Even-sector-specific and documented in detail in Part II [7]; it is listed here because the falsification also restructured the comparative analysis of Branches B and Z.)

Remark 7.1 (CCM-pattern: cancellation, not magnitude). The CCM-side dead routes share a different pattern from the spectral side: each tries to bound a quantity by its individual-term magnitude, and each is defeated by large internal cancellation. The living routes explicitly factor the cancellation through projector identities or scalar secular equations, never through simple norm domination.

8 Methodical Dead Ends

A small number of methodical routes were tested on the spectral side and ruled out on structural grounds.

- **Dobrushin–Zegarliniski uniqueness $\rightarrow$ RH-CAP.** Dobrushin–Zegarliniski uniqueness works only on lattices with localised Gibbs measures. RH has neither; the absence of multiplicative independence enforces a non-local prime-coupling structure that breaks the Dobrushin contraction setting.
- **Sturm node-counting (naive).** Classical Sturm–Liouville node-counting for the eigenfunctions of QW_N fails because the operator is finite-dimensional and the even/odd spectral gap is sub-exponentially small. Node-counting provides no information at the relevant precision scale.
- **Naive Abel summation for far-tail.** Applying Abel summation to the far-tail prime sums $\sum_{p>\lambda} (\log p)^k / p^s$ without controlling the arithmetic oscillation of $\theta(x) - x$ leads to divergent error terms. The PNT transfer lemma (Part II) shows how to do this correctly.
- **Naive superdeterminant form for Möbius- $\mathbb{Z}_2$.** The formal identity $\text{sdet}(I - L_s) = \det(I - L_s^+) / \det(I - L_s^-) = 1/\zeta(s)$ is algebraically correct for finite prime sets but does not converge for the full prime set at $\Re(s) = \frac{1}{2}$ (trace-class obstruction K4 from Part I).
- **Friedrichs + Petersson as Hilbert–Pólya lever.** The plan to use Friedrichs extension of the Petersson inner product to construct a self-adjoint realisation whose spectrum encodes the Riemann zeros fails: the relevant bilinear form is not semi-bounded on the domain needed to capture the critical-line zeros.

9 Summary

Contribution	Status	Location
I-1 Normal-Family Reframe: (ii) $\rightarrow$ (ii-a) + (ii-c)	rigorous	Section 2.2
(ii-c) absorbed into (ii-a) under (Z)	rigorous	Section 2.2
Foundation-Lock: canonical criterion object	rigorous	Section 2.2
Converged empirics: $A_\lambda(0.4) \approx 1.004$ (7 points)	numerical evidence	Table 1
Multiplicative canonical form	empirical (0.012% fit)	Section 2.4
Theorem B-1: $A_\lambda \leq 1.0063$ (conditional)	conditional	Theorem 2.3
c_Ξ Hadamard curvature (RH-neutral)	analytical	Section 4
CCM: three C2ca-inputs identified	conditional reduction	Section 3.2
CCM: effective microcluster gap $g_{\text{eff}} \approx 5\text{--}6$	numerical evidence	Section 3.3
6 excluded spectral paths + 9 excluded CCM paths	documented	Sections 6 and 7
5 methodical dead ends	documented	Section 8

The spectral sector contributes two independent conditional reductions of RH, each with a well-characterised open surface. Branch B has the smallest open surface (a single hypothesis (ii-a)) and the strongest empirical support ($A_\lambda \leq 1.0063$ conditional on uniform log-curvature control). The cross-route synthesis across all three branches is developed in Part IV [6]; the condensed atlas and outlook in Part V [5].

Acknowledgments. Computations were performed on a Mac Studio (Apple M1 Max, 32 GB). The interval arithmetic library mpmath was developed by Fredrik Johansson.

AI Disclosure

This manuscript was developed in extensive collaboration with AI language models (Claude, Anthropic; Copilot, Microsoft/GitHub; Gemini, Google DeepMind). AI contributed to conceptual exploration, analytical work, literature review, writing, and critical review. The author, Lukas Geiger, conceived the research direction, provided domain expertise, and exercised final editorial authority. The author assumes full and sole scholarly responsibility for all content, including any errors.

References

- [1] Alain Connes. The Riemann hypothesis: Past, present and a letter through time. *arXiv preprint*, 2026. Survey on the Riemann Hypothesis: past, present, and a letter to Riemann using only pre-1860 mathematics.
- [2] Alain Connes, Caterina Consani, and Henri Moscovici. Zeta zeros and prolate wave operators. *Annals of Functional Analysis*, 15(87), 2024. Jacobi matrix representation and prolate–zeta connection.
- [3] Alain Connes and Walter D. van Suijlekom. Quadratic forms, real zeros and echoes of the spectral action. *Communications in Mathematical Physics*, 406(12), 2025. Quadratic forms, real zeros, and connections to the spectral action.

- [4] Lukas Geiger. A conditional reduction of the Riemann hypothesis to even dominance of the Weil quadratic form. Zenodo, 2026. Standalone extract from the Landscape series; v1.8 record DOI 10.5281/zenodo.20396370, concept DOI 10.5281/zenodo.19764771.
- [5] Lukas Geiger. The Riemann hypothesis: Conclusio. Preprint, 2026. Part V of a five-part series (in preparation).
- [6] Lukas Geiger. The Riemann hypothesis: Cross-route synthesis. Preprint, 2026. Part IV of a five-part series (in preparation).
- [7] Lukas Geiger. The Riemann hypothesis: Even dominance of the Weil quadratic form. Preprint, 2026. Part II of a five-part series.
- [8] Lukas Geiger. The Riemann hypothesis: Foundations, obstructions, and reorientation. Preprint, 2026. Part I of a five-part series.
- [9] Lukas Geiger. The spectral zookeeper: A conditional reduction toward the Riemann hypothesis via spectral microcluster closure in the Connes–Consani–Moscovici framework. Zenodo, 2026. Companion CCM-route paper; v1.4 record DOI 10.5281/zenodo.20396399, concept DOI 10.5281/zenodo.19673126.
- [10] Hugh L. Montgomery. The pair correlation of zeros of the zeta function. *Proceedings of Symposia in Pure Mathematics*, 24:181–193, 1973.

The Riemann Hypothesis: Cross-Route Synthesis

Part IV of a Five-Part Series

Lukas Geiger

Bernau, Germany

ORCID: [0009-0005-7296-1534](https://orcid.org/0009-0005-7296-1534)

May 31, 2026

Abstract

This paper, Part IV of a five-part series [6, 5, 7, 4], collects three classes of material that do not fit the scope of Parts II (Even Dominance) or III (Spectral Routes): the cross-route synthesis comparing all three active branches A, B, Z; the dormant routes and external lines of the wider landscape; and a self-correcting record of hypothesis corrections during the iterative research process.

The cross-route synthesis identifies a shared Galerkin pattern across branches — each conditional reduction involves a fixed-parameter positive structure that must be promoted to a uniform-limit statement via a growing parameter $N(\lambda)$, $d(\lambda)$, or $m(\lambda)$ — and tabulates the characteristic constants ($C^* = 64\pi^2/3$ as the current Branch A boundary-moment benchmark for $\lambda = 5, 7$, $\kappa_\infty \in [1.6, 1.8] \times 10^{-3}$ for Branch B, $g_0(m) \approx 2\pi/\log \gamma_m$ for Branch Z).

Series status (v3.0, 2026-05-26; criticality-lens update 2026-05-31). The series is restructured from a trilogy to a five-part series; the conditional reduction status from v2.2 is unchanged. A 2026 working phase adds Section 3: a *heuristic* criticality lens ($\text{RH} \Leftrightarrow \Lambda = 0$) that strategically re-weights the routes, a sharper localization of the common wall (the analytic continuation across $\Re s = 1$), a computable diagnostic built from the *classical* Laguerre–Pólya/Herglotz criterion, and four honest closures of *concrete constructions* (not routes, and several of them self-corrections of internally proposed constructions). None of these closes RH, refutes a route, or changes the conditional-reduction status.

MSC 2020: 11M26, 11M36, 47A75

Keywords: Riemann Hypothesis, cross-route synthesis, Galerkin pattern, dormant routes, hypothesis corrections, landscape survey

Contents

1	Introduction	3
2	Cross-Route Synthesis	3
2.1	Characteristic Constants	3
2.2	The Shared Galerkin Pattern	3
2.3	Two Isolated Proof Slots	4
3	The Sharpened Common Wall: A Criticality Lens and an Analytic-Continuation Localization	5
3.1	A criticality lens (heuristic, not a theorem)	5
3.2	The common wall: analytic continuation across $\Re s = 1$	5
3.3	A computable diagnostic from the classical Laguerre–Pólya/Herglotz criterion	6
3.4	Four closures of concrete constructions (not routes)	6
4	Dormant Routes	7
5	External Lines	8
6	Systematically Explored Alternatives	9
7	Independent Results	9
8	Hypothesis Corrections in the Iterative Process	10
8.1	GF1 Correction Chain (Iterations 15–29)	11
8.2	B-1 Corrections and c_{Ξ} RH-Neutral Form	11
8.3	The Relative-Determinant Bypass: Proposed and Refuted (2026)	11
8.4	Value of the Correction Record	11
9	Summary	12

1 Introduction

This paper, Part IV of a five-part series, serves as the “miscellaneous and combined” component of the landscape atlas:

Part I [6]: Foundations and Obstructions.

Part II [5]: Even Dominance (Branch A).

Part III [7]: Spectral Routes (Branches B and Z).

Part IV (this paper): Cross-Route Synthesis, dormant routes, external lines, methodical dead ends, hypothesis corrections.

Part V [4]: Conclusio.

The three active branches A (Even Dominance), B (Hadamard Strip / I-1 Normal-Family), and Z (CCM Microcluster) are developed in detail in Parts II and III. This paper identifies the structural patterns they share, surveys the wider landscape of dormant and external routes, and provides a transparent record of hypothesis corrections during the iterative development.

2 Cross-Route Synthesis

2.1 Characteristic Constants

Each active branch produces a characteristic constant or scaling law whose identification anchors the conditional reduction:

Branch	Constant	Scaling law	Theoretical identification
A (Even Dom.)	$C^* = 64\pi^2/3 \approx 210.55$	sLI $\sim C^* \lambda/(NL)$ for the $\lambda = 5, 7$ benchmark; $\lambda = 3$ separate and $\lambda = 11$ open	$16 \int_0^\infty t^2/\sinh^2(t/2) dt$ (reverse-fit candidate; exact collar kernel known)
B (Hadamard)	$\kappa_\infty \in [1.6, 1.8] \times 10^{-3}$	saturates with λ	finite-part residual curvature of Foundation-Lock
Z (CCM)	$g_0(m) \approx 2\pi/\log \gamma_m$	log-decreasing	Montgomery–Odlyzko (classical)

Remark 2.1 (Identification status). The constant $C^* = 64\pi^2/3$ for Branch A is an identification candidate for the $\lambda = 5, 7$ benchmark. The 27 May audit shows that the simple endpoint-kernel factorisation is false; the exact collar kernel is $K(t) = e^t \cosh(t)/(2 \sinh^2 t)$, so a rigorous derivation must pass through the eigenmode asymptotics. The $\lambda = 3$ sweep is recorded as a low-bandwidth outlier, while the single $\lambda = 11, N = 200$ point ($C_{\text{eff}} \approx 168$) remains an open discriminator (Part II [5], Remark on derivation candidate). The constant κ_∞ for Branch B is established empirically in the conservative range $[1.6, 1.8] \times 10^{-3}$ (Part III [7]). The gap $g_0(m)$ for Branch Z rests on classical pair-correlation heuristics.

2.2 The Shared Galerkin Pattern

Despite their different analytical mechanisms, all three branches share a common structural pattern:

1. **Fixed-parameter positive structure:** At each finite (λ, N) (or (λ, m) for Branch Z), the relevant quantity is computable and has the correct sign/bound.

2. **Uniform-limit requirement:** The conditional reduction requires promoting the fixed-parameter statement to a uniform assertion as $\lambda \rightarrow \infty$ with a parameter $(N(\lambda), d(\lambda), m(\lambda))$ growing alongside.
3. **Transient vs. converged regime:** The warm-up parameter is branch-dependent. For the GF1/ C^* branch the current diagnostic scale is $\xi = NL/\lambda$, not N/L : the $\lambda = 5, 7$ rows enter the benchmark window near $\xi \approx 50$, whereas the $\lambda = 11, N = 200$ point remains an open pre-convergence discriminator. Other Galerkin branches still track their own resolution ratios.

This pattern is consistent with a *Bernstein–Galerkin conjecture*: the Galerkin truncation of the Weil quadratic form at dimension N converges to the infinite-dimensional form in the relevant operator topology, and the convergence is uniform in λ after a λ -dependent warm-up phase.

Remark 2.2 (Status of the conjecture). The Bernstein–Galerkin pattern is an empirical observation from 5–7 data points per branch. It has not been formalised as a mathematical statement. Its role is heuristic: it provides a mechanistic explanation for why all three branches have similar structural gaps (fixed-parameter $\checkmark$, uniform-limit open).

2.3 Two Isolated Proof Slots

The cross-route analysis identifies two concrete, isolated proof slots that, if resolved, would close one or more branches:

1. **B-1: Uniform log-curvature control from Foundation-Lock.** If $H_\lambda(\delta) \leq K\delta^2$ with $K \leq 1.8 \times 10^{-3}$ is derived from the Foundation-Lock structure, then Theorem B-1 (Part III [7]) gives $A_\lambda \leq 1.0063$, closing Branch B.
2. **GF1: Eigenmode asymptotics + factor-16 identification.** If the exact collar kernel and the $q_a(k)$ eigenmode asymptotics rigorously produce the factor 16 in $C^* = 16 \cdot 4\pi^2/3$, and the $\lambda = 11$ discriminator does not persist as a separate branch, then Theorem GF1 (Part II [5]) provides a boundary-moment scaling benchmark for the non-exceptional Branch A regime; the $\lambda = 3$ branch remains a conservative guardrail.

These two slots are *independent*: closing one does not affect the other. Each is a well-posed analytical problem amenable to standard tools (complex analysis for B-1; Fourier analysis and spectral theory for GF1).

Remark 2.3 (Both slots are margin-type — a strategic caveat). Both slots are *margin/coercivity* statements: a uniform bound that closes a gap. They remain well-posed open problems, and nothing below refutes them. Section 3 adds a *heuristic* criticality lens: since RH (if true) sits at the critical value $\Lambda = 0$ with no buffer (Section 3.1), exact-structural mechanisms (identity, symmetry, spectral realization) deserve *additional* strategic weight alongside these margin-type slots. This is a re-reading of priorities, not a theorem and not a refutation (see the caveat in Theorem 3.1).

3 The Sharpened Common Wall: A Criticality Lens and an Analytic-Continuation Localization

A 2026 working phase re-examined all branches through one question — *what kind of argument can close the gap?* — and localized the shared obstruction more sharply. Everything in this section is a **strategic lens and a diagnostic localization**, not a new closure of RH and not a refutation of any route. Several items are honest self-corrections of internally proposed constructions. The conditional-reduction status of Parts II/III is unchanged.

3.1 A criticality lens (heuristic, not a theorem)

The de Bruijn–Newman constant Λ satisfies $\Lambda \leq 0 \Leftrightarrow \text{RH}$; Rodgers and Tao (*Ann. of Math.* 2020) proved $\Lambda \geq 0$. Hence $\text{RH} \Leftrightarrow \Lambda = 0$: if RH holds, it sits *exactly* at the critical value, with no margin to spare.

Remark 3.1 (Criticality lens — and its essential caveat). This invites a *strategic* dichotomy of proof mechanisms: *margin-seeking* (a uniform gap / coercivity / smallness bound) versus *exact-structural* (an identity, a symmetry, or a spectral realization that forces the conclusion without slack). The heuristic reading is that a mechanism producing a strict uniform margin would be establishing something with a buffer, which sits uneasily with a critical truth at $\Lambda = 0$. **Caveat (essential, to avoid over-claiming)**. The step “a uniform margin would force $\Lambda < 0$ ” is an *analogy*, not a theorem. The branch-specific margins — the even-dominance gap $\lambda_1^+ < \lambda_1^-$, the Hadamard constant κ , the CCM residual smallness — are *not* shown to equal a de Bruijn–Newman statement; Λ and these quantities are different objects. Consequently the lens does *not* refute the well-posed proof slots of Section 2 (Theorem 2.3); it only suggests weighting exact-structural routes *in addition*. It is a re-reading of priorities, not a result.

Under this lens the routes re-sort by *type*, not by viability: the margin-type branches A (A7/AVG gap), B ((ii-a) coercivity), Z (residual smallness) keep their exact kernels as assets, while the exact-structural lines — M (Bost–Connes $\leftrightarrow$ Meyer, with the KMS phase transition at $\beta = 1$ playing the role of built-in criticality) and the global adelic trace route (Connes–Consani scaling site / $\mathbb{F}_1$, where the only missing datum is the Hodge-index positivity of an otherwise-available Lorentzian form) — receive additional strategic weight. The empirical signature is consistent: the even-dominance margin grows like a power law $\sim \lambda^{0.72}$ (a criticality signature, not a buffer). This re-weighting is recorded in the project decision log; it does not alter any route’s formal status.

3.2 The common wall: analytic continuation across $\Re s = 1$

Independently of the lens, the correctly-typed routes share a single concrete endpoint. Three lines of work — the Möbius prime-hub, the Wiener–Hopf factorization, and the Sonin/Toeplitz (Connes–Consani) construction — are each *locally/finitely* controllable but break at *exactly* the same place: the global limit, where the local data must assemble into ζ . The unified description is the Weyl–Titchmarsh function M_∞ of the unbounded limit operator; the missing datum is uniform control of that limit from local data *without* importing ζ .

Remark 3.2 (The wall is a classical fact, here sharply localized). A finite Euler product $\prod_{p \in P} (1 - p^{-s})^{-1}$ has, after the substitution $z = -i(s - \frac{1}{2})$, a pole comb in the upper half-plane at $z_{p,k} = 2\pi k / \log p + i/2$ (on $\Im z = \frac{1}{2}$), with a common pole at $z = i/2$ of order $|P|$; hence its logarithmic derivative is not a Herglotz function. The comb is removed only by ζ 's analytic continuation ($\zeta(0)$ finite; the pole at $s = 1$). This is **Euler-product folklore** (the product converges only for $\Re s > 1$). It is not a new obstruction: it is a *precise localization* of the known wall — the missing datum is the analytic continuation across $\Re s = 1$, i.e. a global/functional-equation phenomenon that finite arithmetic-plus-archimedean data structurally do not supply.

3.3 A computable diagnostic from the classical Laguerre–Pólya/Herglotz criterion

Write $\Xi(z) = \xi(\frac{1}{2} + iz)$ and $M_\Xi = -\Xi'/\Xi$. The chain

$$\text{RH} \iff \Xi \in \text{Laguerre-Pólya} \iff M_\Xi \text{ Herglotz} \iff \frac{(M_\Xi(z) - \overline{M_\Xi(w)})}{z - \bar{w}} \succeq 0$$

is **classical** (Hermite–Biehler; the Laguerre–Pólya characterization of real-rooted entire functions; the positivity $\Re \xi'/\xi > 0$ for $\Re s > \frac{1}{2}$ is Voros' criterion). *We claim no novelty for the criterion itself.* The contribution is a *computable, non-circular test protocol* that turns this classical equivalence into a diagnostic instrument: an Euler-anchor giving M_Ξ from ξ'/ξ in the convergent region $\Re w > 1$ (verified numerically to $\sim 10^{-31}$), a ladder of gates (normalization; tail-tightness; a Stieltjes identity; an off-axis Poisson tomography resolving individual residues), and a detector lemma by which an off-line zero forces a negative 1×1 Pick minor just below the migrated pole. The instrument is empirically validated: it detects off-line zeros and discriminates the true zero spectrum from a smooth density-matched control.

Remark 3.3 (What the diagnostic established on a genuine operator). An independent, non-circular discretization of the Connes–Moscovici prolate operator (arXiv:2112.05500) was run through the full instrument. It reproduces Riemann's zero-counting function $N(E)$ to $\sim 3\%$ in the UV — the *archimedean half* (the smooth density, leading term $\frac{1}{2} \log(y/2\pi)$), an independent non-circular confirmation of the CCM density claim. But its poles sit $O(1)$ off the zeros, so the arithmetic fine-structure (the *positions*) is not captured; and the relative-determinant / Krein construction does *not* remove the Euler pole comb, because the prolate background is itself Herglotz (comb-free), leaving nothing to cancel. Both findings re-confirm Theorem 3.2: density yes, positions no; the wall is the analytic continuation. The genuine positive by-product — that the determinant channel $-D'/D$ (residues = multiplicities), not mere eigenvalue convergence, is the object that must converge ($\det_{\text{reg}}(A - z) \rightarrow C \cdot \Xi$) — sharpens the CCM frontier and is recorded among the independent results (Section 7).

3.4 Four closures of concrete constructions (not routes)

The phase produced four genuine negative results, each over a *specific constructed object or technique*. None closes a whole route; each re-types the route or names a successor.

1. **Osterwalder–Schrader reflection, natural form.** The reflection block $Q_+(x) = \int_0^\infty \int_0^\infty x(t) K(t+s) x(s)$ is indefinite, because the prime distributions $\delta(t+s - \log p)$

in the Weil kernel meet the half-axis $t + s = \log p$; the Toeplitz alternative $K(t - s)$ collapses to global Weil positivity = RH. The OS route in its most natural form (inversion reflection on $\{|g| \geq 1\}$) is closed-negative; another half-space choice would be needed.

2. **Möbius $\mathbb{Z}_2$ prime hub, naïve L_s^- .** An obstruction theorem: any continuation of the divergent prime traces $\text{Tr } T_s = P(s)$, $\text{Tr } T_s^2 = P(2s)$ down to $\Re s = \frac{1}{2}$ imports ζ and its zeros via $P(s) = \sum_n \mu(n)/n \cdot \log \zeta(ns)$ — hence circular. (Project-internal; not a literature-verified standard.)
3. **Toeplitz-symbol sub-method (Branch A).** $A_{\sin}$ is not Toeplitz (m -dependent oscillating diagonals $\Rightarrow$ no well-defined symbol $f(\theta)$); a structural defect of the constructed object, not merely “the method equals the open problem”. Successor: the finite Sylvester-inertia bound (done) plus GLT/outlier control.
4. **Relative-determinant / Krein bypass.** The relative determinant $-\partial_z \log[\det(A - z)/\det(A^0 - z)]$ does *not* remove the Euler pole comb: if A^0 is the prolate background it is itself comb-free (nothing to cancel, and $\det(\text{Euler} \cdot \text{prolate})/\det(\text{prolate}) = \det(\text{Euler})$ keeps the comb); the only comb-free alternative needs a self-adjoint prime factor, which is circular (= RH). This is an honest self-correction of an internally proposed “right technology”; the proposed-then-refuted arc is recorded in Section 8.

Crucially, the *determinant channel* $-D'/D$ as a *sharpening* of the CCM frontier remains valid (Section 7); only the *relative determinant* as a *bypass* of the continuation wall is refuted. The two must not be conflated.

4 Dormant Routes

A small number of constructed routes are dormant rather than dead: they are explicit reductions or partial constructions that have not been refuted, but currently lack active development. The internal register (AKTIONSPLAN.md) keeps the full inventory; the four below are summarised here because each has a direct bridge to one of the active branches.

- **D-CCM (Dirichlet–CCM operator $D_{\log}^{(\lambda, N, \chi)}$).** A direct χ -twist of the CCM rank- N perturbation construction. Addresses GRH, not RH directly — the trivial character χ_0 (Riemann case) is in this scaling regime not practically testable. Companion of Branch Z.
- **D-Twist (χ -twisted defect norm).** Diagnostic blueprint with completed Pilot v2 and Milestone 2_χ work; reformulated as λ -asymptotic statement. First tests at χ_{21}, χ_{33} ; Riemann unreachable in this scale.
- **D-PW (Paley–Wiener sector difference $\psi_\chi = \phi_\chi^+ - \phi_\chi^-$).** Classification of the five Paley–Wiener steps complete: S1, S2, S4 mechanical; S3, S5 character-specific. Continuation of the published Dirichlet Character Atlas v2 (Zenodo 20241612).
- **P-M (Möbius-graded $\mathbb{Z}_2$ prime hub).** Conditional lemma proved for finite prime set (Möbius identity exact); full prime set converges only for $\Re(s) > 1$; $\mathbb{Z}_2$ -sector statement for non-trivial zeros open. Structurally appealing alternative

interpretation of the Prime Orbit Axiom (Part I, Definition 2.1). A 2026 obstruction theorem sharpens this (Section 3.4): no non-circular L_s^- exists, since the divergent prime traces continue to $\Re s = \frac{1}{2}$ only by importing ζ and its zeros. This re-types P-M as the *prime half* of a P-M+M fusion — it provably lacks the archimedean-modular half ($s \leftrightarrow 1 - s$) that route M supplies (project-internal).

The remaining dormant entries (P-A, P-C, P-B, M, R, H) are not listed individually here; see `AKTIONSPLAN.md` for the full inventory. Two are *strategically elevated* by the criticality lens of Section 3.1: **M** (Bost–Connes $\leftrightarrow$ Meyer) and the global adelic trace route **P-A** (Connes–Consani scaling site) are of the exact-structural type and now receive additional strategic weight; **P-C** is co-resolved by an active branch; **P-B** (Berry–Keating) is of the right type but unconstructed (self-adjointness is the missing Hilbert–Pólya theorem); **R**, **H** are framework/expository.

5 External Lines

The landscape sits inside a wider research ecosystem. The following lines are cited as anchors; only the directly proof-relevant items appear here.

- **Connes 2026 Mainline** [1] (arXiv:2602.04022). Re-frames the Riemann zeros as the infrared spectrum of a self-adjoint operator and identifies the key open hypothesis MS2 (§6.6(b)). Branches A, B, Z all condition on this line.
- **Connes–van Suijlekom 2025** [3] (arXiv:2511.23257). Theorem 5.6: 1-dim even kernel of the truncated Weil form implies real zeros of $\hat{\eta}$. This is A1 = B1 in the present landscape.
- **Connes–Consani–Moscovici 2025** [2] (arXiv:2511.22755). Foundation of the rank-one perturbation $D_{\log}^{(\lambda, N)}$. Anchors Branch B’s Foundation-Lock and Branch Z’s microcluster construction.
- **Groskin 2026** (arXiv:2605.20224). Independent implementation of the CCM Galerkin matrix at sixteen cutoffs $c = 13 \dots 67, 100$ at high precision (Aitken extrapolation, $\sim 10^{-334}$ depth). Acts as external validator for branches B and Z.
- **Śliwiński 2026** (arXiv:2601.12133). Counter-regime: in the $\lambda = N$ configuration with low precision, proves an inverse-logarithmic dissonance lower bound $\epsilon(\lambda, N) \geq 1/(4 \ln \lambda)$. Does not bound the fixed- λ , $N \gg \lambda$, high-precision setup of this series; stands as a warning against naïve $\lambda = N$ limits.
- **Watson 2026** (arXiv:2602.01248), “The Riemann Ξ -function from primitive Markovian cycles I.” Constructs a Laguerre–Pólya function Ψ from finite reversible Markov dynamics, “entirely Archimedean, using no Euler products, primes, or arithmetic spectral input”; not an RH proof — it isolates “a single remaining analytic problem relating Ψ to $\Xi(2 \cdot)$.” An independent triangulation of the common wall: from the archimedean side it reaches the same *identification* step (relating the constructed object to Ξ) that the routes of this series reach from the arithmetic side (Section 3.2).

- **Franchi Viceré 2026** (Zenodo 19546495). Claims unconditional RH proof via apeirohedron / semi-local spectral descent. Surveyed and found not viable: the stabilisation argument confuses pointwise with strip-uniform convergence; the Apeirohedron limit step is not rigorous; boundary conditions do not replace the even-sector / MS1 input. See `AKTIONSPLAN.md §Klasse 8`.

6 Systematically Explored Alternatives

Eleven alternative strategies (BI-1 through BI-11) were systematically explored during development. This exhaustive search is itself a meta-result: the solution space is mapped.

Approach	Core idea	Result	Status
BI-1: $\det_2$ / Carleman	Extend to Hilbert–Schmidt class	No positivity substitute	Refuted
BI-2: RG fixed point	Weil positivity under scaling	Not formalizable	Discarded
BI-3: Universal property	Physical plausibility $\Rightarrow$ RH	Too abstract	Discarded
BI-5: Lifting programs	Weil $\rightarrow$ Grothendieck $\rightarrow$ RMT	All open except $\mathbb{F}_q$	Not usable
BI-6: Implication ring	$\text{Li} \leftrightarrow \text{Weil} \leftrightarrow \text{NB} \leftrightarrow \text{BD}$	Ring does not exist	Refuted
BI-7: Pöschl–Teller scattering	Gamma structure analog to ξ	No off-line stability	Negative
BI-8: Second-order stability	RH as variational problem	Architecture correct	Realized
BI-9: Soliton models	Topological stability	Coulomb gas unstable	Refuted
BI-10: Building blocks	No-Coordination, shift structure	Successfully reused	Integrated
BI-11: Connes 2026	Even dominance via Q_W	Breakthrough	Core path

The chosen path (BI-8 + BI-10 + BI-11) combines the variational architecture, the arithmetic building blocks, and Connes’ spectral framework. Routes A and B (total positivity, commuting operator) were ruling-outs that are structural discoveries: even dominance holds *despite* their absence, which clarifies that the mechanism is genuinely arithmetic-statistical rather than algebraic (Part II [5]).

7 Independent Results

Several results from this program are of independent interest, regardless of the main proof’s status:

Result	Context	Significance	Status
<i>Positive results</i>			
Shift Parity Lemma	Arithmetic trigonometry	Every prime favors even eigenfunctions	Proved
No-Coordination Lemma	Multiplicative independence	Arithmetic substitute for entropy term	Proved
Leading-Mode Cancellation ($c = 2 + \sqrt{2}$)	Resolvent analysis	Universal pairwise cancellation constant	Proved
Euler–Maclaurin: $\rho^{\text{EM}} < 0$	PNT asymptotics	Global stability anchor	Proved
33 finite-range CAP diagnostics	Interval arithmetic	Conditional bridge	Evidence
Concentration lemma	Even-dominance image measure	No fixed prime set controls a positive share of prime mass	Proved
Finite Sylvester inertia ($\nu_- = 0$)	Interval arithmetic, $\lambda = 100$	Finite even dominance rigorous, no symbol needed	Verified
Residue-calibration channel	Weyl-function diagnostic	Eigenvalue conv. $\neq$ Weyl conv.; channel is $-D'/D$, $\det_{\text{reg}} \rightarrow C\Xi$	Diagnostic
<i>Negative results</i>			
Obstructions K1–K4	Dead-end analysis	Maps impossible strategies	Documented
No absolute PF_∞ for A_λ	Prime shift multiplier	$M(\xi)$ negative on $\sim 49\%$ of domain	Proved
No universal commuting operator	Sturm–Liouville search	Kernel of $[T, D_N(r)]$ is $\{c \cdot I\}$	Proved
<i>Structural insights</i>			
Inverted-Hessian structure	YM/TU/RH comparison	RH is a maximum, not minimum	Confirmed
Cross-disciplinary catalog	Regularization	11 fields share trace-class obstruction	Documented

The negative results protect future researchers from repeating dead-end strategies. The Shift Parity Lemma, the obstruction analysis (K1–K4), and the Mellin decay barrier are publishable independently.

8 Hypothesis Corrections in the Iterative Process

The iterative development produced several hypothesis corrections that are documented here as a record of honest research practice (Sammler-Disziplin).

8.1 GF1 Correction Chain (Iterations 15–29)

Iter	Hypothesis	Status
15	$C_m \approx 8$ universal	falsified (boundary effect)
18	$C_m(\lambda = 5) \approx 4.3$	falsified (no plateau)
19	$\text{sLI} \sim 1/N$	confirmed, but without λ -scaling
20	$N \cdot \text{sLI} \approx 329$ universal	falsified
27	$\text{sLI} = \mathbf{C}^* \lambda / (\mathbf{NL}) + \mathbf{o}(\cdot)$	confirmed for $\lambda = 5, 7$, $C^* = 64\pi^2/3$
28	$\lambda = 3$ all- λ check	low-bandwidth outlier; two-regime guardrail
29	$\lambda = 11$ discriminator and endpoint-kernel audit	$C_{\text{eff}} \approx 168$ at $N = 200$; simple endpoint factorisation falsified

8.2 B-1 Corrections and c_{Ξ} RH-Neutral Form

The Hadamard constant c_{Ξ} was initially written in the RH-specialised form $\sum_{k \geq 1} 1/\gamma_k^2$ (sum over imaginary parts of zeros on the critical line). An external audit corrected this to the RH-neutral form $c_{\Xi} = -\frac{1}{2} \sum_{\rho} 1/(\rho - \frac{1}{2})^2$ (sum over all non-trivial zeros). This is a Sammler-Disziplin correction: the RH-neutral form is mathematically more general and does not assume what is being proved.

8.3 The Relative-Determinant Bypass: Proposed and Refuted (2026)

During the 2026 phase, the relative determinant (Krein spectral shift) $-\partial_z \log[\det(A - z)/\det(A^0 - z)]$ was proposed internally as the “right technology” to remove the upper-half-plane Euler pole comb and so bypass the continuation wall (Theorem 3.2). A subsequent test on the genuine Connes–Moscovici prolate background refuted this as a bypass: the prolate A^0 is itself Herglotz (comb-free), so there is nothing to cancel ($\det(\text{Euler} \cdot \text{prolate})/\det(\text{prolate}) = \det(\text{Euler})$ retains the comb), and the only comb-free alternative needs a self-adjoint prime factor, which is circular (= RH). The proposal is withdrawn as a bypass. The *separate* finding — that the determinant channel $-D'/D$, not mere eigenvalue convergence, is the correct object to converge ($\det_{\text{reg}}(A - z) \rightarrow C \Xi$) — remains valid (Section 7); the two must not be conflated.

A second correction: an internal phase lemma $\Theta_{\sigma} = \xi(\sigma - iz)/\xi(\sigma + iz)$ was sign-reversed (its poles come from the denominator, so under RH it *has* poles); the correct object is the reciprocal phase $\hat{\Theta}_{\sigma} = \xi(\sigma + iz)/\xi(\sigma - iz)$, with $\text{RH} \Leftrightarrow \hat{\Theta}_{\sigma}$ pole-free in the upper half-plane for all $\sigma > \frac{1}{2}$.

8.4 Value of the Correction Record

These corrections demonstrate that the research process is self-correcting. Each correction sharpened the theoretical framework: the GF1 chain moved from a naïve constant ($C_m \approx 8$) through several intermediate forms to the $\lambda = 5, 7$ boundary-moment benchmark $\text{sLI} \sim C^* \lambda / (NL)$; the $\lambda = 3$ sweep and the open $\lambda = 11$ discriminator prevent an all- λ theorem reading, and the c_{Ξ} correction removed a circular RH-dependence from a key constant.

9 Summary

Part IV contributes the cross-route perspective: the three active branches share a Galerkin pattern (fixed-parameter positive, uniform-limit open), and two isolated proof slots (B-1 log-curvature, GF1 factor-16 identification) are the concrete analytical targets, which remain well-posed open problems. Section 3 then adds a 2026 *heuristic* criticality lens ($\text{RH} \Leftrightarrow \Lambda = 0$) that re-weights the routes toward exact-structural mechanisms (without refuting the margin-type slots), localizes the common wall as the analytic continuation across $\Re s = 1$, records a computable Laguerre–Pólya/Herglotz diagnostic (a classical criterion turned into a test instrument), and reports four closures of *concrete constructions* — the OS reflection in natural form, the naïve Möbius L_s^- , the Toeplitz-symbol sub-method, and the relative-determinant bypass — none of which closes a route or RH. The dormant and external routes are catalogued for completeness; the hypothesis correction record provides transparency. The condensed atlas and final assessment follow in Part V [4].

Acknowledgments. Computations were performed on a Mac Studio (Apple M1 Max, 32 GB) and a Hetzner CCX13 cloud server (2 dedicated vCPU, 8 GB). The interval arithmetic library `mpmath` was developed by Fredrik Johansson.

AI Disclosure

This manuscript was developed in extensive collaboration with AI language models (Claude, Anthropic; Copilot, Microsoft/GitHub; Gemini, Google DeepMind). AI contributed to conceptual exploration, analytical work, literature review, writing, and critical review. The author, Lukas Geiger, conceived the research direction, provided domain expertise, and exercised final editorial authority. The author assumes full and sole scholarly responsibility for all content, including any errors.

References

- [1] Alain Connes. The Riemann hypothesis: Past, present and a letter through time. *arXiv preprint*, 2026. Survey on the Riemann Hypothesis: past, present, and a letter to Riemann using only pre-1860 mathematics.
- [2] Alain Connes, Caterina Consani, and Henri Moscovici. Zeta zeros and prolate wave operators. *Annals of Functional Analysis*, 15(87), 2024. Jacobi matrix representation and prolate–zeta connection.
- [3] Alain Connes and Walter D. van Suijlekom. Quadratic forms, real zeros and echoes of the spectral action. *Communications in Mathematical Physics*, 406(12), 2025. Quadratic forms, real zeros, and connections to the spectral action.
- [4] Lukas Geiger. The Riemann hypothesis: Conclusio. Preprint, 2026. Part V of a five-part series (in preparation).
- [5] Lukas Geiger. The Riemann hypothesis: Even dominance of the Weil quadratic form. Preprint, 2026. Part II of a five-part series.

- [6] Lukas Geiger. The Riemann hypothesis: Foundations, obstructions, and reorientation. Preprint, 2026. Part I of a five-part series.
- [7] Lukas Geiger. The Riemann hypothesis: Spectral routes. Preprint, 2026. Part III of a five-part series (in preparation).

The Riemann Hypothesis: Conclusio

Part V of a Five-Part Series

Lukas Geiger

Bernau, Germany

ORCID: [0009-0005-7296-1534](https://orcid.org/0009-0005-7296-1534)

May 31, 2026

Abstract

This paper concludes a five-part investigation of the Riemann Hypothesis through spectral analysis, arithmetic thermodynamics, and the Weil quadratic form:

Part I [5]: Foundations and Obstructions.

Part II [4]: Even Dominance (Branch A).

Part III [6]: Spectral Routes (Branches B and Z).

Part IV [3]: Cross-Route Synthesis.

Part V (this paper): Conclusio.

We present the condensed atlas: three independent conditional reductions of RH (Even Dominance, Hadamard Strip, CCM Microcluster), three structural identifications ($c_{\Xi} \approx 0.0231$, $C^* = 64\pi^2/3$ as a $\lambda = 5, 7$ boundary-moment benchmark with $\lambda = 11$ still open, $g_0(m) \approx 2\pi/\log \gamma_m$), two conditional theorems (B-1, GF1), and two isolated open proof slots. The series title remains “From Landscape to Atlas”; promotion to “From Atlas to Proof” is governed by a five-gate audit protocol triggered only by a substantive mathematical breakthrough.

Series status (v3.0, 2026-05-26; criticality-lens update 2026-05-31). Five-part series; all three conditional reductions unchanged from v2.2. Structural reorganisation only; no new proof claims. A 2026 working phase adds a *heuristic* criticality lens ($\text{RH} \Leftrightarrow \Lambda = 0$), a sharper common-wall localization, and a computable Laguerre–Pólya/Herglotz diagnostic (Part IV [3]); these re-weight priorities and close four *concrete constructions*, but make no new proof claim.

MSC 2020: 11M26, 11M36, 47A75

Keywords: Riemann Hypothesis, conditional reduction, atlas, assessment, open problems, outlook

1 Atlas Overview

Route	Sector	Status	Key result	Paper
A (Even Dom.)	Even (Part II)	conditional	A1–A8 chain + GF1 benchmark with $C^* = 64\pi^2/3$ for $\lambda = 5, 7$; $\lambda = 11$ open	[4]
B (Hadamard)	Spectral (Part III)	conditional	$A_\lambda \leq 1.0063$ under log-curvature control	[6]
Z (CCM)	Spectral (Part III)	conditional	$C2ca.1 + C2ca.2 + C2ca.5$; $g_0(m) = 2\pi/\log \gamma_m$	[6, 7]
Cross-Route	Combined (Part IV)	empirical	shared Galerkin pattern; 2026 criticality lens + sharpened common wall ($\Re s = 1$ continuation)	[3]

2 Result Summary

2.1 Three Structural Identifications

1. Hadamard curvature constant (RH-neutral):

$$c_\Xi = -\frac{1}{2} \sum_{\rho} \frac{1}{(\rho - \frac{1}{2})^2} \approx 0.0231.$$

Under RH, this specialises to $\sum_{k \geq 1} 1/\gamma_k^2$. (Part III [6].)

2. **Boundary-moment constant:** $C^* = 64\pi^2/3 \approx 210.55$, identified as a candidate via $16 \int_0^\infty t^2/\sinh^2(t/2) dt$; the factor 16 is conjectured to arise from the DZT2 ladder normalisation. This is a $\lambda = 5, 7$ benchmark, with $\lambda = 3$ retained as a separate low-bandwidth guardrail and $\lambda = 11$ retained as an open discriminator after the $N = 200$ point. (Part II [4].)

3. **Hecke-normalised gap:** $g_0(m) \approx 2\pi/\log \gamma_m$, classical Montgomery–Odlyzko. (Part III [6].)

2.2 Two Conditional Theorems

- **Theorem B-1** (Part III): Under uniform log-curvature control $H_\lambda(\delta) \leq K\delta^2$ with $K \leq 1.8 \times 10^{-3}$:

$$A_\lambda(\delta) \leq B_\Xi \cdot \exp(K/4) \leq 1.0063,$$

where $B_\Xi = \xi(0)/\xi(\frac{1}{2}) = 1.00579\dots$ is analytical.

- **Theorem GF1** (Part II): Under eigenmode-weighted kernel convergence producing the effective factor 16:

$$\text{sLI} = \frac{64\pi^2}{3} \frac{\lambda}{NL} + o(\cdot).$$

This is formulated for the non-exceptional converged regime; the $\lambda = 3$ sweep and the open $\lambda = 11$ discriminator prevent an all- λ universal reading.

2.3 Two Isolated Open Proof Slots

1. **B-1 slot:** Uniform log-curvature control from Foundation-Lock structure. Closes Branch B.
2. **GF1 slot:** Kernel convergence + factor-16 identification from DZT2 definitions. Closes Branch A’s non-exceptional boundary-moment benchmark, with the low-bandwidth branch kept explicit.

3 Localization Tables

3.1 Li Coefficient Perspective

Step	Statement	Status	Method
S1	Define $\xi(s)$, functional equation	proven	definition
S2	Li coefficients via Bombieri–Lagarias	proven	[1]
S3	Decomposition $\lambda_n = \lambda_n^{(\infty)} + \lambda_n^{(\text{fin})}$	proven	Part I
S4	Archimedean dominance: $\lambda_n^{(\infty)} \sim \frac{n}{2} \log n$	proven	Stirling
S5	$\lambda_n \geq 0$ for all $n \geq 1$	OPEN*	$\iff$ RH
S6	All zeros on $\Re(s) = 1/2$	$\Leftrightarrow$ S5	Li (1997)
S7	Prime counting: $\pi(x) = \text{Li}(x) + O(\sqrt{x} \log x)$	$\Leftarrow$ S6	von Koch

* S5 is equivalent to RH (Li, 1997). The present series gives only conditional reductions.

3.2 Even Dominance Perspective

Step	Statement	Status	Method
E1	Weil explicit formula	proven	Weil (1952)
E2	Weil quadratic form QW_λ	proven	[2]
E3	Even/odd decomposition of QW_λ	proven	Part II
E4	Shift Parity Lemma: each prime favors even	proven	Part II
E5	Gap diagnostics at 33 values, $\lambda \leq 1,300,000$	evidence / conditional	Part II
E5b	Euler–Maclaurin: $\rho^{\text{EM}} < 0$ for $\lambda \geq 1.2\text{M}$	proven	Part II
E6	Cumulative step: even dominance for all λ	conditional reduction	= A6
E7	Connes’ Theorem 6.1 + Hurwitz bridge $\Rightarrow$ RH	conditional on E6	[2]

4 Remaining Open Problems

Remark 4.1 (OP1: Cumulative Even Dominance — reduced to the odd-sector lower bound (v2.2)). The cumulative even dominance problem is reduced to a single spectral question: a quantitative lower bound for $\lambda_1^-(\lambda)$ (Asymptotic Variational Gap Conjecture, Part II).

Remark 4.2 (OP2 resolved: Simplicity of λ_1^+). Certified as simple for all 33 certificate values by interval arithmetic. The intra-even gap grows as $\sim \sqrt{\lambda}$ (Part II).

Conjecture 4.3 (OP3: Analytical Proof of Index Rigidity). *The negative inertia index of K_σ^{sym} is exactly 2 for all N sufficiently large and $\sigma \in (1, \sigma_{\max})$.*

Conjecture 4.4 (OP4: Nuclear Transfer Operator for ξ). *Construct a separable Hilbert space $\mathcal{V}$ and a nuclear family $s \mapsto \mathcal{L}_s$ with $\xi(s)/\xi(0) = \det(I - \mathcal{L}_s)$, spectral radius < 1 for $\Re(s) > 1/2$, and eigenvalue 1 at the zeros.*

Conjecture 4.5 (OP5: Nevanlinna Property of $\log \det_2$). *Is $\log \det_2(I - zR)$ a Nevanlinna function if and only if all zeros lie on the real axis?*

OP2 is resolved. OP1 is reduced to a single asymptotic step. OP3–OP5 remain structural questions.

5 Assessment and Outlook

5.1 What Has Been Achieved

1. A **complete structural analysis** of the gauge-theoretic approach (Part I: R1–R9, K1–K4).
2. The **Shift Parity Lemma** and 33 **computer-assisted finite-range diagnostics** (Part II).
3. The **conditional reduction A1–A8** with analytical backbone (M1'', Euler–Maclaurin, PNT transfer) and the GF1 boundary-moment benchmark with $C^* = 64\pi^2/3$ for $\lambda = 5, 7$ (Part II).
4. Two **independent spectral conditional reductions** (Branches B and Z, Part III).
5. The **cross-route synthesis** identifying shared Galerkin pattern and two isolated proof slots (Part IV).
6. **Fifteen excluded paths** (spectral + CCM) and **five methodical dead ends** (Parts III–IV).
7. A 2026 **criticality lens** ($\text{RH} \Leftrightarrow \Lambda = 0$, heuristic) re-weighting the routes toward exact-structural mechanisms; a sharper **common-wall** localization (the analytic continuation across $\Re s = 1$); a computable **Laguerre–Pólya/Herglotz diagnostic**; and four closures of *concrete constructions* (not routes) — all without changing the conditional-reduction status (Part IV [3]).

5.2 Current Atlas Status

The series title remains “From Landscape to Atlas: Multi-Route Cartography of an Ongoing Expedition Toward the Riemann Hypothesis.” All three living routes are conditional reductions; no single branch has reached conditional-proof level under the project’s audit protocol.

Promotion protocol. The title migration from “Atlas” to “From Atlas to Proof” is governed by a five-gate audit: (1) single named hypothesis H ; (2) every step from H to RH at theorem-level; (3) two independent reviews including a sceptical refuter; (4) 2–3 day cool-off; (5) explicit Decision Record entry. Default: Atlas remains.

5.3 Future Scenarios

- **One branch dies** (its open hypothesis refuted): documented as dead route; the other two continue.
- **One branch matures unconditionally**: that proof stands as RH; other branches retained for independent verification.
- **Both/all three mature**: rare, but mathematically valuable as multi-route verification.
- **None matures, but partial progress continues**: research programme continues; the Landscape preserves the audit. This is the default outlook.

5.4 Outlook

Concrete next steps (not time-bound):

- Rigorous derivation of the two open proof slots (B-1 log-curvature, GF1 factor-16).
- Resolution of the $\lambda = 11$ GF1 discriminator, validation at higher λ , and a theorem-level separation of the $\lambda = 3$ low-bandwidth branch.
- κ_∞ -asymptotics for $\lambda \rightarrow \infty$.
- Promotion-gate audit if a slot closes.
- Develop the exact-structural lines (M = Bost–Connes $\leftrightarrow$ Meyer; P-A = global adelic trace route) that the 2026 criticality lens elevates, in parallel with the two margin-type proof slots (Part IV [3]).

6 Connections to the Wider Landscape

6.1 Connes’ Program

The work in this series fits directly into Connes’ spectral approach to RH [2]. The Shift Parity Lemma (Part II) provides analytical evidence for the even dominance hypothesis (Theorem 6.1). The I-1 Normal-Family reframe (Part III) decomposes Wall (ii) into a single named condition (ii-a). The CCM Microcluster Programme (Part III) attacks the same wall through the Fourier model.

6.2 Nyman–Beurling–Baez-Duarte

The Li criterion, the Weil positivity criterion, the Nyman–Beurling criterion, and the Baez-Duarte criterion are all equivalent to RH. Direct implications between them (without passing through RH) are known only in two cases: Weil $\Rightarrow$ Li ([1]) and Nyman–Beurling $\Rightarrow$ Baez-Duarte (trivial).

6.3 The Trace-Class Obstruction as Universal Pattern

Obstruction K4 (trace-class failure at $\Re(s) = 1/2$, Part I) appears across QFT (UV divergences), statistical mechanics (thermodynamic limit), and NCG (zeta-function regularisation). The Even Dominance approach circumvents this by working with finite-dimensional Galerkin blocks rather than infinite-dimensional transfer operators.

7 Full Summary Table

Contribution	Status	Paper
Thermodynamic formalism for ζ	established	I
9 unconditional structural results (R1–R9)	proven	I
4 excluded paths (K1–K4)	proven	I
Shift Parity Lemma	proven (analytic)	II
Gap diagnostics (33 values, $\lambda \leq 1,300,000$)	evidence / conditional	II
Leading-Mode Cancellation ($c = 2 + \sqrt{2}$)	proven	II
Euler–Maclaurin: $\rho^{\text{EM}}(13) < 0$	proven	II
GF1 benchmark: $\text{sLI} \sim C^* \lambda / (NL)$ for $\lambda = 5, 7$; $\lambda = 3$ separate, $\lambda = 11$ open	conditional (Thm GF1)	II
I-1 Normal-Family Reframe: (ii) $\rightarrow$ (ii-a)	rigorous	III
Foundation-Lock + converged empirics	rigorous + numerical	III
Theorem B-1: $A_\lambda \leq 1.0063$ (conditional)	conditional	III
c_Ξ Hadamard curvature (RH-neutral)	analytical	III
CCM three C2ca-inputs	conditional reduction	III
6 excluded spectral + 9 excluded CCM paths	documented	III
Cross-route Galerkin pattern	empirical	IV
11 systematically explored alternatives	documented	IV
Hypothesis correction chain (GF1 Iter 15–27)	documented	IV
Cumulative step (A6)	conditional reduction (v2.2)	II
Riemann Hypothesis	conditional reduction (v2.2)	I–V

Acknowledgments. Computations were performed on a Mac Studio (Apple M1 Max, 32 GB) and a Hetzner CCX13 cloud server (2 dedicated vCPU, 8 GB). The interval arithmetic library mpmath was developed by Fredrik Johansson.

AI Disclosure

This manuscript was developed in extensive collaboration with AI language models (Claude, Anthropic; Copilot, Microsoft/GitHub; Gemini, Google DeepMind). AI contributed to conceptual exploration, analytical work, literature review, writing, and critical review. The author, Lukas Geiger, conceived the research direction, provided domain expertise, and exercised final editorial authority. The author assumes full and sole scholarly responsibility for all content, including any errors.

References

- [1] Enrico Bombieri and Jeffrey C. Lagarias. Complements to Li’s criterion for the Riemann hypothesis. *Journal of Number Theory*, 77(2):274–287, 1999.
- [2] Alain Connes. The Riemann hypothesis: Past, present and a letter through time. *arXiv preprint*, 2026. Survey on the Riemann Hypothesis: past, present, and a letter to Riemann using only pre-1860 mathematics.
- [3] Lukas Geiger. The Riemann hypothesis: Cross-route synthesis. Preprint, 2026. Part IV of a five-part series (in preparation).

- [4] Lukas Geiger. The Riemann hypothesis: Even dominance of the Weil quadratic form. Preprint, 2026. Part II of a five-part series.
- [5] Lukas Geiger. The Riemann hypothesis: Foundations, obstructions, and reorientation. Preprint, 2026. Part I of a five-part series.
- [6] Lukas Geiger. The Riemann hypothesis: Spectral routes. Preprint, 2026. Part III of a five-part series (in preparation).
- [7] Lukas Geiger. The spectral zookeeper: A conditional reduction toward the Riemann hypothesis via spectral microcluster closure in the Connes–Consani–Moscovici framework. Zenodo, 2026. Companion CCM-route paper; v1.4 record DOI 10.5281/zenodo.20396399, concept DOI 10.5281/zenodo.19673126.