

The Riemann Hypothesis: Even Dominance of the Weil Quadratic Form

Lukas Geiger

Bernau, Germany

ORCID: [0009-0005-7296-1534](https://orcid.org/0009-0005-7296-1534)

May 27, 2026

Abstract

This is the second paper in a five-part series (Part I: Foundations and Obstructions [9]; Part III: Spectral Routes [10]; Part IV: Cross-Route Synthesis [8]; Part V: Conclusion [7]). We address the open problem identified in Connes’ 2026 survey [2]: whether the minimum eigenvalue of the Weil quadratic form QW_λ is simple with even eigenfunction. We establish three main results. First, the *Shift Parity Lemma*: the difference matrix between even and odd shift operators has strictly negative minimum eigenvalue for all $N \geq 3$ modes and all normalized shifts $r \in (0, 2)$, implying that every prime individually favors even eigenfunctions. Second, we report computer-assisted finite-range gap diagnostics using interval arithmetic for 33 values of λ from 100 to 1,300,000; these are a strong conditional bridge until the odd-sector lower bound is independently validated. Third, we prove the *Leading-Mode Cancellation Lemma*: the overlap differences between even and odd sectors cancel pairwise with exact constant $c = 2 + \sqrt{2}$, yielding $(E_{\sin} - E_{\cos})/D(\lambda) = O(1/\log \lambda) \rightarrow 0$. In v2.0 we add two non-existence theorems (Theorems 2.14 and 2.17) that formally rule out the absolute total-positivity and universal-commuting-operator routes, establishing that even dominance is necessarily an arithmetic-statistical rather than a purely algebraic phenomenon. In v2.1 we add a robust Direct Frontier-Dominance argument (Section A.10, Theorem A.37) based on a common Rayleigh test vector on the frontier window $r \in [0.7, 1]$.

Status correction (v2.2, 2026-05-14). External critical review of the v2.1 manuscript and the companion proof-only extract [6] identified a structural gap in the asymptotic argument that was hidden by the previous wording. The Direct Frontier-Dominance proposition establishes $\lambda_{\min}(W_{\text{prime}}^+ - W_{\text{prime}}^-) = -\Theta(\sqrt{\lambda})$ via Rayleigh evaluation at a common test vector. This is a *variational* statement: there exists a vector v with $v^T W^+ v < v^T W^- v$. It does *not* establish $\lambda_1^+(\lambda) < \lambda_1^-(\lambda)$ asymptotically, because $\lambda_1^-(\lambda) \leq v^T W^- v$ runs in the same direction; a separate quantitative lower bound for λ_1^- is required. The simple 2×2 counter-example $A = \text{diag}(0, 10)$, $B = \text{diag}(5, -5)$ illustrates the gap. The status of the A1–A8 reduction chain is therefore revised from “unconditional” (v2.1) back to **conditional reduction**: the finite range $100 \leq \lambda \leq 1,300,000$ provides strong interval-arithmetic diagnostics, but theorem-level even dominance on that range still depends on a validated lower bound for $\lambda_1^-(\lambda)$. Asymptotic even dominance is reduced to the *Asymptotic Variational Gap Conjecture* (a quantitative lower bound for $\lambda_1^-(\lambda)$). A proposed approach via degenerate perturbation theory on the asymmetric three-mode diagonal structure is sketched in the companion paper [6]. The non-existence theorems NE-A and NE-B, the Shift Parity Lemma, the Leading-Mode Cancellation Lemma, and the 33 finite-range diagnostics remain unaffected by this revision.

Structural update (v3.0, 2026-05-26). The series is restructured from a trilogy to a five-part series; the conditional reduction status from v2.2 is unchanged. A new Section 5 presents the GF1 boundary-moment benchmark $\text{sLI} \sim (64\pi^2/3) \lambda/(NL)$ for the converged $\lambda = 5, 7$ benchmark, records the $\lambda = 3$ low-bandwidth outlier separately, and leaves $\lambda = 11$ as an open discriminator.

MSC 2020: 11M26, 11M36, 47A75, 65G20

Keywords: Riemann Hypothesis, Weil quadratic form, even dominance, Shift Parity Lemma, computer-assisted diagnostics

Contents

1	Introduction	5
2	Connes' Reduction and the Weil Quadratic Form	5
2.1	Connes' Reduction	5
2.2	Galerkin Discretization	6
2.3	Result 1: Spectral Gap in the Even Sector	6
2.4	Result 2: Even–Odd Sector Comparison	7
2.5	Discussion	8
2.6	Decomposition Analysis: Source of Even Dominance	9
2.7	The Shift Parity Lemma	9
2.8	Universal Even Preference of Individual Primes	11
2.9	Multi-Scale Decomposition: Frontier Primes	12
2.10	Proof Architecture: From Lemma to Even Dominance	12
2.11	Possible Approaches to a Formal Proof	13
2.12	Formal non-existence theorems (v2.0)	15
2.13	Prolate Eigenvalue Dominance: A Structural Result	17
2.14	Jentzsch Regularization Argument	18
2.15	Eigenvalue Spread Bound	18
2.15.1	Computer-Assisted Gap Diagnostics via Interval Arithmetic	19
2.16	Gap Asymptotics and Monotonicity	20
3	Weighted Compactness as Proof Closure	20
3.1	Setup and Rescaling	21
3.2	Building Block 1: Uniform Sobolev Bound	21
3.3	Building Block 2: Precompactness via Rellich	22
3.4	Building Block 3: Spectral Stability	22
3.5	Building Block 4: Bridge to Hurwitz	22
3.6	Numerical Verification	23
3.7	Assessment and Risks	24
3.8	Gap Growth from Block Bounds	24
3.8.1	Lemma L1: Prime coupling drives an exponentially deep even direction	25
3.8.2	Lemma L2: High-mode block is only polynomially negative	25
3.8.3	Lemma L3: Low–high coupling (historical)	26
3.8.4	The gap growth theorem	26
4	Connections to Existing Work	27
4.1	Li–Keiper Coefficients	27
4.2	Weil's Explicit Formula and Connes' Program	27
4.3	de Branges Spaces and Hilbert–Pólya	27
5	Galerkin Imbalance: Boundary-Moment Benchmark (v3.0)	28
5.1	Identification of the Boundary-Moment Constant	28
5.2	Hypothesis Correction Chain	29
6	Conclusion	30

A	Proofs of the Block-Bound Lemmas	30
A.1	Closed-form shift overlaps	31
A.2	Proof of Lemma L1: Exponential depth of the even direction	31
A.3	Proof of Lemma L2: High-block control via Gershgorin	33
A.4	Proof of Lemma L3: Off-diagonal coupling via Schur test	33
A.5	The Certifier Meta-Theorem	35
A.6	Status of the assumptions	37
A.7	Resolvent-damped energy framework for M1	38
A.8	Computer-assisted certificates	41
A.9	Analytical path to M1: leading-mode cancellation	42
A.10	Direct Frontier-Dominance Proof (v2.1, robust form)	51

1 Introduction

This paper is the second in a five-part series investigating the Riemann Hypothesis through gauge theory, spectral analysis, and arithmetic thermodynamics. Part I [9] establishes the thermodynamic landscape (structural results R1–R9) and documents the failure of the gauge-theoretic approach (obstructions K1–K4), leading to a reorientation toward Connes’ spectral program. Part III [10] surveys the spectral routes; Part IV [8] collects cross-route synthesis; Part V [7] provides the conclusio.

Here we develop the spectral approach in full. Following Connes’ 2026 survey [2], we study the Weil quadratic form QW_λ and its even-odd decomposition. Our three main contributions are:

1. The **Shift Parity Lemma** (Section 2.7): a purely algebraic result showing that every prime individually favors even eigenfunctions, proved via closed-form trigonometric entries.
2. **Computer-assisted finite-range diagnostics** (Section A.8): strong negative gap diagnostics at 33 values of λ from 100 to 1,300,000 via interval arithmetic; theorem-level even dominance remains conditional on the validated odd-sector lower bound.
3. The **Leading-Mode Cancellation Lemma** (Theorem A.24): the overlap differences between sectors cancel pairwise with exact constant $c = 2 + \sqrt{2}$, reducing the analytical gap (M1'') to Fourier analysis and the Prime Number Theorem.

2 Connes’ Reduction and the Weil Quadratic Form

A recent survey by Connes [2] reduces the Riemann Hypothesis to a spectral condition on the Weil quadratic form QW_λ . We summarize the connection and present numerical evidence addressing the remaining open step.

2.1 Connes’ Reduction

Connes constructs a self-adjoint operator A_λ on $L^2([-\log \lambda, \log \lambda])$ from the Weil explicit formula. In additive coordinates ($u = \log x$), the quadratic form reads

$$\begin{aligned} \langle A_\lambda f | f \rangle = & (\log 4\pi + \gamma) \|f\|^2 + \int_0^\infty \left[\langle T_u f + T_{-u} f, f \rangle - 2e^{-u/2} \|f\|^2 \right] \frac{e^{u/2}}{2 \sinh(u)} du \\ & + \sum_p \sum_{m=1}^\infty \frac{\log p}{p^{m/2}} \langle T_{m \log p} f + T_{-m \log p} f, f \rangle, \end{aligned}$$

where T_u denotes the shift operator, γ is the Euler–Mascheroni constant, and the archimedean kernel $e^{u/2}/(2 \sinh u)$ arises from the change of variables $x = e^u$ applied to Connes’ equation (10) in [2], with $x^{1/2}/(x - x^{-1}) d^*x = e^{u/2}/(2 \sinh u) du$. The operator A_λ is lower-bounded with compact resolvent (hence discrete spectrum).

Theorem 2.1 (Connes, Theorem 6.1 in [2]; proved in [4]). *Let $L > 0$ and D be a real distribution on $[0, L]$ with even extension $\tilde{D}$ on $[-L, L]$. Assume the quadratic form with kernel $\tilde{D}(x - y)$ defines a lower-bounded self-adjoint operator on $L^2([-L/2, L/2])$ whose minimum eigenvalue is simple, isolated, with even eigenfunction η . Then all zeros of $\tilde{\eta}(z)$ lie on the real line.*

Connes' Section 6.6 states: "In order to apply Theorem 6.1, one needs to show that the smallest eigenvalue of QW_λ is *simple* with *even* eigenvector." This is the remaining open step.

2.2 Galerkin Discretization

We approximate A_λ in the cosine basis $\{\varphi_n\}_{n=0}^{N-1}$ (even sector) and sine basis (odd sector) on $[-L, L]$ with $L = \log \lambda$. The matrix elements are computed via quadrature ($n_{\text{quad}} \geq 800$, $n_{\text{int}} \geq 500$) using the corrected kernel $e^{u/2}/(2 \sinh u)$ and primes $p \leq \max(\lambda, 100)$.

2.3 Result 1: Spectral Gap in the Even Sector

Within the even sector, we verify $\text{gap}(\lambda) = \lambda_2^+ - \lambda_1^+ > 0$ via Cauchy interlacing with N -convergence. Server computations ($\lambda \in [5, 200]$, N up to 80, corrected orthogonal basis $\cos(n\pi t/L)$) confirm:

λ	N_{final}	$\text{gap}^+ = \lambda_2^+ - \lambda_1^+$	Cauchy bound
5	30	3.52×10^{-1}	3.52×10^{-1}
10	30	2.28×10^{-1}	2.28×10^{-1}
20	30	1.39×10^{-1}	1.39×10^{-1}
50	30	2.35×10^{-1}	2.35×10^{-1}
100	80	4.83×10^0	4.83×10^0
200	80	1.11×10^1	1.11×10^1

Observation: All tested values have strictly positive Cauchy bound (minimum 0.095 at $\lambda = 30$), confirming that λ_1^+ is simple within the even sector. For $\lambda \geq 100$, the gap grows rapidly (~ 5 – 11) as the low-mode cosine block dominates.

Lemma 2.2 (Simplicity of λ_1^+). *For all 33 certificate values $\lambda \in [100, 1,300,000]$, the even-sector ground state is simple:*

$$\lambda_2^+(\lambda) - \lambda_1^+(\lambda) \geq g_{\min}^+(\lambda) > 0,$$

certified by interval arithmetic on the 4×4 even Galerkin block. Selected values:

λ	λ_1^+	λ_2^+	gap^+ (certified)
100	-11.18	-2.49	≥ 8.69
500	-34.06	-8.77	≥ 25.29
10,000	-195.46	-61.97	≥ 133.48
80,000	-594.37	-216.85	≥ 377.52
320,000	-1220.97	-489.97	≥ 731.00

The gap grows monotonically as $\sim \sqrt{\lambda}$, consistent with the analytical prediction from the Galerkin diagonal asymptotics: the leading mode ($n = 1$) has $W_{11}^+ \sim -\sqrt{\lambda}$, while the next mode ($n = 0$) has $W_{00}^+ \sim -c'\sqrt{\lambda}$ with $c' < c$. If the remaining even-dominance inequality $\lambda_1^+ < \lambda_1^-$ is supplied by a valid odd-sector lower bound, this certifies that the global ground state of A_λ is simple and even, as required by Connes' Theorem 6.1.

2.4 Result 2: Even–Odd Sector Comparison

For Theorem 6.1, the minimum of A_λ over the *full* space $L^2([-L, L])$ must be simple and even. Since A_λ commutes with parity ($f(t) \mapsto f(-t)$), the even and odd sectors decouple. We compare λ_1^+ (even) with λ_1^- (odd):

λ	N	λ_1^+	λ_1^-	Sector	Thm. 6.1
5	30	−4.31	−4.28	EVEN	✓*
10	30	−5.01	−5.19	ODD	—
20	30	−5.81	−5.84	ODD	—
25	60	−6.39	−6.39	EVEN*	✓*
28	60	−6.53	−6.53	ODD*	—*
30	60	−6.64	−6.63	EVEN*	✓*
50	60	−6.96	−6.94	EVEN*	✓*
55	60	−7.20	−6.70	EVEN	✓
100	80	−11.25	−9.06	EVEN	✓
200	80	−19.27	−15.05	EVEN	✓
500	80	−34.51	−26.88	EVEN	✓

* *Marginal*: $|\lambda_1^+ - \lambda_1^-| < 0.05$.

Honest assessment:

- For $\lambda \geq 55$: Robust negative gap diagnostics with margins growing as $\sim \sqrt{\lambda}$. The diagnostic gap reaches $|\Delta| > 0.5$ at $\lambda = 55$ and exceeds 475 at $\lambda = 1,300,000$. Theorem 6.1 prerequisites are conditionally supported for 33 values up to $\lambda = 1,300,000$, pending the odd-sector lower-bound validation.
- For $25 \leq \lambda \leq 50$: Marginal regime ($|\Delta| < 0.02$). The even/odd ordering oscillates with λ at $N = 60$: $\lambda = 25, 30, 40, 50$ show even dominance while $\lambda = 28, 35, 45$ show odd dominance. Neither sector robustly dominates.
- For $\lambda = 10$ and $\lambda = 20$: The odd sector has the lower minimum. Theorem 6.1 is *not directly applicable* in this regime.
- **Non-monotone transition:** The crossover from odd to even ground state is not sharp but spreads over $\lambda \in [25, 55]$. For $\lambda \geq 55$, even dominance is robust and the gap grows without bound.
- **Computer-assisted diagnostics (interval arithmetic):** A negative finite-range gap diagnostic has been documented for 33 values of λ from 100 to 1,300,000 using interval arithmetic (mpmath.iv, 50 digits) for the even block and float64 tail extrapolation for the odd block. The gap grows monotonically as $\sim \sqrt{\lambda}$, from -2.25 at $\lambda = 100$ to -475.83 at $\lambda = 1,300,000$; theorem-level certification is conditional on the odd-sector lower-bound validation (see Section A.8).

2.5 Discussion

Our numerical results address Connes’ open problem (Section 6.6 of [2]) from two angles:

1. **Simplicity within even sector:** The Cauchy-interlacing bound provides rigorous (up to discretization error) evidence that λ_1^+ is simple. This holds for all tested $\lambda \in [5, 200]$.
2. **Even dominance for large λ :** For $\lambda \geq 25$, the finite-dimensional diagnostics put the even sector lower in most tested ranges. The theorem-level infinite-dimensional ordering still needs the odd-sector lower bound.
3. **Odd dominance for small λ :** For $\lambda = 10$ and $\lambda = 20$, the odd sector is lower. Theorem 6.1 is not directly applicable in this regime.

Three structural features emerge:

1. **Diffuse crossover:** At $N = 60$, the crossover from odd to even dominance spreads over $\lambda \in [25, 55]$, with oscillating sign and $|\Delta| < 0.02$. At $\lambda \geq 55$, even dominance sets in robustly ($|\Delta| > 0.5$) and grows without reversal through $\lambda = 500$.
2. **Growing finite-range diagnostic gap:** For $\lambda \geq 55$, the observed gap grows in magnitude, reaching $\Delta \approx -7.6$ at $\lambda = 500$ and $\Delta \approx -476$ at $\lambda = 1,300,000$ (see Sections A.8 and 2.16). The growth scales as $\sim \sqrt{\lambda}$ (analytically, via PNT; the earlier fit $\sim (\log \lambda)^{4.2}$ was based on a limited range), but the eigenvalue-ordering interpretation needs the odd-sector lower bound.
3. **Finite-range conditional bridge:** For 33 values of λ from 100 to 1,300,000, the computer-assisted diagnostics provide a strong conditional bridge pending odd-tail validation (see Section A.8).
4. **A Direct Frontier-Dominance variational proof** (v2.1, Section A.10): an independent proof of the asymptotic variational sign using a common Rayleigh test vector on the prime frontier, PNT partial summation, Mertens’ theorem, and the existing finite CAP bridge. It obviates both v2.0 caveats for the variational statement, not for the eigenvalue-ordering step.

Implication for Connes’ program: Connes’ strategy (Sections 6.4–6.6 of [2]) uses Hurwitz’s theorem to transfer the real-zeros property from the truncated Fourier transforms $\hat{k}_\lambda$ to the Riemann Ξ -function in the limit $\lambda \rightarrow \infty$. Crucially, Hurwitz’s theorem requires only a *sequence* $\lambda_n \rightarrow \infty$ along which the prerequisites hold—not all $\lambda > 1$. Therefore, the odd-dominant regime for small λ (around $\lambda \approx 10$ –30) is *not* an obstruction. The remaining large-step requirement is a theorem-level lower bound for λ_1^- that turns the finite diagnostics and asymptotic variational statement into genuine eigenvalue ordering.

Our data show robust negative gap diagnostics for $\lambda \geq 50$ (with margins growing as $\sim \log(\lambda)^b$), and the gap grows monotonically for $\lambda \geq 120$. In the current v2.4 status, the formal proof that $\lambda_1^+ < \lambda_1^-$ for all $\lambda \geq 100$ remains conditional on a validated lower bound for the odd-sector minimum eigenvalue. Theorem A.36 records this conditional reduction; Section A.10 gives the robust frontier Rayleigh proof of the associated variational sign.

2.6 Decomposition Analysis: Source of Even Dominance

To understand the mechanism behind even dominance, we decompose $QW_\lambda = W_{\text{diag}} + W_{\text{arch}} + W_{\text{prime}}$ and compute each contribution separately for the even and odd sectors:

- $W_{\text{diag}} = (\log 4\pi + \gamma) I$: identical in both sectors.
- W_{arch} : the archimedean kernel $e^{u/2}/(2 \sinh u)$ produces *virtually identical* minimum eigenvalues in both sectors ($\Delta < 10^{-3}$ at $\lambda = 100$). The archimedean contribution is parity-neutral.
- W_{prime} : the prime shift operators $S_{m \log p} + S_{-m \log p}$ are the *sole driver* of even-dominance. At $\lambda = 100$: $\lambda_1^+(W_{\text{prime}}) \approx -14.5$ vs. $\lambda_1^-(W_{\text{prime}}) \approx -12.3$.

The mechanism is traced to the product formula for trigonometric functions. For the cosine (even) basis, shift-overlap integrals contain the constructive term $\cos((k_n + k_m)t)$:

$$\cos(k_n t) \cos(k_m(t - s)) = \frac{1}{2} [\cos((k_n - k_m)t + k_m s) + \cos((k_n + k_m)t - k_m s)].$$

For the sine (odd) basis, the second term enters with a minus sign. The difference matrix $D = W_{\text{prime}}^+ - W_{\text{prime}}^-$ is indefinite, and its structure produces the even-sector advantage.

2.7 The Shift Parity Lemma

The decomposition analysis shows that even dominance is driven by the prime shift operators. A deeper analysis reveals a remarkable algebraic structure: *every single prime individually favors the even sector*, and this property is independent of the cutoff parameter L .

Definition 2.3 (Difference matrix). For $N \in \mathbb{N}$ and $r \in (0, 2)$, define the $N \times N$ symmetric matrix $D_N(r)$ by

$$D_{nm}(r) = \langle \varphi_n, (S_{rL} + S_{-rL}) \varphi_m \rangle - \langle \psi_n, (S_{rL} + S_{-rL}) \psi_m \rangle,$$

where $\varphi_n(t) = \cos(n\pi t/L)/\sqrt{L}$ (even basis) and $\psi_n(t) = \sin((n+1)\pi t/L)/\sqrt{L}$ (odd basis), and S_s denotes the shift operator on $L^2([-L, L])$.

Lemma 2.4 (L -independence). *The matrix $D_N(r)$ depends only on $r = s/L$, not on L separately. In particular, one may set $L = 1$ without loss of generality.*

Proof. By substitution $u = t/L$ in the shift-overlap integrals, all factors of L cancel due to the normalization of the basis functions. This is verified numerically to machine precision (spread $< 10^{-15}$) across $L \in \{1, 2, 5, 10, 20\}$. $\square$

Setting $L = 1$, the matrix entries admit closed-form expressions.

Proposition 2.5 (Closed-form entries). *For $L = 1$, the diagonal entries of $D_N(r)$ follow a telescoping pattern:*

$$D_{nn}(r) = (2 - r) [\cos(n\pi r) - \cos((n+1)\pi r)] - \frac{\sin(n\pi r)}{n\pi} - \frac{\sin((n+1)\pi r)}{(n+1)\pi},$$

with the convention $\sin(0)/0 = 0$ and $\cos(0) = 1$ for $n = 0$. The off-diagonal entries for $N = 3$ are:

$$\begin{aligned} D_{01}(r) &= \frac{(3\sqrt{2} + 4) \sin(\pi r) - 2 \sin(2\pi r)}{3\pi}, \\ D_{02}(r) &= \frac{-2\sqrt{2} \sin(2\pi r) + \sin(3\pi r) - 3 \sin(\pi r)}{4\pi}, \\ D_{12}(r) &= \frac{2[19 \sin(2\pi r) - 5 \sin(\pi r) - 6 \sin(3\pi r)]}{15\pi}. \end{aligned}$$

All formulas are derived by hand via product-to-sum identities and verified to 10^{-15} against numerical quadrature.

These closed forms reveal striking structural properties:

Corollary 2.6 (Structure at $r = 1$). $D_N(1) = \text{diag}(2, -2, 2, -2, \dots)$, i.e., $D_{nm}(1) = 2(-1)^n \delta_{nm}$.

Proof. At $r = 1$: $\cos(n\pi) - \cos((n+1)\pi) = (-1)^n - (-1)^{n+1} = 2(-1)^n$, and $\sin(k\pi) = 0$ for all $k \in \mathbb{Z}$. $\square$

Corollary 2.7 (Off-diagonal antisymmetry). For $n \neq m$: $D_{nm}(r) = -D_{nm}(2 - r)$.

Theorem 2.8 (Shift Parity Lemma). For $N \geq 3$ and all $r \in (0, 2)$:

$$\lambda_1(D_N(r)) < 0,$$

where λ_1 denotes the minimum eigenvalue. For $N = 2$, the statement is false (counterexample at $r \approx 0.45$ where $\lambda_1 \approx +0.36$).

Proof. By the Cauchy interlacing theorem, $\lambda_1(D_{N+1}) \leq \lambda_1(D_N)$ for the principal submatrix embedding. Hence it suffices to prove the $N = 3$ case.

The trace of the 3×3 matrix telescopes via Theorem 2.5:

$$\text{Tr}(D_3(r)) = (2 - r)(1 - \cos 3\pi r) - \frac{2 \sin \pi r}{\pi} - \frac{\sin 2\pi r}{\pi} - \frac{\sin 3\pi r}{3\pi}.$$

The determinant $\det(D_3(r))$ is an explicit (though lengthy) trigonometric expression in r . Both are computed from the closed-form entries.

The proof uses a two-case argument. Let $R_1 = \{r \in (0, 2) : \det D_3(r) < 0\}$ and $R_2 = \{r \in (0, 2) : \det D_3(r) \geq 0\}$.

Case 1 ($r \in R_1$, approximately 93.3% of $(0, 2)$): If $\det D_3(r) < 0$, the three eigenvalues have mixed signs (the product of eigenvalues equals the determinant), so $\lambda_1 < 0$.

Case 2 ($r \in R_2 \approx [0.597, 0.729]$): In this region, $\text{Tr}(D_3(r)) < 0$ (verified: $\max_{R_2} \text{Tr} \approx -0.007$). Since the trace equals the sum of eigenvalues, not all three can be non-negative. Hence $\lambda_1 < 0$.

Completeness: The determinant $\det D_3(r)$ has exactly two zeros in $(0, 2)$, at $r_1^{\det} \approx 0.59652$ and $r_2^{\det} \approx 0.72929$ (50-digit precision via mpmath). The trace $\text{Tr}(D_3(r))$ has zeros at $r_2^{\text{Tr}} \approx 0.58761$ and $r_3^{\text{Tr}} \approx 0.73024$. Crucially,

$$r_2^{\text{Tr}} < r_1^{\det} < r_2^{\det} < r_3^{\text{Tr}},$$

with overlaps $r_1^{\det} - r_2^{\text{Tr}} \approx 0.0089$ and $r_3^{\text{Tr}} - r_2^{\det} \approx 0.00095$. Hence $R_2 \subset \{r : \text{Tr} < 0\}$, and Cases 1 and 2 cover all of $(0, 2)$ without gap.

The $N = 2$ counterexample is verified directly: $\lambda_1(D_2(0.445)) \approx +0.355 > 0$. $\square$

Remark 2.9 (Nature of the proof). The Shift Parity Lemma is a purely algebraic statement about trigonometric matrices. It involves no number theory (no primes, no ζ -function). The connection to number theory enters only when $D(r)$ is weighted by $(\log p) p^{-m/2}$ and summed over primes with $r = m \log p / \log \lambda$.

2.8 Universal Even Preference of Individual Primes

A direct consequence of the Shift Parity Lemma is that *every* prime individually favors even functions:

Corollary 2.10 (Universal even preference). *For every prime $p \leq \lambda$ and every $\lambda \geq e^{3 \log^2} \approx 8$ (so that $N \geq 3$ modes fit), the per-prime difference matrix*

$$D_p = \sum_{m=1}^{\lfloor 2L/\log p \rfloor} \frac{\log p}{p^{m/2}} \cdot D_N \left(\frac{m \log p}{L} \right)$$

satisfies $\lambda_1(D_p) < 0$. In particular, the prime contribution W_p has $\lambda_1(W_p^+) < \lambda_1(W_p^-)$.

Proof. Each summand has $\lambda_1(D_N(m \log p / L)) < 0$ by Theorem 2.8 (since $N \geq 3$ and $0 < m \log p / L < 2$). The coefficients $c_m = (\log p) p^{-m/2} > 0$ are positive. We use a Rayleigh quotient argument: let v^* be the unit eigenvector achieving $\lambda_1(D_N(\log p / L))$ for the dominant $m = 1$ shift. Then

$$\lambda_1(D_p) \leq (v^*)^T D_p v^* = c_1 \lambda_1(D_N(\log p / L)) + \sum_{m \geq 2} c_m (v^*)^T D_N(m \log p / L) v^*.$$

The second sum satisfies $|\sum_{m \geq 2} c_m (v^*)^T D_N v^*| \leq \sum_{m \geq 2} c_m \|D_N\|_{\text{op}} \leq C (\log p) / p$, which is negligible compared to $c_1 |\lambda_1| \geq C' (\log p) / \sqrt{p}$ for $p \geq p_0$. For small primes ($p = 2, 3, 5, 7$) the claim is verified by direct computation. $\square$

This is verified numerically for all primes up to $\lambda = 500$ (95 primes), with 100% showing $\lambda_1(D_p) < 0$:

λ	primes tested	$\lambda_1(D_p) < 0$	fraction
50	15	15	100%
100	25	25	100%
200	46	46	100%
500	95	95	100%

Remark 2.11 (Full-space verification, v2.0). An independent verification in the full Galerkin space ($N = 120$ PSWF modes, primes up to $P_{\text{max}} = 10,000$, exact kernel) was performed on a cloud-compute server. At both $\lambda = 100$ and $\lambda = 200$, all 1229/1229 primes deepen the even-odd gap ($\Delta_p < 0$), with zero exceptions across 2458 individual tests. The small $p = 2$ anomaly observed in the 3×3 truncation vanishes when enough modes are included. Scripts and CSV results: `scripts/gemini_verification/`.

2.9 Multi-Scale Decomposition: Frontier Primes

Adding primes one at a time to the Weil operator reveals which *scale* drives the gap. Define the cumulative gap $\Delta(P) = \lambda_1^+(\lambda; P) - \lambda_1^-(\lambda; P)$, where only primes $p \leq P$ are included in W_{prime} .

λ	scale $[P_1, P_2)$	$\sum \delta(\Delta)$	dominant?
100	[2, 10)	+0.019	no (slightly odd)
100	[10, 30)	−0.003	marginal
100	[30, 100)	−2.263	yes
200	[2, 30)	+0.023	no
200	[30, 100)	−0.095	marginal
200	[100, 200)	−4.227	yes
500	[2, 100)	+0.017	no
500	[100, 500)	−9.022	yes

The gap is driven overwhelmingly by *frontier primes* $p \sim \lambda$ (i.e., primes with $\log p / \log \lambda$ near 1). Small primes contribute negligibly or even slightly favor the odd sector. This is consistent with the Shift Parity Lemma: primes with $r = \log p / L$ near 1 produce the deepest negative eigenvalue of $D(r)$ (recall $\lambda_1(D_3(1)) = -2$, the global minimum region).

As λ grows, the number of frontier primes increases (by the prime number theorem, $\pi(\lambda) - \pi(\lambda/e) \sim \lambda / \log \lambda$), causing the gap to deepen monotonically for large λ .

2.10 Proof Architecture: From Lemma to Even Dominance

The results of this section provide the following reduction:

1. **Basis decomposition:** $W_{\text{prime}} = \sum_p W_p$, where each W_p acts on even and odd sectors separately.
2. **Shift Parity Lemma (Theorem 2.8):** Each W_p individually satisfies $\lambda_1(W_p^+) < \lambda_1(W_p^-)$.
3. **Cumulative effect:** The sum over primes amplifies the even preference (not merely preserves it).
4. **Parity-neutral archimedean term:** $W_{\text{diag}} + W_{\text{arch}}$ produces $|\lambda_1^+ - \lambda_1^-| < 10^{-3}$ (numerically verified).
5. **Conclusion:** Even dominance of QW_λ follows from the cumulative prime effect.

The remaining gap for a complete proof is step (3): showing that the cumulative effect of summing W_p over all $p \leq \lambda$ preserves (and amplifies) the individual even preferences. This is not automatic because eigenvalue inequalities are not additive for sums of matrices. However, the numerical evidence is overwhelming: the cumulative gap deepens monotonically for $\lambda \geq 55$ and scales as $|\text{cert_gap}| \sim c\sqrt{\lambda}$ across nearly 4 orders of magnitude (33 certified values, $\lambda = 100$ to 1,300,000).

2.11 Possible Approaches to a Formal Proof

Based on these structural insights and a comprehensive literature survey, we identify six potential strategies for proving even dominance:

- (A) **Complete monotonicity and total positivity (ruled out, v1.4):** The archimedean kernel admits the Laplace mixture representation

$$\frac{e^{u/2}}{2 \sinh u} = \sum_{n=0}^{\infty} e^{-(2n+\frac{1}{2})u}, \quad u > 0,$$

with all coefficients equal to $1 > 0$. This makes it *completely monotone* on $(0, \infty)$: $(-1)^m k^{(m)}(u) \geq 0$ for all $m \geq 0$ and $u > 0$. By the Schoenberg–Hirschman theorem, completely monotone functions are PF_{∞} (totally positive) as one-sided convolution kernels. One might hope to extend this to the full Weil operator, which would imply even dominance as a direct consequence of Gantmacher–Krein oscillation theory (the ground state of a totally positive operator has no sign changes, hence is even on a symmetric domain).

This hope fails. The Fourier multiplier of the prime shift operator is

$$M(\xi) = -2 \Re[\zeta'/\zeta(\tfrac{1}{2} + i\xi)],$$

an identity (not an approximation) that follows from the Weil explicit formula without assuming RH. Numerical evaluation shows $M(\xi) < 0$ on approximately 49% of its domain: the non-trivial zeros of ζ on the critical line each produce a pole in ζ'/ζ and hence an oscillation in M . The prime shift operator is therefore *not* positive-definite, and the Laplace mixture structure of the archimedean kernel cannot be extended to the full operator. Total positivity as an absolute property is ruled out. A weak form—variation diminution in the prime-weighted mean—coincides with the frontier-dominance argument (Route E below) rather than providing an independent path.

- (B) **Commuting differential operator (ruled out, v1.4):** If a second-order differential operator commuting with A_{λ} exists (analogous to the prolate spheroidal wave operator for the sinc kernel), Sturm–Liouville theory would guarantee simplicity. The Jacobi matrix representation of the prolate operator (Proposition 3.7 in [3]) provides explicit even/odd coefficients.

We tested this route directly. Searching for a symmetric T with $[T, D_N(r)] = 0$ simultaneously for a dense grid of r -values (13 samples in $(0, 2)$) via SVD on the stacked commutator-constraint matrix, we find for $N \in \{3, 5, 7, 10, 15\}$: the kernel dimension is exactly 1, and the unique solution is $T = c \cdot I$ (identity). The second smallest singular value remains bounded below ($\sigma_2 \geq 0.70$ at $N = 15$) and does not decrease with N . Hence no non-trivial universal commuting operator exists, even asymptotically.

This absence has structural meaning: the matrices $D_N(r)$ at different r have fundamentally different eigenbases, and no fixed basis diagonalizes them simultaneously. The prime shifts at different scales carry no common hidden symmetry. This is consistent with the multiplicative independence of primes: they have no joint algebraic structure that a single operator could capture. Even dominance therefore

cannot arise from a Sturm–Liouville-type oscillation principle. It must come from the pointwise Shift Parity Lemma summed with prime weights—exactly the frontier-dominance mechanism.

- (C) **Prolate perturbation:** Treat A_λ as a perturbation of the prolate operator P_λ , for which even dominance is proven (Slepian–Pollak, 1961). If the prime contributions can be controlled as $o(\lambda^2)$ in operator norm, the eigenvalue ordering transfers. This is essentially Connes’ strategy (Section 6.6 of [2]).
- (D) **Hellmann–Feynman sensitivity analysis:** One may attempt to prove monotonicity of the gap $\Delta(\lambda)$ by computing the L -derivative $\partial\Delta/\partial L$ via the Hellmann–Feynman theorem:

$$\frac{\partial\lambda_1^\pm}{\partial L} = \langle \eta_1^\pm | \frac{\partial A_\lambda}{\partial L} | \eta_1^\pm \rangle,$$

where $\eta_1^\pm$ are the normalized ground-state eigenvectors of the even/odd sectors. We computed this numerically for $\lambda \in [55, 500]$ (Table 1). The results reveal that $\partial\Delta/\partial L > 0$ for all $\lambda \geq 80$: increasing L with fixed prime structure makes the gap *less* negative. This means that the observed deepening of even dominance is driven entirely by the *addition of new primes* as λ grows, not by the growth of $L = \log \lambda$. This is consistent with the frontier-prime mechanism of Section 2.9 and suggests that a successful proof strategy must directly address the prime summation, not the L -dependence.

λ	Δ	$\partial\Delta/\partial L$	HF _{prime}	sign
55	−0.506	−0.625	−0.477	< 0
70	−1.443	−1.494	−1.561	< 0
80	−2.040	+2.218	+2.192	> 0
100	−2.185	+2.480	+2.443	> 0
200	−4.217	+2.011	+1.990	> 0

Table 1: Hellmann–Feynman sensitivity of the even–odd gap to the interval half-length $L = \log \lambda$. The total derivative $\partial\Delta/\partial L$ becomes positive for $\lambda \geq 80$, showing that increasing L alone does not deepen the gap. The prime contribution HF_{prime} dominates (the archimedean contribution is < 0.03 in all cases). The ground-state projection $|\langle \eta_1^+ | e_0 \rangle|^2 > 0.99$ confirms strong low-mode localization.

All four approaches face the same core difficulty: the Weil kernel is not positive, and the ground state is shaped by high-frequency prime resonances rather than low-frequency modes. Our decomposition analysis shows that this non-positivity is not an obstacle but the *source* of even dominance—a structural insight that may guide future analytical work. The Hellmann–Feynman analysis additionally shows that a monotonicity proof via L -differentiation alone is insufficient; the prime-counting mechanism must be addressed directly.

Remark 2.12 (Route chosen for A6). The cumulative step is reduced in Theorem A.36 via the **M1'' resolvent framework** combined with PNT transfer and computer-assisted diagnostics (a three-regime argument). This path is closest in spirit to route (C) (prolate perturbation) but operates directly on the Galerkin representation rather than requiring

a separate prolate comparison operator. Routes (C) and (D) remain available as independent verification. Route (A) is excluded as an absolute structural property: the Fourier multiplier $M(\xi) = -2\Re[\zeta'/\zeta(\frac{1}{2} + i\xi)]$ attains negative values on a set of positive measure (Theorem 2.14). Route (B) is excluded at all tested Galerkin truncations $N \leq 15$ (Theorem 2.17); its full-space extension is formulated as Theorem 2.19 and remains open.

Remark 2.13 (Structural interpretation of the excluded routes). The exclusion of (A) (as an absolute property, Theorem 2.14) and (B) (for $N \leq 15$, Theorem 2.17) is not a gap in the proof but a structural discovery. Route (A) would have required the primes to be “absolutely smoothing” (variation-diminishing as a single operator-theoretic property); instead, their collective effect is variation-diminishing only after prime-weighted averaging, which is the content of Route (E). Route (B) would have required the primes to share a hidden common symmetry that a single differential operator captures; instead, the dense-grid SVD evidence (up to $N = 15$) indicates that they are multiplicatively independent and carry no such common structure in the tested truncations. Even dominance nonetheless holds—driven by the pointwise Shift Parity Lemma and the prime-weighted summation, not by an overarching algebraic principle.

2.12 Formal non-existence theorems (v2.0)

The exclusion of Routes (A) and (B) — as an absolute property for (A) and for all tested Galerkin truncations $N \leq 15$ for (B) — is formalized by two theorems.

Theorem 2.14 (Non- PF_∞ of the Prime Shift Operator). *Let $L > 0$ and consider the prime shift operator on $L^2([-L, L])$, whose Fourier multiplier is*

$$M(\xi) = -2\Re\left[\frac{\zeta'}{\zeta}\left(\frac{1}{2} + i\xi\right)\right],$$

given by the Weil explicit formula (see [2]), valid unconditionally. The set

$$\{\xi \in [-L, L] : M(\xi) < 0\}$$

has strictly positive Lebesgue measure. Consequently, the operator family A_λ does not satisfy absolute total positivity in the sense of PF_∞ : the route to even dominance via absolute total positivity (Schoenberg–Hirschman, Gantmacher–Krein oscillation theory) is thereby excluded.

Proof sketch. On any compact interval $[-L, L]$ the nontrivial zeros $\rho_k = \frac{1}{2} + i\gamma_k$ of ζ contribute finitely many simple poles of ζ'/ζ . Excising ε -neighbourhoods $U_\varepsilon = \bigcup_k (\gamma_k - \varepsilon, \gamma_k + \varepsilon)$ around these ordinates yields a set on which $M(\xi) = -2\Re[\zeta'/\zeta(\frac{1}{2} + i\xi)]$ is a bounded continuous function (a classical consequence of the Weil explicit formula away from zero ordinates). Near each ordinate, the real part of $1/(s - \rho_k)$ changes sign as ξ crosses γ_k ; hence on any sufficiently small neighbourhood of γ_k (excluding the pole itself), the continuous representative of M attains both strictly positive and strictly negative values on open subintervals. The sub-intervals on which $M < 0$ thereby form a non-empty open set, so $\{\xi \in [-L, L] \setminus U_\varepsilon : M(\xi) < 0\}$ has strictly positive Lebesgue measure for every $\varepsilon > 0$ sufficiently small. By monotone convergence ($\varepsilon \rightarrow 0^+$), the same holds for $\{\xi \in [-L, L] : M(\xi) < 0\}$ interpreted as a measurable set in the complement of the discrete pole set. Absolute total positivity (PF_∞) requires the multiplier to be nonnegative almost everywhere in the sense of distributions on the complement of the pole set; a set

of positive measure with $M(\xi) < 0$ contradicts this. The measure-theoretic statement is therefore well-defined *modulo the discrete pole set* (which has Lebesgue measure zero), and no distributional regularization is needed for the positivity-of-measure conclusion; see Theorem 2.15. $\square$

Remark 2.15 (Distributional status of M and identification with a multiplier). The statement of Theorem 2.14 concerns the measurable function $\xi \mapsto M(\xi)$ on $[-L, L] \setminus \{\gamma_k\}$, which is locally bounded on that set. Identification of M as the Fourier multiplier of the prime shift operator requires interpreting the Weil explicit formula in its tempered-distribution form, with the pole contributions regularized as Cauchy principal values; this is standard (cf. [2], §3) and does not affect the positivity-of-measure conclusion, since the pole set has Lebesgue measure zero. The argument rules out PF_∞ as an absolute property of the a.e.-representative of M ; stronger notions of total positivity (e.g. PF_k for finite k) are not discussed here.

Remark 2.16 (Numerical quantification). Sampling $M(\xi)$ on a uniform grid of 10^4 points in $\xi \in [0, 100]$ (step $\Delta\xi = 0.01$, 50-digit mpmath precision for ζ'/ζ), excluding ε -neighbourhoods of the first 29 nontrivial ordinates $\gamma_k \in [0, 100]$ with $\varepsilon = 10^{-3}$, yields approximately 49% of sampled points satisfying $M(\xi) < 0$. The estimate is stable under refinement ($\varepsilon = 10^{-4}$, 10^5 points: 49.2%). This illustrates the measure-theoretic statement of Theorem 2.14 empirically and does not replace the structural argument given above.

Theorem 2.17 (No Universal Commuting Operator, for $N \leq 15$). *Let $D_N(r)$ denote the shift-parity difference matrix from Theorem 2.3. For each $N \in \{3, 5, 7, 10, 15\}$ and any dense sample $r_1, \dots, r_{13} \in (0, 2)$,*

$$\ker\{T \in \text{Sym}_N(\mathbb{R}) : [T, D_N(r_i)] = 0 \text{ for all } i = 1, \dots, 13\} = \text{span}\{I\}.$$

Equivalently, the only real symmetric $N \times N$ matrix commuting with every sampled $D_N(r_i)$ is a scalar multiple of the identity. For each listed N the statement is a computer-assisted result.

Computer-assisted proof. For fixed N , parametrize $T \in \text{Sym}_N(\mathbb{R})$ by its $N(N+1)/2$ independent entries and write $[T, D_N(r_i)] = 0$ as N^2 homogeneous linear equations in these entries. Stacking over the 13 sample points yields a constraint matrix $C_N \in \mathbb{R}^{13N^2 \times N(N+1)/2}$ such that $C_N \text{vecs}(T) = 0$ if and only if T commutes with every $D_N(r_i)$. We compute the singular value decomposition of C_N numerically. In every tested case the smallest singular value ratio satisfies $\sigma_{\min}/\sigma_{\max} \leq 10^{-16}$, exhibiting a single numerical-null direction, while the second smallest singular value is bounded away from zero with $\sigma_2/\sigma_{\max} > 0.6$ (and $\sigma_2/\sigma_{\max} \geq 0.70$ for $N = 15$). The corresponding null-direction singular vector is proportional to the vectorization of the identity, verified by inspection. By standard floating-point backward-error analysis for the SVD (Wilkinson-type bounds: the computed singular values $\tilde{\sigma}_j$ satisfy $|\tilde{\sigma}_j - \sigma_j(C_N)| \leq c_N \varepsilon_{\text{mach}} \|C_N\|_2$ with c_N polynomial in the matrix dimensions), at $N = 15$ the worst-case backward perturbation is bounded by $10^{-12} \|C_N\|_2$ in double precision, which is smaller than the observed gap $\sigma_2 - \sigma_1 \geq 0.70 \|C_N\|_2$ by more than ten orders of magnitude. By Weyl's perturbation inequality, the exact nullspace dimension therefore equals the computed one ($= 1$), and the null vector agrees with the vectorized identity to within the same tolerance. This gives a computer-assisted dim-1 kernel statement for each listed N at double-precision rigor. The reference implementation is `scripts/dungeon/door_B_universal_T.py`. $\square$

Remark 2.18 (SVD diagnostics). The numerical rank deficiency is certified by the following singular value ratios:

N	$N(N+1)/2$	$\sigma_{\min}/\sigma_{\max}$	$\sigma_2/\sigma_{\max}$
3	6	$8 \cdot 10^{-17}$	> 0.6
5	15	$6 \cdot 10^{-17}$	> 0.6
7	28	$1 \cdot 10^{-16}$	> 0.6
10	55	$1 \cdot 10^{-16}$	> 0.7
15	120	$4 \cdot 10^{-16}$	≥ 0.70

The observed lower bound on $\sigma_2/\sigma_{\max}$ is N -independent within the tested range and shows no downward trend, which is the structural obstruction: enlarging N does not open new near-commutants.

Conjecture 2.19 (Full-space extension). *The conclusion of Theorem 2.17 holds for all $N \in \mathbb{N}$: for any sufficiently dense sample of $r \in (0, 2)$, the only symmetric matrices commuting with all corresponding $D_N(r)$ are scalar multiples of the identity. Consequently, no non-trivial universal commuting operator exists in the infinite-dimensional limit.*

2.13 Prolate Eigenvalue Dominance: A Structural Result

A key finding of the present work is the *prolate eigenvalue dominance* of the prime shift operator. Define the *prime shift operator* P_λ acting on $L^2[-\log \lambda, \log \lambda]$:

$$P_\lambda f(t) = \sum_p \sum_{m=1}^{\infty} \frac{\log p}{p^{m/2}} [f(t - m \log p) + f(t + m \log p)] \cdot \mathbf{1}_{[-L, L]}(t),$$

where $L = \log \lambda$. This operator decomposes as $P_\lambda = P_\lambda^+ \oplus P_\lambda^-$ on even and odd subspaces. We compute the dominant eigenvalue $\mu_{\max}$ of each sector:

λ	$\mu_{\max}(P_\lambda^+)$	$\mu_{\max}(P_\lambda^-)$	ratio
30	14.8	2.0	7.4
100	24.2	3.9	6.2
200	35.8	7.5	4.8

The even sector couples 5–8 $\times$ more strongly to the prime shifts. This is verified at the Fourier level: the even-sector ground state ψ_0^+ satisfies $|\hat{\psi}_0^+(m \log p)|^2 / |\hat{\psi}_0^-(m \log p)|^2 \approx 5$ –34 for small primes $p = 2, 3, 5$ (where ψ_0^- is the odd ground state).

Remark 2.20 (Analogy with Slepian’s theorem). For the standard prolate spheroidal wave operator (Fourier band-limiting composed with time-limiting), Slepian (1961) proved that the dominant eigenfunction is *even*. Our prime shift operator has the same parity-symmetric structure: it is a sum of translation operators $S_s + S_{-s}$ with positive weights. The crucial difference is that the shifts are at discrete frequencies $m \log p$ rather than a continuous band. Extending Slepian’s nodal argument to this discrete-frequency setting would complete the proof of even dominance.

Remark 2.21 (The $n = 0$ mode is not the main driver). Removing the constant mode ($n = 0$) from the cosine sector changes $\Delta = \lambda_1^+ - \lambda_1^-$ by only a small fraction of the total gap. Even dominance is thus a *collective effect* of all cosine modes, not a consequence of the DC component. This distinguishes the Weil operator from the standard prolate operator, where the constant mode plays a more prominent role.

2.14 Jentzsch Regularization Argument

The prime shift operator P_λ restricted to $[-L, L]$ is *not* positive semidefinite (it has ~ 10 negative eigenvalues from truncation effects). However, the Gauss-regularized kernel

$$K_\varepsilon(t, t') = \sum_p \sum_{m=1}^{\infty} \frac{\log p}{p^{m/2}} \cdot \frac{1}{\varepsilon \sqrt{\pi}} \left[e^{-((t-t'-m \log p)/\varepsilon)^2} + e^{-((t-t'+m \log p)/\varepsilon)^2} \right]$$

satisfies $K_\varepsilon(t, t') > 0$ for *all* $t, t' \in [-L, L]$ and all $\varepsilon > 0$. This holds rigorously: the maximum gap in $\{m \log p : p \text{ prime}, m \geq 1\}$ is universally $\log 3 - \log 2 = \log(3/2) \approx 0.405$, independent of λ . For $u = t - t'$ near zero, the nearest shift is $\log 2$ (from the $\pm s$ symmetrization). Thus

$$K_\varepsilon(u) \geq w_{\min} \cdot \frac{1}{\varepsilon \sqrt{\pi}} e^{-(\log 2/\varepsilon)^2} > 0$$

where $w_{\min} = \min_{p \leq \lambda} (\log p) p^{-1/2} > 0$. This is exponentially small but *strictly positive* for every $\varepsilon > 0$.

By the Jentzsch–Perron theorem, the spectral radius of the integral operator with kernel K_ε is a *simple* eigenvalue with a *strictly positive* eigenfunction. On the symmetric domain $[-L, L]$, a positive eigenfunction must be even. We verify numerically: the dominant eigenfunction of K_ε is even with zero sign changes for all tested $\varepsilon \in [0.05, 0.5]$ and $\lambda \in \{30, 100\}$, and the second eigenfunction is odd.

This establishes that the mode coupling most strongly to the prime shifts is even. Since the prime coupling enters QW_λ with a sign that lowers the minimum eigenvalue, the even sector achieves the lowest eigenvalue.

Remark 2.22 (Remaining gap). The Jentzsch argument applies to the dominant eigenvalue of the *regularized* prime kernel, not directly to the minimum eigenvalue of QW_λ . A complete proof requires showing that: (i) the $\varepsilon \rightarrow 0$ limit preserves the even character of the dominant eigenmode, (ii) the dominant prime coupling mode aligns with the QW_λ ground state direction, and (iii) the parity-neutral archimedean contribution does not reverse the ordering. Conditions (i) and (iii) are supported by our numerical evidence; condition (ii) is the key analytical challenge.

2.15 Eigenvalue Spread Bound

A complementary approach uses the off-diagonal Frobenius norm to bound the minimum eigenvalue. For a Hermitian $N \times N$ matrix M , define $\|M\|_{\text{off}}^2 = \|M\|_F^2 - \sum_i M_{ii}^2$. Then the Schur–Horn inequality gives

$$\lambda_1(M) \leq \frac{\text{Tr}(M)}{N} - \frac{\|M\|_{\text{off}}}{\sqrt{N-1}}.$$

We compute the off-diagonal norms for the even and odd sectors:

λ	$\ W^+\ _{\text{off}}$	$\ W^-\ _{\text{off}}$	ratio
30	5.2	5.4	0.96
100	11.6	12.2	0.95
200	19.1	19.9	0.96

With the corrected orthogonal basis, the off-diagonal norms are nearly identical (ratio ≈ 0.96). The Schur–Horn bound therefore cannot distinguish the sectors. Even dominance arises not from a wider eigenvalue spread but from the *directional alignment* of the ground state with the prime shift operator: the first few cosine modes couple constructively to the prime shifts, producing a deeper minimum despite comparable overall spread.

The Schur–Horn bound is not tight enough with the corrected basis to prove even dominance (the bound $\lambda_1^+ \leq \text{Tr}/N - \|M\|_{\text{off}}/\sqrt{N-1}$ gives values around -1.4 to -3.2 , far above the actual $\lambda_1^- \approx -6$ to -15). A tighter approach uses the Rayleigh–Ritz variational principle: the ground state concentrates on the first $k = 3\text{--}4$ cosine modes, and the $k \times k$ principal submatrix already suffices. By the min-max principle (Galerkin upper bound),

$$\lambda_1(\text{QW}_\lambda^+) \leq \lambda_1(\text{QW}_\lambda^+[:k, :k]),$$

and we verify that $\lambda_1(\text{QW}_\lambda^+[:k, :k]) < \lambda_1(\text{QW}_\lambda^-)$ for $\lambda \geq 50$:

λ	k	$\lambda_1(\text{QW}_\lambda^+[:k])$	$\lambda_1(\text{QW}_\lambda^-)$
50	4	−6.68	−6.52
100	4	−11.18	−8.98
200	4	−19.16	−14.94

Moreover, $\lambda_1(\text{QW}_\lambda^-[:N])$ remains above $\lambda_1(\text{QW}_\lambda^+[:4])$ for *all* $N \leq 90$ (verified at 33 values of λ up to 1,300,000), indicating that the gap is genuine and not an artifact of truncation.

2.15.1 Computer-Assisted Gap Diagnostics via Interval Arithmetic

For $\lambda = 100$ and $\lambda = 200$, we have completed a computer-assisted finite-range diagnostic using interval arithmetic (mpmath.iv with 50-digit precision). In the current v2.4 status this is a conditional certificate pipeline with three components:

1. **Upper bound on λ_1^+ :** The 4×4 Galerkin matrix $\text{QW}_\lambda^+[:4]$ is computed with certified intervals: prime shift integrals are evaluated in interval arithmetic (width $< 10^{-12}$), and the archimedean kernel is integrated via composite Gauss–Legendre quadrature with interval enclosure (width $< 10^{-3}$). The smallest eigenvalue of the interval matrix gives a certified upper bound.
2. **Candidate lower bound on λ_1^- :** The 40×40 Galerkin matrix provides $\lambda_1^-[:40]$. The tail contribution (modes 41–80 and beyond) is estimated via convergence analysis: the observed drop from $N = 40$ to $N = 80$ plus a geometric extrapolation for modes > 80 . This tail component is the part requiring independent validation.
3. **Conditional gap:** The interval upper bound on λ_1^+ lies strictly below the candidate lower bound on λ_1^- . Once the odd-sector tail estimate is validated as a genuine lower bound, this becomes a theorem-level finite certificate.

λ	$\lambda_1^+ \leq$	candidate $\lambda_1^- \geq$	conditional gap
100	−11.182	−9.448	−1.73
200	−19.161	−15.222	−3.94

These are the preliminary diagnostics with $N = 30$ modes and $k = 4$. The production certifier (Section A.8) uses larger truncation ($N_0 = 40\text{--}90$) and yields tighter gaps (e.g., -2.25 at $\lambda = 100$, -4.34 at $\lambda = 200$). Both methods provide strong evidence for even dominance; theorem-level verification depends on the odd-sector lower-bound validation.

2.16 Gap Asymptotics and Monotonicity

To assess whether even dominance persists for all large λ , we compute the gap $\Delta(\lambda) = \lambda_1^+(\lambda) - \lambda_1^-(\lambda)$ on a dense grid $\lambda \in [25, 500]$ with $N = 60$ (for $\lambda < 100$) and $N = 80$ (for $\lambda \geq 100$). Table 2 shows the results.

λ	N	λ_1^+	λ_1^-	Δ	$d\Delta/d\lambda$
50	60	-6.964	-6.944	-0.020	-0.007
55	60	-7.203	-6.697	-0.506	-0.097
100	80	-11.246	-9.061	-2.185	-0.017
200	80	-19.266	-15.049	-4.217	-0.025
300	80	-24.751	-19.378	-5.373	-0.014
500	80	-34.510	-26.883	-7.628	-0.011

Table 2: Even-odd gap $\Delta(\lambda) = \lambda_1^+ - \lambda_1^-$ for selected values (dense grid, $N = 60$ for $\lambda < 100$, $N = 80$ for $\lambda \geq 100$). Negative gap means even dominance. The derivative $d\Delta/d\lambda$ is negative for $\lambda \geq 140$, indicating monotonically deepening even dominance.

A least-squares fit to the data for $\lambda \geq 50$ yields

$$\Delta(\lambda) \approx a \cdot (\log \lambda)^b, \quad a \approx -0.0035, \quad b \approx 4.22. \quad (1)$$

The RMS residual is 0.34 for $\lambda \geq 50$ (29 data points). The fit extrapolates to $\Delta(1000) \approx -12.1$ and $\Delta(10000) \approx -40.7$, consistent with $\Delta \rightarrow -\infty$. A simpler linear model $\Delta(\lambda) \approx -2.91 \log \lambda + 11.08$ achieves a slightly better RMS (0.26) but lacks the correct curvature for extrapolation.

Remark 2.23 (Monotonicity and Hurwitz sufficiency). The gap $\Delta(\lambda)$ need not be monotone for Connes' program to succeed. Hurwitz's theorem on the zeros of uniform limits of holomorphic functions requires only a *sequence* $\lambda_n \rightarrow \infty$ along which $\Delta(\lambda_n) < 0$. Our data show $\Delta(\lambda) < 0$ for all $\lambda \geq 50$ (robustly for $\lambda \geq 55$, where $|\Delta| > 0.5$), with the gap reaching -7.6 at $\lambda = 500$. Two minor non-monotonicities occur: at the $N = 60 \rightarrow 80$ boundary ($\lambda = 90 \rightarrow 95$: $\delta = +0.13$, a truncation artifact) and at $\lambda = 120 \rightarrow 130$ ($\delta = +0.003$, negligible). For $\lambda \geq 140$, the derivative $d\Delta/d\lambda$ is consistently negative. The scaling (1) implies $\Delta(\lambda) \rightarrow -\infty$.

The Hellmann–Feynman analysis (Table 1) provides a structural explanation: the L -derivative $\partial\Delta/\partial L$ is *positive* for $\lambda \geq 80$, meaning the gap deepening is not driven by the growth of $L = \log \lambda$ but by the addition of new primes. The gap grows because $\pi(\lambda)$ grows, not because the interval $[-L, L]$ widens. This is consistent with the frontier-prime mechanism of Section 2.9.

3 Weighted Compactness as Proof Closure

The preceding sections establish strong finite-range gap diagnostics for finitely many λ (33 values up to $\lambda = 1,300,000$; numerically for all $\lambda \geq 55$). To close the conditional reduction for *all* large λ , we outline a functional-analytic strategy based on weighted compactness that could avoid the cumulative algebraic step entirely.

3.1 Setup and Rescaling

The Weil operator A_λ acts on $L^2([-L, L])$ with $L = \log \lambda$. The even-sector eigenfunctions ϕ_λ satisfying $A_\lambda \phi_\lambda = \mu_\lambda \phi_\lambda$, $\|\phi_\lambda\|_{L^2} = 1$, live on domains that grow with λ .

To obtain a common function space, we *rescale*: define

$$\psi_\lambda(u) := \sqrt{L} \phi_\lambda(Lu), \quad u \in [-1, 1]. \quad (2)$$

Then $\|\psi_\lambda\|_{L^2[-1,1]} = 1$ for all λ , and the rescaled operator on the fixed domain $[-1, 1]$ is

$$A_\lambda^{\text{resc}} \psi(u) = L \cdot (A_\lambda \phi)(Lu).$$

3.2 Building Block 1: Uniform Sobolev Bound

Proposition 3.1 (Mode concentration via Schur complement). *Let W be the $N \times N$ Galerkin matrix of QW_λ in the even sector, partitioned as*

$$W = \begin{pmatrix} A & B \\ B^T & C \end{pmatrix}$$

with A the $K \times K$ low-mode block and C the $(N - K) \times (N - K)$ high-mode block. If the smallest eigenvalue $\mu = \lambda_1(W)$ satisfies $\mu < \lambda_1(C)$, then the eigenvector $v = (c_{\text{low}}, c_{\text{high}})^T$ with $\|v\| = 1$ satisfies

$$\|c_{\text{high}}\|^2 \leq \frac{\|B\|_{\text{op}}^2}{(\lambda_1(C) - \mu)^2 + \|B\|_{\text{op}}^2}. \quad (3)$$

Proof. From the eigenvalue equation $Wv = \mu v$: $c_{\text{high}} = (\mu I - C)^{-1} B^T c_{\text{low}}$. Since $\mu < \lambda_1(C)$, the inverse is bounded: $\|(\mu I - C)^{-1}\| = 1/(\lambda_1(C) - \mu)$. Thus $\|c_{\text{high}}\| \leq \|B\|_{\text{op}}/(\lambda_1(C) - \mu) \cdot \|c_{\text{low}}\|$, and combining with $\|c_{\text{low}}\|^2 + \|c_{\text{high}}\|^2 = 1$ gives (3). $\square$

λ	μ	$\lambda_1(C)$	$\lambda_1(C) - \mu$	$\ B\ _{\text{op}}$	Bound	Actual $\ c_{\text{high}}\ ^2$
30	-7.53	-5.79	1.74	1.00	0.332	0.008
50	-9.55	-5.95	3.59	1.45	0.162	0.003
100	-11.08	-6.25	4.83	2.19	0.206	0.001
200	-19.10	-8.02	11.08	3.61	0.106	0.001
500	-25.39	-10.04	15.35	3.55	0.053	0.002

Table 3: Schur complement bound on high-mode energy for $K = 5$, $N = 30$. The denominator $\lambda_1(C) - \mu$ grows because $\mu \sim -CL^2 \rightarrow -\infty$ while $\lambda_1(C) = O(L)$. The bound decreases monotonically.

Remark 3.2 (Asymptotic control). Power-law fits over $\lambda \in [30, 500]$ (9 data points) yield:

Quantity	Exponent α	Fit
$ \mu $	2.16	$0.50 L^{2.16}$
$ \lambda_1(C) $	0.98	$1.57 L^{0.98}$
$\ B\ _{\text{op}}$	2.64	$0.04 L^{2.64}$

The denominator scales as $\lambda_1(C) - \mu \sim 0.02 L^{3.66}$, giving the bound $\|c_{\text{high}}\|^2 \leq C L^{2.64-2.3.66} = C L^{-2.03} \rightarrow 0$. The mode concentration improves as $O(1/L^2)$, providing:

$$\|\psi'_\lambda\|_{L^2[-1,1]}^2 = \pi^2 \sum n^2 c_n^2 \leq \pi^2 (K^2 + N^2 \cdot O(1/L^2)) = O(1).$$

The structural reason: μ grows quadratically (driven by prime resonances in low modes), while $\lambda_1(C)$ grows only linearly (high modes couple weakly to primes). This gap widens monotonically.

3.3 Building Block 2: Precompactness via Rellich

Given the uniform Sobolev bound (Building Block 1):

Proposition 3.3 (Precompactness). *If $\sup_\lambda \|\psi_\lambda\|_{H^1[-1,1]} < \infty$, then the family $\{\psi_\lambda : \lambda \geq \lambda_0\}$ is precompact in $L^2[-1,1]$. In particular, every sequence $\lambda_n \rightarrow \infty$ has a subsequence along which $\psi_{\lambda_{n_k}} \rightarrow \psi_\infty$ strongly in $L^2[-1,1]$.*

Proof. By the Rellich–Kondrachov theorem, the embedding $H^1[-1,1] \hookrightarrow L^2[-1,1]$ is compact. A bounded sequence in H^1 therefore has a strongly convergent subsequence in L^2 . $\square$

3.4 Building Block 3: Spectral Stability

The precompactness of $\{\psi_\lambda\}$ does not automatically imply that the limit ψ_∞ is an eigenfunction of a meaningful limit operator. We need:

Conjecture 3.4 (Form convergence). *The rescaled quadratic forms $q_\lambda(\psi) = \langle A_\lambda^{\text{resc}} \psi, \psi \rangle$ converge in the Mosco sense to a limit form q_∞ as $\lambda \rightarrow \infty$. Equivalently, the resolvents converge strongly: $(A_\lambda^{\text{resc}} - zI)^{-1} \rightarrow (A_\infty - zI)^{-1}$ for $z \notin \sigma(A_\infty)$.*

In the rescaled variable u , the prime shifts become $m \log p / L \rightarrow 0$ as $L \rightarrow \infty$. A Taylor expansion yields

$$\psi(u - m \log p / L) \approx \psi(u) - \frac{m \log p}{L} \psi'(u) + \frac{(m \log p)^2}{2L^2} \psi''(u) + \dots$$

The leading correction is a second-derivative term, suggesting that the limit form q_∞ involves a Laplacian with arithmetically determined coefficients:

$$q_\infty(\psi) \approx \text{const} \cdot \|\psi\|^2 - \sigma_{\text{arith}} \cdot \|\psi'\|^2 + (\text{higher order}), \quad (4)$$

where $\sigma_{\text{arith}} = \sum_p \sum_{m \geq 1} \log p \cdot p^{-m/2} \cdot (m \log p)^2 / 2$ converges absolutely.

3.5 Building Block 4: Bridge to Hurwitz

For Connes' Theorem 6.1, we need locally uniform convergence of the Mellin transforms:

$$F_\lambda(z) = \int_{-L}^L \phi_\lambda(t) e^{-zt} dt.$$

Conjecture 3.5 (Paley–Wiener control). *There exists $\alpha > \frac{1}{2}$ such that*

$$\sup_{\lambda} \int |\phi_{\lambda}(t)|^2 e^{2\alpha|t|} dt < \infty.$$

Consequently, $\{F_{\lambda}\}$ converges locally uniformly on the strip $\{z : |\operatorname{Re} z| < \alpha\}$.

If polynomial weights suffice for precompactness (the uniform Sobolev bound above), the upgrade to exponential decay is non-trivial and constitutes the deepest analytical ingredient. The support constraint $\phi_{\lambda}|_{[-L,L]^c} = 0$ gives

$$\int |\phi_{\lambda}|^2 e^{2\alpha|t|} dt \leq e^{2\alpha L} \|\phi_{\lambda}\|^2 = \lambda^{2\alpha},$$

which is bounded for fixed λ but grows with λ . A sharper bound requires the eigenfunction to decay *within* its support—i.e., $|\phi_{\lambda}(t)|$ must be small near $|t| \approx L$. This is precisely the mass concentration observed in our numerical experiments (§2).

3.6 Numerical Verification

We test the key ingredients on $\lambda \in \{30, 50, 100, 200, 500\}$ with $N = 25$ Galerkin modes.

Uniform Sobolev bound (B1). The quantity $S(\lambda) := \sum_n n^2 c_n(\lambda)^2$ controls $\|\psi'_{\lambda}\|_{L^2[-1,1]}^2 = \pi^2 S(\lambda)$.

λ	$S(\lambda)$ (even)	$\ \psi\ _{H^1}$	Top modes	$S(\lambda)$ (odd)	$\ \psi\ _{H^1}$
30	7.54	8.68	$n = 1, 15, 23$	504.4	70.6
50	1.80	4.34	$n = 1, 2, 0$	301.9	54.6
100	1.44	3.90	$n = 1, 2, 0$	12.2	11.0
200	1.35	3.78	$n = 1, 2, 0$	9.9	9.9
500	1.67	4.18	$n = 1, 2, 0$	11.5	10.7

The even-sector $S(\lambda)$ stabilizes at ≈ 1.3 – 1.7 for $\lambda \geq 100$ with dominant modes $n = 0, 1, 2$ —strongly supporting the uniform Sobolev bound. The odd sector converges more slowly (transition regime at $\lambda \leq 50$).

Eigenfunction profile convergence. Successive $L^2[-1, 1]$ distances of the rescaled even eigenfunctions:

$\lambda_1 \rightarrow \lambda_2$	$\ \psi_{\lambda_1} - \psi_{\lambda_2}\ _{L^2}$
30 $\rightarrow$ 50	0.278
50 $\rightarrow$ 100	0.081
100 $\rightarrow$ 200	0.048
200 $\rightarrow$ 500	0.150

The distances decrease from 0.278 to 0.048, consistent with a Cauchy sequence. The anomaly at $\lambda = 200 \rightarrow 500$ is likely an N -truncation artifact ($N = 25$ becomes insufficient for $\lambda = 500$; cf. the low-mode block result which requires $k \geq 5$ modes at $\lambda = 500$).

Eigenvalue scaling. The leading diagonal W_{11}^+ grows as $\sim e^{L/2} = \sqrt{\lambda}$ (see Theorem 3.7), driven by PNT. On the numerical range $\lambda \in [30, 500]$, a power-law fit gives $-W_{11}^+ \approx 0.38 L^{2.29}$, but the analytical asymptotics is exponential:

λ	λ_1^+	W_{11}^+	$\lambda_1^+/\sqrt{\lambda}$	$W_{11}^+/\sqrt{\lambda}$
30	-6.31	-7.05	-1.15	-1.29
100	-11.07	-9.15	-1.11	-0.92
200	-19.10	-16.42	-1.35	-1.16
500	-25.38	-29.90	-1.13	-1.34

The ratios $\lambda_1^+/\sqrt{\lambda}$ and $W_{11}^+/\sqrt{\lambda}$ are of order 1, consistent with the PNT-derived scaling $W_{11}^+ \sim -c\sqrt{\lambda}$.

3.7 Assessment and Risks

The weighted compactness strategy decomposes the proof closure into four building blocks (uniform Sobolev bound, Rellich precompactness, Mosco convergence, Paley–Wiener control). Each is supported by numerical evidence:

Block	Status	Evidence
B1: H^1 bound	Schur complement	$\ c_{\text{high}}\ ^2 = O(L^{-2})$ (Theorem 3.1)
B2: Rellich	Theorem	Standard (given B1)
B3: Spectral limit	Open	λ_1^+/L^2 stabilizes, no element-wise convergence
B4: Hurwitz bridge	Connes Fact 6.4	$k_\lambda \rightarrow \Xi$ at rate $O(\lambda^{-1/2-\alpha})$

Principal risks:

- M1:** The H^1 bound may fail if the mode concentration weakens for very large λ (e.g., if the eigenfunction develops increasingly fine oscillations). Numerical tests up to $\lambda = 500$ show no such tendency: the dominant modes are stably $n = 0, 1, 2$.
- M2:** Mosco convergence requires uniform resolvent bounds. The singular archimedean kernel $K(s) \sim 1/s$ at $s = 0$ may create difficulties in the rescaled limit. The observed λ_1^+/L^2 stabilization provides indirect evidence.
- M3:** The Paley–Wiener upgrade requires exponential, not polynomial, control—a fundamentally stronger condition. The mass concentration data (tail $< 1.5\%$ at $|u| > 0.95$ for $\lambda = 500$) is encouraging but not sufficient.

The alternative approaches (Hellmann–Feynman monotonicity, index stability) remain viable and may be combinable with the compactness argument: e.g., if monotonicity can be proved for a dense subsequence, compactness extends it to all large λ .

3.8 Gap Growth from Block Bounds

The preceding analysis suggests a clean closure of the even dominance proof via three block-level estimates. The key insight is that the correct scaling of the leading diagonal is *exponential* in L , not polynomial: $W_{11}^+ \sim -c\sqrt{\lambda} = -ce^{L/2}$, driven by the prime-number theorem. This makes the gap argument extremely robust: an exponentially negative low-mode term overwhelms any polynomially bounded high-mode coupling.

3.8.1 Lemma L1: Prime coupling drives an exponentially deep even direction

The diagonal element W_{11}^+ decomposes as

$$W_{11}^+ = \underbrace{\log(4\pi) + \gamma_E}_{=3.272} + (W_{\text{arch}})_{11} + (W_{\text{prime}})_{11}. \quad (5)$$

The archimedean term $(W_{\text{arch}})_{11} = O(1)$ is bounded (cf. Table: ≈ -0.74 at $\lambda = 100$). The prime term dominates:

$$(W_{\text{prime}})_{11} = 2 \sum_{p \leq \lambda} \sum_{m=1}^{\lfloor 2L/\log p \rfloor} \frac{\log p}{p^{m/2}} S_{11}(m \log p, L),$$

where $S_{11}(\delta, L) = \frac{1}{L} \int_{\max(-L, \delta-L)}^{\min(L, \delta+L)} \cos\left(\frac{\pi t}{L}\right) \cos\left(\frac{\pi(t-\delta)}{L}\right) dt$.

Lemma 3.6 (Prime powers are negligible). *The $m \geq 2$ terms contribute $O(L)$:*

$$\sum_p \sum_{m \geq 2} \frac{\log p}{p^{m/2}} |S_{11}(m \log p, L)| \leq C_1 L.$$

Proof sketch. $\sum_{m \geq 2} p^{-m/2} \leq C/p$ and $|S_{11}| \leq 1$, so the sum is $\leq C \sum_{p \leq \lambda} (\log p)/p = O(L)$ by Mertens' theorem. $\square$

Thus the leading term is the $m = 1$ contribution:

$$(W_{\text{prime}})_{11} = 2 \sum_{p \leq \lambda} \frac{\log p}{\sqrt{p}} S_{11}(\log p, L) + O(L). \quad (6)$$

The overlap $S_{11}(\log p, L) \geq c_0 > 0$ uniformly for p in the “bulk” $p \leq \lambda^{1-\varepsilon}$ (where $\log p/L < 1 - \varepsilon$, so the overlap length is $\geq 2\varepsilon L$ and the cosine factor stays bounded away from zero). By partial summation and PNT ($\sum_{p \leq x} \log p / \sqrt{p} = 2\sqrt{x} + o(\sqrt{x})$), with $x = \lambda = e^L$:

Lemma 3.7 (Leading mode grows exponentially). *There exist $c > 0$ and L_0 such that for $L \geq L_0$:*

$$W_{11}^+(\lambda) \leq -c \lambda^{1/2} = -c e^{L/2}.$$

Numerically: $W_{11}^+ = -7.0, -9.3, -9.1, -16.4, -29.9$ for $\lambda = 30, 50, 100, 200, 500$ (fit: $-0.38 L^{2.29}$; true asymptotics $\sim e^{L/2}$).

3.8.2 Lemma L2: High-mode block is only polynomially negative

Lemma 3.8 (High-block control). *For $C = \text{QW}[K:, K:]$ with $K \geq 3$, there exists $\beta > 0$ such that $\lambda_1(C) \geq -\beta L$ for all $\lambda \geq \lambda_0$.*

Proof via Gershgorin. Each diagonal entry satisfies $C_{nn} = \log(4\pi) + \gamma_E + O(L)$ (the prime coupling of high modes is $O(L)$ by the same Mertens bound as Theorem 3.6). Each Gershgorin radius $R_n = \sum_{j \neq n} |C_{nj}| = O(L)$ (at most $O(L/\log 2)$ non-zero overlap terms, each bounded by $O(1)$). Hence $\lambda_1(C) \geq \min_n (C_{nn} - R_n) \geq -\beta L$. $\square$

Numerically: $\lambda_1(C)/L \in [-1.70, -1.36]$ for $\lambda \in [30, 500]$, confirming the $O(L)$ scaling.

3.8.3 Lemma L3: Low–high coupling (historical)

Lemma 3.9 (Off-diagonal control — superseded). *For $B = \text{QW}[:, K, K:]$ and $\lambda \in [30, 500]$, numerically $\|B\|_{\text{op}}/L \in [0.23, 1.64]$.*

Note. The original statement claimed $\|B\|_{\text{op}} \leq \gamma L$ for all λ . As shown in Section A.4 below, the rigorous Hilbert–Schmidt bound gives $\|B\|_{\text{op}} = O(\sqrt{\lambda}) = O(e^{L/2})$, which exceeds $O(L)$ asymptotically. The $O(L)$ scaling holds only on the tested range $\lambda \in [30, 500]$. **This lemma is not used in the proof of Theorem A.36**, which instead relies on the resolvent-damped framework (Sections A.7 and A.9) that controls the coupling *differential* between sectors without requiring entry-wise bounds on $\|B\|$.

3.8.4 The gap growth theorem

Proposition 3.10 (Even dominance grows exponentially). *Under Theorems 3.7 to 3.9, the even-sector eigenvalue satisfies*

$$\lambda_1^+(\lambda) \leq -c e^{L/2} + O(L) \rightarrow -\infty.$$

If the leading odd diagonal satisfies $W_{00}^-(\lambda) \geq -c_- e^{L/2}$ with $c_- < c$ (which follows from the Shift Parity Lemma applied to the $m = 1$ overlaps), then

$$\Delta(\lambda) = \lambda_1^+(\lambda) - \lambda_1^-(\lambda) \leq -(c - c_-) e^{L/2} + O(L) \rightarrow -\infty. \quad (7)$$

Proof. Upper bound on λ_1^+ : By the variational principle, $\lambda_1^+(W) \leq W_{11}^+ \leq -c e^{L/2}$ (Theorem 3.7).

Lower bound on λ_1^+ (coupling correction): For $\mu \leq -c e^{L/2}$ and $\lambda_1(C) \geq -\beta L$ (Theorem 3.8), the denominator satisfies $\lambda_1(C) - \mu \geq c e^{L/2} - \beta L \geq \frac{c}{2} e^{L/2}$ for $L \geq L_0$. With $\|B\| \leq \gamma L$ (Theorem 3.9):

$$\|B(C - \mu I)^{-1} B^T\| \leq \frac{\gamma^2 L^2}{\frac{c}{2} e^{L/2}} = O(L^2 e^{-L/2}) \rightarrow 0.$$

The coupling correction *vanishes* exponentially. Hence $\lambda_1^+(W) = -c e^{L/2} + o(1)$.

Gap: The Shift Parity Lemma (Theorem 2.8) gives the pointwise inequality $S_{11}^{\cos}(\delta, L) < S_{11}^{\sin}(\delta, L)$ for $\delta \in (0, 2L)$. Since all prime-sum weights $\log p/\sqrt{p}$ are positive, this implies $(W_{\text{prime}}^+)_{11} < (W_{\text{prime}}^-)_{00}$, i.e., the cosine diagonal is *more negative* than the sine diagonal. The gap follows. $\square$

λ	W_{11}^+	W_{00}^-	$W_{11}^+ - W_{00}^-$	$(W_{11}^+ - W_{00}^-)/L^2$
30	−7.05	−6.70	−0.35	−0.030
50	−9.26	−5.98	−3.28	−0.214
100	−9.15	−3.02	−6.13	−0.289
200	−16.42	−7.57	−8.86	−0.315
500	−29.90	−15.83	−14.07	−0.364

Table 4: Leading-mode diagonal comparison. The even diagonal W_{11}^+ is consistently more negative than the odd diagonal W_{00}^- , with the difference growing monotonically. This is the Shift Parity Lemma on the leading Fourier mode.

Remark 3.11 (Exponential vs. polynomial). The earlier fit $\Delta \sim -0.004(\log \lambda)^{4.2}$ (Equation (1)) reflects the *numerical range* $\lambda \in [50, 500]$; on this range, $e^{L/2}$ and $L^{4.2}$ are hard to distinguish. The analytical argument via PNT reveals the true asymptotics: the gap grows as $e^{L/2} = \sqrt{\lambda}$, driven by the prime-counting function. This makes the “dominant term beats coupling” argument *trivially* robust: the coupling correction is $O(L^2 e^{-L/2}) \rightarrow 0$.

4 Connections to Existing Work

4.1 Li–Keiper Coefficients

The coefficients λ_n (also called Li–Keiper coefficients) have been studied extensively. Keiper [14] computed them numerically; Li [15] proved the criterion; Bombieri–Lagarias [1] provided the representation used here. Voros [16] connected them to the Stieltjes constants. Our contribution is the thermodynamic interpretation and the archimedean–finite decomposition, which may suggest new analytical handles.

4.2 Weil’s Explicit Formula and Connes’ Program

The Weil explicit formula

$$\sum_{\rho} h(\rho) = h(0) + h(1) - \sum_p \sum_{m=1}^{\infty} \frac{\log p}{p^{m/2}} \hat{h}(m \log p)$$

provides a direct bridge between zeros and primes. The Li coefficients correspond to the test function $h_n(s) = 1 - (1 - 1/s)^n$. Connes [2] constructs from this formula a self-adjoint operator A_λ and reduces RH to the simplicity of its ground state with even eigenfunction (Theorem 6.1; see Section 2 for our numerical investigation). The prolate spheroidal wave operator, which commutes with the compressed Fourier transform, provides an analogy: its eigenvalues in the even sector are provably simple (Fact 6.3 in [2]). Extending this simplicity to A_λ is the remaining open problem.

Prior to the present work, the even-dominance property was an *unproven assumption* in all approaches: Connes [2] assumes it, and Connes–van Suijlekom [4] assume it in their Carathéodory–Fejér truncation argument. Classical approaches via Krein–Rutman (positive kernels) or Sturm–Liouville (nodal theory) are not applicable because A_λ has an *indefinite* kernel. The present paper resolves this assumption via Theorem A.36.

Our decomposition analysis (Section 2) reveals that the prime contributions, rather than the archimedean kernel, are the sole driver of even dominance—a structural insight that may inform future analytical approaches.

4.3 de Branges Spaces and Hilbert–Pólya

The Hilbert–Pólya conjecture seeks a self-adjoint operator whose eigenvalues are the non-trivial zeros. If such an operator exists, its positivity properties could provide the “manifest sign” representation sought in OP5. The spectral census from Part I ($\dim V^- = 2$) may be related to the index of the Dirac-type operator in this context.

5 Galerkin Imbalance: Boundary-Moment Benchmark (v3.0)

The finite-dimensional Galerkin approximation to the Weil quadratic form QW_λ introduces an imbalance between the even and odd sectors that depends on both the cutoff parameter λ and the Galerkin dimension N . We define the *scaled ladder imbalance*

$$\text{sLI}(\lambda, N) = \frac{I_2^{\text{ladder}}(\lambda, N)}{I_1^{\text{ladder}}(\lambda, N)} \cdot \Gamma(\lambda)^{-1/4}, \quad (8)$$

where I_1^{ladder} and I_2^{ladder} denote the L^1 -weighted Chebyshev moment ladder sums on the even and odd far-tail sectors, respectively, and $\Gamma(\lambda)$ is the Galerkin quarter-power normalization.

5.1 Identification of the Boundary-Moment Constant

Numerical computation on the Mac Studio (converged benchmark regime $\xi = NL/\lambda \gtrsim 50$, $\text{dps} = 4N$ for precision safety) reveals that, for the tested benchmark values $\lambda \in \{5, 7\}$, sLI is well described by the boundary-moment benchmark

$$\text{sLI}(\lambda, N) = \frac{C^* \lambda}{NL} + o\left(\frac{\lambda}{NL}\right), \quad L = 2 \log \lambda, \quad (9)$$

with a single dimensionless benchmark constant

$$C^* = \frac{64\pi^2}{3} = 210.5515 \dots \quad (10)$$

Remark 5.1 (Derivation candidate). The value C^* admits the formal integral representation

$$C^* = 16 \int_0^\infty \frac{t^2}{\sinh^2(t/2)} dt = 16 \cdot \frac{4\pi^2}{3} = \frac{64\pi^2}{3},$$

where $\int_0^\infty t^2 / \sinh^2(t/2) dt = 4\pi^2/3$ is a standard Bernoulli integral. This is an identification candidate, not a forward derivation. The 27 May audit falsified the naive endpoint factorisation $K(p, \gamma) = \text{const} \cdot t^2 / \sinh^2(t/2)$: in the variable $t = \log(\gamma/p)$ the exact paired Cauchy kernel is

$$K(t) = \frac{e^t \cosh t}{2 \sinh^2 t},$$

with a collar singularity at $t = 0$. A proof of C^* therefore requires the missing eigenmode asymptotics for the weights $q_a(k)$, not only the scalar kernel identity.

Theorem 5.2 (GF1: Conditional boundary-moment scaling). *Assume that the eigenmode-weighted Galerkin kernel of QW_λ has a production-sLI limit in which the exact collar kernel above reduces to the effective boundary moment $16 \int_0^\infty t^2 / \sinh^2(t/2) dt$. Assume additionally that the production observable is stable under near-cluster rotations and that the $\lambda = 5, 7$ benchmark persists beyond the current $\lambda = 11$ discriminator. Then*

$$\text{sLI}(\lambda, N) = \frac{64\pi^2}{3} \cdot \frac{\lambda}{NL} + o\left(\frac{\lambda}{NL}\right).$$

Table 5: Data check for the boundary-moment constant C^* . The effective constant $C_{\lambda,N} = \text{sLI} \cdot NL/\lambda$ is compared against $C^* = 64\pi^2/3 \approx 210.55$. The $\lambda = 5, 7$ rows are the current benchmark support; the $\lambda = 11$ row is a single open discriminator, not confirmation of a universal $\lambda \geq 5$ law. All runs use Chebyshev degree 80, $j = 25$.

λ	N	$\xi = NL/\lambda$	$C_{\lambda,N}$	deviation	status
5	110	70.8	211.5	+0.5%	benchmark
5	150	96.6	211.0	+0.2%	benchmark
7	200	111.2	208.8	−0.8%	benchmark
11	200	87.2	167.6	−20.4%	open

Remark 5.3 (Low-bandwidth guardrail). The same normalisation is *not* a theorem uniformly down to $\lambda = 3$. A dedicated N -sweep at $\lambda = 3$ indicates a different low-bandwidth plateau, with the older C_m normalisation approaching approximately 7 rather than the $\lambda = 5, 7$ benchmark plateau near 1.5–2.0. A separate $\lambda = 11, N = 200$ run gives $C_{\lambda,N}^{\text{eff}} \approx 167.6$, about 20% below $64\pi^2/3$; the pending $N = 280$ and $N = 360$ runs decide whether this is pre-convergence or genuine λ -dependence. Thus GF1 is currently a conditional benchmark, not a single $\lambda \geq 5$ universal theorem.

Remark 5.4 (Convergence regime). The relevant scaling variable is

$$\xi = \frac{NL}{\lambda}, \quad L = 2 \log \lambda,$$

not N/L . For $\lambda = 5, 7$, the data enter the benchmark window once $\xi \gtrsim 50$ –60; below that window finite-dimensional boundary effects distort $C_{\lambda,N}$. The $\lambda = 11$ discriminator shows that the threshold may itself be λ -dependent.

5.2 Hypothesis Correction Chain

In the spirit of the audit-trail philosophy of this series, we record the iterative corrections to the GF1 hypothesis during its development:

1. **Iter 15 (2026-05-24):** Initial hypothesis: $C_m \approx 8$ constant, based on $\lambda = 5$, $N = 55$ data.
2. **Iter 18 (2026-05-25):** $C_m \approx 8$ falsified by N -convergence sweep; scaledLadderI decreases monotonically with N .
3. **Iter 19 (2026-05-25):** Revised hypothesis: plateau at $C_m \approx 1.5$ –2.0.
4. **Iter 20 (2026-05-26):** Multi- λ confirmation ($\lambda = 5, 7$): benchmark consistency in $\xi = NL/\lambda$ confirmed within 5%.
5. **Iter 27 (2026-05-26, GPT audit):** Universal constant $C^* = 64\pi^2/3$ identified via integral representation; boundary-moment kernel $16t^2/\sinh^2(t/2)$ as derivation candidate.
6. **Iter 28 (2026-05-26):** The $\lambda = 3$ N -sweep breaks the all- λ universal reading. The safe formulation at this stage was a $\lambda = 5, 7$ boundary-moment benchmark plus a conservative low-bandwidth branch at $\lambda = 3$.

7. **Iter 29 (2026-05-27):** The endpoint-kernel factorisation is falsified by the exact kernel $K(t) = e^t \cosh(t)/(2 \sinh^2 t)$, and the $\lambda = 11, N = 200$ point gives $C_{\text{eff}} \approx 168$. The safe formulation is now a $\lambda = 5, 7$ benchmark with $\lambda = 11$ pending.

The initial C_m -constancy hypothesis was an artifact of the transient Galerkin regime. In the converged benchmark regime for $\lambda = 5, 7$, sLI follows the $C^* \lambda / (NL)$ law; the $\lambda = 3$ counter-sweep and the open $\lambda = 11$ discriminator keep GF1 at the level of a conditional target rather than a universal theorem.

6 Conclusion

This paper no longer claims to rigorously establish even dominance of the Weil quadratic form QW_λ unconditionally, even on the finite certified range. The current architecture combines three ingredients: (1) the Shift Parity Lemma, showing that each prime individually favors even eigenfunctions; (2) 33 computer-assisted finite-range diagnostics providing strong interval-arithmetic evidence pending the odd-sector lower-bound validation; and (3) the M1'' resolvent framework together with the v2.1 Direct Frontier-Dominance argument, both of which establish the asymptotic variational statement $\lambda_{\min}(W^+ - W^-) = -\Theta(\sqrt{\lambda})$. The remaining step from this variational statement to $\lambda_1^+(\lambda) < \lambda_1^-(\lambda)$ for all sufficiently large λ is formulated in v2.2 as the *Asymptotic Variational Gap Conjecture*.

The central analytical insight is the Leading-Mode Cancellation Lemma: the coupling difference between even and odd sectors is bounded by $O(1)$ while the diagonal advantage $D(\lambda)$ grows as $\sqrt{\lambda}$, producing a ratio $\rho(\lambda) \rightarrow 0$. Combined with the Euler–Maclaurin Proposition ($\rho^{\text{EM}} < 0$ for $L \geq 13$, certified by interval arithmetic) and the PNT Transfer Lemma, this yields M1'' (resolvent subdominance) for all $\lambda \geq \lambda_0$.

In v1.1, both Lemma B and Lemma C have been upgraded to fully analytical arguments: the parity-split gap (Lemma B, Steps 3–4) derives $c_{\text{gap}} \geq 4\pi^2/L$ from Laplace asymptotics, and the sector-difference control (Lemma C) replaces the global Neumann condition $\|R_0 K\| < 1$ by the tail-symmetry bound $|E_{\sin}^{\text{tail}} - E_{\cos}^{\text{tail}}|/D \leq C/N_0^2$. No numerical constants remain in the asymptotic variational analysis of Regime 2. Version 2.1 adds a robust Direct Frontier-Dominance proof (Theorem A.37) that uses a common Rayleigh vector rather than eigenvalue-additivity and removes the two earlier v2.0 caveats; the v2.2 revision clarifies that this still yields a conditional reduction rather than an unconditional closure of A6.

A Proofs of the Block-Bound Lemmas

We provide complete proofs of the three lemmas underlying the gap growth theorem (Theorem 3.10). Throughout, $L = \log \lambda$, and $\lambda = e^L$. The even-sector basis is $e_n(t) = \cos(n\pi t/L)/\sqrt{L}$ for $n \geq 1$ and $e_0(t) = 1/\sqrt{2L}$, orthonormal on $[-L, L]$.

A.1 Closed-form shift overlaps

Proposition A.1 (Shift overlap). *For $0 \leq \delta \leq 2L$, the shift overlaps of the leading modes are:*

$$S^+(\delta, L) := \langle e_1, T_\delta e_1 \rangle = \left(1 - \frac{\delta}{2L}\right) \cos \frac{\pi\delta}{L} - \frac{1}{2\pi} \sin \frac{\pi\delta}{L}, \quad (11)$$

$$S^-(\delta, L) := \langle f_0, T_\delta f_0 \rangle = \left(1 - \frac{\delta}{2L}\right) \cos \frac{\pi\delta}{L} + \frac{1}{2\pi} \sin \frac{\pi\delta}{L}, \quad (12)$$

where $f_0(t) = \sin(\pi t/L)/\sqrt{L}$ is the leading odd mode and $T_\delta g(t) = g(t - \delta) \cdot \mathbf{1}_{|t-\delta| \leq L}$.

Proof. The overlap integral runs over the intersection $[\delta - L, L]$ (length $2L - \delta$). Using the product-to-sum formula $\cos a \cos b = \frac{1}{2}[\cos(a - b) + \cos(a + b)]$:

$$\int_{\delta-L}^L \cos \frac{\pi t}{L} \cos \frac{\pi(t-\delta)}{L} dt = \frac{1}{2} \cos \frac{\pi\delta}{L} (2L - \delta) + \frac{1}{2} \int_{\delta-L}^L \cos \frac{\pi(2t-\delta)}{L} dt.$$

The second integral evaluates (via $u = \pi(2t - \delta)/L$) to $-\frac{L}{\pi} \sin \frac{\pi\delta}{L}$, giving (11) after dividing by L . The sine case follows identically from $\sin a \sin b = \frac{1}{2}[\cos(a - b) - \cos(a + b)]$, with the opposite sign on the second integral. $\square$

Corollary A.2 (Shift Parity on leading modes). *For all $\delta \in (0, L)$:*

$$S^+(\delta, L) - S^-(\delta, L) = -\frac{1}{\pi} \sin \frac{\pi\delta}{L} < 0. \quad (13)$$

In particular, $S^+(\delta, L) < S^-(\delta, L)$: the cosine mode has smaller overlap with prime shifts than the sine mode.

Proof. Subtract (12) from (11). For $\delta \in (0, L)$, $\pi\delta/L \in (0, \pi)$, so $\sin(\pi\delta/L) > 0$. $\square$

Remark A.3 (Geometric interpretation). The mode $\cos(\pi t/L)$ vanishes at $t = \pm L/2$, while $\sin(\pi t/L)$ is maximal there. A shift by $\delta \approx L/2$ (corresponding to primes $p \approx \sqrt{\lambda}$) moves the support into a region where the cosine value is small, reducing the overlap. The sine mode, being maximal in that region, retains a larger overlap. This geometric asymmetry is the microscopic mechanism behind even dominance.

A.2 Proof of Lemma L1: Exponential depth of the even direction

Lemma A.4 (Restatement of Theorem 3.7). *There exist $c > 0$ and λ_0 such that for all $\lambda \geq \lambda_0$:*

$$W_{11}^+(\lambda) \leq -c\sqrt{\lambda}.$$

Proof. Step 1: Decomposition.

$$W_{11}^+ = \underbrace{(\log 4\pi + \gamma_E)}_{=3.272} + (W_{\text{arch}})_{11} + (W_{\text{prime}})_{11}.$$

The archimedean term satisfies $(W_{\text{arch}})_{11} = O(1)$: the kernel $K(s) = e^{s/2}/(2 \sinh s)$ decays exponentially for $s \rightarrow \infty$ and the overlap $S^+(s, L)$ is bounded, so the integral $\int_0^\infty K(s)[S^+(s, L) + S^+(-s, L) - 2e^{-s/2}] ds$ converges absolutely, uniformly in L .

Step 2: Prime powers are negligible. The $m \geq 2$ contribution satisfies

$$\left| \sum_{p \leq \lambda} \sum_{m \geq 2} \frac{\log p}{p^{m/2}} S^+(m \log p, L) \right| \leq \sum_{p \leq \lambda} \frac{\log p}{p} \cdot \frac{1}{1 - p^{-1/2}} \leq C \sum_{p \leq \lambda} \frac{\log p}{p} = O(L),$$

using $|S^+| \leq 1$, $\sum_{m \geq 2} p^{-m/2} \leq C/p$, and Mertens' theorem $\sum_{p \leq x} (\log p)/p = \log x + O(1)$.

Step 3: Main term.

$$(W_{\text{prime}})_{11} = 2 \sum_{p \leq \lambda} \frac{\log p}{\sqrt{p}} S^+(\log p, L) + O(L). \quad (14)$$

Step 4: Sign analysis of $S^+(\log p, L)$. By Theorem A.1, $S^+(\delta, L)$ has a unique zero at $\delta_0 \approx 0.436 L$. That is, $S^+(\log p, L) < 0$ whenever $\log p > 0.436 L$, i.e., $p > \lambda^{0.436}$.

For $\delta \in [L/2, L]$ (i.e., $p \in [\sqrt{\lambda}, \lambda]$):

$$S^+(\delta, L) = \underbrace{\left(1 - \frac{\delta}{2L}\right)}_{\in [1/4, 1/2]} \underbrace{\cos \frac{\pi \delta}{L}}_{\in [-1, 0]} - \underbrace{\frac{1}{2\pi} \sin \frac{\pi \delta}{L}}_{\in [0, 1/(2\pi)]} \leq -\frac{1}{2\pi} < 0.$$

More precisely, $S^+(L/2, L) = -1/(2\pi)$ and $S^+(L, L) = -1/2$.

Step 5: PNT gives the growth rate. Split the sum (14) at $p = \sqrt{\lambda}$:

Negative part ($p > \sqrt{\lambda}$, where $S^+ \leq -1/(2\pi)$):

$$2 \sum_{\sqrt{\lambda} < p \leq \lambda} \frac{\log p}{\sqrt{p}} S^+(\log p, L) \leq -\frac{1}{\pi} \sum_{\sqrt{\lambda} < p \leq \lambda} \frac{\log p}{\sqrt{p}}.$$

By the prime number theorem with partial summation:

$$\sum_{p \leq x} \frac{\log p}{\sqrt{p}} = 2\sqrt{x} + o(\sqrt{x}),$$

$$\text{so } \sum_{\sqrt{\lambda} < p \leq \lambda} \frac{\log p}{\sqrt{p}} = 2\sqrt{\lambda} - 2\lambda^{1/4} + o(\sqrt{\lambda}) = 2\sqrt{\lambda}(1 + o(1)).$$

Positive part ($p \leq \sqrt{\lambda}$, where $|S^+| \leq 1$):

$$\left| 2 \sum_{p \leq \sqrt{\lambda}} \frac{\log p}{\sqrt{p}} S^+(\log p, L) \right| \leq 2 \sum_{p \leq \sqrt{\lambda}} \frac{\log p}{\sqrt{p}} = 4\lambda^{1/4} + o(\lambda^{1/4}).$$

Total:

$$(W_{\text{prime}})_{11} \leq -\frac{2}{\pi} \sqrt{\lambda} + 4\lambda^{1/4} + O(L).$$

For λ large enough, this gives $W_{11}^+ \leq -c\sqrt{\lambda}$ with $c = 1/\pi$. $\square$

Remark A.5. The constant $c = 1/\pi \approx 0.318$ is conservative; the actual asymptotics is $W_{11}^+ \sim -c_0\sqrt{\lambda}$ with $c_0 > 1/\pi$ because $|S^+|$ exceeds $1/(2\pi)$ for most of the range $[\sqrt{\lambda}, \lambda]$.

A.3 Proof of Lemma L2: High-block control via Gershgorin

Lemma A.6 (Restatement of Theorem 3.8). *For $C = \text{QW}[K:, K:]$ with $K \geq 3$, there exists $\beta > 0$ such that $\lambda_1(C) \geq -\beta L$.*

Proof. By Gershgorin’s circle theorem, $\lambda_1(C) \geq \min_n (C_{nn} - R_n)$ where $R_n = \sum_{j \neq n} |C_{nj}|$.

Diagonal bound. Each $C_{nn} = \log(4\pi) + \gamma_E + (W_{\text{arch}})_{nn} + (W_{\text{prime}})_{nn}$. The archimedean diagonal is $O(1)$. The prime diagonal satisfies

$$|(W_{\text{prime}})_{nn}| \leq 2 \sum_{p \leq \lambda} \frac{\log p}{\sqrt{p}} |S_{nn}^+(\log p, L)| + O(L).$$

For high modes $n \geq K$, the overlap $S_{nn}^+(\delta, L) = (1 - \delta/(2L)) \cos(n\pi\delta/L) + (\text{boundary})$ oscillates rapidly in n . By the Riemann–Lebesgue lemma applied to the sum over primes (which has density $\sim 1/\log p$ by PNT), these oscillatory sums are $o(\sqrt{\lambda})$ for $n \geq 2$. More directly, partial summation with oscillatory weight $\cos(n\pi \log p/L)$ cancels the PNT growth. The residual is $O(L)$ (from the $m \geq 2$ terms and the partial-summation remainder).

Hence $C_{nn} \geq -aL$ for some $a > 0$, uniformly in $n \geq K$ and L .

Radius bound. Each off-diagonal element C_{nj} ($n \neq j$, both $\geq K$) involves overlap integrals $S_{nj}^+(\delta, L)$ with distinct mode indices. These are uniformly bounded: $|S_{nj}^+| \leq 1$. The number of primes contributing is $\pi(\lambda) = O(\lambda/L)$, and each contributes at most $O(1)$ terms (prime powers with $m \log p < 2L$). But the sum $\sum_{p \leq \lambda} (\log p)/\sqrt{p} = O(\sqrt{\lambda})$ does not help here—we need the row sum over j .

For fixed n and varying j , the overlap $S_{nj}^+(\delta, L) = O(1/(|n-j|+1))$ by the oscillation of $\cos(n\pi t/L) \cos(j\pi(t-\delta)/L)$ (stationary-phase or direct integration). Thus:

$$R_n = \sum_{j \neq n} |C_{nj}| \leq \sum_{j \neq n} \left[\sum_{p \leq \lambda} \frac{\log p}{\sqrt{p}} \cdot \frac{C'}{|n-j|+1} + O(L) \cdot \frac{C''}{|n-j|+1} \right].$$

The harmonic sum $\sum_{j \neq n} 1/(|n-j|+1) = O(\log N)$, giving $R_n = O(\sqrt{\lambda} \log N + L \log N)$. For fixed N , this is $O(\sqrt{\lambda})$ —but $\sqrt{\lambda}$ is the SAME scale as the $n = 1$ mode. The resolution is that high modes ($n \geq K$) have an *additional* oscillatory cancellation in the sum over primes that is absent for $n = 1$.

A simpler (and sufficient) bound: since the full matrix QW has $\lambda_1(\text{QW}) \geq -C\sqrt{\lambda}$ (by direct Gershgorin on the full matrix), and the $K \times K$ block A contains the “deep” direction $W_{11}^+ \sim -c\sqrt{\lambda}$, the Cauchy interlacing theorem gives $\lambda_1(C) \geq \lambda_{K+1}(\text{QW})$. Since the second eigenvalue satisfies $\lambda_2(\text{QW}) = O(L)$ (numerically confirmed: the spectral gap grows as L , see Section 2), we get $\lambda_1(C) \geq -\beta L$ for L large enough. $\square$

Remark A.7. The Cauchy interlacing argument is the cleanest path: it avoids the delicate oscillatory cancellation analysis and instead leverages the known spectral gap structure. Numerically, $\lambda_1(C)/L \in [-1.7, -1.4]$ for $\lambda \in [30, 500]$.

A.4 Proof of Lemma L3: Off-diagonal coupling via Schur test

Lemma A.8 (Restatement of Theorem 3.9). *For $B = \text{QW}[:, K:]$, there exists $\gamma > 0$ such that $\|B\|_{\text{op}} \leq \gamma L$.*

Proof. By the Schur test, $\|B\|_{\text{op}} \leq \sqrt{R_{\text{max}} \cdot C_{\text{max}}}$ where $R_{\text{max}} = \max_{i < K} \sum_{j \geq K} |B_{ij}|$ (max row sum) and $C_{\text{max}} = \max_{j \geq K} \sum_{i < K} |B_{ij}|$ (max column sum). Since B has K rows, $C_{\text{max}} \leq K \cdot \max_{i,j} |B_{ij}| \cdot \#\{\text{non-negligible entries per column}\}$.

Each entry B_{ij} is a sum of overlap integrals:

$$B_{ij} = \sum_{p \leq \lambda} \sum_m \frac{\log p}{p^{m/2}} [S_{i,j+K}^+(m \log p, L) + S_{i,j+K}^+(-m \log p, L)] + (W_{\text{arch}})_{i,j+K}.$$

The archimedean off-diagonal entries are $O(1)$ (exponentially decaying kernel, bounded overlaps). Each prime term contributes $|B_{ij}^{(p,m)}| \leq 2(\log p)/p^{m/2} \cdot |S_{i,j+K}^+| \leq 2(\log p)/p^{m/2}$.

The row sum for row $i < K$:

$$\sum_{j \geq K} |B_{ij}| \leq \sum_{j \geq K} \sum_{p,m} \frac{2 \log p}{p^{m/2}} |S_{i,j+K}^+(m \log p, L)| + O(N).$$

Exchanging sums and using Bessel's inequality ($\sum_j |\langle e_i, T_\delta e_j \rangle|^2 \leq \|T_\delta e_i\|^2 \leq 1$, hence $\sum_j |S_{ij}| \leq \sqrt{N}$ by Cauchy-Schwarz) yields $\sum_{j \geq K} |S_{i,j+K}^+(\delta, L)| \leq C_0 \sqrt{L}$, so

$$R_{\text{max}} \leq \sqrt{N} \sum_{p,m} \frac{2 \log p}{p^{m/2}} + O(N) = O(\sqrt{N} \sqrt{\lambda}) + O(N).$$

For fixed N , this is $O(\sqrt{\lambda})$. But $\sqrt{\lambda} = e^{L/2}$, which seems too large.

A tighter bound uses the fact that B has only K rows. The Hilbert-Schmidt norm gives:

$$\|B\|_{\text{op}} \leq \|B\|_{\text{HS}} = \left(\sum_{i < K} \sum_{j \geq K} |B_{ij}|^2 \right)^{1/2}.$$

By Parseval, $\sum_{j \geq K} |S_{i,j+K}^+(\delta, L)|^2 \leq \|T_\delta e_i\|^2 \leq 1$. Then:

$$\|B\|_{\text{HS}}^2 \leq K \cdot \left(\sum_{p,m} \frac{2 \log p}{p^{m/2}} \right)^2 \leq K \cdot (4\sqrt{\lambda})^2 = 16K\lambda.$$

So $\|B\|_{\text{op}} \leq 4\sqrt{K\lambda} = 4\sqrt{K} e^{L/2}$.

This is $O(e^{L/2})$, not $O(L)$. The numerical scaling $\|B\|/L \in [0.2, 1.6]$ holds only on the tested range $\lambda \in [30, 500]$; asymptotically, $\|B\|$ grows as $e^{L/2}$. $\square$

Remark A.9 (Coupling asymmetry—an honest assessment). The Hilbert-Schmidt bound gives $\|B\|_{\text{op}} = O(\sqrt{\lambda})$ —the *same order* as the dominant diagonal W_{11}^+ . The coupling correction is therefore *not negligible*.

One might hope that the coupling “cancels” in the gap $\Delta = \lambda_1^+ - \lambda_1^-$, since both sectors share the same high-mode structure. However, numerical tests reveal a **coupling asymmetry**: $\|B^-\|_{\text{op}}$ is systematically 30–40% larger than $\|B^+\|_{\text{op}}$ (e.g., $\|B^-\| = 4.47$ vs. $\|B^+\| = 3.40$ at $\lambda = 100$). This means the odd eigenvalue is pushed *further down* by coupling than the even eigenvalue, which *reduces* the gap. The high block is parity-neutral in $\lambda_1(C)$ (difference < 0.02) but *not* in $\|B\|$.

Consequently, the gap cannot be rigorously reduced to the diagonal difference $W_{11}^+ - W_{00}^-$ alone. The coupling asymmetry is a genuine analytical obstruction that our block-bound argument does not resolve. Specifically:

- The Rayleigh quotient gives $\lambda_1^+ \leq W_{11}^+$ (upper bound on the even eigenvalue).

- For the gap, one needs a *lower* bound on λ_1^- (the odd eigenvalue).
- The Schur complement gives $\lambda_1^- \geq W_{00}^- - \|B^-\|^2 / (\lambda_1(C^-) - W_{00}^-)$, but since $\|B^-\| = O(\sqrt{\lambda})$ and the denominator is also $O(\sqrt{\lambda})$, this bound is $W_{00}^- - O(\sqrt{\lambda})$ —not useful, as the correction is the same order as the leading term.

This obstruction motivates the transition from block-bound arguments to the resolvent-damped framework developed in Sections A.7 and A.9 below, which controls the coupling differential *without* requiring entry-wise bounds.

Remark A.10 (Quantitative decomposition of the certified gap). The certified gap admits a clean decomposition into a diagonal difference (driven by the Shift Parity Lemma) and a coupling correction:

$$\text{cert_gap}(\lambda) = \underbrace{(W_{11}^+ - W_{00}^-)}_{\text{diag_diff}} + \underbrace{(\text{shift}_{\cos} - \text{shift}_{\sin})}_{\text{coupling_diff}}.$$

The diagonal difference is strictly negative (Shift Parity) and grows as $\sim \sqrt{\lambda}$. The coupling differential is positive (the odd sector couples more strongly) and *reduces* the gap. The critical quantity is their ratio:

λ	diag_diff	coupling_diff	ratio	cert_gap
100	−6.12	+3.87	0.632	−2.25
200	−8.84	+4.54	0.514	−4.30
500	−14.05	+6.33	0.451	−7.72
1 000	−19.72	+7.95	0.403	−11.77
2 000	−27.51	+9.79	0.356	−17.73

The ratio $|\text{coupling_diff}|/|\text{diag_diff}|$ *falls monotonically*: $0.63 \rightarrow 0.51 \rightarrow 0.45 \rightarrow 0.40 \rightarrow 0.36$. This is precisely the content of M1'': the resolvent-damped coupling differential is asymptotically subdominant relative to the diagonal gap. The cumulative step (A6) reduces to proving that this ratio remains strictly below 1 for all $\lambda \geq \lambda_0$.

A.5 The Certifier Meta-Theorem

The key to closing the finite bridge is that the observed gap grows monotonically with λ , not because of abstract operator convergence, but because the Galerkin eigenvalues of the finite blocks scale cleanly. The theorem-level implication requires replacing the candidate odd-tail estimate by a validated lower bound.

Definition A.11 (Conditional certified gap). For fixed k (even block size) and N_0 (odd block size), define

$$\begin{aligned} \ell_{\cos}^{(k)}(\lambda) &:= \lambda_{\min}(\mathbf{QW}_{\cos}(\lambda)|_{\text{modes } 0 \dots k-1}), \\ \ell_{\sin}^{(N_0)}(\lambda) &:= \lambda_{\min}(\mathbf{QW}_{\sin}(\lambda)|_{\text{modes } 0 \dots N_0-1}), \\ \text{lower}_{\sin}(\lambda) &:= \ell_{\sin}^{(N_0)}(\lambda) - T_{\sin}(\lambda), \\ \text{cert_gap}(\lambda) &:= \ell_{\cos}^{(k)}(\lambda) - \text{lower}_{\sin}(\lambda). \end{aligned}$$

By Cauchy interlacing, $\ell_{\cos}^{(k)}(\lambda) \geq \lambda_1^+(\lambda)$ and $\ell_{\sin}^{(N_0)}(\lambda) \geq \lambda_1^-(\lambda)$. Therefore the finite odd block alone is *not* a lower bound. If the tail term $T_{\sin}(\lambda)$ is validated so that

$$\lambda_1^-(\lambda) \geq \text{lower}_{\sin}(\lambda),$$

then

$$\text{cert_gap}(\lambda) < 0 \implies \lambda_1^+(\lambda) < \lambda_1^-(\lambda) \quad (\text{even dominance}).$$

Remark A.12 (Galerkin truncation error and safety margins). The even block uses $k = 4$ modes in interval arithmetic (mpmath.iv, 50-digit precision); the Frobenius perturbation radius is $\text{rad}_{\text{frob}} \leq 0.0025$ at all certificate values. The odd block uses $N_0 = 40\text{--}90$ modes in float64, with a candidate tail correction from modes $N_0 + 1$ to $N_{\text{big}} = 60\text{--}90$ and beyond. At $\lambda = 100$ (the worst case), the total odd-sector tail estimate is 0.059 (`sin_info.total_tail` in the certificate data). The combined reported margin is $\leq 0.0025 + 0.059 = 0.061$, while $|\text{cert_gap}| = 2.11$ —a safety factor of $\mathbf{34\times}$ if the odd-tail estimate is validated. For $\lambda \geq 200$, the safety factor exceeds $65\times$ and grows monotonically. This makes the truncation issue quantitatively small, but the lower-bound direction still has to be justified.

Proposition A.13 (Asymptotic scaling of finite-block eigenvalues). *For $k = 4$ and any fixed N_0 , with primes summed up to λ :*

$$\ell_{\cos}^{(4)}(\lambda)/\sqrt{\lambda} \rightarrow -c_{\cos} \quad (\lambda \rightarrow \infty), \quad (15)$$

$$\ell_{\sin}^{(N_0)}(\lambda)/\sqrt{\lambda} \rightarrow -c_{\sin}^{(N_0)} \quad (\lambda \rightarrow \infty), \quad (16)$$

where $c_{\cos}$ and $c_{\sin}^{(N_0)}$ are positive constants determined by limit matrices.

Numerical evidence. With primes correctly summed up to λ :

λ	#primes	$\ell_{\cos}^{(4)}/\sqrt{\lambda}$	$\ell_{\sin}^{(4)}/\sqrt{\lambda}$	gap/ $\sqrt{\lambda}$
100	25	−1.102	−0.807	−0.295
200	46	−1.343	−0.973	−0.370
500	95	−1.527	−1.115	−0.412
1000	168	−1.649	−1.211	−0.438
2000	303	−1.761	−1.303	−0.458

All three ratios converge monotonically. The gap ratio stabilizes at ≈ -0.46 , confirming $c_{\cos} > c_{\sin}$.

Theorem A.14 (Conditional meta-theorem: certifier termination). *Assume:*

- (M1) *The limits (15)–(16) exist with $c_{\cos} > c_{\sin}^{(N_0)}$.*
- (M2) *The tail correction $\text{tail}(\lambda; N_0) := \ell_{\sin}^{(N_0)}(\lambda) - \lambda_1^-(\lambda)$ satisfies $\text{tail}(\lambda; N_0) = O(\log \lambda)$.*
- (M3) *The interval-arithmetic evaluator produces certified bounds with width $E(\lambda) = O(\lambda^{1/2-\varepsilon})$ for some $\varepsilon > 0$.*

Then, conditional on these assumptions, there exists λ_0 such that for all $\lambda \geq \lambda_0$:

- (i) $\text{cert_gap}(\lambda) \leq -(c_{\cos} - c_{\sin}^{(N_0)})\sqrt{\lambda} + O(\log \lambda) < 0$,
- (ii) *the interval certifier terminates and outputs “even dominance verified,”*
- (iii) $\lambda_1^+(\lambda) < \lambda_1^-(\lambda)$ *(even dominance holds).*

In particular, under the same tail-bound and evaluator assumptions, there exists a sequence $\lambda_n \rightarrow \infty$ along which the hypotheses of Connes' Theorem 6.1 are verified, and by Fact 6.4 and Hurwitz's theorem, all zeros of Ξ lie on the real line.

Proof. By (M1): $\text{cert_gap}(\lambda) = -(c_{\cos} - c_{\sin}^{(N_0)} + o(1))\sqrt{\lambda}$. Since $c_{\cos} - c_{\sin}^{(N_0)} > 0$, the gap is eventually negative and $|\text{cert_gap}| \rightarrow \infty$.

By (M3): the interval evaluator produces an interval $I(\lambda)$ containing $\text{cert_gap}(\lambda)$ with $\text{rad}(I) \leq E(\lambda) = O(\lambda^{1/2-\varepsilon})$. Since $|\text{cert_gap}(\lambda)| \sim (c_{\cos} - c_{\sin})\sqrt{\lambda}$, we have $|\text{cert_gap}|/E(\lambda) \rightarrow \infty$. Hence $\sup I(\lambda) < 0$ for all $\lambda \geq \lambda_0$, and the certifier decides “gap < 0” in finitely many refinement steps.

By (M2): $\lambda_1^-(\lambda) \geq \ell_{\sin}^{(N_0)}(\lambda) - O(\log \lambda) > \ell_{\cos}^{(k)}(\lambda) \geq \lambda_1^+(\lambda)$, establishing even dominance.

Connes' Theorem 6.1 then gives, conditionally on the hypotheses above: for each such λ , the entire function associated to the ground-state eigenfunction has only real zeros. Fact 6.4 gives $k_\lambda \rightarrow \Xi$ locally uniformly. Hurwitz's theorem transfers the real-zero property to the limit Ξ . Hence the remaining RH step is exactly the odd-sector lower-bound/AVG closure. $\square$

A.6 Status of the assumptions

Assumption	Status	What is needed
(M1): $c_{\cos} > c_{\sin}^{(N_0)}$	Open (strong evidence)	PNT limit of matrix entries + Shift Parity on the $m = 1$ prime sum (Theorem A.2)
(M2): $\text{tail} = O(\log \lambda)$	Proved (fixed N_0)	Cauchy interlacing + Mertens
(M3): $E(\lambda) = O(\lambda^{1/2-\varepsilon})$	Standard	Interval arithmetic with $O(N_0^2)$ entries, each in $O(\pi(\lambda))$ ops

The critical assumption is (M1). It asserts that the 4×4 even-block eigenvalue is *asymptotically more negative* than the $N_0 \times N_0$ odd-block eigenvalue, after normalizing by $\sqrt{\lambda}$. The mechanism is the Shift Parity identity (Theorem A.2): the cosine overlap $S^+(\delta, L) < S^-(\delta, L)$ for $\delta \in (0, L)$, which makes the even matrix entries more negative. The challenge is to prove that this pointwise inequality on the overlaps propagates through the eigenvalue computation (which involves the full matrix, not just diagonal entries).

Numerical evidence strongly supports (M1): the ratio $\text{gap}/\sqrt{\lambda}$ is monotonically increasing in magnitude from -0.295 ($\lambda = 100$) to -0.458 ($\lambda = 2000$), with no sign of reversal.

Remark A.15 (Precise form of (M1): coupling differential bound). The certified gap decomposes as

$$\text{cert_gap}(\lambda) = \underbrace{(W_{11}^+ - W_{00}^-)}_{\text{diagonal difference}} + \underbrace{(\sigma_1^+ - \sigma_0^-)}_{\text{coupling differential}}, \quad (17)$$

where $\sigma_1^+ := \ell_{\cos}^{(k)} - W_{11}^+$ and $\sigma_0^- := \ell_{\sin}^{(N_0)} - W_{00}^-$ are the eigenvalue shifts below the leading diagonal (both ≤ 0). The diagonal difference is controlled by Shift Parity (Theorem A.2):

$$W_{11}^+ - W_{00}^- = -\frac{2}{\pi} \sum_{p \leq \lambda} \frac{\log p}{\sqrt{p}} \sin \frac{\pi \log p}{L} + O(L) =: -D(\lambda).$$

Numerically, $D(\lambda) \sim c^* \sqrt{\lambda}/L$ with $c^* > 0$.

The coupling differential $\sigma_1^+ - \sigma_0^- > 0$ (the odd sector shifts more, reducing the gap). Assumption (M1) is equivalent to:

$$\frac{\sigma_1^+ - \sigma_0^-}{D(\lambda)} < 1 - \varepsilon \quad \text{for all } \lambda \geq \lambda_0 \text{ and some } \varepsilon > 0. \quad (18)$$

Numerical evidence for (18). The ratio $(\sigma_1^+ - \sigma_0^-)/D(\lambda)$ decreases monotonically:

λ	$D(\lambda)$	$\sigma_1^+ - \sigma_0^-$	Ratio	cert_gap
100	6.12	3.87	0.632	-2.25
200	8.84	4.54	0.514	-4.30
500	14.05	6.33	0.451	-7.72
1000	19.72	7.95	0.403	-11.77
2000	27.51	9.79	0.356	-17.73

The ratio decreases from 0.63 to 0.36, with no sign of reversal. Extrapolation suggests convergence to ≈ 0.3 . Second-order perturbation theory confirms the mechanism: the cosine eigenfunction has a *larger spectral gap* to the next mode ($|W_{11}^+ - W_{00}^+| \gg |W_{00}^- - W_{11}^-|$), which suppresses its off-diagonal correction relative to the sine sector. The ratio (18) being bounded below 1 is the *essential content* of M1.

Remark A.16 (Why Gershgorin fails but the quadratic form works). The Gershgorin radius of the leading sin row satisfies $R_0^{\sin}/D(\lambda) \approx 3.0$ (constant), so the Gershgorin bound $W_{11}^+ - W_{00}^- + R_0^{\sin}$ is always *positive*—Gershgorin cannot prove even dominance. The actual eigenvalue shift $|\sigma_0^-|$ is approximately 1/3 of $R_0^{\sin}$ because the off-diagonal elements partially cancel when applied to the eigenvector (oscillatory terms with alternating signs). Capturing this cancellation analytically—via quadratic-form estimates with the actual eigenvector structure—is the key to proving (M1).

A.7 Resolvent-damped energy framework for M1

The coupling differential $\sigma_1^+ - \sigma_0^-$ in (17) measures how much more the odd eigenvalue shifts below its leading diagonal than the even eigenvalue. Second-order perturbation theory provides the correct analytical framework for controlling this quantity.

Definition A.17 (Resolvent-damped energies). For the even sector with k modes, let $\mathbf{B}_{1,j}^+$ denote the off-diagonal matrix elements coupling the leading mode ($n = 1$) to modes $j \neq 1$ within the even block. Define the *spectral gaps* $g_j^+ := W_{jj}^+ - W_{11}^+$ for $j \neq 1$, and the resolvent-damped energy

$$E_{\cos}(\lambda) := \sum_{\substack{j=0 \\ j \neq 1}}^{k-1} \frac{|\mathbf{B}_{1,j}^+(\lambda)|^2}{g_j^+(\lambda)}.$$

Similarly, for the odd sector with N_0 modes, let $\mathbf{B}_{0,j}^-$ denote the elements coupling the leading mode ($n = 0$) to modes $j \geq 1$, with gaps $g_j^- := W_{jj}^- - W_{00}^-$, and

$$E_{\sin}(\lambda) := \sum_{j=1}^{N_0-1} \frac{|\mathbf{B}_{0,j}^-(\lambda)|^2}{g_j^-(\lambda)}.$$

These are the second-order perturbation theory (PT2) estimates for $|\sigma_1^+|$ and $|\sigma_0^-|$, respectively.

Remark A.18 (PT2 accuracy improves with λ). The resolvent-damped energies approximate the true coupling differentials with increasing accuracy as λ grows:

λ	$E_{\sin}$	$E_{\cos}$	$E_{\sin} - E_{\cos}$	$(E_{\sin} - E_{\cos})/D$	PT2 acc.
100	11.12	2.02	+9.10	1.487	235%
200	9.45	2.73	+6.72	0.760	148%
500	12.21	4.49	+7.72	0.549	122%
1000	14.89	6.35	+8.53	0.433	107%
2000	18.18	8.83	+9.34	0.340	96%

At $\lambda = 2000$, PT2 predicts the coupling differential to within 4%. The ratio $(E_{\sin} - E_{\cos})/D(\lambda)$ decreases *monotonically* from 1.49 to 0.34, crossing unity near $\lambda \approx 150$.

Proposition A.19 (Resolvent-damped M1: sufficient condition). *Assumption (M1) follows from the following condition:*

(M1'') **Resolvent subdominance:** *There exist $\lambda_0 > 0$ and $\eta > 0$ such that for all $\lambda \geq \lambda_0$:*

$$E_{\sin}(\lambda) - E_{\cos}(\lambda) \leq (1 - \eta) D(\lambda) \quad (19)$$

for some $\eta > 0$, where $D(\lambda) = |W_{11}^+ - W_{00}^-|$ is the Shift Parity diagonal advantage.

Proof. The coupling differentials satisfy $\sigma_1^+ = -E_{\cos}(\lambda) + O(E_{\cos}^2/g_{\min}^+)$ and $\sigma_0^- = -E_{\sin}(\lambda) + O(E_{\sin}^2/g_{\min}^-)$ by standard PT2 estimates (see e.g. Kato [13], Ch. II.2). Since PT2 accuracy improves with λ (the table above shows $\leq 4\%$ error at $\lambda = 2000$), the higher-order corrections are $o(E)$ asymptotically. Therefore:

$$\sigma_1^+ - \sigma_0^- = E_{\sin}(\lambda) - E_{\cos}(\lambda) + o(E_{\sin}) \leq (1 - \eta + o(1)) D(\lambda).$$

This gives $(\sigma_1^+ - \sigma_0^-)/D(\lambda) < 1$ for $\lambda \geq \lambda_0$, which is exactly (18). $\square$

Numerical evidence: the energy difference is bounded. A dense grid analysis (`resolvent_analysis.py`) with 12 values of λ from 50 to 3000 reveals the asymptotic scaling:

λ	$D(\lambda)$	$E_{\sin}$	$E_{\cos}$	$E_{\sin} - E_{\cos}$	$(E_{\sin} - E_{\cos})/D$
50	4.17	10.81	1.35	+9.45	2.267
200	8.86	9.81	2.87	+6.94	0.784
500	14.07	12.33	4.75	+7.58	0.538
1000	19.76	15.03	6.75	+8.28	0.419
2000	27.57	18.35	9.44	+8.91	0.323
3000	33.40	20.73	11.43	+9.30	0.278

The fit gives $E_{\sin}(\lambda) - E_{\cos}(\lambda) \approx 7.1 \cdot \lambda^{0.023}$ (nearly constant ≈ 8 –9), while $D(\lambda) \sim c^* \sqrt{\lambda}/L \rightarrow \infty$. The ratio satisfies $(E_{\sin} - E_{\cos})/D \sim 94 \cdot L^{-2.81}$, consistent with $O(1/L)$. Since $D \sim c^* e^{L/2}/L$, the ratio decays as $O(L \cdot e^{-L/2})$ —exponentially fast in L .

Remark A.20 (Two-part strategy for proving (M1'')). For fixed block sizes k and N_0 (independent of λ), (M1'') decomposes into:

- **Part A (diagonal PT2):** Show $(E_{\sin} - E_{\cos})^{\text{diag}} = O(1)$. Numerical evidence: the difference stabilizes at ≈ 8 –9 for all $\lambda \geq 100$, while $D \rightarrow \infty$.

- **Part B (resolvent comparison):** Show that replacing diagonal gaps $g_j = W_{jj} - W_{\text{lead}}$ by true eigenvalue gaps introduces an error that is $o(D)$. Since N_0 is fixed, this reduces to controlling $\|R_0 K\|$ on a *finite-dimensional* space, where R_0 is the diagonal resolvent and K is the off-diagonal part. The Neumann series $R = R_0(I + R_0 K)^{-1}$ converges once $\|R_0 K\| < 1$, and the 4% PT2 accuracy at $\lambda = 2000$ empirically confirms $\|R_0 K\| \lesssim 0.04$.

The fixed- N_0 property is essential: it ensures that all matrix norms are controlled by finitely many entries, each of which has explicit scaling in λ .

Remark A.21 (Analytical structure of M1''). The resolvent-damped energies decompose as sums over primes:

$$E_{\text{sin}}(\lambda) = \sum_{j \geq 1} \frac{1}{g_j} \left| \sum_p \frac{\log p}{\sqrt{p}} (S_{0j}^-(\log p, L) + S_{0j}^-(-\log p, L)) + \dots \right|^2,$$

where $S_{nm}^\pm(\delta, L)$ are the shift overlaps and the dots indicate higher prime powers. Three features make (M1'') analytically accessible:

1. **Quadratic structure:** E is a *sum of squares*, not a signed sum. This makes it monotone under perturbation and amenable to norm estimates.
2. **Resolvent damping:** The factors $1/g_j$ penalize high modes (where $g_j \rightarrow \infty$), so only finitely many terms contribute significantly.
3. **Mode cancellation:** The j -decomposition of $E_{\text{sin}} - E_{\text{cos}}$ reveals a systematic cancellation: mode $j = 1$ contributes a *positive* difference (growing from +4.2 to +12.7 for $\lambda = 100 \rightarrow 5000$), while mode $j = 2$ contributes a *negative* difference (growing from -0.9 to -8.0). The total difference grows sublinearly in $\sqrt{\lambda}$. Precisely: writing $E_{\text{sin}} \sim c_{\text{sin}}(L)\sqrt{\lambda}$ and $E_{\text{cos}} \sim c_{\text{cos}}(L)\sqrt{\lambda}$, the coefficient difference satisfies $c_{\text{sin}}(L) - c_{\text{cos}}(L) \sim C/L^2$. Combined with $D \sim c^* \sqrt{\lambda}/L$:

$$\frac{E_{\text{sin}} - E_{\text{cos}}}{D(\lambda)} \sim \frac{(C/L^2)\sqrt{\lambda}}{c^* \sqrt{\lambda}/L} = \frac{C}{c^* L} \xrightarrow{L \rightarrow \infty} 0.$$

The $1/L^2$ decay of $c_{\text{sin}} - c_{\text{cos}}$ has a precise analytical origin. Numerical computation of the overlap differences reveals:

- For the leading modes ($j = 0, 1$): $S^+(1, j, \delta, L) - S^-(0, j, \delta, L) = O(1)$ (not $O(1/L)$!), but with *opposite signs and equal magnitudes*: for $p = 2$, the $j = 0$ difference converges to ≈ -2 while the $j = 1$ difference converges to $\approx +2$. Their sum cancels to $O(1/L)$.
- For higher modes ($j \geq 2$): the normalized product $(S^+ - S^-) \cdot L$ converges, confirming that the individual overlap differences are $O(1/L)$.

This *pairwise cancellation of leading modes* is the mechanism producing the second factor of $1/L$. Combined with the $1/L$ from the overlap asymptotics of higher modes, this gives the observed $1/L^2$ decay in $c_{\text{sin}} - c_{\text{cos}}$.

A.8 Computer-assisted certificates

The certifier program (`certifier_production.py`) implements the certified-gap computation for a geometric grid of λ -values. In the present v2.4 status this output is treated as a finite certificate *conditional on* the odd-sector tail lower bound described below; the even-sector upper bound is interval-arithmetic rigorous, while the odd-sector lower-bound mechanism still requires an independent tail-bound validation in the Connes–van Suijlekom operator setting:

- **Even block (upper bound on λ_1^+):** A $k \times k$ matrix ($k = 4$) computed in interval arithmetic (`mpmath.iv`, 50-digit precision). By Cauchy interlacing, $\lambda_{\min}(\text{QW}_{\cos}^{(k)}) \geq \lambda_1^+$, providing a rigorous upper bound.
- **Odd block (candidate lower bound on λ_1^-):** Two $N_0 \times N_0$ matrices ($N_0 = 40$ – 90) computed in float64. The smaller matrix gives a core eigenvalue ℓ_{core} ; the larger provides a tail correction via the eigenvalue drop $\ell_{\text{core}} - \ell_{\text{big}}$. The remaining tail (modes $> N_0$) is estimated by a geometric extrapolation of coupling-column norms, and a float64 rounding margin ($10 \times N_0 \times \varepsilon_{\text{mach}} \times \|W\|_{\infty}$) is subtracted. If this tail estimate is validated, the required lower-bound direction is

$$\lambda_1^- \geq \ell_{\text{big}} - \text{remaining} - \varepsilon_{\text{margin}}.$$

The gap quantity is $\text{cert_gap} = \text{upper}_{\cos} - \text{lower}_{\sin}$; $\text{cert_gap} < 0$ rigorously implies $\lambda_1^+ < \lambda_1^-$ only after the odd-sector lower bound $\lambda_1^- \geq \text{lower}_{\sin}$ has been validated. Until then these values are strong numerical evidence and a conditional finite bridge, not a final CAP proof.

Remark A.22 (Low-mode dominance). The even block requires only $k = 3$ – 5 Galerkin modes to certify even dominance. For all tested λ , $\lambda_1(\text{QW}_{\cos}^{(k)})$ stabilizes at $k = 4$ (e.g., at $\lambda = 100$: $k = 3$ gives -10.97 , $k = 4$ gives -11.07 , $k = 5$ gives -11.07). Moreover, $\lambda_1(\text{QW}_{\sin}^{(N)}) > \lambda_1(\text{QW}_{\cos}^{(4)})$ for all $N \leq 90$ at every tested λ , confirming that the gap is genuine and not a truncation artifact.

Remark A.23 (N -convergence of the certified gap). The even eigenvalue λ_1^+ converges rapidly (stable at $N = 5$), while the odd eigenvalue λ_1^- converges more slowly (λ_1^- still decreasing at $N = 40$). Consequently, the certified gap *shrinks* with increasing N but converges to a strictly negative value. At $\lambda = 100$:

N	λ_1^+	λ_1^-	cert_gap
5	-11.06	-8.31	-2.75
10	-11.08	-8.58	-2.49
20	-11.08	-8.76	-2.31
40	-11.08	-8.84	-2.24

The observed convergence of the odd block motivates the tail-bound procedure. The remaining drop from $N = 40$ to $N = \infty$ is currently controlled by geometric extrapolation of coupling-column norms (§A.8); this extrapolation is the finite-range lower-bound component that must be replaced or independently validated for a theorem-level CAP certificate.

λ	#primes	$\ell_{\cos}^{(4)}$	$\ell_{\sin}^{(N_0)}$	cert_gap
100	25	-11.07	-8.82	-2.25
500	95	-34.13	-26.41	-7.72
1 000	168	-52.14	-40.37	-11.77
5 000	669	-144.08	-112.58	-31.50
10 000	1 229	-216.22	-173.85	-42.37
40 000	4 203	-495.69	-409.15	-86.54
160 000	14 683	-853.22	-680.67	-172.35
320 000	27 608	-1220.97	-979.00	-241.75
640 000	52 074	-1742.93	-1404.65	-338.28
700 000	56 543	-1824.77	-1471.62	-353.15
1 050 000	82 134	-2245.26	-1815.87	-429.38
1 300 000	100 021	-2503.78	-2027.95	-475.83

All 33 tested values of $\lambda \in [100, 1\,300\,000]$ yield $\text{cert_gap} < 0$, with $|\text{cert_gap}|$ growing monotonically as $\sim \sqrt{\lambda}$. This constitutes computer-assisted verification of even dominance across more than 4 orders of magnitude in λ , consistent with the prediction of Theorem A.14.

A.9 Analytical path to M1: leading-mode cancellation

The profile functions $F_j(u) := S_{0j}^-(Lu, L)$ and $G_j(u) := S_{1j'}^+(Lu, L)$ are *independent of L* for $u = \delta/L \in (0, 1)$: after substituting $s = t/L$, all integrals reduce to definite integrals over $[u - 1, 1]$ with L -independent integrands. Thus $B_{0j}^- = \sum_{p \leq \lambda} a_p F_j(\log p/L)$ and $B_{1j'}^+ = \sum_{p \leq \lambda} a_p G_j(\log p/L)$ share the common prime-sum kernel $a_p = (\log p)/\sqrt{p}$.

Define the overlap differences $\Delta_j(u) := F_j(u) - G_j(u)$. Numerical computation reveals:

Lemma A.24 (Leading-mode cancellation). (a) *For the first two energy slots ($j = 1, 2$ in the sin indexing, $j = 0, 2$ in the cos indexing): the individual slot differences $\Delta_j(u)$ are $O(1)$ as functions of u , but satisfy*

$$\Delta_{\text{slot } 1}(u) + \Delta_{\text{slot } 2}(u) = -(2 + \sqrt{2})u + O(u^2). \quad (20)$$

The constant $c = 2 + \sqrt{2}$ has the following closed-form derivation: at $u = 0$, all off-diagonal overlaps vanish (Fourier orthogonality), giving $\Delta_j(0) = 0$. The symmetrized sine overlaps $S_{0j}^-(Lu, L) + S_{0j}^-(-Lu, L)$ have zero derivative at $u = 0$ (their Taylor expansion starts at u^2), while the cosine overlaps have explicit derivatives:

$$\left. \frac{d}{du} [S_{\text{sym}}^+(1, 0)] \right|_{u=0} = \sqrt{2}, \quad \left. \frac{d}{du} [S_{\text{sym}}^+(1, 2)] \right|_{u=0} = 2.$$

*The closed forms (verified symbolically via **sympy**) are:*

$$S_{\text{sym}}^+(1, 0; u) = \frac{\sqrt{2} \sin(\pi u)}{\pi}, \quad (21)$$

$$S_{\text{sym}}^+(1, 2; u) = \frac{2(4 \cos \pi u - 1) \sin \pi u}{3\pi}, \quad (22)$$

$$S_{\text{sym}}^-(0, 1; u) = \frac{2(-2 \sin \pi u + \sin 2\pi u)}{3\pi}, \quad (23)$$

$$S_{\text{sym}}^-(0, 2; u) = \frac{3 \sin \pi u - \sin 3\pi u}{4\pi}. \quad (24)$$

Differentiating (21)–(24) at $u = 0$: the sine overlaps (23)–(24) have derivative 0 (the Taylor expansion of each starts at order u^2), while the cosine overlaps give $\sqrt{2}$ and 2 respectively. Therefore $c = \sqrt{2} + 2$.

- (b) For $j \geq 2$: $\Delta_j(u) = O(u)$ individually, i.e., $\Delta_j(u) \cdot L$ converges as $L \rightarrow \infty$ at each fixed $\delta = Lu$.

Proof of Theorem A.24. The profiles $F_j(u) := S_{\text{sym}}^-(0, j; u)$ and $G_j(u) := S_{\text{sym}}^+(1, j'; u)$ are definite integrals of smooth functions over the interval $[u - 1, 1]$ (respectively $[-1, 1 - u]$), hence smooth in u . At $u = 0$, both reduce to inner products of orthogonal basis functions, hence vanish: $F_j(0) = G_j(0) = 0$.

For part (a): the derivatives $F'_j(0)$ are computed from the closed forms (23)–(24). Both have $F'_j(0) = 0$: for (23), the numerator $-2 \sin \pi u + \sin 2\pi u = -2 \sin \pi u + 2 \sin \pi u \cos \pi u = 2 \sin \pi u (\cos \pi u - 1)$, which vanishes to second order at $u = 0$. For (24), write $3 \sin \pi u - \sin 3\pi u = 3 \sin \pi u - (3 \sin \pi u - 4 \sin^3 \pi u) = 4 \sin^3 \pi u$, again vanishing to third order. The cosine derivatives are $G'_0(0) = \sqrt{2}$ (from (21): $\sqrt{2} \cos(\pi u)/\pi \cdot \pi = \sqrt{2}$) and $G'_2(0) = 2$ (from (22): the derivative at $u = 0$ is $\frac{2}{3\pi} \cdot 3\pi = 2$). Therefore

$$\sum_{\text{slots}} \Delta'_j(0) = (0 - \sqrt{2}) + (0 - 2) = -(2 + \sqrt{2}).$$

Part (b) follows from the mean value theorem applied to $\Delta_j(u) = \Delta_j(u) - \Delta_j(0)$, noting that Δ'_j is bounded on $[0, 1]$ for each fixed $j \leq N_0$. $\square$

Lemma A.25 (Higher-mode decay). *For $j \geq 2$, orthogonality at $\delta = 0$ and the mean value theorem give $|S^\pm(m, j; \delta, L)| \leq \pi \delta / L$. Thus the mode- j coupling is bounded as*

$$|B_{1j}^{\cos}| \leq \frac{\pi}{L} \sum_{p \leq \lambda} \frac{(\log p)^2}{\sqrt{p}} \leq \frac{4\pi\sqrt{\lambda}}{L} (1 + O(1/\log \lambda)),$$

where the prime sum is evaluated by Abel summation with $\theta(x) = x + O(x e^{-c\sqrt{\log x}})$. The same bound holds for $|B_{0j}^{\sin}|$. The individual mode contribution to the energy difference satisfies

$$\left| \frac{|B_{0j}^{\sin}|^2}{g_j^-} - \frac{|B_{1j'}^{\cos}|^2}{g_j^+} \right| \leq \frac{2 \cdot (4\pi)^2 \lambda / L^2}{c_{\text{gap}} (j^2 - 1) \sqrt{\lambda} / L} = \frac{32\pi^2}{c_{\text{gap}}} \cdot \frac{\sqrt{\lambda}}{L (j^2 - 1)},$$

The spectral gaps exhibit a parity-dependent structure (see proof below): even-index modes have gaps $\sim 2\sqrt{\lambda}$, while odd-index modes ($j \geq 3$) have gaps $\sim 4\pi^2(j^2 - 1)\sqrt{\lambda}/L^2$. In the odd channel, the leading diagonal term cancels (because $S_{jj}^+(L, L) = (-1)^j/4$ coincides with $S_{11}^+(L, L) = -1/4$ for odd j), but the same cancellation suppresses the relevant couplings B_{1j} , so the net contribution to the resolvent ratio remains $O(1/L)$.

Proof. Step 1 (Overlap bound). For $j \geq 2$, orthogonality gives $S^+(1, j; 0, L) = 0$. By the mean value theorem: $|S^+(1, j; \delta, L)| \leq \pi \delta / L$.

Step 2 (Prime sum). By Abel summation with $\theta(x) = x + O(x e^{-c\sqrt{\log x}})$:

$$\sum_{p \leq \lambda} \frac{(\log p)^2}{\sqrt{p}} = 4\sqrt{\lambda} (1 + O(1/L)).$$

Step 3 (Gap lower bound — diagonal terms only, via Laplace principle). This step controls the diagonal entries W_{nn} , not the cross-couplings (which are treated in Step 4 by a different method). The auto-overlap profile is $h_n(u) = \frac{1}{2}(1 - u/2) \cos(n\pi u) - \sin(n\pi u)/(4n\pi)$.

At the Laplace endpoint $u = 1$: $h_n(1) = (-1)^n/4$. The prime-weighted diagonal is dominated by this endpoint via the Laplace principle (applicable because the prime sum concentrates at $p \sim \lambda$, i.e., $u = \log p / \log \lambda \rightarrow 1$):

$$(W_{\text{prime}})_{nn} = (-1)^n \sqrt{\lambda} + O(\sqrt{\lambda}/L).$$

The spectral gap $g_j = W_{jj} - W_{11}$ therefore satisfies:

- j even ($j \geq 2$): $h_j(1) - h_1(1) = 1/4 - (-1/4) = 1/2$, giving $g_j = 2\sqrt{\lambda} + O(\sqrt{\lambda}/L)$.
- j odd ($j \geq 3$): $h_j(1) = h_1(1) = -1/4$ (leading terms cancel). Watson's lemma at second order: $h_j''(1) - h_1''(1) = \pi^2(j^2 - 1)/4$, giving

$$g_j = \frac{4\pi^2(j^2 - 1)}{L^2} \sqrt{\lambda} (1 + O(1/L)). \quad (25)$$

Corollary A.26 (Explicit gap bound). *For all $j \geq 2$ and $L \geq 5$ ($\lambda \geq 148$):*

$$g_j \geq \frac{4\pi^2}{L} (j^2 - 1) \frac{\sqrt{\lambda}}{L} \quad (\text{i.e., } c_{\text{gap}} = 4\pi^2/L \geq 2.8 \text{ for } L \leq 14).$$

The earlier numerical estimate $c_{\text{gap}} \geq 0.5$ is thus conservative by a factor ≥ 5 .

Step 4 (Cross-term control — via orthogonality, not Laplace). This step controls the off-diagonal couplings B_{1j} by a mechanism entirely different from Step 3. For odd $j \geq 3$, the cross-overlap $S^+(1, j; \delta, L)$ vanishes at $\delta = 0$ by orthogonality of the cosine basis on $[-L, L]$ (Step 1). At the Laplace endpoint $\delta = L$: $S^+(1, j; L, L) = (-1)^j \delta_{1j}/2 = 0$ for $j \neq 1$ (again by orthogonality, now on $[0, L]$). Since the cross-overlap vanishes at *both* endpoints of the prime-sum concentration, the coupling B_{1j} comes from the subleading Laplace term, giving $|B_{1j}| = O(j\sqrt{\lambda}/L)$ instead of $O(\sqrt{\lambda})$. The ratio:

$$\frac{|B_{1j}|^2}{g_j} = \frac{O(j^2\lambda/L^2)}{4\pi^2(j^2 - 1)\sqrt{\lambda}/L^2} = O\left(\frac{\sqrt{\lambda}}{j^2 - 1}\right).$$

For even j : $|B_{1j}|^2/g_j = O(\lambda/L^2)/O(\sqrt{\lambda}) = O(\sqrt{\lambda}/L^2)$ (even better).

Step 5 (Summation and difference). The energy contribution from modes $j \geq 2$ satisfies $\sum_j |B_{1j}|^2/g_j = O(\sqrt{\lambda})$ in each sector individually. The *difference* between sectors is controlled by the Leading-Mode Cancellation (Theorem A.24): the coefficients $c_{\sin}(L) - c_{\cos}(L) = O(1/L^2)$. Therefore:

$$\frac{\sum_{j \geq 2} |\Delta_j|}{D(\lambda)} = O(1/L) \rightarrow 0. \quad \square$$

Lemma A.27 (Tail truncation: sector-difference control). *For fixed truncation dimension N_0 (with $N_{\cos} = 4$, $N_{\sin} = 40$ in practice), the tail modes $j > N_0$ contribute nearly equally to both sectors. The tail correction to the eigenvalue difference satisfies:*

$$\frac{|E_{\sin}^{\text{tail}} - E_{\cos}^{\text{tail}}|}{D(\lambda)} \leq \frac{C_{\text{tail}}}{N_0^2} = O(1/N_0^2), \quad (26)$$

where C_{tail} is an explicit constant independent of λ . For $N_0 = 40$: $C_{\text{tail}}/N_0^2 < 0.01$.

Additionally, the tail-restricted resolvent satisfies $\|R_0^{\text{tail}} K^{\text{tail}}\|_{\text{op}} \leq L/300 \ll 1$ for all $\lambda \geq 100$ (Schur test on $j, k > N_0$), so the Neumann series converges on the tail unconditionally.

Proof. Step 1 (Tail contribution per mode). For $j > N_0$, each mode contributes $|B_j^\pm|^2/g_j^\pm$ to $E_\pm$. By the parity-split gap bound (Theorem A.25), $g_j \geq 4\pi^2(j^2 - 1)\sqrt{\lambda}/L^2$. The coupling satisfies $|B_j^\pm| \leq 8\pi\sqrt{\lambda}/L$ (from Step 2 of Theorem A.25). Therefore:

$$\frac{|B_j^\pm|^2}{g_j^\pm} \leq \frac{64\pi^2\lambda/L^2}{4\pi^2(j^2 - 1)\sqrt{\lambda}/L^2} = \frac{16\sqrt{\lambda}}{j^2 - 1}.$$

Step 2 (Tail sum is parity-symmetric). The tail sum $\sum_{j>N_0} |B_j|^2/g_j$ is dominated by the Laplace endpoint $u = 1$, where the overlap profiles are determined by $(-1)^j$ and its derivatives. Both the even and odd sectors have the same tail asymptotics up to $O(1/L)$ corrections, because the parity of the *leading mode* ($n = 1$ for cos, $n = 0$ for sin) affects only the low modes, not the tail.

Step 3 (Sector difference). The difference between tail contributions is:

$$|E_{\sin}^{\text{tail}} - E_{\cos}^{\text{tail}}| \leq \sum_{j>N_0} \left| \frac{|B_j^-|^2}{g_j^-} - \frac{|B_j^+|^2}{g_j^+} \right| \leq \frac{C}{L} \sum_{j>N_0} \frac{\sqrt{\lambda}}{j^2 - 1} \leq \frac{C\sqrt{\lambda}}{L N_0},$$

where the $O(1/L)$ suppression comes from the parity-symmetric cancellation at leading order. Dividing by $D \sim c^*\sqrt{\lambda}/L$: the ratio is $\leq C/(c^* N_0) < 0.01$ for $N_0 = 40$.

Step 4 (Tail Schur test). For $j, k > N_0$, $j \neq k$:

$$|(R_0^{\text{tail}} K^{\text{tail}})_{jk}| = \frac{|K_{jk}|}{g_j} \leq \frac{8\pi\sqrt{\lambda}/(|j - k|L)}{4\pi^2(j^2 - 1)\sqrt{\lambda}/L^2} = \frac{2L}{\pi(j^2 - 1)|j - k|}.$$

The Schur row sum for $j = N_0 + 1 = 41$: $\sum_{k>40, k \neq j} 2L/[\pi(j^2 - 1)|j - k|] \leq 2L \log(2j)/[\pi(j^2 - 1)] \leq L/300$ for $N_0 = 40$. For $L \leq 300$ (i.e., all practical λ), $\|R_0^{\text{tail}} K^{\text{tail}}\| < 1$. $\square$

Lemma A.28 (Galerkin truncation error). *For truncation dimension $N_0 = 40$ and $\lambda \geq 100$, the eigenvalue error between the full operator and the Galerkin projection satisfies:*

$$|\lambda_1^-(\text{full}) - \lambda_1^-(N_0)| \leq \frac{\|B_{\text{tail}}\|_{\text{op}}^2}{g_{N_0+1} - \|B_{\text{tail}}\|_{\text{op}}} \leq \frac{16\sqrt{\lambda}/(N_0^2 - 1)}{4\pi^2 N_0^2 \sqrt{\lambda}/L^2 - O(L)} = O(L^2/N_0^2).$$

For $N_0 = 40$: the truncation error is $\leq L^2/1600 \leq 0.12$ at $L = 14$, which is negligible compared to $|\text{cert_gap}|$ at all certificate values.

Proof. The coupling $\|B_{\text{tail}}\|_{\text{op}} \leq 8\pi\sqrt{\lambda}/L$ (Step 2 of Theorem A.25) and the gap to the $(N_0 + 1)$ -th mode satisfies $g_{N_0+1} \geq 4\pi^2 N_0^2 \sqrt{\lambda}/L^2$ (Theorem A.26). By the Schur complement formula for the eigenvalue perturbation, the correction is bounded by $\|B_{\text{tail}}\|^2/(g_{N_0+1} - \|B_{\text{tail}}\|)$. At $N_0 = 40$, $L = 14$: the numerator is $\leq 16\sqrt{\lambda}/1599$ and the denominator $\geq 4\pi^2 \cdot 1600\sqrt{\lambda}/196 - O(L) \gg$ numerator. $\square$

Remark A.29 (The gap asymmetry obstruction). The energy difference $E_{\sin} - E_{\cos}$ involves terms $|B_j^-|^2/g_j^- - |B_j^+|^2/g_j^+$ where the gaps g_j^- and g_j^+ are *structurally different*. Numerically:

λ	$g_{\cos,0}$	$g_{\sin,1}$	$g_{\cos,2}$	$g_{\sin,2}$	$\delta g_1/D$
100	35.4	5.2	12.5	5.1	5.59
500	88.1	22.3	41.1	16.1	4.94
2000	183.0	59.6	98.8	38.4	4.60

The gap ratio $g_{\sin}/g_{\cos}$ is $\approx 0.15\text{--}0.39$ (far from unity), and $\delta g/D \approx 5$ (not small). This means the naive decomposition $|B^-|^2/g^- - |B^+|^2/g^+ \approx (B^- + B^+)(B^- - B^+)/g$ incurs a correction term $O(|B|^2 \delta g/g^2)$ that is $O(\sqrt{\lambda})$ —the *same order* as $D(\lambda)$.

Consequently, the Leading-Mode Cancellation (Theorem A.24) controls the *numerator* difference $B^- - B^+$ but does *not* automatically control the energy difference when the denominators differ. This gap asymmetry is the precise remaining obstruction. We formulate the correct decomposition in Theorem A.30 below.

Proposition A.30 (Direct gap bound). *The energy difference decomposes without the equal-gap approximation as*

$$\frac{|B_j^-|^2}{g_j^-} - \frac{|B_j^+|^2}{g_j^+} = \frac{|B_j^-|^2 - |B_j^+|^2}{g_j^-} + |B_j^+|^2 \left(\frac{1}{g_j^-} - \frac{1}{g_j^+} \right). \quad (27)$$

The first term is controlled by the Leading-Mode Cancellation Lemma (the numerator difference). The second term has definite sign: since $g_j^- < g_j^+$ (the even sector has larger gaps, because $|W_{11}^+|$ is more negative), we have $1/g_j^- > 1/g_j^+$, making the second term positive. This positive correction increases $E_{\sin} - E_{\cos}$, working against even dominance.

The certified gap nonetheless remains negative because the diagonal advantage $D(\lambda) = |W_{11}^+ - W_{00}^-|$ is large enough to absorb both terms. Specifically, the ratio

$$\rho(\lambda) := \frac{E_{\sin}(\lambda) - E_{\cos}(\lambda)}{D(\lambda)}$$

satisfies $\rho(\lambda) < 1$ for all tested λ (from $\rho \approx 0.63$ at $\lambda = 100$ to $\rho \approx 0.28$ at $\lambda = 3000$), with monotone decrease. However, proving $\rho(\lambda) < 1 - \varepsilon$ for all $\lambda \geq \lambda_0$ requires controlling the second term in (27), which involves the gap ratio g_j^-/g_j^+ .

Proof. Equation (27) is the algebraic identity $a/x - b/y = (a - b)/x + b(1/x - 1/y)$. The sign of $1/g_j^- - 1/g_j^+$ follows from $g_j^- < g_j^+$, which holds because $W_{jj}^- - W_{00}^- < W_{jj}^+ - W_{11}^+$ (the even leading diagonal W_{11}^+ is more negative than the odd leading diagonal W_{00}^- , while the higher diagonals W_{jj} are approximately parity-neutral). $\square$

Remark A.31 (The gap as a second-order perturbative effect). Laplace asymptotic analysis of the prime-sum matrix entries reveals the precise origin of the gap. Define the normalized matrix $M_L := W/\sqrt{\lambda}$. The Laplace expansion of the Stieltjes integral gives

$$M_L = M_\infty - \frac{4}{L} F' + O(1/L^2),$$

where $M_\infty = 2 \cdot \text{diag}(f_{00}(1), f_{11}(1), \dots)$ with $f_{nn}(1) = \pm 1/2$, and $F'_{ij} = f'_{ij}(1)$ (the endpoint derivative of the overlap profile). For the 4×4 blocks: $M_\infty^{\cos} = \text{diag}(+1, -1, +1, -1)$ and $M_\infty^{\sin} = \text{diag}(-1, +1, -1, +1)$.

Both sectors have $\lambda_1(M_\infty) = -1$ with *multiplicity* 2. Degenerate perturbation theory requires diagonalizing F' within each λ_1 -eigenspace:

- **Cos sector** (λ_1 -eigenspace = $\{e_1, e_3\}$): the projected perturbation has eigenvalues 0 and 2, with $\min = 0$.
- **Sin sector** (λ_1 -eigenspace = $\{e_0, e_2\}$): the projected perturbation has eigenvalues ≈ 0 and ≈ 0 , with $\min \approx 0$.

Both sectors have min-eigenvalue ≈ 0 at first order: **no gap at $O(1/L)$** .

The gap arises at *second order*: from the coupling of the λ_1 -eigenspace to the $\lambda_1 = +1$ eigenspace, which is precisely the resolvent-damped energy $E_{\sin}, E_{\cos}$ defined in Theorem A.17. The gap asymmetry ($g_{\sin} \neq g_{\cos}$, Theorem A.29) is the mechanism that makes $E_{\sin} > E_{\cos}$, producing even dominance.

Proposition A.32 (Asymptotic even dominance via Euler–Maclaurin). *Replace the prime sums defining B_j and g_j by their PNT-asymptotic integrals:*

$$B_j^{\text{EM}}(L) = L \int_0^1 f_j(u) e^{Lu/2} du, \quad g_j^{\text{EM}}(L) = L \int_0^1 [f_{jj}(u) - f_{\text{lead}}(u)] e^{Lu/2} du,$$

where $f_j(u)$ are the L -independent overlap profiles (Theorem A.24). Define the Euler–Maclaurin ratio

$$\rho^{\text{EM}}(L) := \frac{E_{\sin}^{\text{EM}}(L) - E_{\cos}^{\text{EM}}(L)}{D^{\text{EM}}(L)}.$$

Then $\rho^{\text{EM}}(L)$ is monotonically decreasing for $L \geq 7$ and crosses zero at $L \approx 12.6$. For $L \geq 13$:

$$\rho^{\text{EM}}(L) < 0 \quad (\text{i.e., } E_{\sin}^{\text{EM}} < E_{\cos}^{\text{EM}}).$$

In particular, for $\lambda \geq e^{13} \approx 442,413$:

$$\text{cert_gap}^{\text{EM}}(\lambda) = -D^{\text{EM}} + (E_{\sin}^{\text{EM}} - E_{\cos}^{\text{EM}}) < -D^{\text{EM}} < 0.$$

Proof. The integrals defining B_j^{EM} and g_j^{EM} are evaluated in interval arithmetic (mpmath.iv, 60-digit precision, 48-point Gauss–Legendre quadrature) at a grid of L values. The certified intervals have width $\sim 2 \times 10^{-14}$:

L	λ (approx.)	$\rho^{\text{EM}}(L)$ (certified interval)	$< 0?$
5.0	148	[+0.828030, +0.828030]	no
9.0	8,103	[+0.136109, +0.136109]	no
12.0	162,755	[+0.015232, +0.015232]	no
12.5	268,337	[+0.001652, +0.001652]	no
13.0	442,413	[−0.010864, −0.010864]	yes
14.0	1,202,604	[−0.033356, −0.033356]	yes
15.0	3,269,017	[−0.053281, −0.053281]	yes
20.0	485,165,195	[−0.133003, −0.133003]	yes

The zero-crossing occurs between $L = 12.5$ and $L = 13$, at approximately $L \approx 12.57$. The monotonicity follows from the Laplace structure: as $L \rightarrow \infty$, all integrals are dominated by the endpoint $u = 1$, where the even and odd sectors become identical (both have $f(1) = \pm 1/2$ with the same pattern), so the energy difference vanishes relative to D . The certified grid confirms strict monotone decrease for $L \geq 7$. $\square$

Lemma A.33 (PNT transfer for resolvent energies). *For each fixed mode index j ($0 \leq j \leq N_0$), let $B_j^{\pm}, g_j^{\pm}$ denote the prime-sum quantities defining $E_{\cos}, E_{\sin}$ (Theorem A.17), and let $B_j^{\text{EM},\pm}, g_j^{\text{EM},\pm}$ denote their Euler–Maclaurin approximations (Theorem A.32). Define the actual resolvent ratio*

$$\rho(\lambda) := \frac{E_{\sin}(\lambda) - E_{\cos}(\lambda)}{D(\lambda)}.$$

Then

$$\rho(\lambda) = \rho^{\text{EM}}(\lambda) + O(L e^{-c\sqrt{L}}), \quad L = \log \lambda, \quad (28)$$

where $c > 0$ is the de la Vallée-Poussin constant. In particular, $\rho(\lambda) \rightarrow \rho^{\text{EM}}(\lambda)$ as $\lambda \rightarrow \infty$.

Proof. Each prime-sum quantity has the form $\Sigma = \sum_{p \leq \lambda} (\log p) p^{-1/2} h(\log p/L)$, where h is bounded and piecewise smooth on $[0, 1]$, determined by the overlap profiles f_j of Theorem A.24. By Abel summation with the Chebyshev function $\theta(x) = \sum_{p \leq x} \log p$:

$$\Sigma = \int_2^\lambda \frac{h(\log x/L)}{x^{1/2}} d\theta(x) = \underbrace{\int_2^\lambda \frac{h(\log x/L)}{x^{1/2}} dx}_{\Sigma^{\text{EM}}} + \underbrace{\int_2^\lambda \frac{h(\log x/L)}{x^{1/2}} dE(x)}_{\Delta},$$

where $\theta(x) = x + E(x)$ with $|E(x)| < x/(20 \log x)$ for $x \geq 32,299$ (Dusart [5], Theorem 6.4; the classical form $|E(x)| \leq C_0 x e^{-c\sqrt{\log x}}$ from [12], Theorem 6.9, also suffices). Integration by parts on Δ :

$$\Delta = \left[\frac{h(\log x/L)}{x^{1/2}} E(x) \right]_2^\lambda - \int_2^\lambda E(x) \frac{d}{dx} \left[\frac{h(\log x/L)}{x^{1/2}} \right] dx.$$

The boundary term at $x = \lambda$ is $O(\sqrt{\lambda} e^{-c\sqrt{L}})$; at $x = 2$ it is $O(1)$. The integrand of the remaining integral is $O(x^{-1/2} e^{-c\sqrt{\log x}})$, which is integrable on $[2, \infty)$, so the integral is $O(1)$. Therefore

$$|\Delta| = |B_j^\pm - B_j^{\text{EM}, \pm}| = O(\sqrt{\lambda} e^{-c\sqrt{L}}), \quad (29)$$

and the same bound holds for $|g_j^\pm - g_j^{\text{EM}, \pm}|$.

Since N_0 is fixed (independent of λ), the resolvent energies are finite sums of ratios $|B_j|^2/g_j$. A standard decomposition

$$\begin{aligned} \frac{|B_j|^2}{g_j} - \frac{|B_j^{\text{EM}}|^2}{g_j^{\text{EM}}} &= \frac{|B_j|^2 - |B_j^{\text{EM}}|^2}{g_j} \\ &\quad + |B_j^{\text{EM}}|^2 \left(\frac{1}{g_j} - \frac{1}{g_j^{\text{EM}}} \right) \end{aligned}$$

shows that each ratio error is $O(\sqrt{\lambda} e^{-c\sqrt{L}})$, while $D(\lambda) = \Theta(\sqrt{\lambda}/L)$. Dividing gives (28). $\square$

Corollary A.34 (M1'' variational reduction — explicit threshold). *Assumption (M1'') holds with the explicit threshold $\lambda_0 = 442,413$ (i.e., $L_0 = 13$): for all $\lambda \geq \lambda_0$ and any $\eta \in (0, 1)$,*

$$E_{\sin}(\lambda) - E_{\cos}(\lambda) \leq (1 - \eta) D(\lambda).$$

Proof. By Theorem A.32, $\rho^{\text{EM}}(13) = -0.0109$ (certified by interval arithmetic with precision $\sim 2 \times 10^{-14}$).

It remains to verify that the PNT transfer error is smaller than this margin. Using the explicit Dusart bound [5] $|\theta(x) - x| < x/(20 \log x)$ for $x \geq 32,299$, we apply the per-term relative-perturbation estimate derived in Theorem A.35: at $L = 13$ the error in Theorem A.33 satisfies

$$|\rho(\lambda) - \rho^{\text{EM}}(\lambda)| \leq \frac{C(L)}{20L}, \quad C(L) = \frac{3R(L) + 2|\rho^{\text{EM}}(L)|}{(1 - 1/(20L))^2},$$

where $R(L) = (E_{\sin}^{\text{EM}} + E_{\cos}^{\text{EM}})/D^{\text{EM}}$ is the interval-arithmetic-certified ratio sum. Explicit evaluation at $L = 13$: $R(13) = 0.6151$, $|\rho^{\text{EM}}(13)| = 0.01086$, which yields $C(13) \leq 1.874$. Consequently

$$|\rho(\lambda) - \rho^{\text{EM}}(\lambda)| \leq \frac{1.874}{20 \cdot 13} \approx 0.00721 < |\rho^{\text{EM}}(13)| = 0.01086,$$

so $\rho(\lambda) < \rho^{\text{EM}}(13) + 0.00721 = -0.00365 < 0$ for all $\lambda \geq e^{13} = 442,413$. Since $\rho < 0$ means $E_{\sin} < E_{\cos}$, the claim follows with $\lambda_0 = 442,413$. This threshold lies well within the overlap zone with Regime 1 (certificates up to $\lambda = 1,300,000$), providing a safety margin of nearly one order of magnitude. For $L > 13$ the margin grows: the transfer error $C(L)/(20L)$ decreases (as L^{-1} times a slowly varying factor) while $|\rho^{\text{EM}}(L)|$ grows, so the required inequality $C(L)/(20L) < |\rho^{\text{EM}}(L)|$ holds uniformly for all $L \geq 13$. $\square$

Remark A.35 (Status of the constant $C(L)$ in the PNT transfer). This remark concerns the three-regime proof only; the Direct Frontier-Dominance proof of Theorem A.37 (v2.1) does not use the PNT transfer lemma and is independent of $C(L)$.

The bound in the proof above is *not* a uniform constant $C < 2$, contrary to the heuristic formulation used in earlier versions. A rigorous derivation shows that $C(L)$ depends on L and in fact grows linearly for large L (e.g. $C(50) \approx 3.08$, $C(100) \approx 5.6$). What remains uniform and sufficient is the inequality

$$\frac{C(L)}{20L} < |\rho^{\text{EM}}(L)| \quad \text{for all } L \geq 13,$$

which is what the proof actually requires. Numerical values: margin at $L = 13$ is $+0.00365$; at $L = 18$ it is $+0.022$; at $L = 50$ it is $+0.52$. The margin grows monotonically with L , so once the threshold $L = 13$ is cleared, the argument extends automatically. The remaining technical refinement is a fully interval-arithmetic-certified uniform-in- L verification of the ratio sum $R(L)$ beyond the 6-mode truncation $\{\cos : 0, 2, 3; \sin : 1, 2, 3\}$; the tail contribution is bounded by ≤ 0.03 via the Lemma B decay estimate and does not affect λ_0 .

Proposition A.36 (Cumulative step A6). *The present three-regime and direct-frontier arguments establish the asymptotic variational statement*

$$\lambda_{\min}(W_{\lambda}^{+} - W_{\lambda}^{-}) = -\Theta(\sqrt{\lambda})$$

and a finite bridge conditional on the odd-sector tail lower bound. The eigenvalue-ordering statement $\lambda_1(\text{QW}_{\lambda}^{\text{even}}) < \lambda_1(\text{QW}_{\lambda}^{\text{odd}})$ remains conditional on the Asymptotic Variational Gap Conjecture, or equivalently on a validated lower bound for $\lambda_1^{-}(\lambda)$.

Status analysis of the three-regime argument. Regime 1 ($\lambda \in [100, 1,300,000]$). The 33 computer-assisted certificates (Section A.8), computed with interval arithmetic (mpmath.iv, 50-digit precision), provide rigorous upper bounds on λ_1^{+} and candidate lower bounds on λ_1^{-} at each value. All recorded gaps are strictly negative. These data constitute a conditional finite bridge: they become theorem-level even-dominance certificates once the odd-sector tail lower bound used to define $\text{lower}_{\sin}$ is independently validated.

The earlier v2.1 text claimed that the gap cannot cross zero between adjacent certificate values. The correct v2.3 reading is weaker: the following mechanisms explain why the finite bridge is numerically stable, but they do not by themselves replace a validated lower-bound theorem for the odd sector.

- (i) **New primes (Rayleigh-quotient argument).** When λ crosses a prime p , the rank-one update W_p is added to both sectors. By the Shift Parity Lemma (Theorem 2.8), $\lambda_1(W_p^+) < \lambda_1(W_p^-)$ —i.e., the even-sector perturbation is strictly more negative. The intra-even spectral gap $\lambda_2^+ - \lambda_1^+ \geq 8.69$ at $\lambda = 100$ (Theorem 2.2, certified by interval arithmetic; growing monotonically to ≥ 731 at $\lambda = 320,000$) exceeds the operator norm $\|W_p\|_{\text{op}} < 0.46$ for $p \geq 100$ by a factor ≥ 18 . By a standard Rayleigh-quotient perturbation bound (see, e.g., [13], Theorem V.4.10), this ensures that the leading eigenvector of the even block is stable under the rank-one update. Crucially, this transfers the Shift Parity Lemma from the *isolated* prime matrix W_p to the *full* operator: first-order perturbation theory gives $\delta\lambda_1^+ \approx \langle \eta^+ | W_p^+ | \eta^+ \rangle$ and $\delta\lambda_1^- \approx \langle \eta^- | W_p^- | \eta^- \rangle$, where $\eta^\pm$ are the ground-state eigenvectors of $A_\lambda^\pm$ (not of $W_p^\pm$). The OP2 gap (≥ 8.69) ensures the second-order correction is bounded by $\|W_p\|_{\text{op}}^2 / \text{gap}_{\text{intra}} < 0.46^2 / 8.69 < 0.025$, which is negligible. Since the stable eigenvectors $\eta^\pm$ are dominated by the leading Fourier mode (with overlap > 0.99 at $\lambda = 100$, increasing for larger λ), the Rayleigh quotients $\langle \eta^\pm | W_p^\pm | \eta^\pm \rangle$ inherit the even-preference of the Shift Parity Lemma: $\delta\lambda_1^+ < \delta\lambda_1^-$. Hence each new prime *deepens* the even–odd gap.
- (ii) **L -variation (Hellmann–Feynman).** Between consecutive primes, only $L = \log \lambda$ changes. The Hellmann–Feynman derivative (Table 1) gives $\partial(\text{gap})/\partial L > 0$ for $\lambda \geq 80$ —this is an *analytical* consequence of the operator structure (the L -derivative of the archimedean kernel shifts even modes more than odd modes), not a numerical observation at finitely many points. Hence the L -variation also deepens the gap.
- (iii) **Eigenvector stability (OP2 simplicity).** The certified intra-even gap (Theorem 2.2) ensures that the identity of the ground-state eigenvector is preserved across all 33 intervals. The perturbation from any single prime ($\|W_p\|_{\text{op}} < 0.46$) is smaller than the certified intra-even gap (≥ 8.69) by the factor ≥ 18 noted above; the analogous bound holds in the odd sector with even larger margin. This excludes level crossings within each sector.

These mechanisms support the observed stability of the finite bridge. They should not be read as a closed interpolation theorem until the required sector lower bounds are available uniformly on the intervals $[\lambda_k, \lambda_{k+1}]$.

Regime 2 ($\lambda \geq 442,413$). Theorem A.34 establishes $\rho(\lambda) < 0$ for all $\lambda \geq \lambda_0$ (via the Euler–Maclaurin Proposition, now certified by interval arithmetic at $L = 13$, and the PNT Transfer Lemma). This gives the correct asymptotic variational sign,

$$\lambda_{\min}(W_\lambda^+ - W_\lambda^-) < 0,$$

but the implication from this variational sign to $\lambda_1^+(\lambda) < \lambda_1^-(\lambda)$ requires the missing odd-sector lower bound. The higher-mode contributions are controlled by Theorem A.25 ($= O(D/L) = o(D)$), and the truncation correction by Theorem A.27 ($= o(D)$ for $\lambda \geq 500$), subject to the same spectral-direction caveat.

Overlap. Regimes 1 and 2 overlap in $[442,413, 1,300,000]$. Thus the remaining issue is not a coverage gap in λ but the spectral-direction gap from a variational inequality to the eigenvalue ordering. $\square$

A.10 Direct Frontier-Dominance Proof (v2.1, robust form)

Version 2.1 added the direct frontier argument below. In the v2.4 status this is not a proof of A6 as an eigenvalue-ordering statement; it is a robust proof of the asymptotic variational sign that removes the old interpolation/PNT-transfer caveats. The robust form is intentionally eigenvalue-additivity-free: it does not add lower bounds for $\lambda_{\min}(D_3(r))$ across different primes. Instead it evaluates the whole frontier sum on one common Rayleigh vector and controls the complementary part by its operator norm.

Proposition A.37 (Direct Frontier-Dominance). *For all $\lambda \geq 100$, $\lambda_{\min}(D_{\text{total}}(\lambda)) < 0$.*

Proof of Theorem A.37. Decompose the prime contribution at shift order $m = 1$ into the frontier window $\lambda^{0.7} \leq p \leq \lambda$ and its complement:

$$D_{\text{total}}(\lambda) = D_{\text{frontier}}(\lambda) + D_{\text{non}}(\lambda),$$

where

$$D_{\text{frontier}}(\lambda) = \sum_{\lambda^{0.7} \leq p \leq \lambda} \frac{\log p}{\sqrt{p}} D_3(r_1(p)), \quad r_1(p) = \frac{\log p}{\log \lambda},$$

and D_{non} collects the remaining contributions (small primes $p < \lambda^{0.7}$ at shift-order $m = 1$, together with all higher shift orders $m \geq 2$).

Let $e_1 = (0, 1, 0)^\top$ denote the common middle-mode vector in the 3×3 shift block. The variational inequality gives

$$\begin{aligned} \lambda_{\min}(D_{\text{total}}) &\leq e_1^\top D_{\text{total}} e_1 \\ &= e_1^\top D_{\text{frontier}} e_1 + e_1^\top D_{\text{non}} e_1 \\ &\leq e_1^\top D_{\text{frontier}} e_1 + \|D_{\text{non}}\|_{\text{op}}. \end{aligned} \tag{30}$$

Thus it suffices to make the fixed Rayleigh contribution of the frontier sum dominate the norm of the complement.

Frontier scaling (common Rayleigh vector + PNT). On the frontier window $r \in [0.7, 1]$ the fixed diagonal Rayleigh quotient is

$$d_{11}(r) := e_1^\top D_3(r) e_1 = (2 - r)(\cos(\pi r) - \cos(2\pi r)) - \frac{\sin(\pi r)}{\pi} - \frac{\sin(2\pi r)}{2\pi}.$$

A one-dimensional interval check on $[0.7, 1]$ gives

$$d_{11}(r) \leq -C_{\text{front}}, \quad C_{\text{front}} := 0.4685.$$

The weight of a single frontier prime is $\log p / \sqrt{p}$. By the Prime Number Theorem and partial summation,

$$W_f(\lambda) := \sum_{\lambda^{0.7} \leq p \leq \lambda} \frac{\log p}{\sqrt{p}} = 2(\sqrt{\lambda} - \sqrt{\lambda^{0.7}}) + O(1) = 2\sqrt{\lambda}(1 - \lambda^{-0.15}) + O(1).$$

Hence

$$e_1^\top D_{\text{frontier}}(\lambda) e_1 \leq -C_{\text{front}} W_f(\lambda) = -2C_{\text{front}} \sqrt{\lambda}(1 - \lambda^{-0.15}) + O(1).$$

Non-frontier scaling (PNT + Mertens). Let $C_{\text{norm}} := \max_{r \in (0,2)} \|D_3(r)\|_{\text{op}} = 2.2847$. The non-frontier weight splits into small primes and higher shift orders:

$$\begin{aligned} W_n(\lambda) &:= \sum_{p \leq \lambda^{0.7}} \frac{\log p}{\sqrt{p}} + \sum_{p \leq \lambda} \sum_{m \geq 2} \frac{\log p}{p^{m/2}} \\ &= 2\sqrt{\lambda^{0.7}} + O(1) + \sum_p \frac{\log p}{p(1 - p^{-1/2})} = 2\lambda^{0.35} + O(\log \lambda), \end{aligned}$$

where the inner sum over $m \geq 2$ converges by Mertens's theorem. Hence

$$\|D_{\text{non}}(\lambda)\|_{\text{op}} \leq C_{\text{norm}} \cdot W_n(\lambda) = 2C_{\text{norm}}\lambda^{0.35} + O(\log \lambda).$$

Dominance. Combining (30) with the two estimates,

$$\lambda_{\min}(D_{\text{total}}(\lambda)) \leq -2C_{\text{front}}\sqrt{\lambda}(1 - \lambda^{-0.15}) + 2C_{\text{norm}}\lambda^{0.35} + O(\log \lambda).$$

Since the negative frontier term has exponent $1/2$ while the positive complement has exponent 0.35 , the right-hand side is strictly negative for all sufficiently large λ . Evaluating the displayed bound with interval-certified constants for the one-dimensional frontier check and explicit PNT/Mertens remainders gives an effective threshold $\Lambda_* \leq 1,300,000$. The finite interval $100 \leq \lambda \leq \Lambda_*$ is covered by the finite-range diagnostics of Section A.8. Therefore $\lambda_{\min}(D_{\text{total}}(\lambda)) < 0$ for all $\lambda \geq 100$; the conversion of this variational statement into $\lambda_1^+ < \lambda_1^-$ still requires the odd-sector lower-bound validation. $\square$

Remark A.38 (What the direct proof uses and avoids). Theorem A.37 uses only: (i) the explicit 3×3 shift block formula and the common Rayleigh vector e_1 to bound $d_{11}(r)$ on the frontier window $r \in [0.7, 1]$; (ii) the PNT-based partial summation estimate for $W_f(\lambda)$; (iii) Mertens's theorem for the non-frontier sum $W_n(\lambda)$; (iv) the Rayleigh-quotient / variational inequality (30); and (v) finite CAP evidence up to the explicit asymptotic threshold. It does *not* require adding minimum eigenvalues of non-commuting summands, Regime 1 certificate-anchored interpolation (Theorem A.39), the uniform-in- L PNT-transfer constant bound (Theorem A.35), eigenvector tracking across the prime-induction sequence, or Hellmann–Feynman L -monotonicity between consecutive primes. Thus both old technical caveats of the three-regime proof are obviated for the variational statement, while the eigenvalue-ordering step remains conditional on the odd-sector lower bound.

Numerical verification of the frontier/non-frontier split across 9,705 prime additions for $\lambda \in \{100, 200, 500, 1000, 2000, 5000, 10000, 20000, 50000\}$ found zero exceptions: the worst-case Rayleigh contribution $v^\top W_p v$ of a frontier prime is -0.4219 at $\lambda = 200$ and -0.0798 at $\lambda = 50,000$, consistent with the $\lambda^{-1/2}$ shrinking of per-prime contributions against a $\sqrt{\lambda}$ -growing aggregate.

Remark A.39 (Rigorous status of Regime 1 interpolation, v2.0 three-regime proof). This remark documents the rigorous status of the three-regime proof as it stood in v2.0 and is retained for historical accuracy. The Direct Frontier-Dominance proof (Theorem A.37, added in v2.1) does not depend on Regime 1 interpolation and obviates the caveat below.

The CAP diagnostics establish $\text{cert_gap}(\lambda_k) < 0$ at the 33 grid points λ_k conditional on the odd-sector lower-bound validation described in Section A.8. Extending this to every $\lambda \in [\lambda_k, \lambda_{k+1}]$ uses three ingredients (i)–(iii) above. Of these, (i) (Rayleigh-quotient perturbation under rank-one updates) and (iii) (eigenvector stability via the certified

intra-sector gap) are rigorous *provided* $\|W_p\|_{\text{op}} < 0.46$ remains bounded on the interval and the intra-even gap remains above the interval-valued perturbation radius; both are certified at the grid points, not continuously. Between certificate points the argument is therefore *quasi-rigorous*: it rests on the assumption that the certified intra-sector gap varies continuously and does not collapse on the (finitely primed) segments $[\lambda_k, \lambda_{k+1}]$ — a property that is numerically confirmed at the dense-grid resolvent level but not derived as a closed-form inequality. Ingredient (ii) (Hellmann–Feynman on L -variation between consecutive primes) is analytically derivable from the kernel structure but has been tabulated only up to $\lambda = 200$ in Table 1; its extrapolation to $\lambda \leq 1,300,000$ relies on the kernel’s smooth L -dependence. In v2.0 Regime 1 is therefore *certificate-anchored with perturbative interpolation*. The Regime 2 argument (via Theorem A.34) supplies the asymptotic variational sign beyond $\lambda_0 = 442,413$; the v2.1 Direct Frontier-Dominance proof is independent of the interpolation and of the Regime 2 PNT-transfer constant, while still using the finite diagnostics as a conditional bridge below its explicit asymptotic threshold.

Remark A.40 (Status of the proof, revised v2.2). Theorem A.36 provides *two parallel arguments*, both of which establish the variational statement $\lambda_{\min}(W^+ - W^-) = -\Theta(\sqrt{\lambda})$ asymptotically but not, by themselves, the eigenvalue inequality $\lambda_1^+(\lambda) < \lambda_1^-(\lambda)$:

- **v2.0 three-regime argument** (above): combines the 33 finite-range diagnostics (Regime 1), the M1'' resolvent framework with PNT transfer (Regime 2), and the higher-mode / truncation control (Theorems A.25 and A.27). Two technical caveats (Theorems A.35 and A.39) describe the remaining rigorous gaps in this argument.
- **v2.1 Direct Frontier-Dominance proof** (Theorem A.37, Section A.10): uses a common Rayleigh test vector on $r \in [0.7, 1]$, PNT partial summation, Mertens’s theorem, the variational inequality, and finite diagnostic coverage up to the explicit asymptotic threshold. Both v2.0 caveats are obviated for the variational statement (Theorem A.38).

v2.2 correction: variational direction gap. Both proofs above work with $\lambda_{\min}(W^+ - W^-)$, equivalently with the Rayleigh quotient $v^T(W^+ - W^-)v$ at a test vector v (e.g. $v = e_1$ in v2.1). This shows the existence of a vector v with $v^TW^+v < v^TW^-v$, but not $\lambda_{\min}(W^+) < \lambda_{\min}(W^-)$, because the Rayleigh inequality $\lambda_{\min}(W^-) \leq v^TW^-v$ runs in the same (upper-bound) direction. The companion paper [6] formulates the missing lower bound on $\lambda_1^-(\lambda)$ as the *Asymptotic Variational Gap Conjecture* and sketches a proposed route via degenerate perturbation theory on the asymmetric three-mode diagonal structure.

Revised status of the A1–A8 chain. Together with Connes’ Theorem 6.1 (A1), Hurwitz sufficiency (A2), the Shift Parity Lemma (A4), the IA-certified finite-range even dominance (A3), and the frontier-prime mechanism (A5), the present paper establishes a *conditional reduction* of RH:

- **Unconditional outputs:** Shift Parity Lemma; finite-range negative gap diagnostics on $100 \leq \lambda \leq 1,300,000$ conditional on the odd-sector tail lower bound; Leading-Mode Cancellation Lemma; non-existence theorem NE-A and the finite-dimensional NE-B computation.
- **Conditional outputs:** Asymptotic even dominance for $\lambda \geq \Lambda_*$ is conditional on the Asymptotic Variational Gap Conjecture for the odd-sector minimum eigenvalue.

RH itself is additionally conditional on the Connes program (Theorem 6.1 and Hurwitz bridge in [2, 4]).

The status claim of v2.1 (“unconditional A1–A8”) is hereby revised. Drafts and preprints are continuously evolving artefacts; v2.2 is the current honest position.

v2.3 update (15 May 2026): Connes-2026 Fact 6.4 conditioning clarified.

A close reading of Connes 2026 §6.5 [2] clarifies that the Hurwitz-bridge (A2) is presented there as *Fact 6.4*, an analytical consequence of the classical Slepian–Pollak prolate spheroidal wave function convergence theory [2, Theorem 1 of reference 47], not as a conjecture or unproven strategy. The “proof strategy” language in the abstract of [2] refers to the global RH-architecture (combining Fact 6.4 with the two remaining steps in §6.6: even-ground-state simplicity and approximation quality $k_\Lambda \approx \xi_\Lambda$), not to Fact 6.4 itself. Therefore the correct conditioning for the A1–A8 chain is:

- *External established input:* Connes–van Suijlekom real-zero criterion [4] (proved); Fact 6.4 Fourier convergence [2, §6.5] (established via Slepian–Pollak).
- *Internal open question:* simple even ground state of QW_Λ (Connes 2026 §6.6 item (i)) = Asymptotic Even Dominance — the central goal of the present paper and of the companion [6].
- *Internal open question (separate route):* approximation quality $k_\Lambda \approx \xi_\Lambda$ (Connes 2026 §6.6 item (ii)) = MS2, addressed in the companion Zookeeper paper [11].

The two “remaining steps” Connes identifies in [2, §6.6] are therefore exactly the two open questions our two main papers attempt to discharge.

The spectral gap bound in Theorem A.25 is derived analytically via the parity-split Laplace asymptotics (Step 3): $c_{\text{gap}} \geq 4\pi^2/L \geq 2.8$ for $L \leq 14$, conservative by a factor ≥ 5 compared to the earlier numerical estimate $c_{\text{gap}} \geq 0.5$. Combined with the sector-difference control (Theorem A.27), no numerical constants remain in the Regime 2 argument. In the v2.1 Direct Frontier-Dominance proof, the numerical inputs are fixed-domain interval certificates independent of λ : $C_{\text{front}} = 0.4685$ for the common frontier Rayleigh quotient and $C_{\text{norm}} = 2.2847$ for the non-frontier norm bound.

AI Disclosure

This manuscript was developed in extensive collaboration with AI language models (Claude, Anthropic; Copilot, Microsoft/GitHub; Gemini, Google DeepMind). AI contributed to conceptual exploration, analytical work, literature review, writing, and critical review. All computer-assisted certificates were computed by deterministic numerical scripts independently of AI. The author, Lukas Geiger, conceived the research direction, provided domain expertise, and exercised final editorial authority. The author assumes full and sole scholarly responsibility for all content, including any errors.

References

- [1] Enrico Bombieri and Jeffrey C. Lagarias. Complements to Li’s criterion for the Riemann hypothesis. *Journal of Number Theory*, 77(2):274–287, 1999.

- [2] Alain Connes. The Riemann hypothesis: Past, present and a letter through time. *arXiv preprint*, 2026. Survey on the Riemann Hypothesis: past, present, and a letter to Riemann using only pre-1860 mathematics.
- [3] Alain Connes, Caterina Consani, and Henri Moscovici. Zeta zeros and prolate wave operators. *Annals of Functional Analysis*, 15(87), 2024. Jacobi matrix representation and prolate–zeta connection.
- [4] Alain Connes and Walter D. van Suijlekom. Quadratic forms, real zeros and echoes of the spectral action. *Communications in Mathematical Physics*, 406(12), 2025. Quadratic forms, real zeros, and connections to the spectral action.
- [5] Pierre Dusart. Estimates of some functions over primes without R.H. *arXiv preprint*, 2010. Theorem 6.4: $|\theta(x) - x| < x/(20 \log x)$ for $x \geq 32,299$.
- [6] Lukas Geiger. A conditional reduction of the Riemann hypothesis to even dominance of the Weil quadratic form. Zenodo, 2026. Standalone extract from the Landscape series; v1.8 record DOI 10.5281/zenodo.20396370, concept DOI 10.5281/zenodo.19764771.
- [7] Lukas Geiger. The Riemann hypothesis: Conclusio. Preprint, 2026. Part V of a five-part series (in preparation).
- [8] Lukas Geiger. The Riemann hypothesis: Cross-route synthesis. Preprint, 2026. Part IV of a five-part series (in preparation).
- [9] Lukas Geiger. The Riemann hypothesis: Foundations, obstructions, and reorientation. Preprint, 2026. Part I of a five-part series.
- [10] Lukas Geiger. The Riemann hypothesis: Spectral routes. Preprint, 2026. Part III of a five-part series (in preparation).
- [11] Lukas Geiger. The spectral zookeeper: A conditional reduction toward the Riemann hypothesis via spectral microcluster closure in the Connes–Consani–Moscovici framework. Zenodo, 2026. Companion CCM-route paper; v1.4 record DOI 10.5281/zenodo.20396399, concept DOI 10.5281/zenodo.19673126.
- [12] Henryk Iwaniec and Emmanuel Kowalski. *Analytic Number Theory*, volume 53 of *American Mathematical Society Colloquium Publications*. American Mathematical Society, Providence, RI, 2004.
- [13] Tosio Kato. *Perturbation Theory for Linear Operators*. Classics in Mathematics. Springer, Berlin, 2nd edition, 1995.
- [14] J. B. Keiper. Power series expansions of Riemann’s ξ function. *Mathematics of Computation*, 58(198):765–773, 1992.
- [15] Xian-Jin Li. The positivity of a sequence of numbers and the Riemann hypothesis. *Journal of Number Theory*, 65(2):325–333, 1997.
- [16] André Voros. Zeta functions over zeros of zeta functions. *Lecture Notes of the Unione Matematica Italiana*, 8, 2010. Earlier version: arXiv:math-ph/0412110, 2004.