

The Riemann Hypothesis: Conclusio

Lukas Geiger

Bernau, Germany

ORCID: [0009-0005-7296-1534](https://orcid.org/0009-0005-7296-1534)

May 15, 2026

Abstract

This final paper in a three-part series draws together the results of the preceding papers into a unified picture. We catalog the proven structural results (R1–R9), the excluded paths (K1–K4), and the new contributions: the Shift Parity Lemma, the computer-assisted finite-range gap diagnostics for 33 values of λ up to 1,300,000, the resolvent-damped M1'' framework, the Leading-Mode Cancellation Lemma with exact constant $c = 2 + \sqrt{2}$, and the Euler–Maclaurin Proposition establishing $\rho^{\text{EM}}(L) < 0$ for $L \geq 14$ ($\lambda \geq 1,200,000$). The finite-range diagnostics are theorem-level only conditional on a validated odd-sector lower bound. Part III also accounts for the v2.0 non-existence theorems (NE-A, NE-B in Part II), which restructure the comparative analysis of proof strategies and confirm that even dominance is not reducible to total-positivity or Sturm–Liouville-type arguments. We present the proof architecture from Connes’ Theorem 6.1 through even dominance to the Riemann Hypothesis as a conditional reduction, formulate open problems, and discuss connections to the wider landscape of analytic number theory.

Status correction (v2.2, 2026-05-14). The v2.1 claim that the Direct Frontier-Dominance argument renders the A1–A8 reduction unconditional has been revised after external critical review. The Direct Frontier-Dominance argument establishes $\lambda_{\min}(W^+ - W^-) = -\Theta(\sqrt{\lambda})$ asymptotically via a common Rayleigh test vector, which is a variational statement, not $\lambda_1^+(\lambda) < \lambda_1^-(\lambda)$. A separate quantitative lower bound on $\lambda_1^-(\lambda)$ is required and is currently formulated as the *Asymptotic Variational Gap Conjecture* in the companion paper [3]; see Part II “Status of the proof” remark (v2.2 revision) for the revised wording. Status of the A1–A8 chain: **conditional reduction**. The finite range $100 \leq \lambda \leq 1,300,000$ supplies strong numerical evidence and a conditional bridge, but theorem-level finite even dominance still requires the odd-sector tail lower bound. Asymptotic even dominance is conditional on the Asymptotic Variational Gap Conjecture; RH itself is therefore conditional on this internal eigenvalue-ordering step.

v2.3 (15 May 2026): Conditioning clarified. A close reading of Connes 2026 §6.5 [4] shows that the Hurwitz-bridge (step A2) is presented there as *Fact 6.4*, an analytical consequence of the classical Slepian–Pollak prolate wave function theory ([Theorem 1 of Ref. 47]Connes2026), not as a conjecture. Connes’ own “remaining steps” (§6.6) are: (i) simple even ground state of QW_Λ (= Asymptotic Even Dominance, addressed in Part II and in the companion proof-only paper); (ii) approximation quality $k_\Lambda \approx \xi_\Lambda$ (= MS2, addressed in the companion Zookeeper paper). Drafts and preprints in this series evolve continuously; the present v2.3 abstract is the current honest position.

MSC 2020: 11M26, 11M36, 47A75, 65G20

Keywords: Riemann Hypothesis, Li coefficients, even dominance, Euler–Maclaurin, proof architecture, open problems

Contents

1	Introduction	4
2	Catalog of Proven Results	4
2.1	Structural Results (R1–R9)	4
2.2	New Results from Parts I–II	5
3	Catalog of Excluded Paths	6
4	The Complete Proof Architecture	7
4.1	The Reduction Chain	8
4.2	The Cumulative Step (A6) — Conditional Reduction	8
4.3	Proof Strategies for A6: Comparative Analysis	9
5	The Localization Table	11
5.1	Li Coefficient Perspective	11
5.2	Even Dominance Perspective	11
6	Systematically Explored Alternatives	12
7	Independent Results	13
8	Remaining Open Problems	14
9	Connections to the Wider Landscape	14
9.1	Connes’ Program	14
9.2	Nyman–Beurling–Baez–Duarte	15
9.3	The Trace-Class Obstruction as Universal Pattern	15
9.4	Selberg Zeta and Yang–Mills	15
10	Assessment and Outlook	15
10.1	What We Have Achieved	15
10.2	The Gap-Closure Path	16
10.3	Analytical Status of the Control Lemmas (v1.1)	17
10.4	The Remaining Distance	17
10.5	Philosophical Reflection: Genesis of the Proof Architecture	17
11	Summary	19

1 Introduction

This paper concludes a three-part investigation of the Riemann Hypothesis through gauge theory, spectral analysis, and arithmetic thermodynamics:

Part I [1]: Foundations and Obstructions — structural results (R1–R9), four critical obstructions (K1–K4), reorientation to Connes’ spectral program.

Part II [2]: Even Dominance — Shift Parity Lemma, computer-assisted finite-range diagnostics (33 values, λ up to 1,300,000), resolvent-damped M1'' framework, Leading-Mode Cancellation ($c = 2 + \sqrt{2}$), Euler–Maclaurin Proposition ($\rho^{\text{EM}} < 0$ for $\lambda \geq 1,200,000$).

Part III (this paper): Conclusio — synthesis, proof architecture, open problems.

The key message is that the Riemann Hypothesis, when approached through the Li criterion and the Weil quadratic form, reduces to a well-characterized analytical condition: an asymptotic variational/gap structure of the coupling differential. The Euler–Maclaurin Proposition (Part II, Proposition 4.11) identifies the correct variational sign for large λ ; combined with 33 computer-assisted finite-range diagnostics up to $\lambda = 1,300,000$, it gives a strong conditional bridge. The theorem-level transfer from this variational diagnostic to the eigenvalue ordering $\lambda_1^+ < \lambda_1^-$ remains the open odd-sector lower-bound step.

2 Catalog of Proven Results

2.1 Structural Results (R1–R9)

The following results are unconditional (they do not assume RH):

R1 Gibbs Framework. The prime hub graph G_σ with Ruelle transfer operator L_σ defines a well-posed Gibbs measure μ_σ for all $\sigma > 1$. The partition function $P(\sigma) = -\log \zeta(\sigma)^{-1}$ diverges as $\sigma \rightarrow 1^+$ (Part I, §2).

R2 Bombieri–Lagarias Decomposition. $\lambda_n = \lambda_n^{(\infty)} + \lambda_n^{(\text{fin})}$, where the archimedean part $\lambda_n^{(\infty)} \sim \frac{n}{2} \log n$ and the finite part satisfies $|\lambda_n^{(\text{fin})}| \leq C\sqrt{n} \log n$ unconditionally (Part I, §3).

R3 Non-Identification. The excess free energy $I_n = \text{KL}(\mu_\sigma^{(n)} \parallel \mu_\sigma)$ satisfies $I_n \geq 0$ unconditionally (as a KL divergence), but $I_n \neq \lambda_n$. The Gibbs normalization annihilates the finite-part contribution (Part I, §5).

R4 Variance Bound. $\text{Var}_{\mu_\sigma}[X_{n,\sigma}] \leq \sum_p (\log p)^{2n} p^{-\sigma}$ for $\sigma > 1$. This bounds the fluctuations of the thermodynamic Li coefficient approximation (Part I, §5).

R5 Arithmetic Uncertainty Relation. For any compactly supported test function h :

$$\text{Var}[h(t)] \cdot \text{Var}[\hat{h}(\omega)] \geq \frac{1}{4} |\langle h, h' \rangle|^2,$$

connecting the spectral and spatial uncertainties of functions tested against the Weil quadratic form (Part I, §5).

- R6 Convexity.** $\frac{d^2}{ds^2} \log \xi(\sigma) > 0$ near $\sigma = 1$. The function $\log \xi$ is convex, not concave, implying λ_1 is the minimum (not maximum) of the derivative profile. This overturns the original monotonicity argument (Part I, §§5 and 11).
- R7 OS Reflection Positivity.** The Osterwalder–Schrader reflection $\theta : t \mapsto -t$ satisfies $\langle f, \theta f \rangle_{\mu_\sigma} \geq 0$ for all test functions f supported on $t \geq 0$. This provides a positivity structure in the thermodynamic formalism (Part I, §5).
- R8 Spectral Census.** The arithmetic kernel K_σ^{sym} on $\ell^2(\mathcal{P}_N)$ has negative inertia index 2 for all computed N and $\sigma \in (1, 1.5)$. Exactly two eigenvalues are negative; the remaining $N - 2$ are positive (Part I, §5).
- R9 Gauge Equivalence.** For any admissible gauge g , the kernel $K_\sigma^g = K_\sigma^{\text{sym}} + \Delta K_g$ satisfies $\text{sign}(\det K_\sigma^g) = \text{sign}(\det K_\sigma^{\text{sym}})$. The gauge freedom is tautological and cannot change the sign structure (Part I, §6).

2.2 New Results from Parts I–II

- N1 Shift Parity Lemma** (Part II, §2). For all $N \geq 3$ and all $r \in (0, 2)$:

$$\lambda_1(D_N(r)) < 0,$$

where $D_N(r) = S_N^+(r) - S_N^-(r)$ is the difference between even- and odd-sector shift matrices. Proved analytically (two-case det/Tr argument for $N = 3$, monotonicity reduction for $N \geq 4$) and certified by interval arithmetic (55,000 subintervals, 60-digit precision, Gershgorin disk bounds).

- N2 Computer-assisted finite-range diagnostics** (Part II). For 33 values of λ from 100 to 1,300,000, interval arithmetic documents a strong negative gap diagnostic (mpmath.iv, 50-digit precision). The reported gap grows monotonically as $\sim \sqrt{\lambda}$, from -2.25 at $\lambda = 100$ to -475.83 at $\lambda = 1,300,000$. As theorem-level even dominance this remains conditional on a validated odd-sector lower bound.

- N3 Universal Even Preference of Individual Primes** (Part II, Corollary 2.5). For every prime p with $r = \log p / \log \lambda \in (0, 2)$:

$$\lambda_1(W_p^+) \leq \lambda_1(W_p^-) + \lambda_1(D_N(r)) < \lambda_1(W_p^-),$$

by the Shift Parity Lemma and Weyl’s inequality.

- N4 Frontier-Prime Mechanism** (Part II, §2). Even dominance is driven by primes $p \sim \lambda$ (normalized shift $r \approx 1$), not by small primes. As λ grows, the number of frontier primes increases as $\pi(\lambda) - \pi(\lambda/e) \sim \lambda / \log \lambda$.

- N5 Hellmann–Feynman Analysis** (Part II, §2). The L -derivative $\partial \Delta / \partial L > 0$ for all $\lambda \geq 80$ through $\lambda = 500$. The gap deepening is driven by new primes entering the sum (prime-counting mechanism), not by the growth of $L = \log \lambda$. This rules out L -monotonicity as a proof strategy.

- N6 Gap Asymptotics** (Part II). The even–odd gap scales as $\Delta(\lambda) \sim -c\sqrt{\lambda}$ (exponential in $L = \log \lambda$), robustly negative for all $\lambda \geq 55$, with $|\Delta| \rightarrow \infty$.

N7 Resolvent-Damped M1'' Framework (Part II). The coupling differential decomposes via second-order perturbation theory. The resolvent-damped energy difference satisfies $(E_{\sin} - E_{\cos})/D(\lambda) = O(1/L) \rightarrow 0$, reducing the remaining analytical step to Fourier analysis.

N8 Leading-Mode Cancellation Lemma (Part II). The overlap differences between even and odd sectors cancel pairwise at $u = 0$ (Fourier orthogonality), with the derivative giving the exact constant $c = 2 + \sqrt{2} \approx 3.414$. The closed-form profiles $S_{\text{sym}}^+(1, 0; u) = \sqrt{2} \sin(\pi u)/\pi$ and $S_{\text{sym}}^+(1, 2; u) = \frac{2}{3\pi}(4 \cos \pi u - 1) \sin \pi u$ are verified symbolically.

N9 Euler–Maclaurin Proposition (Part II, Proposition 4.11). Replace the prime sums defining B_j and g_j by their PNT-asymptotic integrals to obtain the Euler–Maclaurin ratio $\rho^{\text{EM}}(L)$. Then $\rho^{\text{EM}}(L)$ is monotonically decreasing for $L \geq 7$ and crosses zero at $L \approx 13.5$. For $L \geq 14$ (i.e., $\lambda \geq e^{14} \approx 1,200,000$):

$$\rho^{\text{EM}}(L) < 0 \quad (\text{i.e., } E_{\sin}^{\text{EM}} < E_{\cos}^{\text{EM}}),$$

so that $\text{cert_gap}^{\text{EM}}(\lambda) < 0$.

N10 PNT Transfer Lemma and M1'' Reduction (Part II). By Abel summation with the Chebyshev function $\theta(x) = x + O(x e^{-c\sqrt{\log x}})$, the resolvent ratio satisfies $\rho(\lambda) = \rho^{\text{EM}}(\lambda) + O(L e^{-c\sqrt{L}})$. Since $\rho^{\text{EM}} < 0$ for $L \geq 14$ and the error vanishes, this gives an asymptotic variational reduction for the resolvent subdominance mechanism. The remaining theorem-level step is still the odd-sector lower bound needed to convert the variational sign into the eigenvalue inequality $\lambda_1^+ < \lambda_1^-$.

3 Catalog of Excluded Paths

Four independent obstructions kill the original gauge-theoretic proof strategy:

K1 Gibbs Normalization Gap. The $1/P(\sigma)$ -normalization annihilates the finite-part contribution $\lambda_1^{(\text{fin})} = \gamma$. The thermodynamic Li coefficient $m_n(\sigma)$ does not converge to λ_n (Part I, §9).

K2 Sign Change of $F(\sigma)$. The target functional $F(\sigma) = A(\sigma)B(\sigma) - P(\sigma)[C(\sigma) + D(\sigma)]$ changes sign at $\sigma_0 \approx 1.14$. No admissible gauge can restore positivity beyond the boundary layer (1, 1.14) (Part I, §10).

K3 Wrong Monotonicity Direction. The convexity of $\log \xi$ near $s = 1$ means λ_1 is the minimum of the derivative profile, not the maximum. The monotonicity argument runs in the wrong direction (Part I, §11).

K4 Trace-Class Obstruction. The Euler-product operator $\mathcal{L}_s = \text{diag}(p^{-s})$ is trace-class only for $\Re(s) > 1$ ($\sum_p p^{-1/2} = +\infty$). The Fredholm determinant $\det(I - \mathcal{L}_s)$ does not exist at the critical line. This is the fundamental density gap between primes and geodesics (Part I, §12).

Remark 3.1 (The obstructions are independent). Each of K1–K4 is individually fatal. K1 kills the Gibbs identification, K2 kills the gauge approach, K3 reverses the monotonicity argument, K4 kills the spectral determinant construction. Even if three were resolved, the fourth would suffice to block the original strategy.

Additionally, the following paths have been explored and found insufficient:

- **de Branges positivity:** Conrey and Li (2000) showed that the kernel positivity condition is *not* satisfied for the relevant Hilbert spaces associated with ζ .
- **Stieltjes realization:** The sequence $c_n = \lambda_n/n$ is strictly increasing, violating complete monotonicity. The spectral measure under RH lives on $|w| = 1$, not $[0, 1]$.
- **L -monotonicity of the gap:** The Hellmann–Feynman analysis (N5) shows $\partial\Delta/\partial L > 0$ for $\lambda \geq 80$, ruling out a proof of gap deepening via L -differentiation alone.
- **Absolute total positivity (PF_∞) for A_λ (v1.4):** The Fourier multiplier of the prime shift operator is $M(\xi) = -2 \Re[\zeta'/\zeta(\frac{1}{2} + i\xi)]$, an identity from the Weil explicit formula valid without assuming RH. Since M is negative on approximately 49% of its domain (the non-trivial zeros of ζ each generate an oscillation), the prime shift operator is not positive-semidefinite, and A_λ is not PF_∞ . Weak variation diminution in the prime-weighted mean holds and coincides with the frontier-dominance argument (N4 + N6), but there is no absolute “primes smooth everything” principle.
- **Universal commuting operator (v1.4):** An SVD search for symmetric T satisfying $[T, D_N(r)] = 0$ simultaneously on a dense r -grid in $(0, 2)$ shows, for $N \in \{3, 5, 7, 10, 15\}$, that the kernel is one-dimensional and consists solely of multiples of the identity. The second smallest singular value stays bounded below by ~ 0.70 at $N = 15$ and does not decrease with N . The prime shifts at different scales therefore carry no common hidden symmetry. This rules out any Sturm–Liouville-type simplicity argument and confirms that the proof mechanism is arithmetic-statistical (prime-weighted summation of pointwise Shift Parity), not algebraic.

4 The Complete Proof Architecture

We present the logical chain from known mathematics to the Riemann Hypothesis, identifying exactly where each component fits and where the gap remains.

4.1 The Reduction Chain

Step	Statement	Status	Source
A1	Connes' Theorem 6.1: even dominance $\Rightarrow$ zeros of $\hat{\eta}$ real	proven	[5, 4]
A2	Hurwitz: even dominance for $\lambda_n \rightarrow \infty$ suffices	proven	[4]
A3	Even dominance at 33 values, $\lambda \leq 1,300,000$	numerical evidence / conditional finite bridge	Part I
A4	Shift Parity Lemma: each prime favors even	proven	Part I
A5	Frontier-prime mechanism	proven	Part I
A5b	Euler–Maclaurin: $\rho^{\text{EM}}(13) < 0$ (certified by interval arithmetic)	proven	Part I
A6	Cumulative step: $\sum_p W_p$ preserves even preference	variational, v2.2[‡]	Part II
A7	Even dominance for $\lambda_n \rightarrow \infty$ (eigenvalue inequality)	conditional (v2.2) [‡]	Part II
A8	All zeros on $\Re(s) = 1/2$ (RH)	conditional reduction (v2.2)	A1+.

[‡] **v2.2 revision (2026-05-14).** The v2.1 Direct Frontier-Dominance argument establishes $\lambda_{\min}(W^+ - W^-) = -\Theta(\sqrt{\lambda})$ asymptotically (variational), but not the eigenvalue inequality $\lambda_1^+ < \lambda_1^-$. A separate lower bound on $\lambda_1^-(\lambda)$ is needed and is currently the *Asymptotic Variational Gap Conjecture* [3]. Step A4 (Shift Parity Lemma) remains unconditional. Step A3 remains a conditional finite bridge until the odd-sector lower-bound construction is validated.

Historical note on Proposition A6. The v2.1 draft presented the Direct Frontier-Dominance argument as an unconditional closure of A6. The v2.2 revision withdraws that claim: the frontier argument proves the asymptotic variational statement $\lambda_{\min}(W^+ - W^-) = -\Theta(\sqrt{\lambda})$, while the eigenvalue inequality $\lambda_1^+ < \lambda_1^-$ beyond the certified range remains conditional on the *Asymptotic Variational Gap Conjecture*; see Part II, the revised “Status of the proof” remark.

4.2 The Cumulative Step (A6) — Conditional Reduction

The cumulative step was the central challenge: showing that the sum $\sum_{p \leq \lambda} W_p$ preserves the individual-prime even preference established by the Shift Parity Lemma. This is nontrivial because eigenvalue inequalities are not additive for sums of matrices.

Part II (Proposition A6) now separates this challenge into a rigorous finite bridge and an asymptotic variational problem:

- Regime 1** ($\lambda \in [100, 1,300,000]$): 33 computer-assisted finite-range diagnostics, computed with interval arithmetic (50-digit precision), provide strong negative gap evidence at the certified values, theorem-level only after validation of the odd-sector lower bound for λ_1^- .

2. **Regime 2** (asymptotic): the M1'' resolvent framework and the Direct Frontier-Dominance route identify the correct asymptotic variational sign, i.e. $\lambda_{\min}(W^+ - W^-) = -\Theta(\sqrt{\lambda})$. Three subsidiary results control the approximation quality:
 - *Higher-mode decay* (Part II, Lemma B): modes $j \geq 2$ contribute $O(D/L) = o(D)$, via the mean value theorem applied to overlap functions plus PNT-controlled prime sums.
 - *Resolvent truncation* (Part II, Lemma C): the Neumann series converges ($\|R_0 K\| < 1$ for both sectors when $\lambda \geq 500$), so the truncation error is $o(D)$.
 - *PNT transfer* (Part II, Lemma PNT): $|\rho(\lambda) - \rho^{\text{EM}}(\lambda)| = O(L e^{-c\sqrt{L}})$, exponentially small relative to $|\rho^{\text{EM}}|$.
3. **Overlap**: the certified range overlaps the explicit asymptotic threshold, so the remaining gap is no longer a coverage gap but a spectral-direction gap for λ_1^- .

The key structural ingredients that make A6 work are:

- The *Shift Parity Lemma* (A4): provides the diagonal advantage $D(\lambda) = \Theta(\sqrt{\lambda})$.
- The *Leading-Mode Cancellation* (constant $c = 2 + \sqrt{2}$): ensures that the coupling difference is $O(1)$, vastly subdominant to D .
- The *Frontier-prime mechanism* (A5): new primes entering at each λ add fresh even-favoring contributions, ensuring the gap deepens monotonically ($|\text{cert_gap}| \sim c\sqrt{\lambda}$ across 4 orders of magnitude).

v2.1 addition — Direct Frontier-Dominance. In v2.1 (Part II, Section on Direct Frontier-Dominance) a second asymptotic route for A6 is given. Its robust form uses the variational inequality $\lambda_{\min}(D_{\text{total}}) \leq e_1^\top D_{\text{frontier}} e_1 + \|D_{\text{non}}\|$ with the common vector $e_1 = (0, 1, 0)^\top$, rather than adding individual minimum eigenvalues. On $r \in [0.7, 1]$ the scalar frontier quotient satisfies $e_1^\top D_3(r) e_1 \leq -0.4685$; PNT partial summation gives a negative frontier contribution of order $\sqrt{\lambda}$, while Mertens-type estimates bound the non-frontier norm by $O(\lambda^{0.35} + \log \lambda)$. The asymptotic threshold lies within the existing CAP range, so the finite bridge plus the frontier estimate establishes the variational statement $\lambda_{\min}(W^+ - W^-) < 0$ beyond the threshold and obviates the two earlier v2.0 caveats. What remains is the separate odd-sector lower bound needed to upgrade this to $\lambda_1^+ < \lambda_1^-$.

4.3 Proof Strategies for A6: Comparative Analysis

Four candidate strategies for the cumulative step were identified during the research. The following table compares their target mechanism, strengths, key bottleneck, and the role they ultimately played:

	Target	Strength	Bottleneck	Decision
A	Monotone eigenvalue ordering via total positivity (PF_∞)	If it holds: A6 is automatic (order-preserving sums)	Fourier multiplier $M(\xi) = -2\Re[\zeta'/\zeta(\frac{1}{2} + i\xi)]$ is negative on $\sim 49\%$ of its domain; PF_∞ fails as an absolute property	Excluded (Thm. NE-A)
B	Sturm–Liouville structure from a commuting operator	If found: ground state rigidity makes A6 mechanical	SVD over 13 samples of r shows the universal commutator kernel is trivial ($T = c \cdot I$) for N up to 15; σ_2 bounded below, independent of N	Excluded for $N \leq 15$ (Thm. NE-B); full-space open
C	Stability of even ground state as perturbation of prolate operator	Connes’ natural route; “primes = small structured perturbation”	Quantitative perturbation control in the relevant spectral gap regime	Backup
D	Global bound on odd spectral mass via $\text{Tr}(M^k)$	Traces are additive — fits the cumulative nature of A6	From moments to lowest eigenvalue requires sharp concentration	Backup
M1''	Resolvent-damped energy decomposition + PNT transfer	Directly leverages Shift Parity, Leading-Mode Cancellation, and 33 finite-range diagnostics	Requires higher-mode decay (Lemma B), truncation control (Lemma C), and the odd-sector lower bound	Chosen

Why M1'' + PNT + CAP was chosen. Routes A and B would each be “jackpot” solutions — if they work, A6 is immediate — but v1.4 establishes that neither exists as a closed structural mechanism. Route A (total positivity) fails because the Fourier multiplier of the prime shift operator, $M(\xi) = -2\Re[\zeta'/\zeta(\frac{1}{2} + i\xi)]$, is negative on roughly half its domain by the explicit formula; the non-trivial zeros of ζ are precisely the oscillations of M . Route B (commuting operator) fails because an SVD search over a dense r -grid shows that the only operator commuting universally with $D_N(r)$ is the identity, for all $N \leq 15$ —the prime shifts at different scales have structurally different eigenbases and share no common algebraic symmetry.

These ruling-outs are not failures of the proof but structural discoveries: even dominance holds *despite* the absence of these two properties, which clarifies that the mechanism is genuinely arithmetic-statistical rather than algebraic. Route C (prolate perturbation) is the closest surviving alternative: the M1'' framework is structurally similar to treating prime contributions as a perturbation. Route D (trace moments) provides a natural backup, since traces are additive and bypass the eigenvalue non-additivity obstruction. The chosen path combines the strongest available tools:

- *Analytical backbone:* M1'' shows $\rho(\lambda) \rightarrow 0$ via the Leading-Mode Cancellation (Lemma A) and PNT Transfer.
- *Finite-range numerics:* 33 diagnostics cover the finite range conditionally, pending the odd-sector lower-bound validation.
- *Error control:* Lemma B (higher-mode decay) and Lemma C (Neumann convergence) ensure that the resolvent approximation is faithful.

Routes C and D remain available as independent verification paths. In particular, a trace-moment proof of Lemma C would remove the reliance on the Neumann convergence condition $\|R_0 K\| < 1$, providing a fully “soft” analytical argument for Regime 2.

5 The Localization Table

We present two perspectives on how the difficulty of RH is distributed.

5.1 Li Coefficient Perspective

From the Bombieri–Lagarias decomposition and the thermodynamic analysis:

Step	Statement	Status	Method
S1	Define $\xi(s)$, functional equation	proven	definition
S2	Li coefficients via Bombieri–Lagarias	proven	[6]
S3	Decomposition $\lambda_n = \lambda_n^{(\infty)} + \lambda_n^{(\text{fin})}$	proven	Part I
S4	Archimedean dominance: $\lambda_n^{(\infty)} \sim \frac{n}{2} \log n$	proven	Stirling
S5	$\lambda_n \geq 0$ for all $n \geq 1$	OPEN¹	$\iff$ RH
S6	All zeros on $\Re(s) = 1/2$	$\iff$ S5	Li (1997)
S7	Prime counting: $\pi(x) = \text{Li}(x) + O(\sqrt{x} \log x)$	$\Leftarrow$ S6	von Koch

5.2 Even Dominance Perspective

From the Connes program and Part II:

Step	Statement	Status	Method
E1	Weil explicit formula	proven	Weil (1952)
E2	Weil quadratic form QW_λ	proven	Connes (2026)
E3	Even/odd decomposition of QW_λ	proven	Part II
E4	Shift Parity Lemma: each prime favors even	proven	Part II
E5	Gap diagnostics at 33 values, $\lambda \leq 1,300,000$	evidence / conditional	Part II (finite-ra
E5b	Euler–Maclaurin: $\rho^{\text{EM}} < 0$ for $\lambda \geq 1.2\text{M}$	proven	Part II
E6	Cumulative step: even dominance for all λ	conditional reduction	= A6
E7	Connes’ Theorem 6.1 + Hurwitz bridge $\Rightarrow$ RH	conditional on E6	[4]

The two perspectives are related but *not equivalent*: E6 (cumulative even dominance) would *imply* S5 (and hence RH) via Connes’ Theorem 6.1 and Hurwitz’s theorem. However, the converse is unknown — RH could hold for reasons unrelated to even dominance. E6 is a *sufficient* condition, not a characterization. The Shift Parity Lemma bridges the perspectives by showing that the individual-prime structure underlying E6 is algebraically controlled.

6 Systematically Explored Alternatives

During the development of the proof, eleven alternative strategies (BI-1 through BI-11) were systematically explored and evaluated. This exhaustive search is itself a meta-result: *the solution space is mapped, and the chosen path is the only consistent one.*

Approach	Core idea	Result	Status
BI-1: $\det_2$ / Carleman	Extend to Hilbert–Schmidt class	No positivity substitute	Refuted
BI-2: RG fixed point	Weil positivity under scaling	Not formalizable	Discarded
BI-3: Universal property	Physical plausibility $\Rightarrow$ RH	Too abstract	Discarded
BI-5: Lifting programs	Weil $\rightarrow$ Grothendieck $\rightarrow$ RMT	All open except $\mathbb{F}_q$	Not usable
BI-6: Implication ring	$\text{Li} \leftrightarrow \text{Weil} \leftrightarrow \text{NB} \leftrightarrow \text{BD}$	Ring does not exist	Refuted
BI-7: Pöschl–Teller scattering	Gamma structure analog to ξ	No off-line stability	Negative
BI-8: Second-order stability	RH as variational problem	Architecture correct	Realized
BI-9: Soliton models	Topological stability	Coulomb gas unstable	Refuted
BI-10: Building blocks	No-Coordination, shift structure	Successfully reused	Integrated
BI-11: Connes 2026	Even dominance via Q_W	Breakthrough	Core path

The key insight from this exploration: BI-8 (second-order stability) was structurally correct — positivity on a restricted subspace suffices, and the decision falls at second order — but mechanistically wrong: the “entropy” in RH is not thermodynamic (KL divergence) but arithmetic-geometric (trigonometric structure of prime shifts, No-Coordination Lemma). The Shift Parity Lemma replaces the variance/KL term entirely.

7 Independent Results

Several results obtained during this program are of independent interest, regardless of the status of the main proof:

Result	Context	Significance	Status
<i>Positive results (new mathematics)</i>			
Shift Parity Lemma	Arithmetic trigonometry	Every prime individually favors even eigenfunctions	Proved
No-Coordination Lemma	Multiplicative independence	Arithmetic substitute for entropy term	Proved
Leading-Mode Cancellation ($c = 2 + \sqrt{2}$)	Resolvent analysis	Universal pairwise cancellation constant	Proved
OS Reflection Positivity	Spectral theory	Bochner + contractive semigroup	Proved as structural input
Euler–Maclaurin: $\rho^{\text{EM}} < 0$	PNT asymptotics	Global stability anchor	Proved
Hellmann–Feynman: $\partial\Delta/\partial L > 0$	Sensitivity analysis	Gap deepening driven by new primes, not L	Proved
Decomposition: primes drive variational even preference	Spectral decomposition	Archimedean kernel is parity-neutral	Variational / conditional
33 finite-range CAP diagnostics	Interval arithmetic	Conditional bridge pending odd lower bound	Evidence / conditional
<i>Negative results (obstructions and impossibilities)</i>			
Obstructions K1–K4	Dead-end analysis	Maps the space of impossible strategies	Documented
Convexity Trap	Variational methods	KL positivity $\nrightarrow$ Li positivity	Proved
Weyl-law obstruction	Operator theory	No naive Hilbert–Pólya operator exists	Proved
Mellin decay barrier	Functional analysis	Li polynomials cannot approximate Weil class	Proved
Implication ring non-existence	Criterion comparison	Li $\rightarrow$ Weil direct implication unknown	Documented
Coulomb-gas instability	Statistical mechanics	Pure position models unstable off-line	Proved
No absolute PF_∞ for A_λ	Prime shift multiplier	$M(\xi) = -2 \Re[\zeta'/\zeta(\frac{1}{2} + i\xi)]$ negative on $\sim 49\%$ of domain	Proved (v1.4)
No universal commuting operator	Sturm–Liouville search	Kernel of $[T, D_N(r)]$ is $\{c \cdot I\}$ for all tested N	Proved (v1.4)
<i>Structural insights (meta-results)</i>			
Inverted-Hessian structure	YM/TU/RH comparison	$\partial^2 \lambda_n / \partial \delta^2 < 0$: RH is a maximum	Confirmed
Cross-disciplinary catalog	Regularization	11 fields share trace-class obstruction	Documented

These 17 results demonstrate that the program produced substantial contributions independent of the main proof. In particular, the Shift Parity Lemma, the obstruction analysis (K1–K4), and the Mellin decay barrier are publishable independently. The negative results are arguably the most valuable: they protect future researchers from repeating dead-end strategies.

8 Remaining Open Problems

With OP1 sharply reduced but not yet fully closed, we formulate four concrete open problems:

Remark 8.1 (OP1: Cumulative Even Dominance — reduced to the odd-sector lower bound (v2.2)). The cumulative even dominance problem is no longer viewed as a missing coverage argument. Proposition A6 (Part II) now isolates the remaining gap to a single spectral question: after the certified range, both the v2.0 three-regime route and the v2.1 Direct Frontier-Dominance route establish the correct asymptotic variational sign, but a separate quantitative lower bound for $\lambda_1^-(\lambda)$ is still needed to conclude $\lambda_1^+(\lambda) < \lambda_1^-(\lambda)$ for all sufficiently large λ . This remaining step is formulated in Part II and in the companion extract [3] as the *Asymptotic Variational Gap Conjecture*.

Remark 8.2 (OP2 resolved: Simplicity of λ_1^+). The minimum eigenvalue λ_1^+ of QW_λ^+ is certified as simple for all 33 certificate values $\lambda \in [100, 1,300,000]$ by interval arithmetic on the 4×4 even Galerkin block (Part II, Simplicity Lemma). The intra-even gap $\lambda_2^+ - \lambda_1^+$ grows monotonically from 8.69 ($\lambda = 100$) to ≥ 731 ($\lambda = 320,000$), scaling as $\sim \sqrt{\lambda}$. For the asymptotic regime, the analytical gap follows from the Galerkin diagonal asymptotics: W_{11}^+ is the unique most-negative diagonal entry. On the certified range this identifies the global ground state of A_λ as simple and even; asymptotically it supplies the even-sector half of the remaining spectral gap problem.

Conjecture 8.3 (OP3: Analytical Proof of Index Rigidity). *The negative inertia index of K_σ^{sym} is exactly 2 for all N sufficiently large and $\sigma \in (1, \sigma_{\max})$.*

Conjecture 8.4 (OP4: Nuclear Transfer Operator for ξ). *Construct a separable Hilbert space $\mathcal{V}$ and a nuclear family $s \mapsto \mathcal{L}_s$ with $\xi(s)/\xi(0) = \det(I - \mathcal{L}_s)$, spectral radius < 1 for $\Re(s) > 1/2$, and eigenvalue 1 at the zeros. This connects to Deninger’s foliated spaces and the Hilbert–Pólya program.*

Conjecture 8.5 (OP5: Nevanlinna Property of $\log \det_2$). *Is $\log \det_2(I - zR)$ a Nevanlinna function (positive imaginary part in the upper half-plane) if and only if all zeros lie on the real axis? This would be a new RH-equivalent criterion in the $\det_2$ -framework (Part I, §6).*

OP2 is resolved and OP1 has been reduced to a single asymptotic odd-sector lower-bound problem. OP3–OP5 remain structural questions that may either close this last spectral step or provide a deeper conceptual route.

9 Connections to the Wider Landscape

9.1 Connes’ Program

The work in Part II fits directly into Connes’ spectral approach to RH [4]. The Shift Parity Lemma provides analytical evidence for the even dominance hypothesis that Connes uses in Theorem 6.1. The computer-assisted diagnostics at 33 values of λ up to 1,300,000 provide strong finite-range evidence across more than 4 orders of magnitude, conditional on the odd-sector lower-bound validation. The resolvent $M1''$ framework, the Leading-Mode Cancellation Lemma, and the PNT Transfer Lemma establish the asymptotic variational sign; the remaining challenge is the eigenvalue-ordering/lower-bound step.

9.2 Nyman–Beurling–Baez-Duarte

The Li criterion, the Weil positivity criterion, the Nyman–Beurling criterion, and the Baez-Duarte criterion are all equivalent to RH. However, direct implications between them (without passing through RH) are known only in two cases: Weil $\Rightarrow$ Li (Bombieri–Lagarias, 1999) and Nyman–Beurling $\Rightarrow$ Baez-Duarte (trivial). A closed implication ring would constitute a proof.

9.3 The Trace-Class Obstruction as Universal Pattern

The obstruction K4 (trace-class failure at $\Re(s) = 1/2$) is not unique to the Riemann zeta function. The same density gap between primes ($\pi(x) \sim x/\log x$) and hyperbolic geodesics ($\#\{\gamma : l_\gamma \leq T\} \sim e^T/T$) appears in:

- **QFT:** UV divergences regularized by cutoff/renormalization.
- **Statistical mechanics:** Thermodynamic limit requiring finite-volume approximation.
- **NCG:** Connes’ noncommutative geometry uses zeta-function regularization of spectral triples.

In each case, an external parameter (temperature, cutoff, scale) provides regularization. Number theory has no such parameter — only the gamma function, which formally resolves the divergence but has not been realized operator-theoretically.

9.4 Selberg Zeta and Yang–Mills

The Selberg zeta function $Z_\Gamma(s)$ for a compact hyperbolic surface $\Gamma \backslash \mathbb{H}$ has the same Euler-product structure as Riemann’s $\zeta(s)$, but with “primes” replaced by primitive geodesics. The trace-class obstruction K4 does *not* apply: geodesic lengths grow exponentially ($\#\{\gamma : l_\gamma \leq T\} \sim e^T/T$), ensuring trace-class convergence for all $\Re(s) > 1/2$. The Selberg zeta RH is proved for compact surfaces (Selberg, 1956).

This contrast illuminates the arithmetic difficulty: the Riemann zeta function’s prime counting ($\pi(x) \sim x/\log x$) produces exactly the “wrong” density for trace-class arguments. The Even Dominance approach circumvents this obstruction by working with the Weil quadratic form (finite-dimensional Galerkin blocks) rather than infinite-dimensional transfer operators.

A Yang–Mills analogy exists: Kirk’s zero-free strip for Selberg zeta (via Ruelle’s theorem) mirrors the nullstellenfreie Region in prime number theory. The spectral gap in the Laplacian on $\Gamma \backslash \mathbb{H}$ plays the role of the Even–Odd gap in the Weil form. Whether the Shift Parity mechanism has a geometric counterpart for Selberg zeta is an open question of independent interest.

10 Assessment and Outlook

10.1 What We Have Achieved

1. A **complete structural analysis** of the gauge-theoretic approach to RH, with nine unconditional results (R1–R9) and four conclusively excluded paths (K1–K4).

2. The **Shift Parity Lemma**, a purely algebraic result about trigonometric shift matrices, establishing universal even preference of individual primes.
3. A **computer-assisted finite-range diagnostic** of even dominance at 33 values ($\lambda = 100$ to $1,300,000$), providing strong evidence for Connes' spectral hypothesis across more than 4 orders of magnitude, conditional on the odd-sector lower-bound validation.
4. The **conditional reduction of the cumulative step (A6/E6)**: (i) the v2.0 three-regime argument combines computer-assisted finite-range diagnostics ($\lambda \leq 1,300,000$), the M1'' resolvent framework with PNT transfer, and three control lemmas (higher-mode decay, resolvent truncation, PNT error transfer); (ii) the v2.1 Direct Frontier-Dominance argument (Part II, Proposition *prop:direct-frontier*) uses a common Rayleigh test vector, PNT+Mertens partial summation, and finite diagnostic coverage up to the explicit asymptotic threshold. Together they isolate the remaining asymptotic gap to the *Asymptotic Variational Gap Conjecture* for λ_1^- .
5. The **Euler–Maclaurin Proposition**, establishing that $\rho^{\text{EM}} < 0$ for $\lambda \geq 1,200,000$, providing the analytical backbone for Regime 2 of the bridge argument.
6. **Structural insights** (frontier-prime mechanism, Hellmann–Feynman analysis, gap asymptotics) that illuminate the arithmetic mechanism behind even dominance.

10.2 The Gap-Closure Path

The Euler–Maclaurin Proposition (Part II, Proposition 4.11) establishes that the PNT-continuous approximation yields $\rho^{\text{EM}}(L) < 0$ for $L \geq 13$, i.e., for $\lambda \geq e^{13} \approx 442,413$. This is certified by interval arithmetic (mpmath.iv, 60-digit precision); the zero-crossing occurs at $L \approx 12.57$. This means that in the continuous (prime-replaced-by-integral) limit, the leading variational balance favors the even sector for all sufficiently large λ . To transfer this to the actual prime sums, two components are needed:

1. **Computer-assisted diagnostics** ($\lambda \leq 1,300,000$): Strong conditional bridge at 33 values, pending the odd-sector lower-bound validation (Part II, §3).
2. **PNT Transfer Lemma (Part II)**: By Abel summation with the Chebyshev function error $\theta(x) - x = O(x \exp(-c\sqrt{\log x}))$, the actual resolvent ratio satisfies $\rho(\lambda) = \rho^{\text{EM}}(\lambda) + O(L e^{-c\sqrt{L}}) \rightarrow \rho^{\text{EM}}(\lambda)$ as $\lambda \rightarrow \infty$. Since $\rho^{\text{EM}} < 0$ for $L \geq 13$ (interval-arithmetic certified), there exists λ_0 such that $\rho(\lambda) < 0$ for all $\lambda \geq \lambda_0$.

The status of M1'' (resolvent subdominance) is therefore: **proved as an asymptotic variational framework**. In v2.1 the robust Direct Frontier-Dominance proof also bypasses the M1'' route for the variational statement by using a common frontier Rayleigh vector plus finite diagnostic coverage. Combined with the higher-mode decay bound (Lemma B), the resolvent truncation control (Lemma C), and the 33 computer-assisted diagnostics, this yields the current Part II status: a conditional finite bridge and a conditional asymptotic reduction pending the odd-sector lower bound.

10.3 Analytical Status of the Control Lemmas (v1.1)

In v1.0, Lemma B and Lemma C relied on numerically verified constants ($c_{\text{gap}} \geq 0.5$ and $\|R_0 K\| < 1$). In v1.1, both have been replaced by fully analytical arguments:

1. **Lemma B (parity-split gap):** The spectral gap has a two-tier structure derived from the Laplace asymptotics of the auto-overlap profiles. Even-index modes have gaps $\sim 2\sqrt{\lambda}$ (the endpoint values differ), while odd-index modes have gaps $\sim 4\pi^2(j^2 - 1)\sqrt{\lambda}/L^2$ (the endpoint values coincide, so the gap comes from the second derivative). In the odd channel, the same cancellation that reduces the gap also suppresses the couplings B_{1j} , preserving $\sum |\Delta_j| = o(D)$.
2. **Lemma C (sector-difference control):** The global condition $\|R_0 K\| < 1$ has been replaced by the weaker and directly relevant statement that the tail contributions to $E_{\sin}$ and $E_{\cos}$ are nearly equal (by parity symmetry at leading Laplace order), so their difference is $O(D/N_0^2) \ll D$. The tail-restricted Schur test gives $\|R_0^{\text{tail}} K^{\text{tail}}\| \leq L/300 \ll 1$.

No numerical constants remain in the proof of Regime 2. For explicit references: Lemma B, Step 3 (Part II, Lemma 4.10) derives the parity-split gap with $c_{\text{gap}} \geq 4\pi^2/L$ (Corollary 4.10a); Lemma C (Part II, Lemma 4.12) replaces the global Neumann condition by a sector-difference bound of $O(1/N_0^2)$.

10.4 The Remaining Distance

The reduction chain A1–A8 (Section 4) is now sharply localized. Steps A1–A5 are established, and A3 gives a strong finite-range diagnostic bridge at the 33 interval-arithmetic values, but it is theorem-level only after an odd-sector lower-bound validation. The remaining distance is no longer a broad proof-architecture problem but a single asymptotic odd-sector lower-bound problem: upgrading the variational negativity of $W^+ - W^-$ to the eigenvalue inequality $\lambda_1^+ < \lambda_1^-$. In v2.1 an independent Direct Frontier-Dominance proof (Part II, Proposition *prop:direct-frontier*) was added that bypasses the interpolation and PNT-transfer caveats by using a common frontier Rayleigh vector and finite diagnostic coverage; the v2.2 revision clarifies exactly what this argument proves and what it still leaves open.

10.5 Philosophical Reflection: Genesis of the Proof Architecture

The proof architecture did not arise from a linear deduction but from a systematic exploration of structurally related stability problems. The starting point (BI-8) was the observation that three *a priori* unrelated contexts — Yang–Mills theory, turbulence theory, and the Riemann Hypothesis — share the same formal architecture: in each case, the positivity of a second variation decides stability or regularity. In Yang–Mills and turbulence, this positivity follows from an entropy or variance term. In the arithmetic context, no such probabilistic degree of freedom exists.

Early numerical experiments revealed that the naïve transfer to RH appears in *inverted form*: the Li coefficients λ_n have a *maximum* at the critical configuration ($\delta = 0$, i.e., RH), not a minimum. The correct functional is therefore $F_{\text{RH}} = -\lambda_n$, with minimum at RH. This inversion pointed to a deeper structural mechanism: stability here cannot come from convexity in the classical (KL-divergence) sense.

Soliton and Coulomb-gas models (BI-9) were tested as topological alternatives. All failed: the Coulomb gas showed $\partial^2 E / \partial \delta^2 < 0$ universally, meaning logarithmic repulsion *destabilizes* the on-line configuration. The decisive insight was that models capturing only the *positions* of zeros, without their analytic structure (entirety, functional equation, Euler product), cannot provide stability. Local potentials are insufficient.

A critical review of earlier work (BI-10) identified three durable building blocks: the Weyl-law obstruction (no naïve Hilbert–Pólya operator), the Convexity Trap (KL positivity does not imply Li positivity), and the No-Coordination Lemma (multiplicative independence of primes prevents coherent phase synchronization). The last of these proved to be the arithmetic replacement for the missing entropy term.

The breakthrough came with Connes’ 2026 survey [4] (BI-11), which provided a completely different framework. Instead of density arguments or global positivity, the Weil form is treated as a finite operator on a Paley–Wiener space. Connes’ Theorem 6.1 reduces RH to an even-dominance condition, and a Hurwitz argument shows this condition is needed only along a sequence $\lambda_n \rightarrow \infty$. This precisely realized the “positivity on a restricted subspace” anticipated by BI-8.

The proof architecture of Part II is the direct implementation of these insights. The Shift Parity Lemma formalizes the constructive interference of even modes and destructive interference of odd modes under prime shifts. The resolvent analysis (M1’’) and the subsidiary lemmas on mode decay and truncation show that this local even preference has the correct asymptotic variational sign under summation. Proposition A6 isolates the finite bridge and the remaining asymptotic spectral step; it does not by itself close the odd-sector eigenvalue-ordering problem.

In retrospect, BI-8 correctly predicted the architecture but misidentified the mechanism. The “entropy” of the Riemann Hypothesis is not thermodynamic but geometric-arithmetic: it arises from the trigonometric structure of prime shifts and the non-coordinability of their phases. The present work can therefore be understood as a distillation of this development, in which the heuristic elements have been separated from the formal, verifiable arguments and the remaining odd-sector lower-bound problem is now explicit.

The v1.4 ruling-outs of Routes A and B sharpen this picture further. One might have expected the primes, acting through A_λ , to obey either an absolute smoothing principle (total positivity) or a hidden symmetry principle (commuting operator). Neither holds: the Fourier multiplier $M(\xi)$ of the prime shift operator oscillates through zero because the non-trivial zeros of ζ are its poles, and no fixed algebraic structure coordinates the shifts at different scales. What remains is the pointwise Shift Parity Lemma—an algebraic identity of trigonometric overlaps—combined with the statistical concentration of prime weight at the frontier $p \sim \lambda$. The absence of a global smoothing or symmetry principle is not a gap to be filled; it is the correct structure. Even dominance arises because the primes are multiplicatively independent, not despite that fact.

The spectral gap bound (Lemma B) and the tail truncation control (Lemma C) have been derived analytically in v1.1 of Part II. Lemma B now uses a parity-split gap structure: even-index modes have gaps $\sim 2\sqrt{\lambda}$ while odd-index modes have gaps $\sim 4\pi^2(j^2 - 1)\sqrt{\lambda}/L^2$ (the leading Laplace term cancels for same-parity modes, but the couplings are equally suppressed). Lemma C replaces the global $\|R_0 K\| < 1$ condition with a sector-difference argument: the tail contributions to $E_{\sin}$ and $E_{\cos}$ are nearly equal, so their difference is $O(D/N_0^2) \ll D$. Both arguments are fully analytical and require no numerical input beyond the PNT error bound.

11 Summary

Contribution	Status	Pap
Thermodynamic formalism for ζ	established	I
9 unconditional structural results (R1–R9)	proven	I
4 excluded paths (K1–K4)	proven	I
Localization: all of RH $\equiv$ Step S5	proven	I
Shift Parity Lemma	proven (analytic)	II
Gap diagnostics (33 values, $\lambda \leq 1,300,000$)	evidence / conditional	II
Frontier-prime mechanism	proven	II
Hellmann–Feynman: prime-counting drives gap	proven	II
Leading-Mode Cancellation ($c = 2 + \sqrt{2}$)	proven	II
Euler–Maclaurin: $\rho^{\text{EM}}(13) < 0$ (interval-arithmetic certified)	proven	II
Resolvent M1'': $(E_{\sin} - E_{\cos})/D = O(1/L)$	proven	II
Gap asymptotics: $ \text{cert_gap} \sim c\sqrt{\lambda}$	numerical	II
M1'' (resolvent subdominance)	variational reduction	II
Higher-mode decay (Lemma B)	proved	II
Tail truncation (Lemma C)	proved (analytical)	II
PNT transfer error (Lemma PNT)	proved with effective C	II
Cumulative step (A6)	conditional reduction (v2.2)	II
Riemann Hypothesis	conditional reduction (v2.2)	I+II

The current status of A6 and RH rests on the three-regime bridge argument (Part II, Proposition A6) together with the v2.1 Direct Frontier-Dominance analysis: computer-assisted certificates cover $\lambda \leq 1,300,000$ (Regime 1), and the asymptotic M1''/PNT and frontier analyses identify the correct variational sign beyond the explicit threshold. In v1.1, Lemma B was upgraded to a parity-split gap bound (analytical) and Lemma C was reformulated as a sector-difference control (replacing the stronger but unnecessary $\|R_0K\| < 1$ condition). In v2.1 the robust Direct Frontier-Dominance proof removed the two older v2.0 caveats. The v2.2 revision clarifies that the remaining task is a separate lower bound for $\lambda_1^-(\lambda)$; RH is therefore currently a *conditional reduction*, not an unconditional closure.

Acknowledgments. Computations were performed on a Hetzner CCX13 cloud server (2 dedicated vCPU, 8 GB RAM). The interval arithmetic library mpmath was developed by Fredrik Johansson.

AI Disclosure

This manuscript was developed in extensive collaboration with AI language models (Claude, Anthropic; Copilot, Microsoft/GitHub; Gemini, Google DeepMind). AI contributed to conceptual exploration, analytical work, literature review, writing, and critical review.

The author, Lukas Geiger, conceived the research direction, provided domain expertise, and exercised final editorial authority. The author assumes full and sole scholarly responsibility for all content, including any errors.

References

- [1] L. Geiger, *The Riemann Hypothesis: Foundations, Obstructions, and Reorientation*, 2026.
- [2] L. Geiger, *The Riemann Hypothesis: Even Dominance of the Weil Quadratic Form*, 2026.
- [3] L. Geiger, *A conditional proof of the Riemann hypothesis via even dominance of the Weil quadratic form*, Zenodo, 2026, DOI: 10.5281/zenodo.20011910.
- [4] A. Connes, *The Riemann Hypothesis: Past, Present and a Letter Through Time*, arXiv:2602.04022v1, 2026.
- [5] A. Connes and W. D. van Suijlekom, *Quadratic Forms, Real Zeros and Echoes of the Spectral Action*, arXiv:2511.23257, 2025.
- [6] E. Bombieri and J. C. Lagarias, *Complements to Li’s criterion for the Riemann hypothesis*, J. Number Theory **77** (1999), 274–287.
- [7] X.-J. Li, *The positivity of a sequence of numbers and the Riemann hypothesis*, J. Number Theory **65** (1997), 325–333.
- [8] A. Weil, *Sur les “formules explicites” de la théorie des nombres premiers*, Comm. Sém. Math. Univ. Lund (1952), 252–265.
- [9] J. B. Keiper, *Power series expansions of Riemann’s ξ function*, Math. Comp. **58** (1992), 765–773.
- [10] C. Deninger, *On the nature of the “explicit formulas” in analytic number theory*, in: Number Theoretic Methods, Developments in Mathematics, vol. **8**, Kluwer Academic Publishers (2002), 97–118.
- [11] A. Connes, *Trace formula in noncommutative geometry and the zeros of the Riemann zeta function*, Selecta Math. **5** (1999), 29–106.
- [12] J. B. Conrey and X.-J. Li, *A note on some positivity conditions related to zeta and L-functions*, Int. Math. Res. Not. **2000**, no. 18, 929–940.